


# Implementing Processes

## Chapter 5

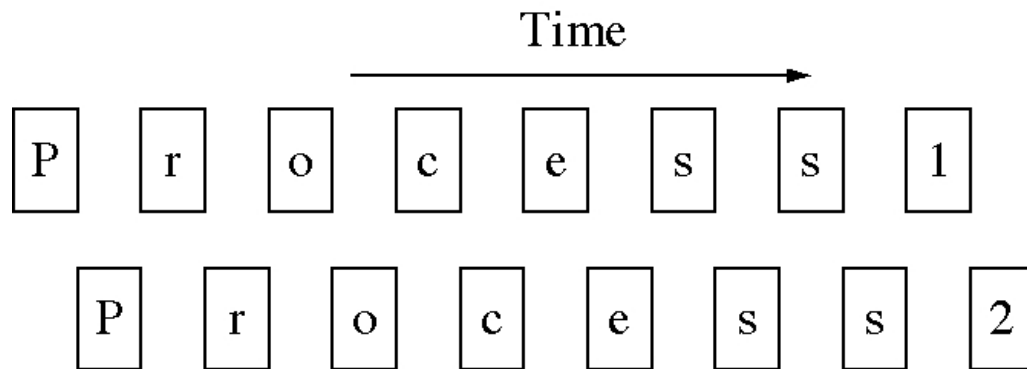
# Key concepts in chapter 5

- Simple operating systems (SOS)
  - Implementation of processes
  - System initialization
  - Process switching
  - System call handling
  - Waiting in the OS
- Operating systems as table and event managers

# Implementing processes by interleaving the processor



**(a)** Serial execution of processes 1 and 2

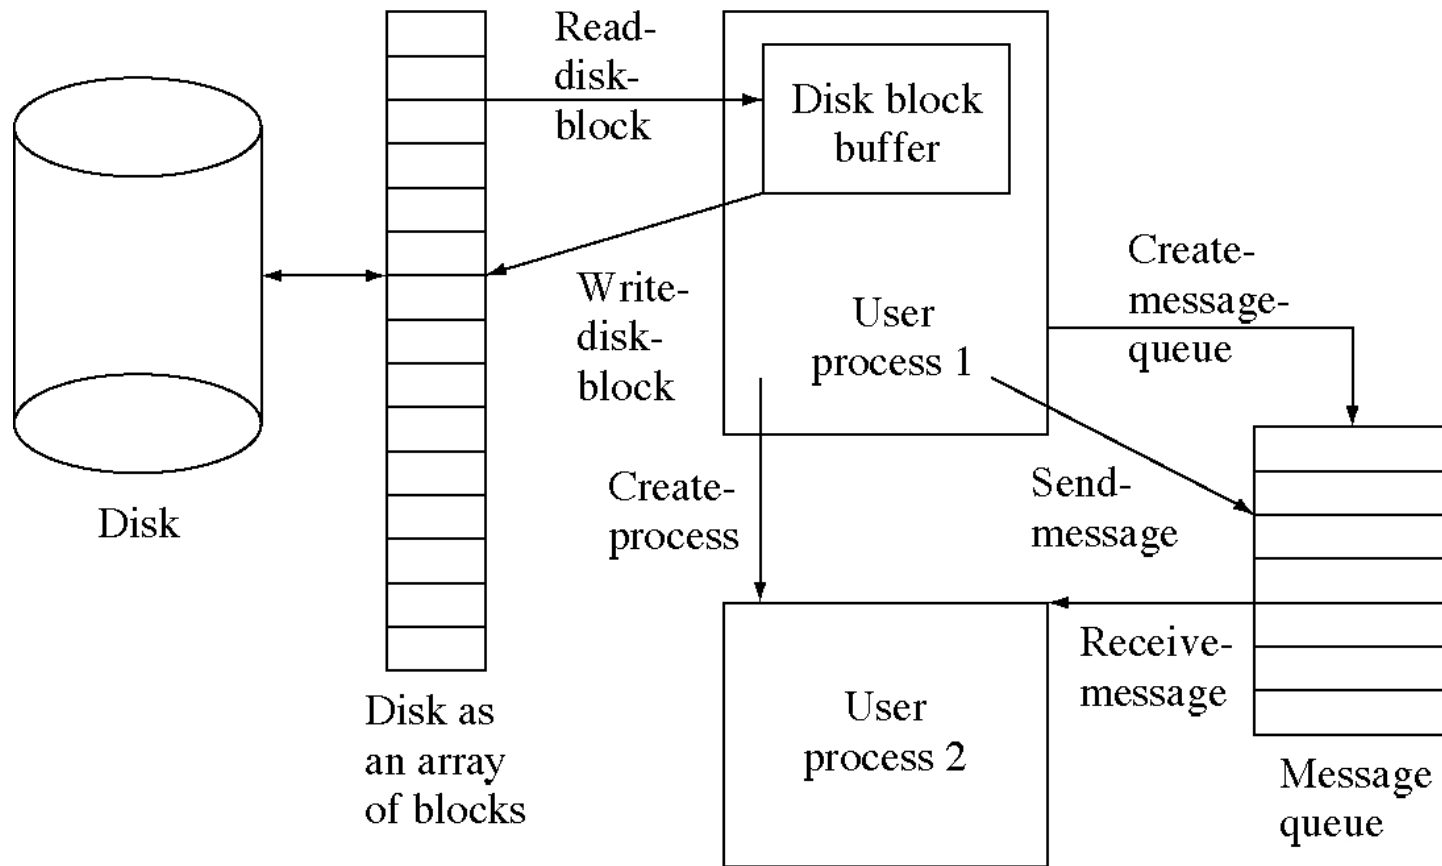


**(b)** Interleaved execution of processes 1 and 2

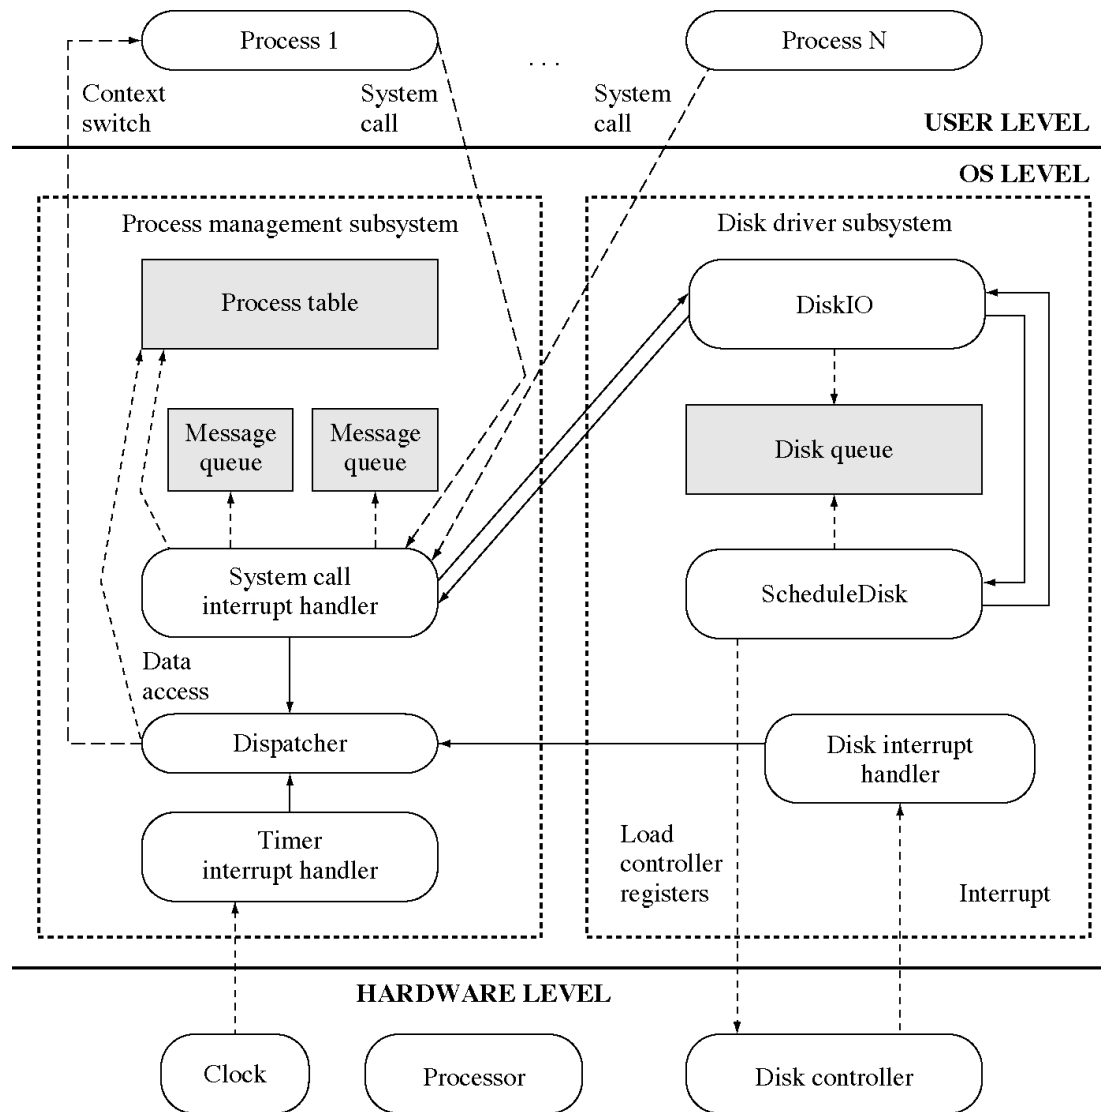
# System call interface of the Simple OS (SOS)

- `CreateProcess(int firstBlock, int nBlocks)`
- `ExitProcess(int exitCode)`
- `CreateMessageQueue()`
- `SendMessage(int mqid, int *msg)`
- `ReceiveMessage(int mqid, int *msg)`
- `ReadDiskBlock(int block, char *buffer)`
- `WriteDiskBlock(int block, char *buffer)`

# SOS objects and operations



# SOS architecture



# System constants (1 of 2)

- ```
// Boolean values
enum { False=0, True=1 };

// hardware constants (determined by the hardware)
const int DiskBlockSize      = 4096;
const int NumberOfRegisters = 32;

// system constants (we can change these constants
//   to tune the operating system)
const int SystemStackSize = 4096; // bytes
const int ProcessSize = 512*1024; // bytes
const int TimeQuantum = 100000;
    // 100000 microseconds = 100 milliseconds
const int MessageSize = 8; // 8 words = 32 bytes
const int InitialProcessDiskBlock = 4341; //disk block #
const int EndOfFreeList = -1;
```

# System constants (2 of 2)

- ```
// system limits (we can change these)
const int NumberOfProcesses = 20;
const int NumberOfMessageQueues = 20;
// The total number of message buffers
const int NumberOfMessageBuffers = 100;
// event handler offsets (determined by the hardware)
const int SystemCallHandler = 0;
const int TimerHandler = 4;
const int DiskHandler = 8;
const int ProgramErrorHandler = 12;
// system call numbers (arbitrary numbers,
//   as long as they are all different)
const int CreateProcessSystemCall = 1;
const int ExitProcessSystemCall = 2;
const int CreateMessageQueueSystemCall = 3;
const int SendMessageSystemCall = 4;
const int ReceiveMessageSystemCall = 5;
const int DiskReadSystemCall = 6;
const int DiskWriteSystemCall = 7;
```



# Process global data

- ```
struct SaveArea {
    int ia, psw, base, bound,
    reg[NumberOfRegisters];
};
enum ProcessState { Ready, Running, Blocked };
typedef int Pid;

struct ProcessDescriptor {
    int slotAllocated;
    int timeLeft; // time left from the last time
    slice
    ProcessState state; // ready, running or blocked
    SaveArea sa; // register save area
};

int current_process
int SystemStack[SystemStackSize];
ProcessDescriptor pd[NumberOfProcesses];
    // pd[0] is the system
```

# Message global data

- `typedef int MessageBuffer[MessageSize];`

```
MessageBuffer
```

```
message_buffer[NumberOfMessageBuffers];
```

```
int free_message_buffer;
```

```
int message_queue_allocated[NumberOfMessageQueues];
```

```
Queue<int> * message_queue[NumberOfMessageQueues];
```

```
struct WaitQueueItem {
```

```
    Pid pid;
```

```
    char * buffer;
```

```
};
```

```
Queue<WaitQueueItem *> *
```

```
wait_queue[NumberOfMessageQueues];
```

# Interrupt vector area

- `char * SystemCallVector`  
    `= &SystemCallInterruptHandler;`
- `char * TimerVector`  
    `= &TimerInterruptHandler;`
- `char * DiskVector`  
    `= &DiskInterruptHandler;`
- `char * ProgramErrorVector`  
    `= &ProgramErrorInterruptHandler`

# Disk global data

- `int process_using_disk;`

```
struct DiskRequest {  
    int command;  
    int disk_block;  
    char * buffer;  
    int pid;  
};
```

```
Queue<DiskRequest *> disk_queue;
```

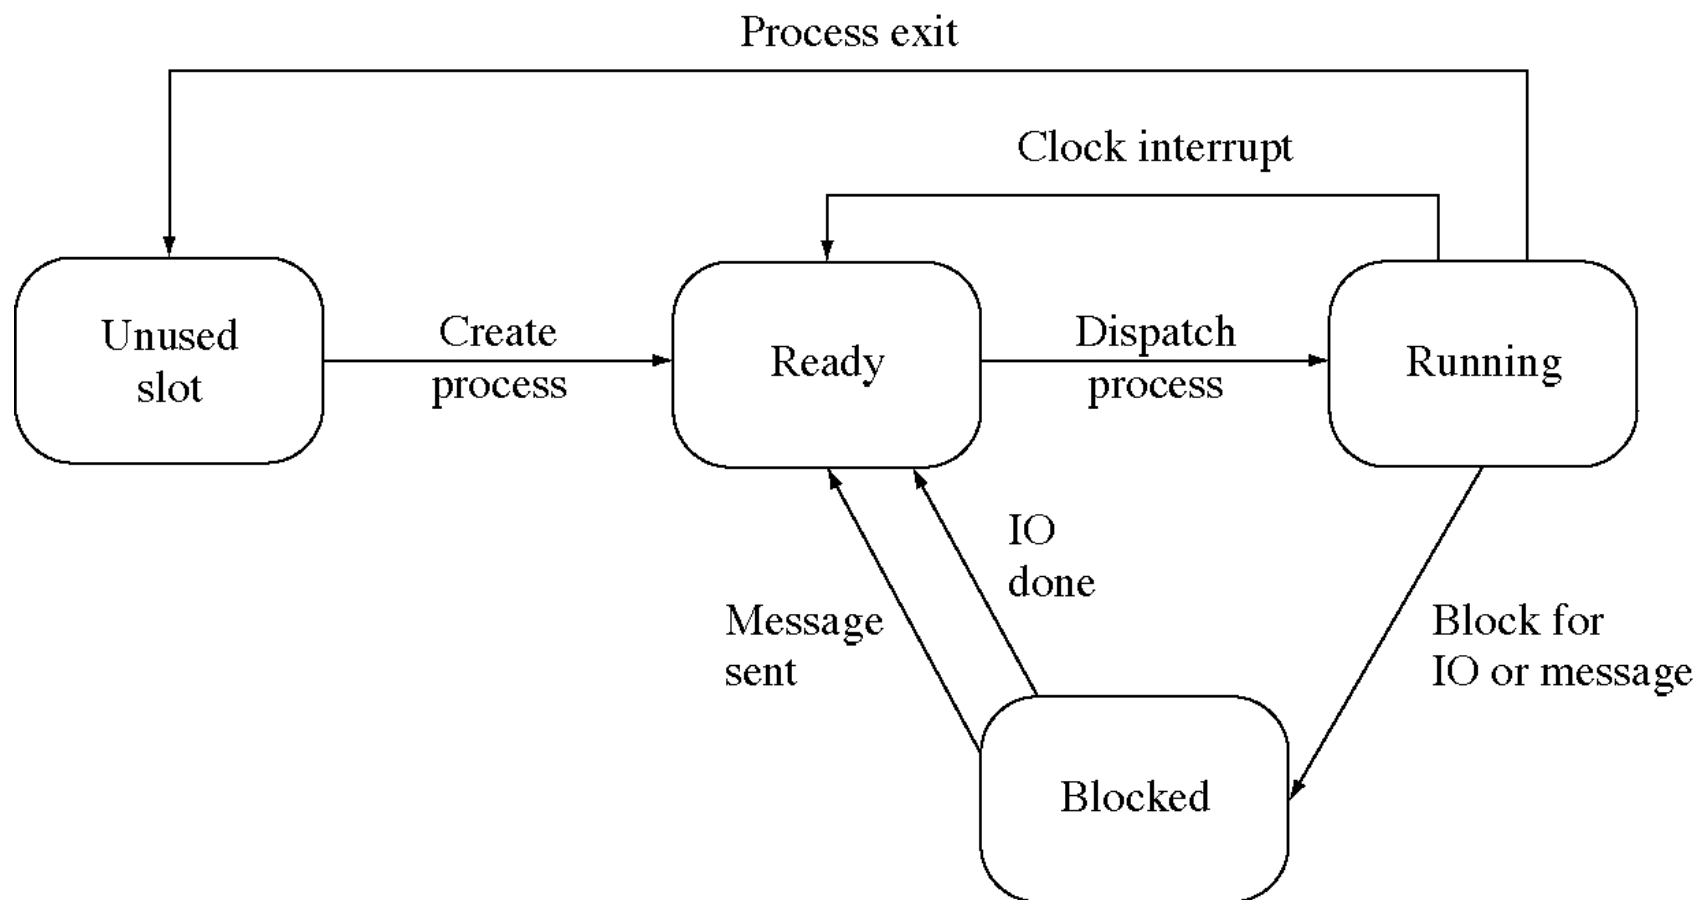
# Message buffer allocation procedures

- ```
int GetMessageBuffer( void ) {  
    // get the head of the free list  
    int msg_no = free_message_buffer;  
    if( msg_no != EndOfFreeList ) {  
        // follow the link to the next buffer  
        free_message_buffer = message_buffer[msg_no][0];  
    }  
    return msg_no;  
}  
  
void FreeMessageBuffer( int msg_no ) {  
    message_buffer[msg_no][0] = free_message_buffer;  
    free_message_buffer = msg_no;  
}
```

# Process creation

- ```
int CreateProcessSysProc(int first_block,int n_blocks) {  
    int pid;  
    for( pid = 1; pid < NumberOfProcesses; ++pid ) {  
        if( pd[pid].slotAllocated ) break;  
    }  
    if( pid >= NumberOfProcesses ) { return -1; }  
    pd[pid].slotAllocated = True;  
    pd[pid].state = Ready;  
    pd[pid].sa.base = pid * ProcessSize;  
    pd[pid].sa.bound = ProcessSize;  
    pd[pid].sa.psw = 3; // user mode, interrupts enabled  
    pd[pid].sa.ia = 0;  
    char * addr = (char *)(pd[pid].sa.base);  
    for( i = 0; i < n_blocks; ++i ) {  
        while( DiskBusy() ) ;  
        IssueDiskRead( first_block + i, addr, 0/*no int*/);  
        addr += DiskBlockSize;  
    }  
    return pid;  
}
```

# Process states



# Dispatcher

- ```
void Dispatcher( void ) {  
    current_process = SelectProcessToRun();  
    RunProcess( current_process );  
}
```



# Select a process to run

- ```
int SelectProcessToRun( void ) {  
    static int next_proc = NumberOfProcesses;  
    if( current_process > 0  
        && pd[current_process].state == Ready  
        && pd[current_process].timeLeft > 0 ) {  
        pd[current_process].state = Running;  
        return current_process;  
    }  
    for( int i = 1; i < NumberOfProcesses; ++i ) {  
        if( ++next_proc >= NumberOfProcesses )  
            next_proc = 1;  
        if( pd[next_proc].slotAllocated  
            && pd[next_proc].state == Ready ) {  
            pd[next_proc].timeLeft = TimeQuantum;  
            pd[next_proc].state = Running;  
            return next_proc;  
        }  
    }  
    return -1;  
}
```

# Run a process

- ```
void RunProcess( int pid ) {  
    if( pid >= 0 ) {  
        SaveArea * savearea = &(pd[pid].sa);  
        int quantum = pd[pid].timeLeft;  
        asm {  
            load    savearea+0, iia  
            load    savearea+4, ipsw  
            load    savearea+8, base  
            load    savearea+12, bound  
            loadall savearea+16  
            load    quantum, timer  
            rti  
        }  
    } else {  
        waitLoop: goto waitLoop;  
    }  
}
```

# The system stack

- All code needs a stack
  - the compiler expects to have one for us by the running program
- We play tricks on the C++ compiler and fiddle with its stack

# Timer interrupt handler

- ```
void TimerInterruptHandler( void ) {
    if( current_process > 0 ) {
        SaveArea * savearea = &(pd[current_process].sa);
        asm {
            store    iia, savearea+0
            store    ipsw, savearea+4
            store    base, savearea+8
            store    bound, savearea+12
            storeall  savearea+16
            load      SystemStack+SystemStackSize, r30
        }
    }
    pd[current_process].timeLeft = 0;
    pd[current_process].state = Ready;
    Dispatcher();
}
```

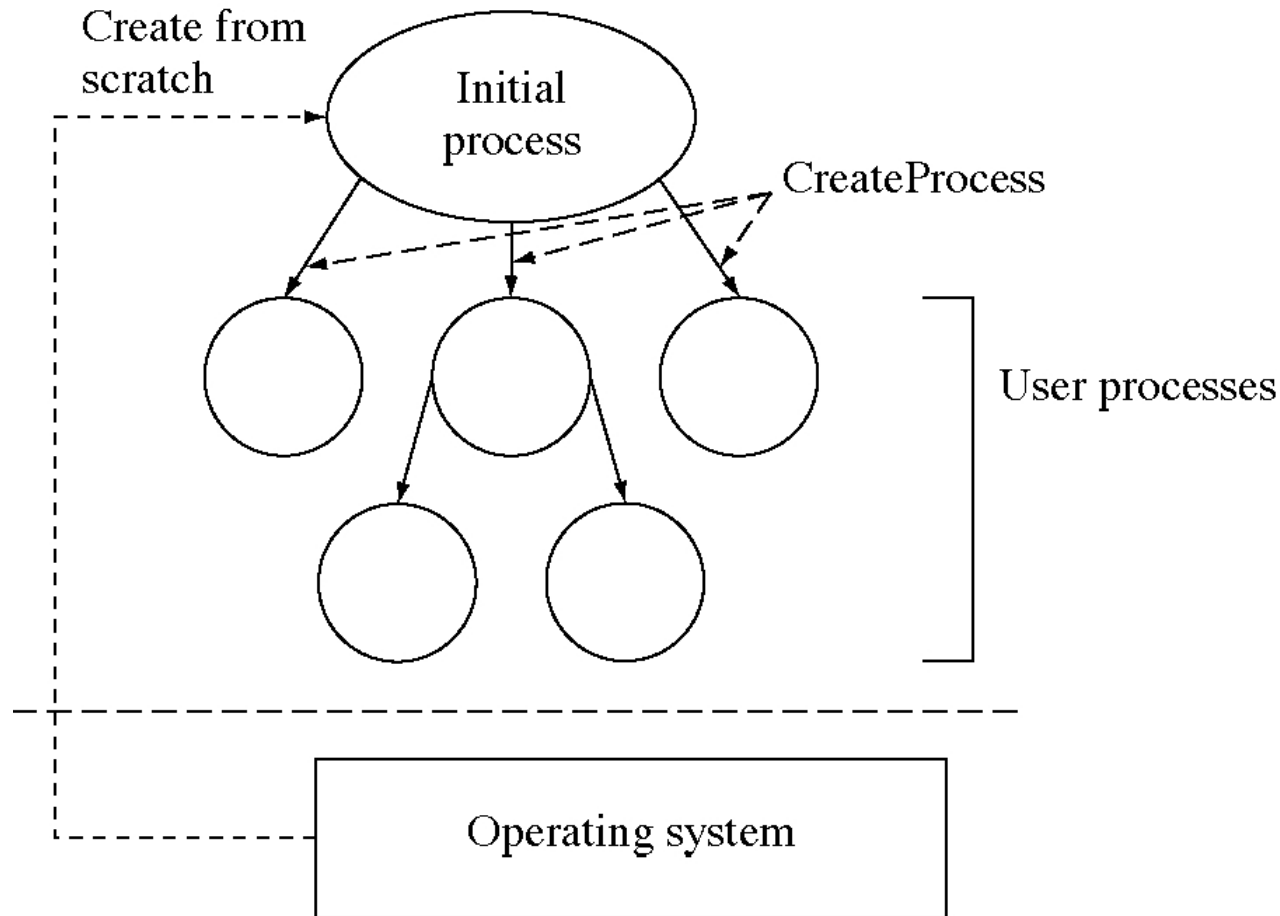
# System initialization

- ```
int main( void ) {  
    asm{ load  SystemStack+SystemStackSize,r30 }  
    asm{ load  &SystemCallVector,iva }  
    pd[0].slotAllocated = True;  
    pd[0].state = Blocked;  
    for( i = 1; i < NumberOfProcesses; ++i )  
        pd[i].slotAllocated = False;  
    (void)CreateProcessSysProc(  
        InitialProcessDiskBlock, 1 );  
    for( i = 0; i < (NumberOfMessageBuffers-1); ++i )  
        message_buffer[i][0] = i + 1;  
    message_buffer[NumberOfMessageBuffers-1][0]  
        = EndOfFreeList;  
    free_message_buffer = 0;  
    for( i = 0; i < NumberOfMessageQueues; ++i )  
        message_queue_allocated[i] = False;  
    process_using_disk = 0;  
    Dispatcher();  
}
```

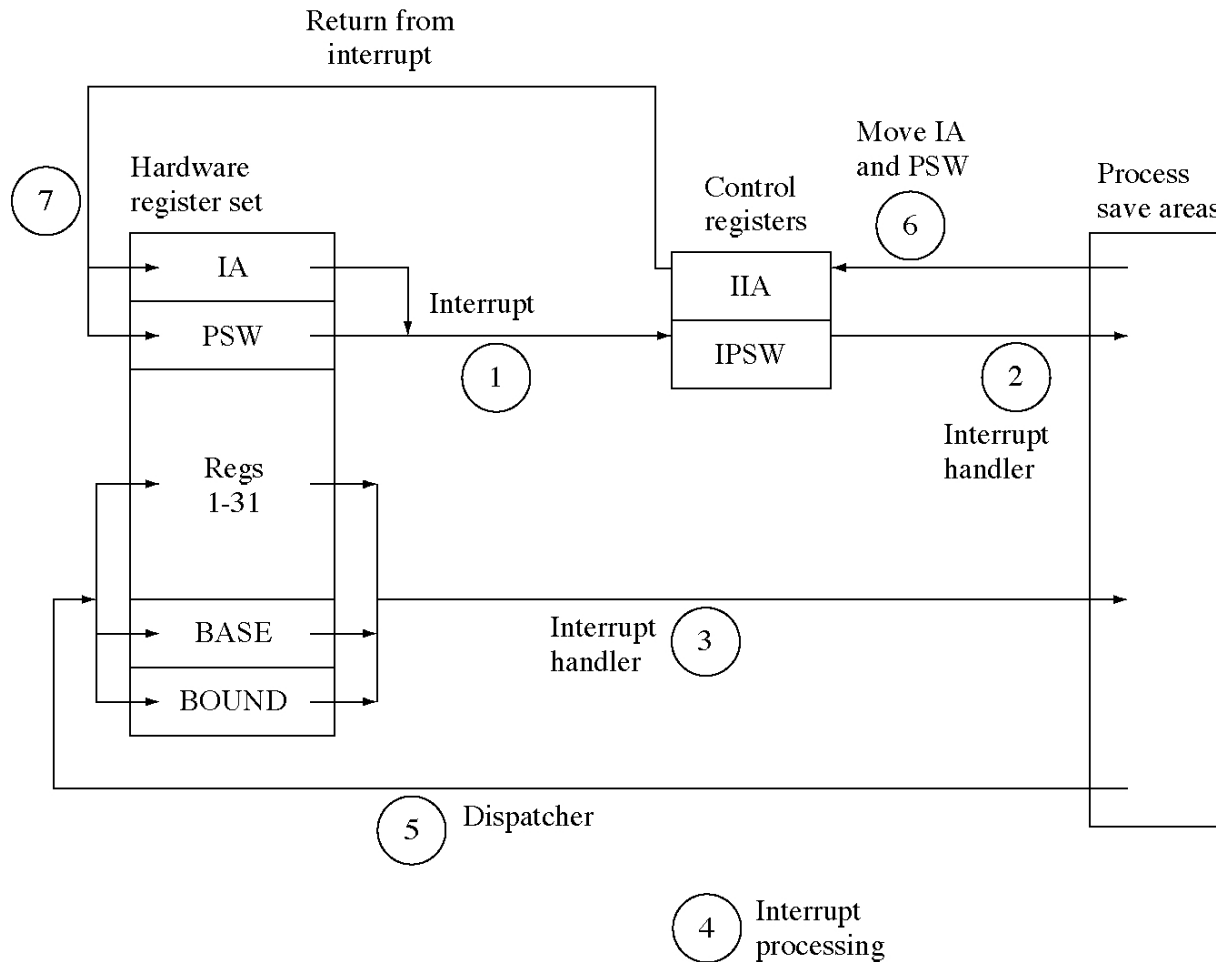
# The initial process

- ```
void main() {  
    // start the two counting processes  
    (void)CreateProcess(  
        UserProcessA, UserProcessASize );  
    (void)CreateProcess(  
        UserProcessB, UserProcessBSize );  
    // Nothing else for this process to do.  
    // We haven't implemented a Wait system call,  
    // so just exit.  
    ExitProcess( 0 );  
}
```

# Initial process creates other processes

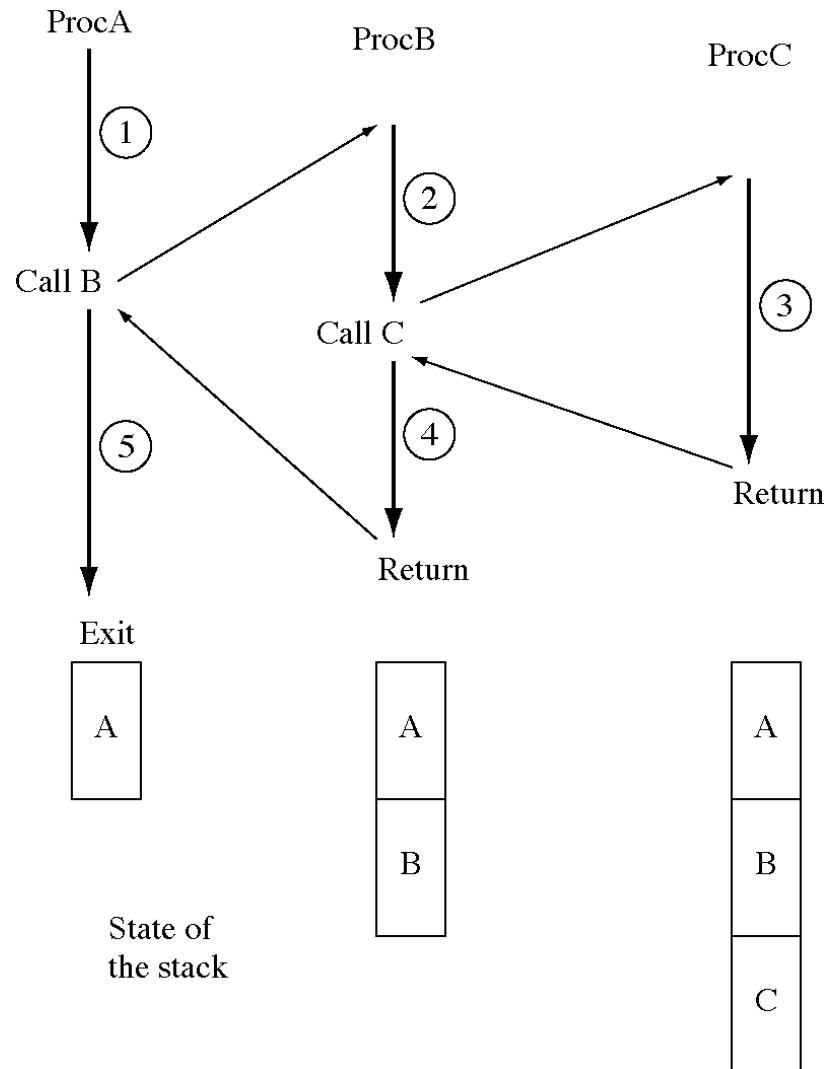


# Process switching

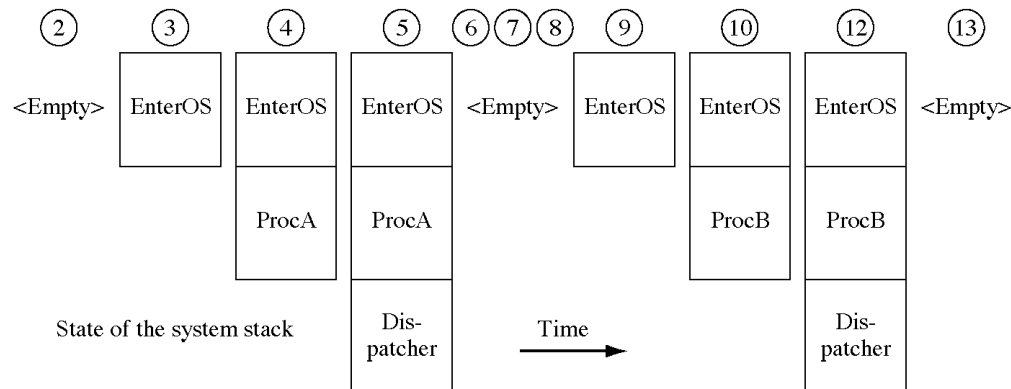
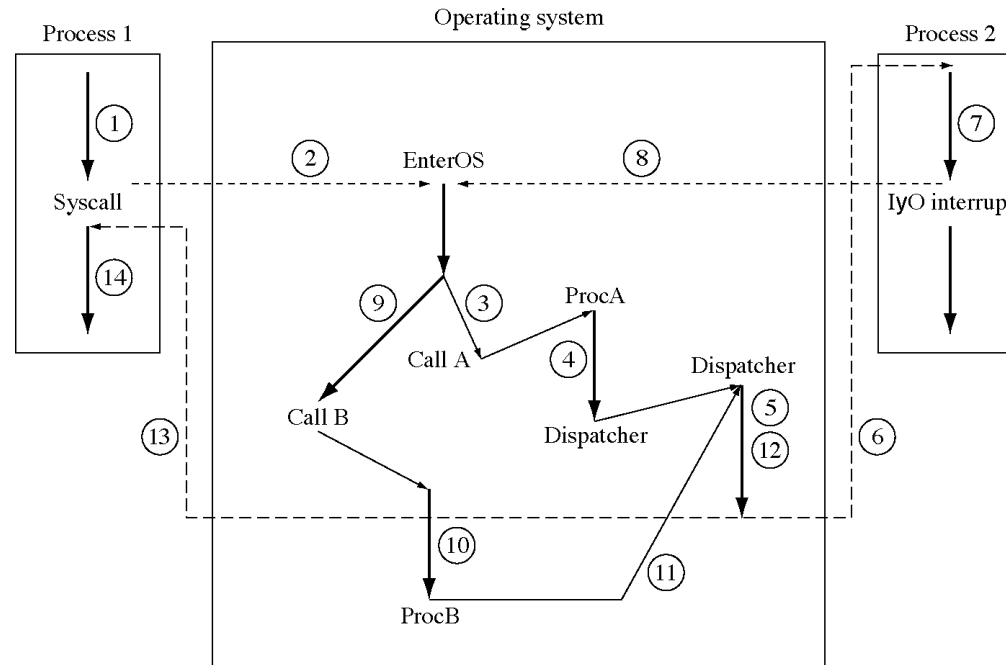




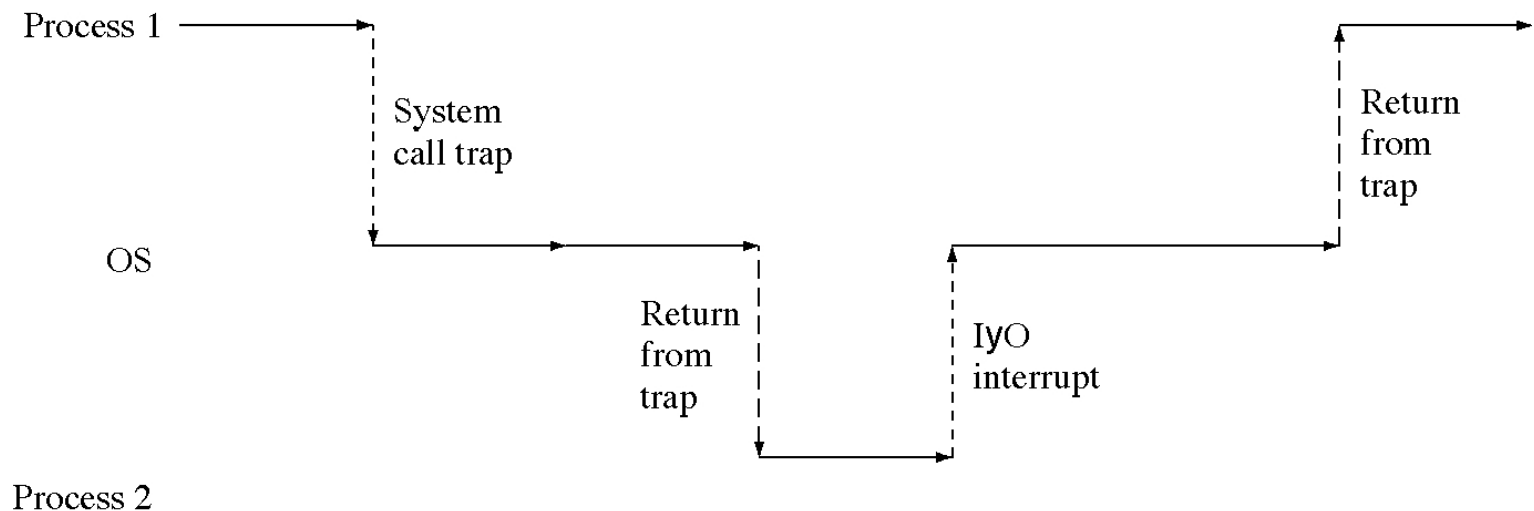
# Flow of control within a process



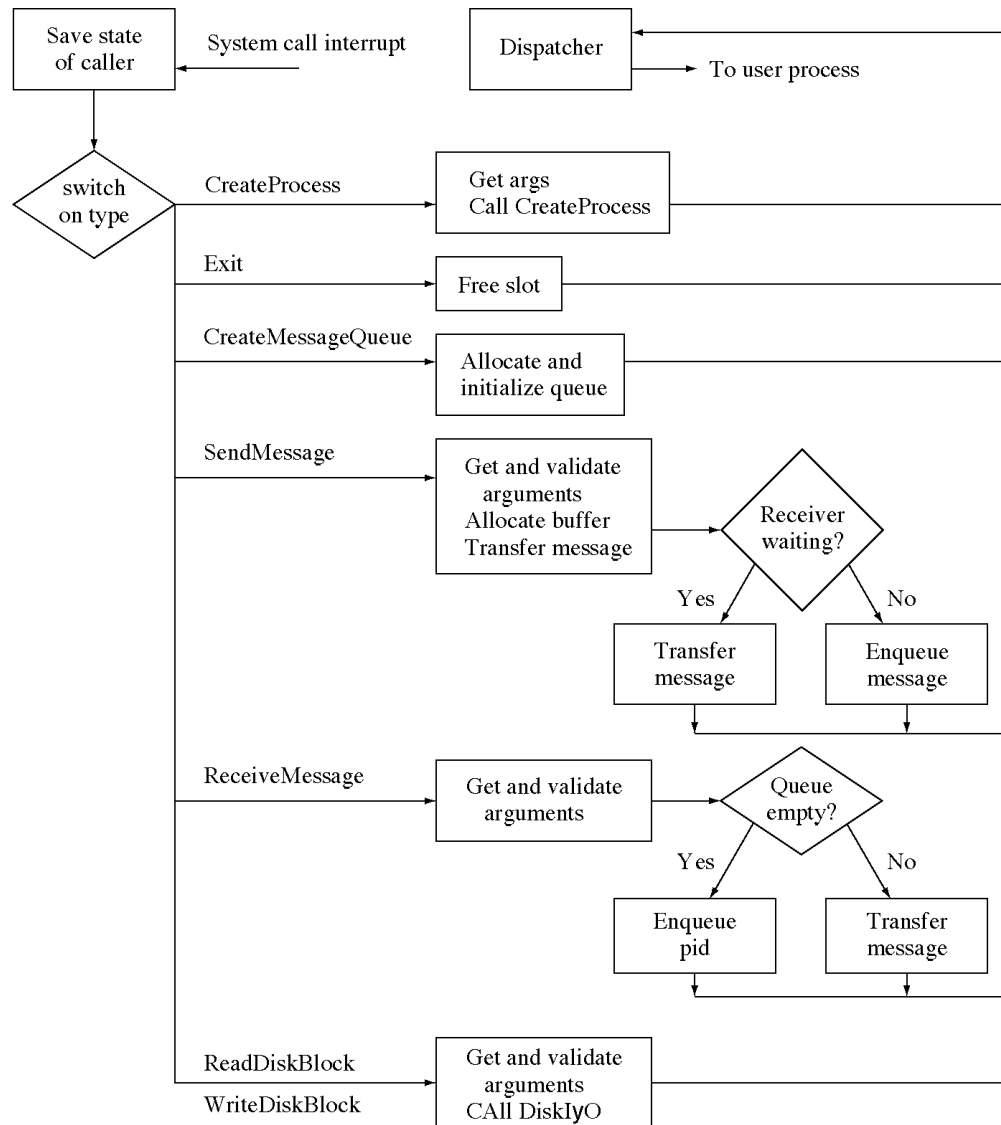
# Process switching control flow



# Flow of control during process switching (another view)



# System call interrupt handler



# System call interrupt handler (1 of 6)

- ```
void SystemCallInterruptHandler( void ) {  
    SaveArea * savearea = &(pd[current_process].sa);  
    int saveTimer;  
    asm {  
        store    timer, saveTimer  
        load     #0, timer  
        store    iia, savearea+0  
        store    ipsw, savearea+4  
        store    base, savearea+8  
        store    bound, savearea+12  
        storeall  savearea+16  
        load     SystemStack+SystemStackSize, r30  
    }  
    pd[current_process].timeLeft = saveTimer;  
    pd[current_process].state = Ready;  
    int system_call_number;  
    asm { store r8, system_call_number }  
    switch( system_call_number ) {
```

# System call interrupt handler (2 of 6)

- `case CreateProcessSystemCall:`  
    `// get the system call arguments from the`  
    `registers`  
    `int block_number; asm { store r9,block_number }`  
    `int number_of_blocks;`  
    `asm { store r10,number_of_blocks }`  
    `// put the return code in R1`  
    `pd[current_process].sa.reg[1]`  
    `= CreateProcessSysProc(`  
        `block_number,number_of_blocks);`  
    `break;`  
`case ExitProcessSystemCall:`  
    `char * return_code; asm { store r9,return_code }`  
    `// we don't save the return code in this OS so`  
    `// just free up the pd slot`  
    `pd[current_process].slotAllocated = False;`  
    `break;`

# System call interrupt handler (3 of 6)

- case CreateMessageQueueSystemCall:  
    // find a free message queue  
    int i;  
    for( i = 0; i < NumberOfMessageQueues; ++i ) {  
        if( !message\_queue\_allocated[i] ) {  
            break;  
        }  
    }  
    if( i >= NumberOfMessageQueues ) {  
        // signal the error, message queue overflow  
        // return a value that is invalid  
        pd[current\_process].sa.reg[1] = -1;  
        break;  
    }  
    message\_queue\_allocated[i] = True;  
    message\_queue[i] = new Queue<int>;  
    wait\_queue[i] = new Queue<WaitQueueItem \*>;  
    pd[current\_process].sa.reg[1] = i;  
    break;

# System call interrupt handler (4 of 6)

- case SendMessageSystemCall:

```
int * user_msg; asm { store r9,user_msg }
int to_q; asm { store r10,to_q }
if( !message_queue_allocated[to_q] ) {
    pd[current_process].sa.reg[1] = -1;
    break; }
int msg_no = GetMessageBuffer();
if( msg_no == EndOfFreeList ) {
    pd[current_process].sa.reg[1] = -2;
    break
}
CopyToSystemSpace( current_process, user_msg,
    message_buffer[msg_no], MessageSize );
if( !wait_queue[to_q].Empty() ) {
    WaitQueueItem item = wait_queue[to_q].Remove();
    TransferMessage( msg_no, item.buffer );
    pd[item.pid].state = Ready;
} else
    message_queue[to_q].Insert( msg_no );
pd[current_process].sa.reg[1] = 0;
break;
```



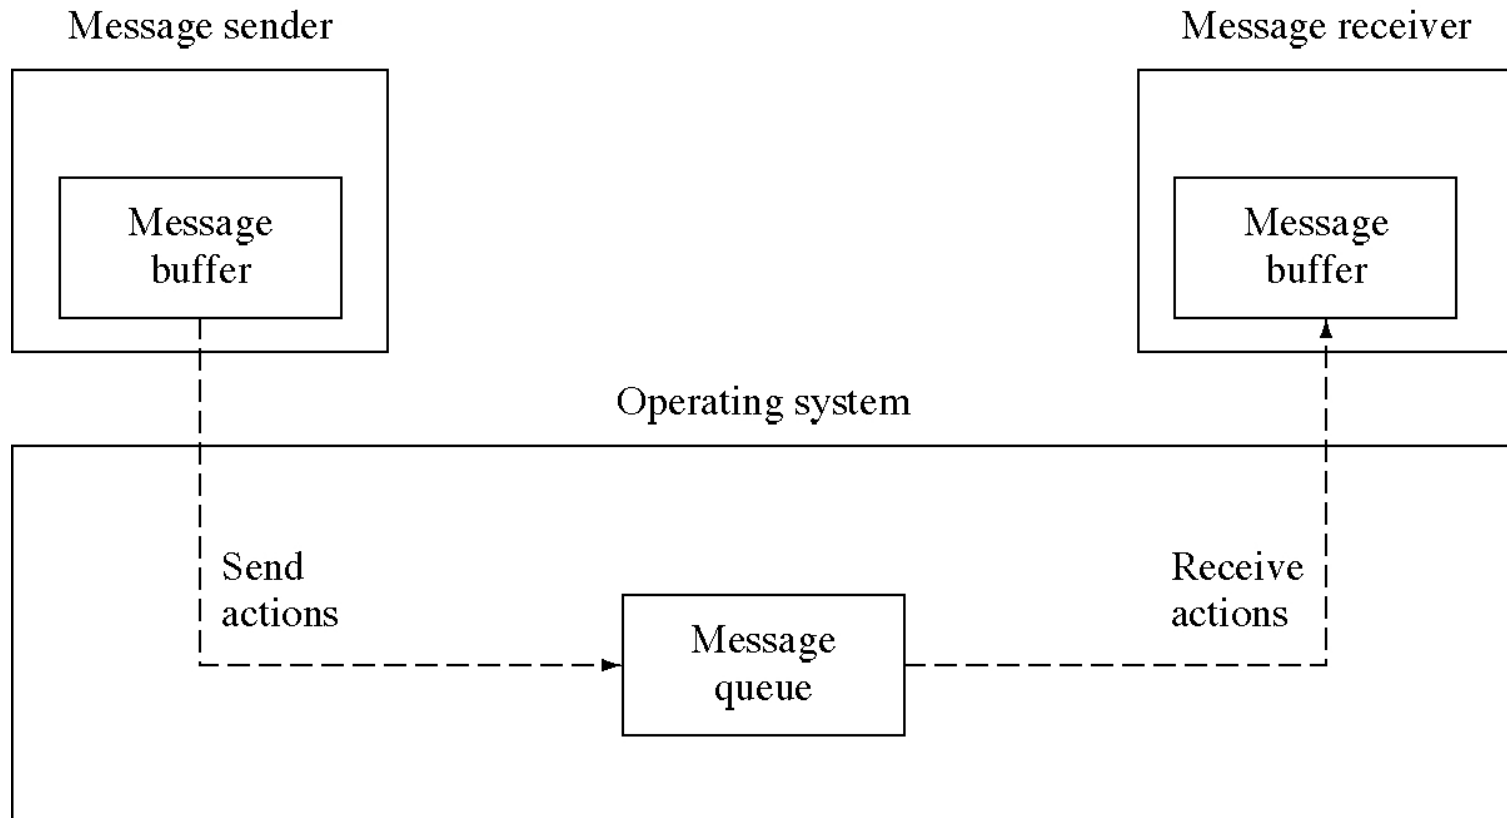
# System call interrupt handler (5 of 6)

- case ReceiveMessageSystemCall:  
    int \* user\_msg; asm { store r9,user\_msg }  
    int from\_q; asm { store r10,from\_q }  
    // check for an invalid queue identifier  
    if( !message\_queue\_allocated[from\_q] ) {  
        pd[current\_process].sa.reg[1] = -1;  
        break;  
    }  
    if( message\_queue[from\_q].Empty() ) {  
        pd[current\_process].state = Blocked;  
        WaitQueueItem item;  
        item.pid = current\_process;  
        item.buffer = user\_msg;  
        wait\_queue[from\_q].Insert( item );  
    } else {  
        int msg\_no = message\_queue[from\_q].Remove();  
        TransferMessage( msg\_no, user\_msg );  
    }  
    pd[current\_process].sa.reg[1] = 0;  
    break;

# System call interrupt handler (6 of 6)

- ```
    case DiskReadSystemCall:
    case DiskWriteSystemCall:
        char * buffer; asm { store r9,buffer }
        buffer += pd[current_process].sa.base;
        // convert to physical address
        int disk_block; asm { store r10,disk_block }
        DiskIO( system_call_number, disk_block,
buffer );
        pd[current_process].sa.reg[1] = 0;
        break;
    }
    Dispatcher();
}
```

# Send and receive actions



# Transfer between system and user memory

- ```
void CopyToSystemSpace(
    int pid, char * from, char * to, int len ) {
    from += pd[pid].sa.base;
    while( len-- > 0 )
        *to++ = *from++;
}
```

```
void CopyFromSystemSpace(
    int pid, char * to, char * from, int len ) {
    to += pd[pid].sa.base;
    while( len-- > 0 )
        *to++ = *from++;
}
```

# Program error interrupt handler

- ```
void ProgramErrorInterruptHandler( void ) {  
    asm {  
        // stop the interval timer  
        //    and clear any pending timer interrupt  
        load    #0,timer  
        // no need to save the processor state  
        //  
        // set up the stack  
        load    SystemStack+SystemStackSize,r30  
    }  
    pd[current_process].slotAllocated = False;  
    Dispatcher();  
}
```

# Disk I/O

- ```
void DiskIO(
    int command, int disk_block, char * buffer ) {
    // Create a new disk request
    // and fill in the fields.
    DiskRequest * req = new DiskRequest;
    req->command = command;
    req->disk_block = disk_block;
    req->buffer = buffer;
    req->pid = current_process;
    // Then insert it on the queue.
    disk_queue.Insert( req );
    pd[current_process].state = Blocked;
    // Wake up the disk scheduler if it is idle.
    ScheduleDisk();
}
```

# Disk scheduling

- ```
void ScheduleDisk( void ) {  
    // If the disk is already busy  
    if( DiskBusy() ) return;  
    DiskRequest * req = disk_queue.Remove();  
    // no disk request to service so return.  
    if( req == 0 ) return;  
    // remember process waiting for the disk operation  
    process_using_disk = req->pid;  
    // issue read or write, disk interrupt enabled  
    if( req->command == DiskReadSystemCall )  
        IssueDiskRead( req->disk_block, req->buffer,  
1 );  
    else  
        IssueDiskWrite(  
            req->disk_block, req->buffer, 1 );  
}
```

# Disk interrupt handler

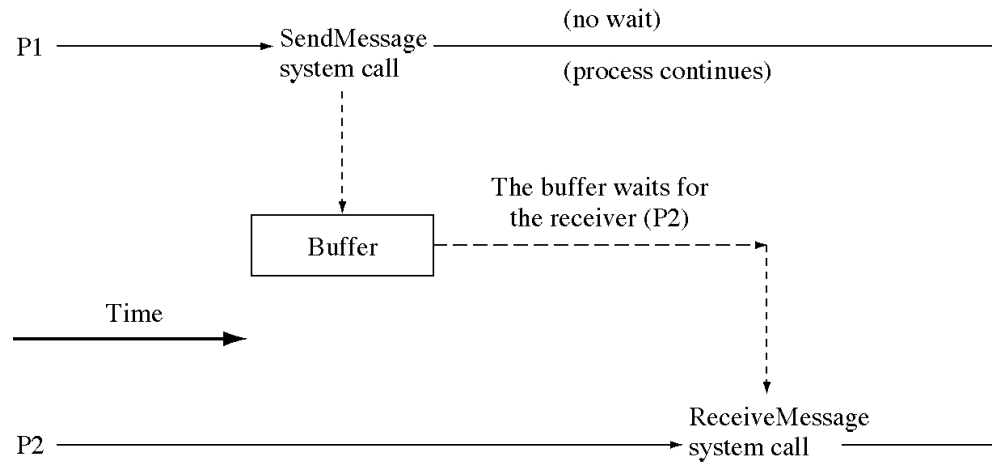
- ```
void DiskInterruptHandler( void ) {
    if( current_process > 0 ) {
        SaveArea * savearea = &(pd[current_process].sa);
        int saveTimer;
        asm { store    timer, saveTimer
              load     #0, timer
              store    iia, savearea+0
              store    ipsw, savearea+4
              store    base, savearea+8
              store    bound, savearea+12
              storeall savearea+16
              load     SystemStack+SystemStackSize, r30 }
        pd[current_process].timeLeft = saveTimer;
        pd[current_process].state = Ready; }
    pd[process_using_disk].state = Ready;
    process_using_disk = 0;
    ScheduleDisk();
    Dispatcher();
}
```



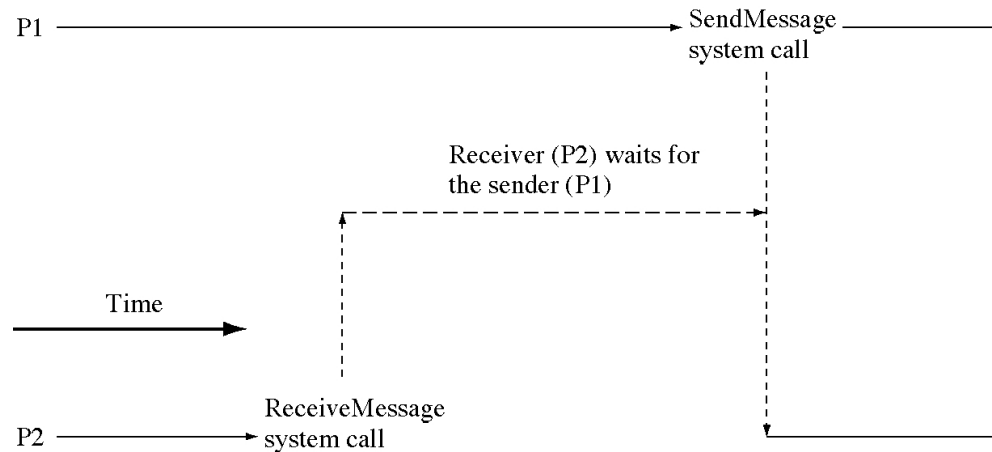
# Disk interface implementation

- ```
int DiskBusy( void ) {  
    disk_status_reg stat = *Disk_status;  
    return stat.busy;  
}  
void IssueDiskRead( int block_number, char * buffer,  
    int enable_disk_interrupt ) {  
    disk_control_reg control_reg;  
    // assemble the necessary control word  
    control_reg.command = 1;  
    control_reg.disk_block = block_number;  
    control_reg.interrupt_enabled  
        = enable_disk_interrupt;  
    // store the control words  
    //    in the disk control register  
    *Disk_memory_addr = buffer;  
    *Disk_control_reg = control_reg;  
}
```

# Waiting for messages



(a) `SendMessage` occurs before the `ReceiveMessage`



(b) `ReceiveMessage` occurs before the `SendMessage`

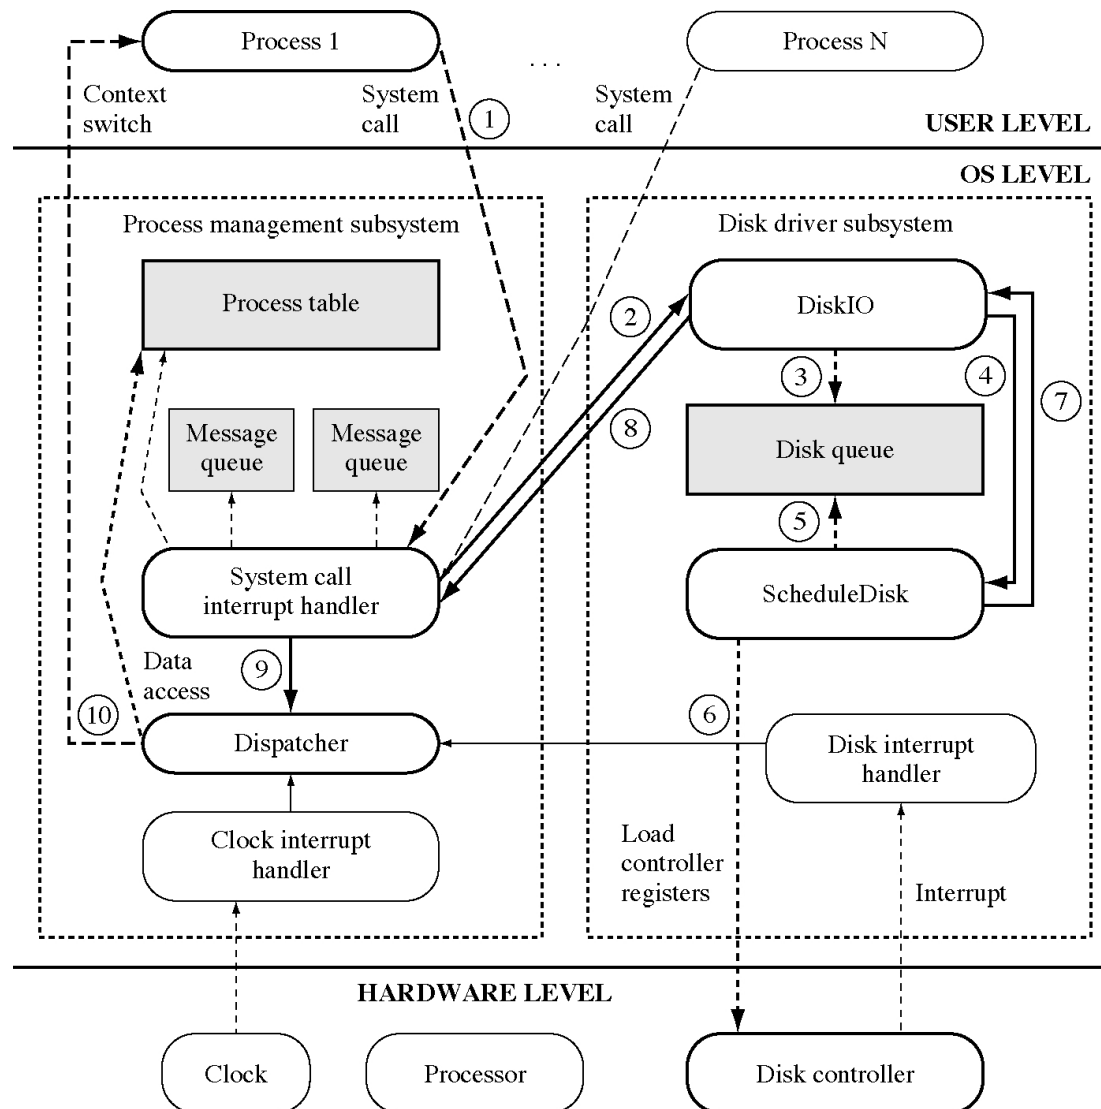
# Waiting inside a system call

- Some systems calls must wait
  - E.g. ReceiveMessage, ReadDiskBlock
- The OS suspends the process and saves its state
  - but how does the state of the OS processing the system call get saved?
  - Special provision must be made for this

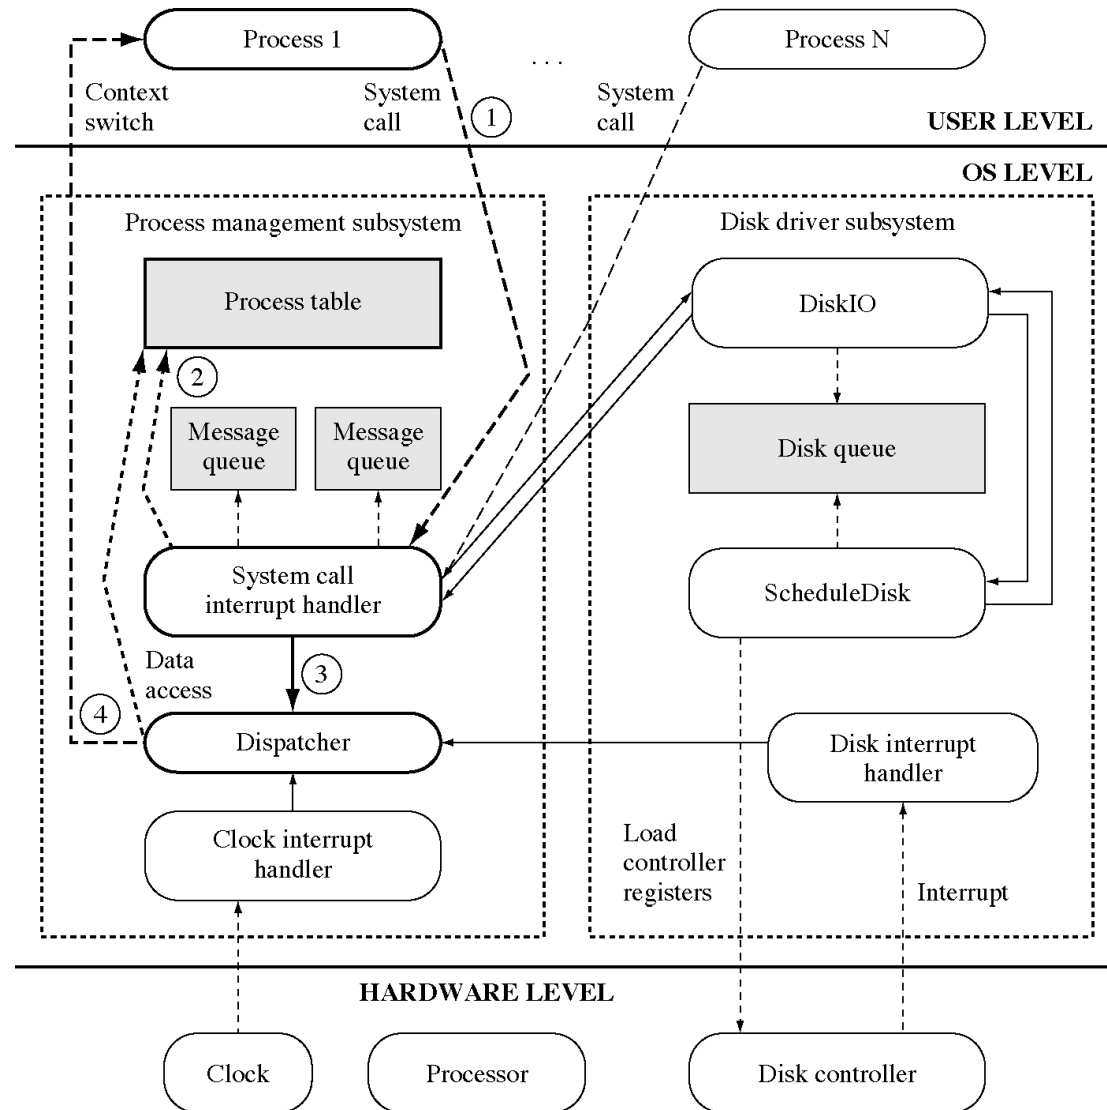
# Suspending system calls

- Find a place to save the state that will be needed when the system call resumes
  - usually this is in a waiting queue
- Arrange to be resumed when the event you are waiting for occurs
  - the OS component that handles the event will also handle this duty

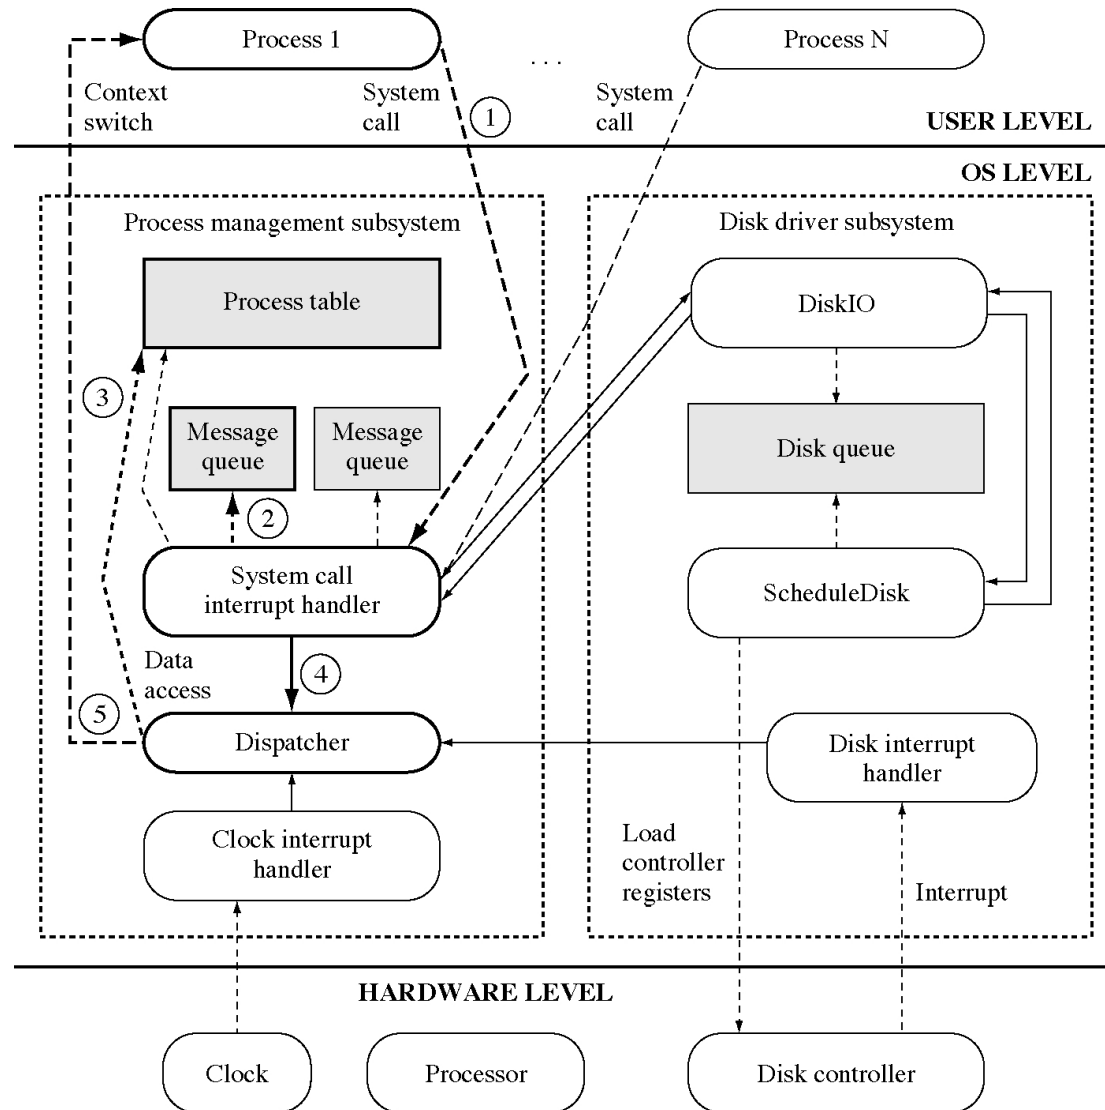
# Disk read flow of control



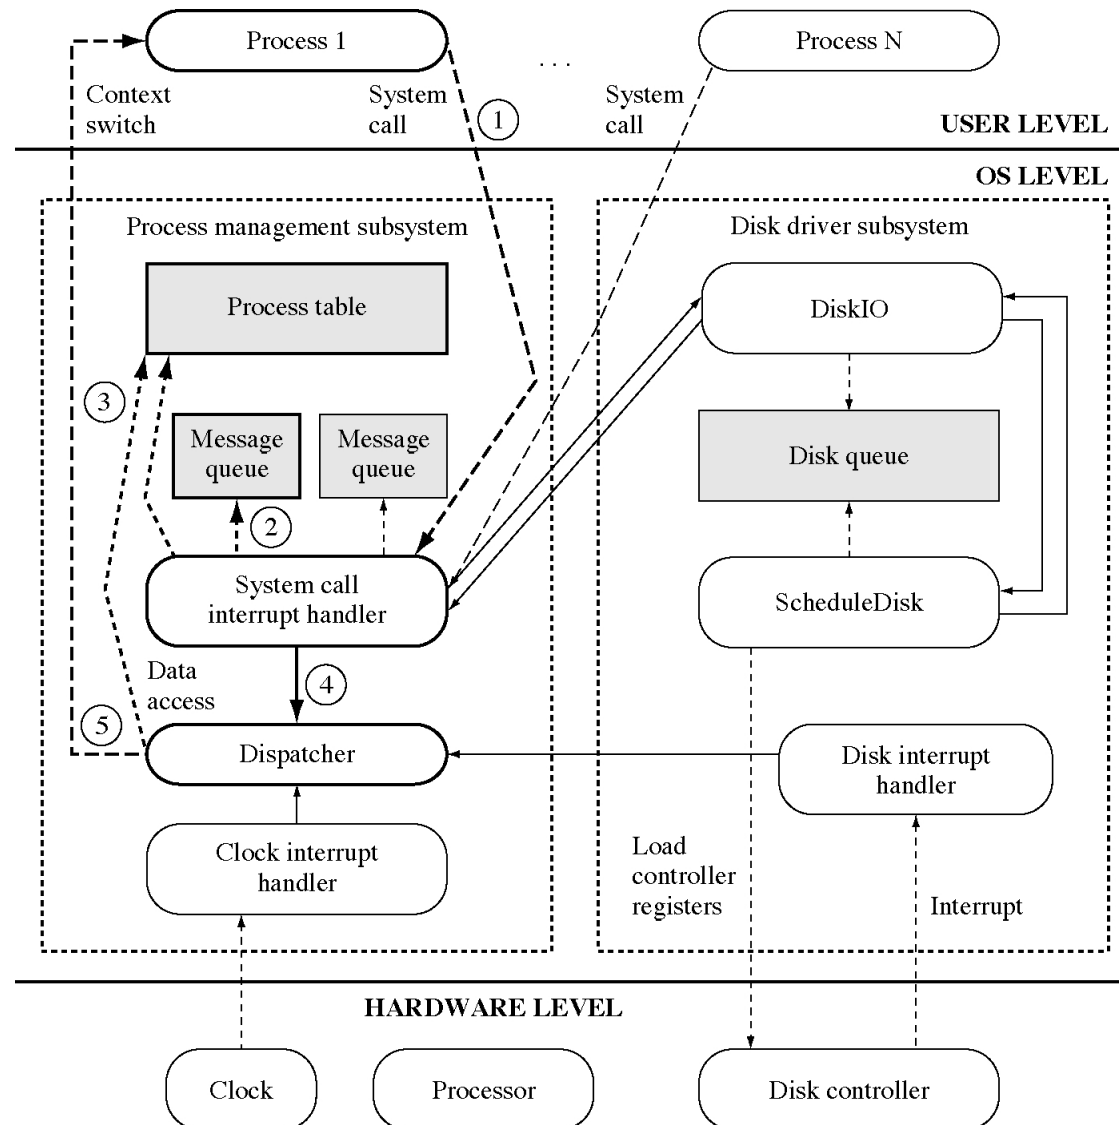
# Create process flow of control



# Create message queue control flow



# Send message flow of control





# Interrupts in the OS

- The OS is not set up to handle this
  - data save areas will be destroyed
  - so we don't allow interrupts in system code
- Chapter 6 shows how to handle this problem

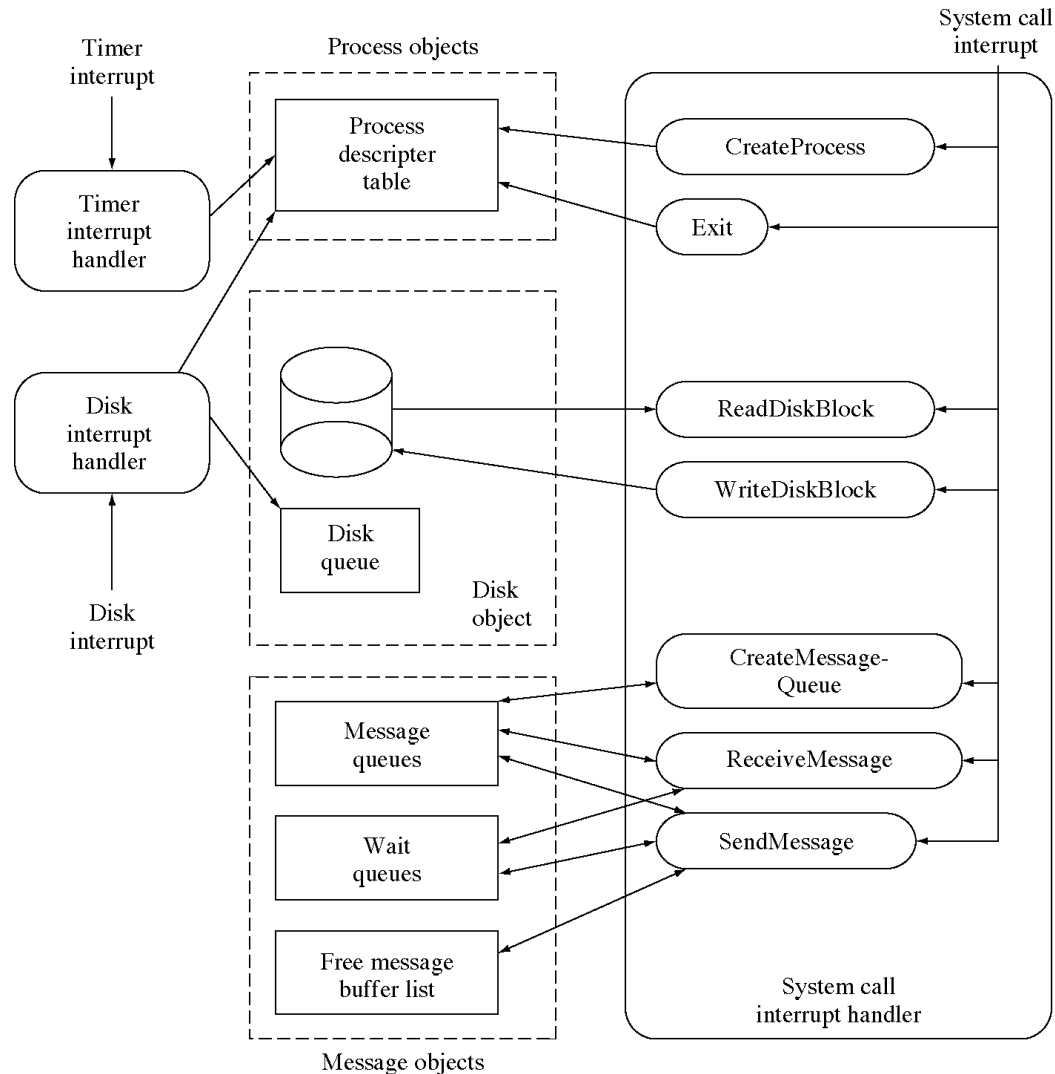
# The OS as an event and table manager

- Interrupt
 

|         |                     |                                    |
|---------|---------------------|------------------------------------|
|         | <i>Data Updated</i> | Processing Done                    |
| – Timer | <i>pd</i>           | Switch processes                   |
| – Disk  | <i>pd</i>           | Unblock process and start next I/O |
- System Call
 

|                  |                                             |                               |
|------------------|---------------------------------------------|-------------------------------|
|                  | <i>Data Updated</i>                         | Processing Done               |
| – Createprocess  | <i>pd</i>                                   | Initialize process table slot |
| – Exit           | <i>pd</i>                                   | Free process table slot       |
| – CreateMsgQueue | <i>pd, message_queue</i>                    | Initialize message queue      |
| – SendMessage    | <i>pd, message_queue<br/>message_buffer</i> | Queue or transfer message     |
| – ReceiveMessage | <i>pd, message_queue<br/>message_buffer</i> | Block or transfer message     |
| – ReadDiskBlock  | <i>pd, disk_queue</i>                       | Queue disk request            |
| – WriteDiskBlock | <i>pd, disk_queue</i>                       | Queue disk request            |

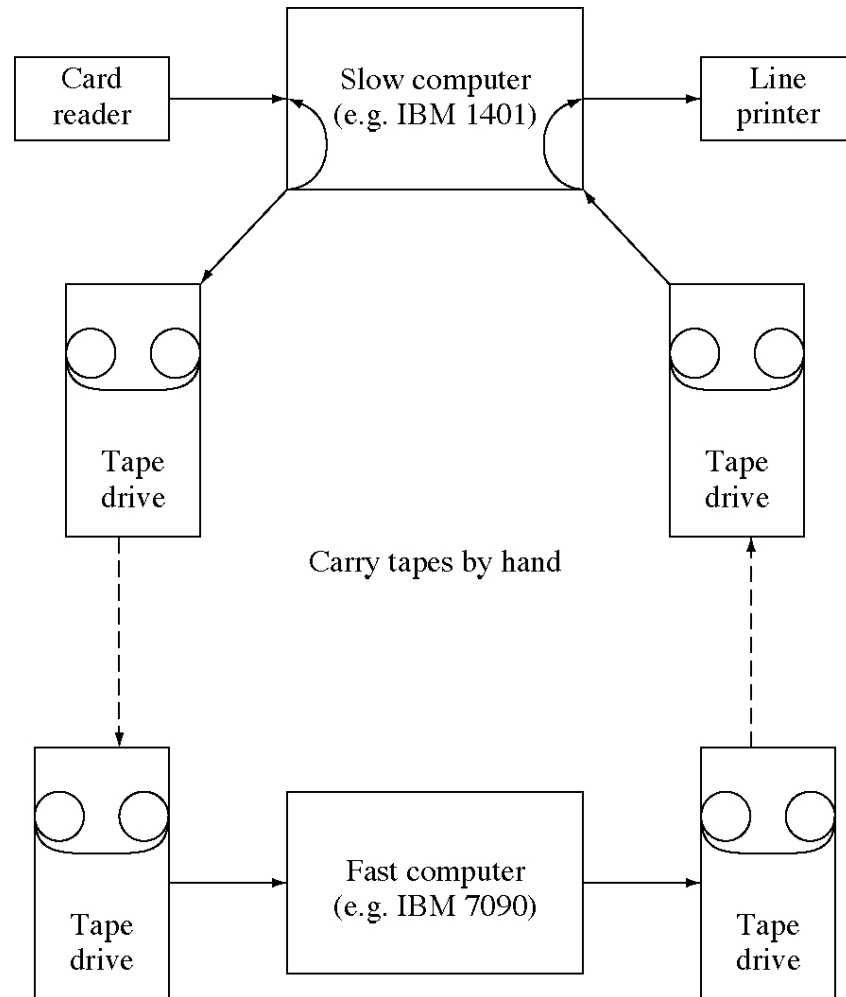
# Interrupt event handling




# Typical process descriptor fields

- Process ID and name
- Memory allocated to the process
- Open files
- Process state
- User name and protection privileges
- Register save area
- CPU time user
- Pending software interrupts
- Parent process
- User ID


# A batch operating system




# Speeding up I/O (3 methods)



(a) Monoprogramming






(c) Multiprogramming




(b) Faster I/O


# I/O overlap



# TSRs in PCs



**(a)** PC with monoprogramming—slow task switching



**(b)** PC with TSRs