

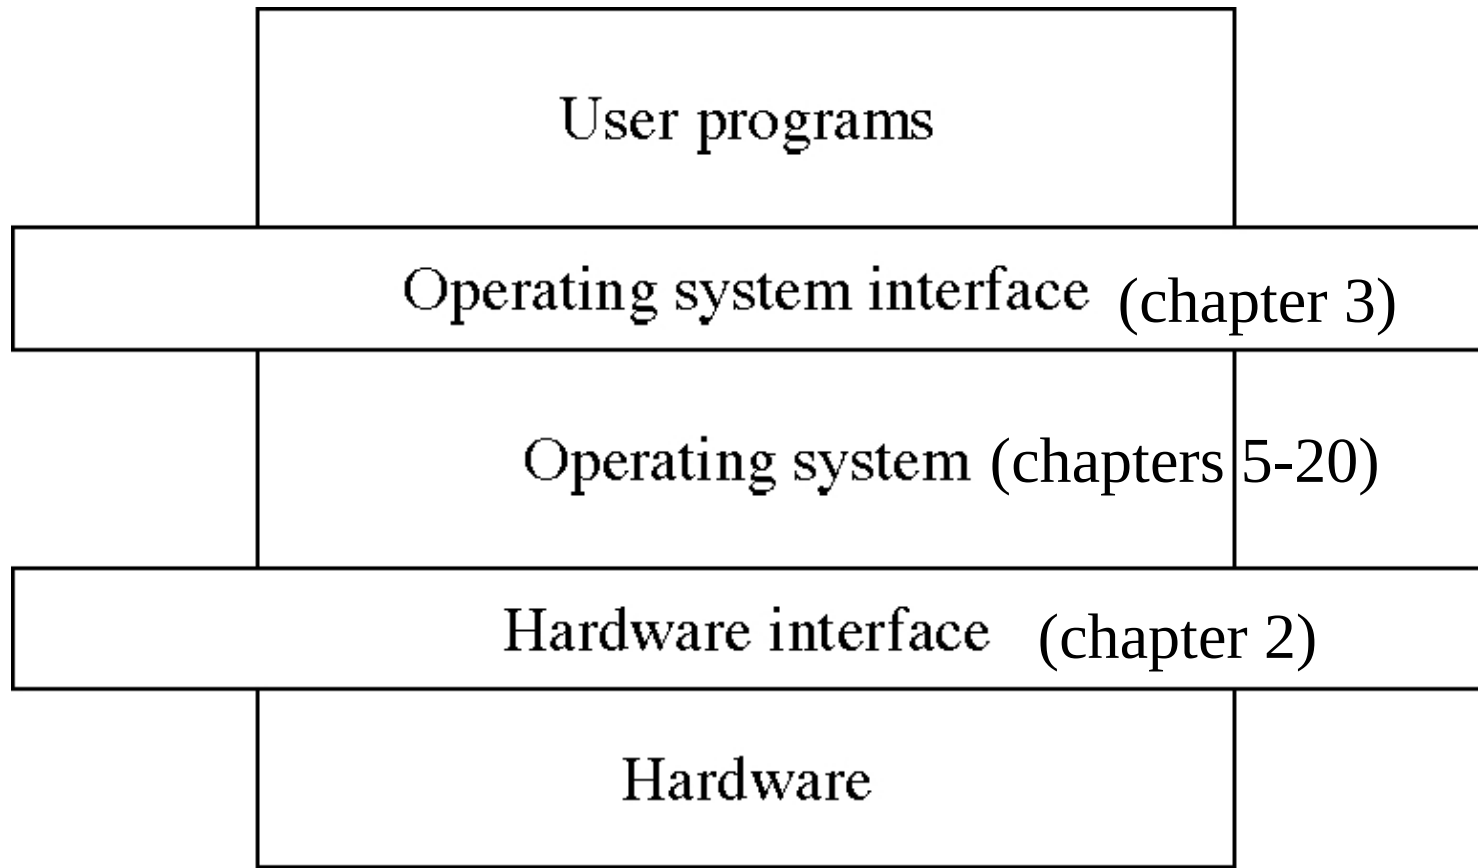
# The Operating System Interface

## Chapter 3

# Key concepts in chapter 3

- System calls
- File and I/O system
  - hierarchical file naming
  - file interface: open, read, write, lseek, close
  - file versus open file
  - devices as files (in naming and in interface)
- Process
  - operations: create, exit, wait
- Shell

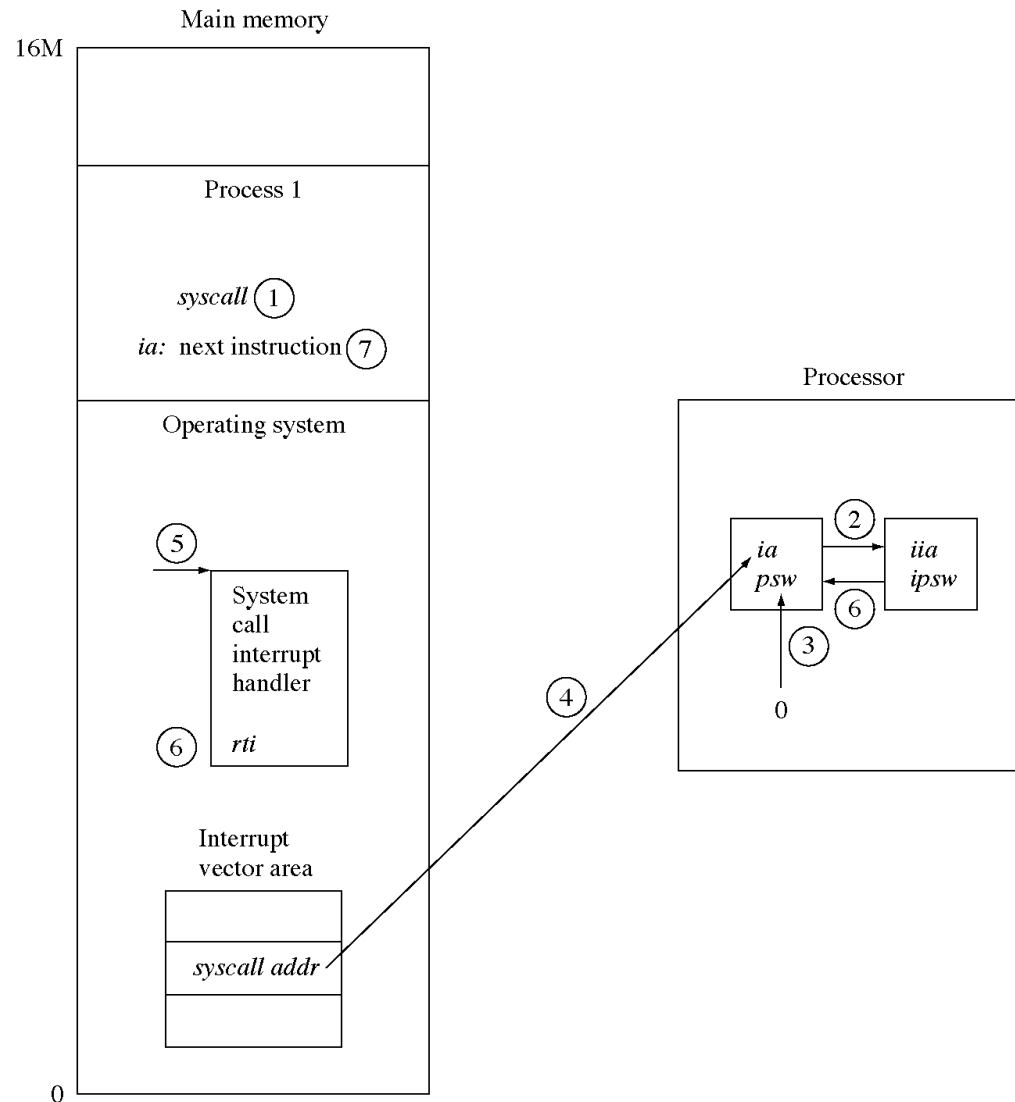
# The OS Level Structure



# System calls

- A special machine instruction
  - that causes an interrupt
  - various names: syscall, trap, svc
- Usually not generated by HLLs
  - but in assembly language functions

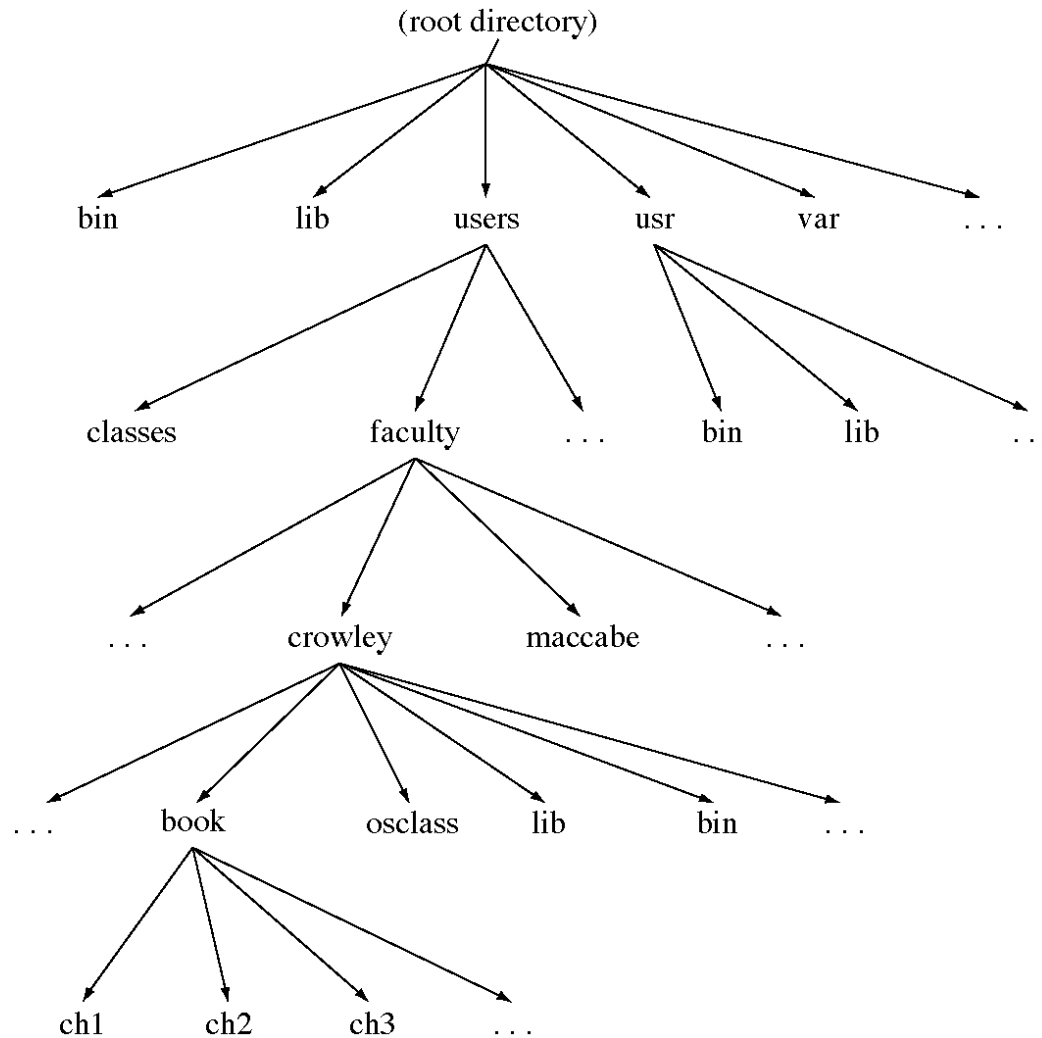
# System call flow of control



# Hierarchical file naming systems

- A tree of directories and files
  - directory: contains file and directory names
- Objects (files and directories) are named with path names
  - later: other kinds of objects (e.g. devices)
- Path names contains a component name for each directory in the path

# A file naming system

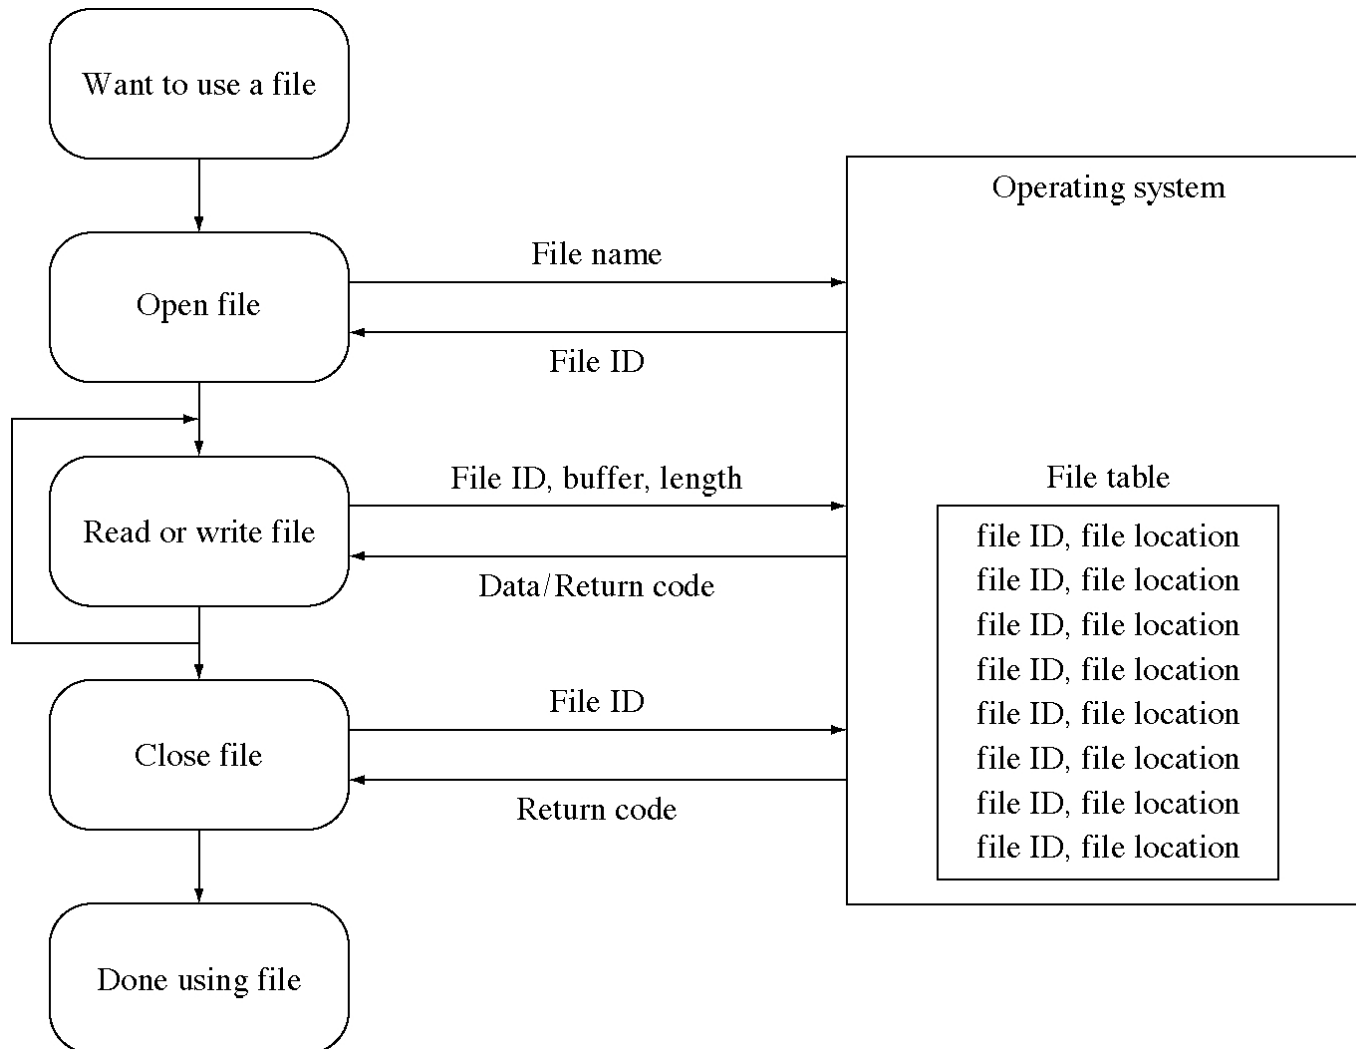


# File and I/O system calls

- `int open(char *name, int flags)`
- `int read(int fid, char *buffer, int count)`
- `int write(int fid, char *buffer, int count)`
- `int lseek(int fid, int offset, int from)`
- `int close(int fid)`
- `int unlink(char *name)`




# Steps in using a file



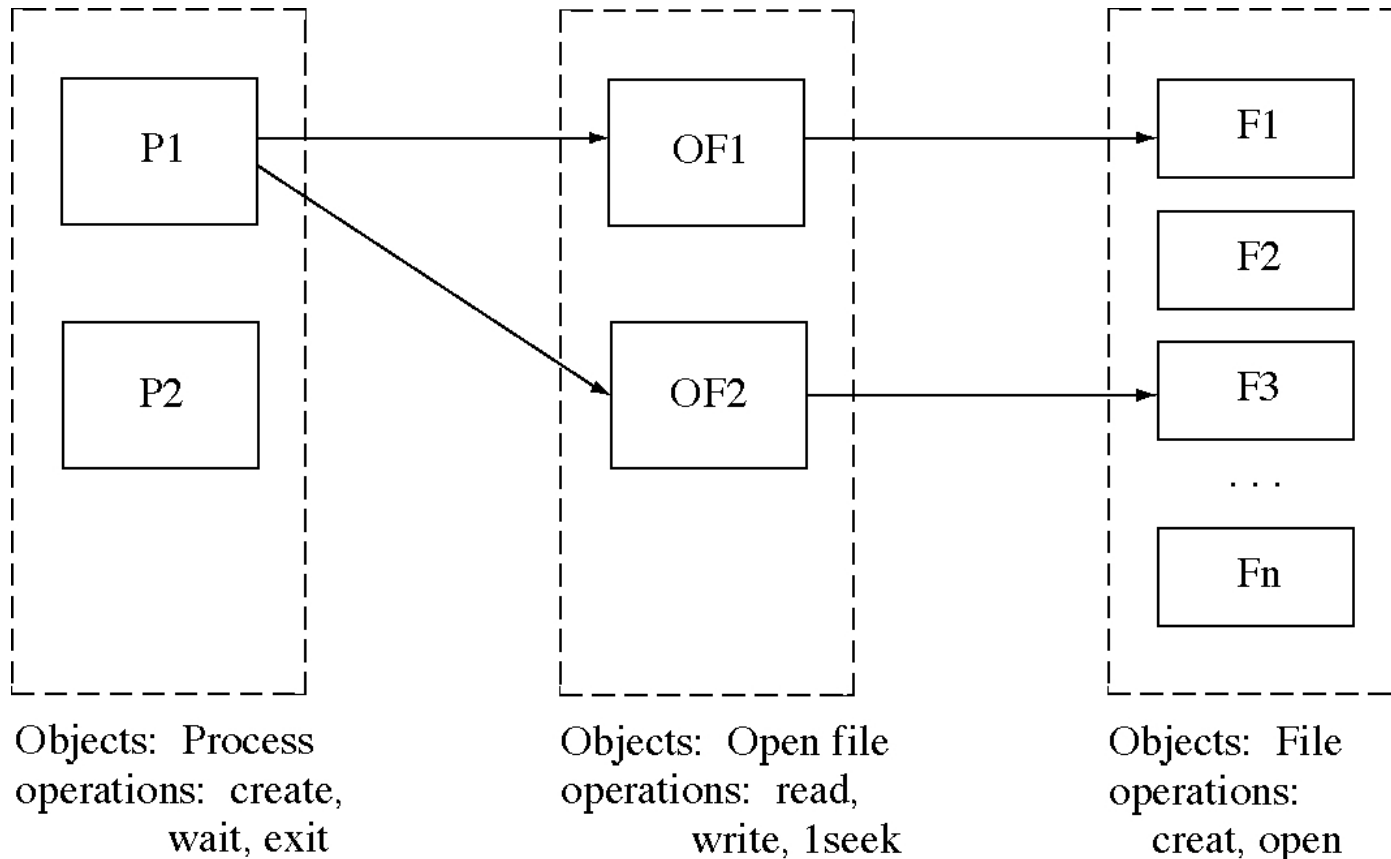
# Files versus open files

- *File*: passive container of bytes on disk
- *Open file*: active source (or sink) of bytes in a running program
  - usually connected to a file
  - but can be connected to a device or another process

# Files and open files



# OS objects and operations



# File copy

- enum { Reading=0, Writing=1, ReadAndWrite=2, ReadWriteFile=0644 };

```
void FileCopy( char * fromFile, char * toFile ) {
    int fromFD = open( fromFile, Reading );
    if( fromFD < 0 ) {
        cerr << "Error opening " << fromFile << endl;
        return; }
    int toFD = creat( toFile, ReadWriteFile );
    if( toFD < 0 ) {
        cerr << "Error opening " << toFile << endl;
        close( fromFD ); return; }
    while( 1 ) {
        char ch; int n = read( fromFD, &ch, 1 );
        if( n <= 0 ) break;
        n = write( toFD, &ch, 1 );
        if( n < 0 ) {
            cerr << "Error writing " << toFile << endl;
            return; }
    }
    close( fromFD ); close( toFD );
}
```

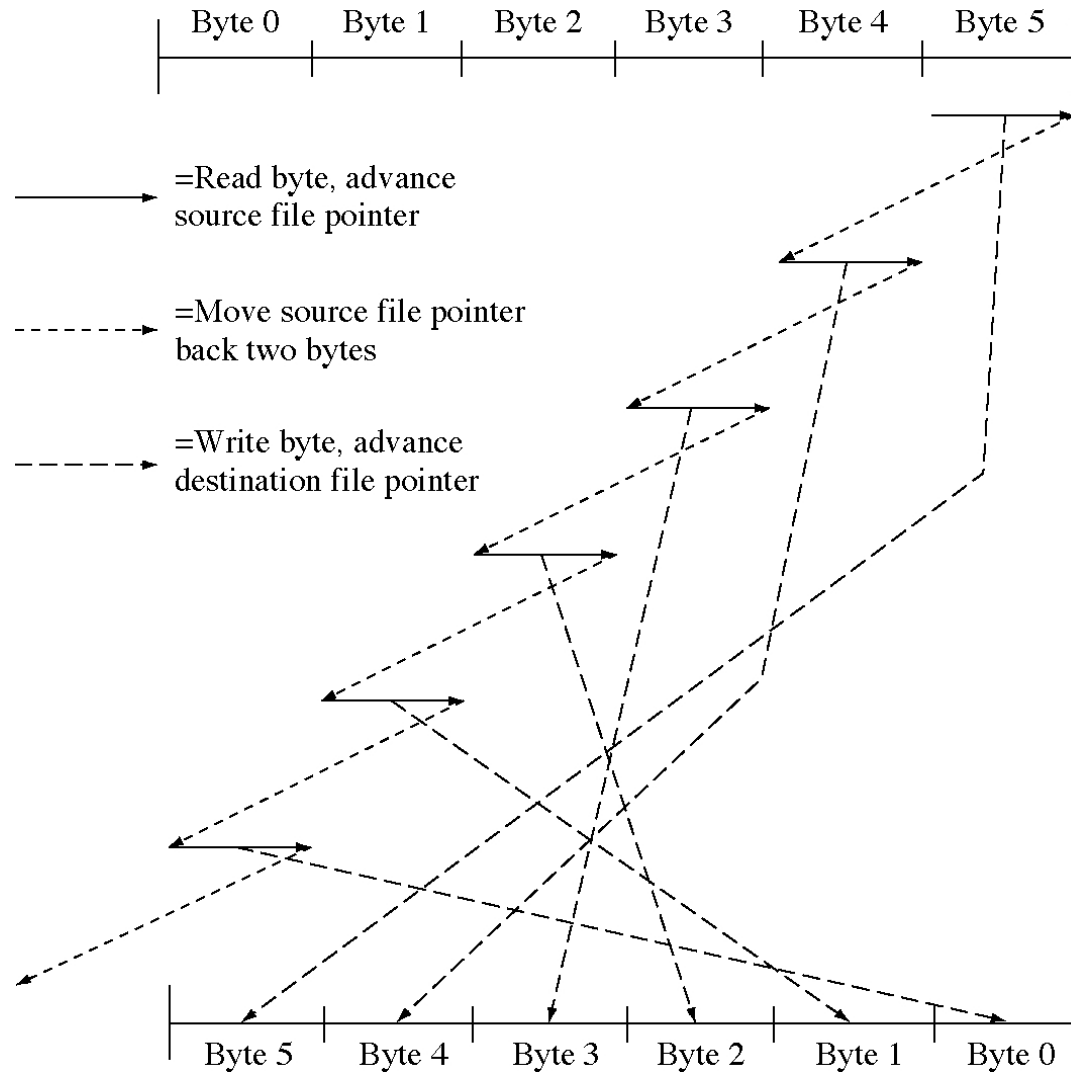
# File reverse (1 of 2)

- ```
enum { Reading=0, Writing=1, ReadAndWrite=2 };
enum { SeekFromBeginning=0, SeekFromCurrent=1, SeekFromEnd=2 };
void Reverse( char * fromFile, char * revFile ) {
    int fromFD = open( fromFile, Reading );
    if( fromFD < 0 ) {
        cerr << "Error opening " << fromFile << endl;
        return;
    }
    // move the internal file pointer so the next character
    // read will be the last character of the file
    int ret lseek( fromFD, -1, SeekFromEnd );
    if( ret < 0 ) {
        cerr << "Error seeking on " << fromFile << endl;
        close( fromFD );
        return;
    }
    int revFD = creat( revFile, 0 );
    if( revFD < 0 ) {
        cerr << "Error creating " << revFile << endl;
        close( fromFD );
        return;
    }
}
```

# File reverse (2 of 2)

- ```
while( 1 ) {
    char ch;
    int n = read( fromFD, &ch, 1 );
    if( n < 0 ) {
        cerr << "Error reading " << fromFile << endl;
        return;
    }
    n = write( revFD, &ch, 1 );
    if( n < 0 ) {
        cerr << "Error writing " << revFile << endl;
        return;
    }
    // exit the loop if lseek returns an error.
    // The expected error is that the computed offset will
    // be negative.
    if( lseek(fromFD, -2, SeekFromCurrent) < 0 )
        break;
}
close( fromFD );
close( revFD );
}
```

# Reversing a file





# Design technique: Interface design

- There are many different sets of system calls with the same functionality
  - which one is best depends on how they will be used
  - we try to make them easy to use and efficient (minimize the number of system calls necessary to get the job done)
- One should always consider several design alternatives and evaluate them

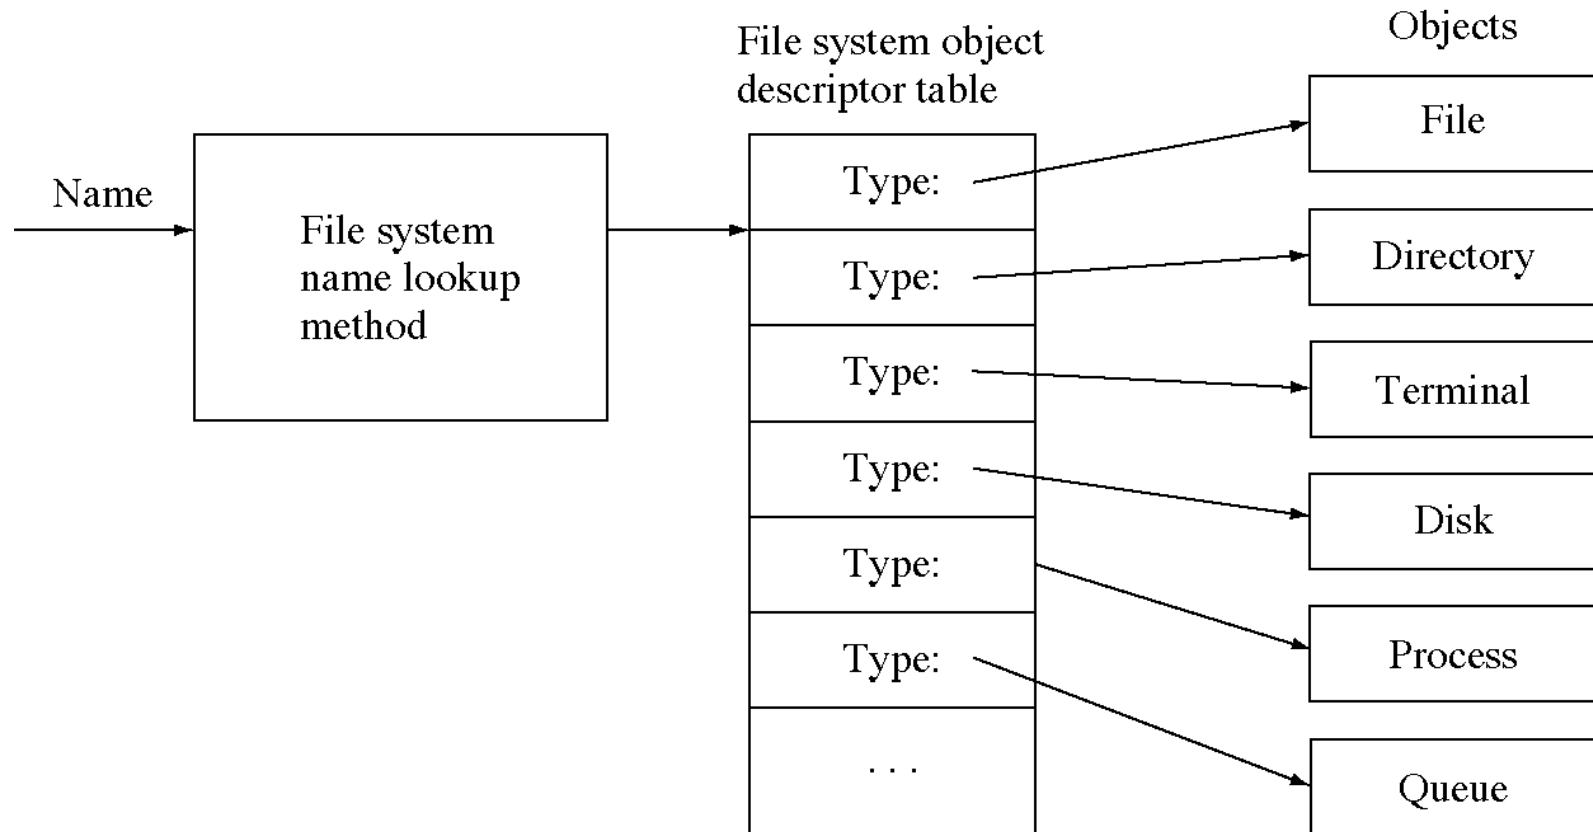
# Meta-data

- Meta-data describes the file rather than being the data in file itself
  - also called meta-information
- Examples of meta-data
  - Who owns the file
  - Who can use the file and how
  - When the file was created, last used, last modified
- `int stat(char * name, StatInfo *statInfo)`
  - this calls returns the file meta-data

# Naming OS objects

- File naming system names files (and directories)
  - but why limit it to that
- Other OS objects need names:
  - processes
  - devices
  - IPC: message queues, pipes, semaphores

# Mapping names to objects



# Devices as files

- Devices are named as files
  - they can be opened as files, to create open files
  - they can be used as byte streams: sources of bytes and sinks for bytes
- Examples
  - `copy someFile /dev/tty17`
  - `copy aFile /dev/tape01`

# The process concept

- *Program*: a static, algorithmic description, consists of instructions

```
- int main() {  
    int i, prod=1;  
    for(i=0; i<100; ++i)  
        prod = prod*i;  
}
```

- *Process*: dynamic, consists of instruction executions

# Simple create process

- ```
void CreateProcess1( void ) {
    int pid1 = SimpleCreateProcess( "compiler" );
    if( pid1 < 0 ) {
        cerr << "Could not create process \"compiler\""
              << endl;
        return; }
    int pid2 = SimpleCreateProcess( "editor" );
    if( pid2 < 0 ) {
        cerr << "Could not create process \"editor\""
              << endl;
        return; }
    // Wait until they are both completed.
    SimpleWait( pid1 );
    SimpleWait( pid2 );
    // "compiler" and "editor" also end by making
    // SimpleExit system calls
    SimpleExit();
}
```

# Process system calls

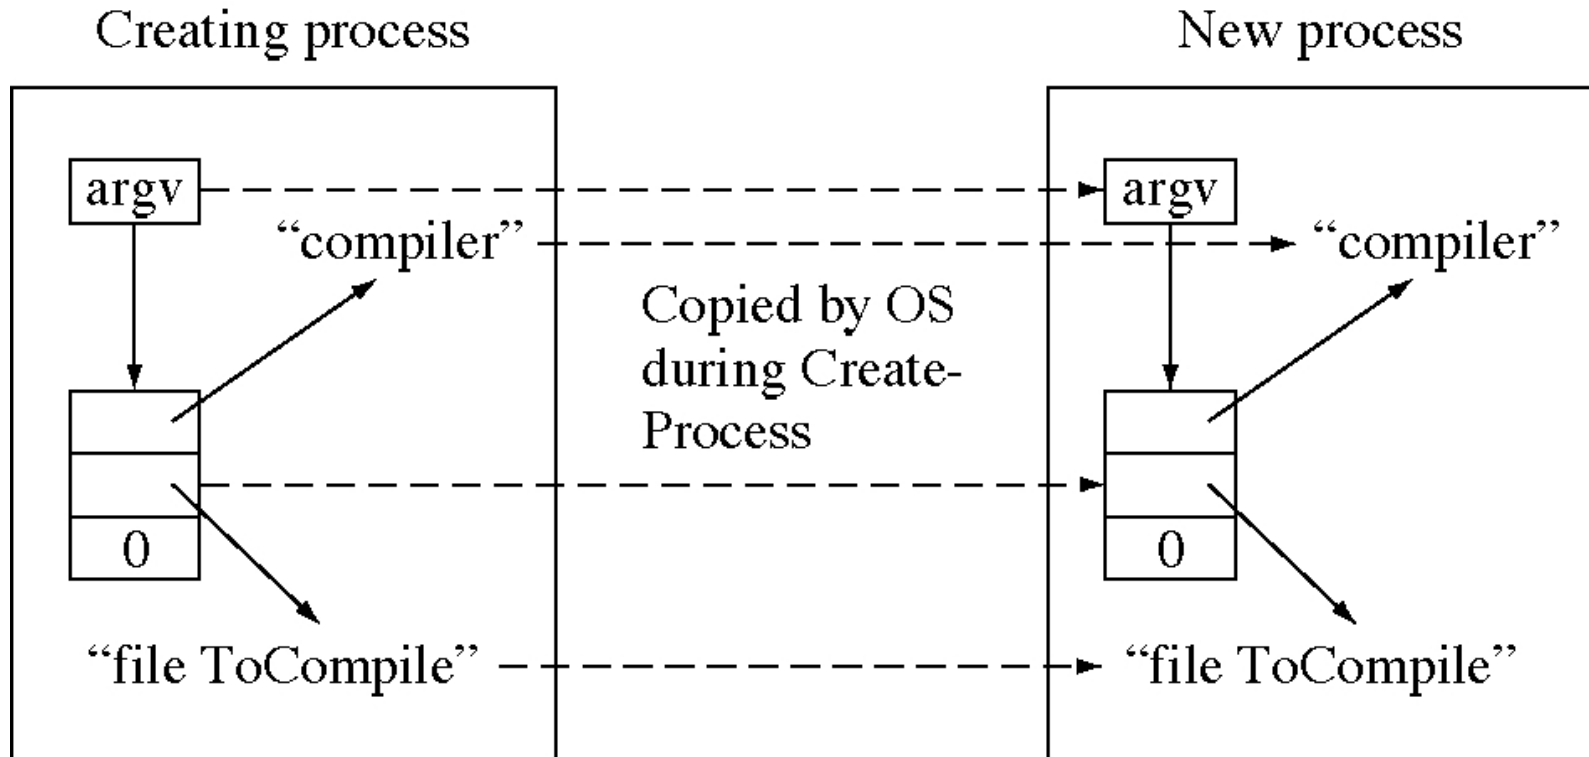
- `int CreateProcess(  
    char *progName, int argc, char *argv[ ])`
  - progName is the program to run in the process
  - returns a process identifier (pid)
- `void Exit(int returnCode)`
  - exits the process that executes the exit system calls
- `int Wait(int pid)`
  - waits for a child process to exit



# Create process

- ```
void CreateProcess2( void ) {
    static char * argv[3]
        = { "compiler", "fileToCompile", (char *) 0 };
    int pid1 = CreateProcess( "compiler", 3, argv );
    if( pid1 < 0 ) {
        cerr << "Could not create process \"compiler\""
            << endl;
        return;
    }
    char * argv[3];
    argv[0] = "editor";
    argv[1] = "fileToEdit";
    argv[2] = (char *) 0;
    int pid2 = CreateProcess( "editor", 3, argv );
    if( pid2 < 0 ) {
        cerr << "Could not create process \"compiler\""
            << endl;
        return;
    }
    (void) Wait( pid1 );
    (void) Wait( pid2 );
    Exit( 0 );
}
```

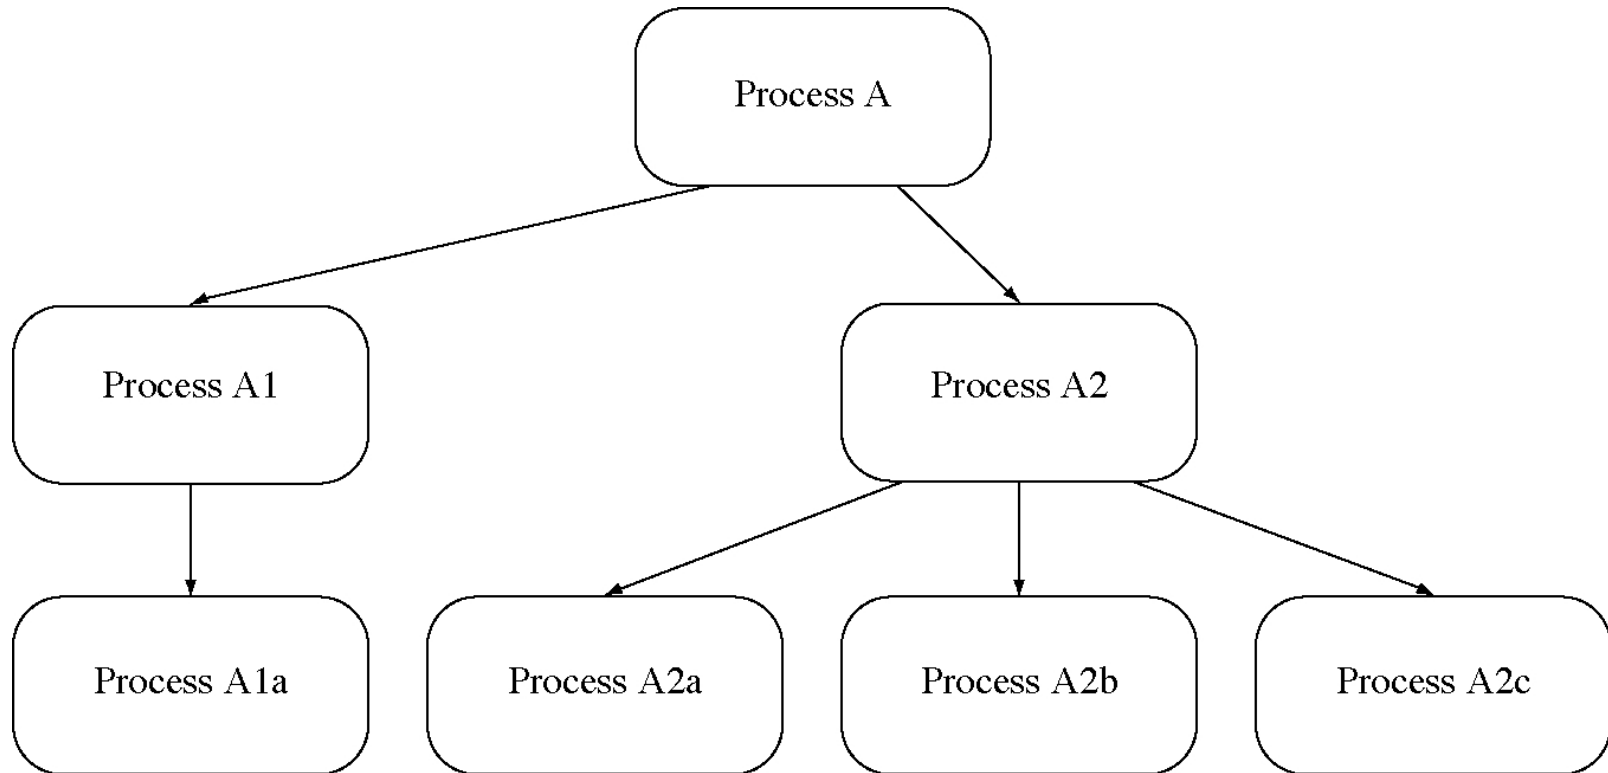
# How arguments are passed



# Print arguments

- ```
// This program writes out its arguments.
#include <iostream.h>
void main( int argc, char * argv[ ] ) {
    int i;
    for( i = 0; i < argc; ++i ) {
        cout << argv[i] << " ";
    }
    cout << "\n";
}
```

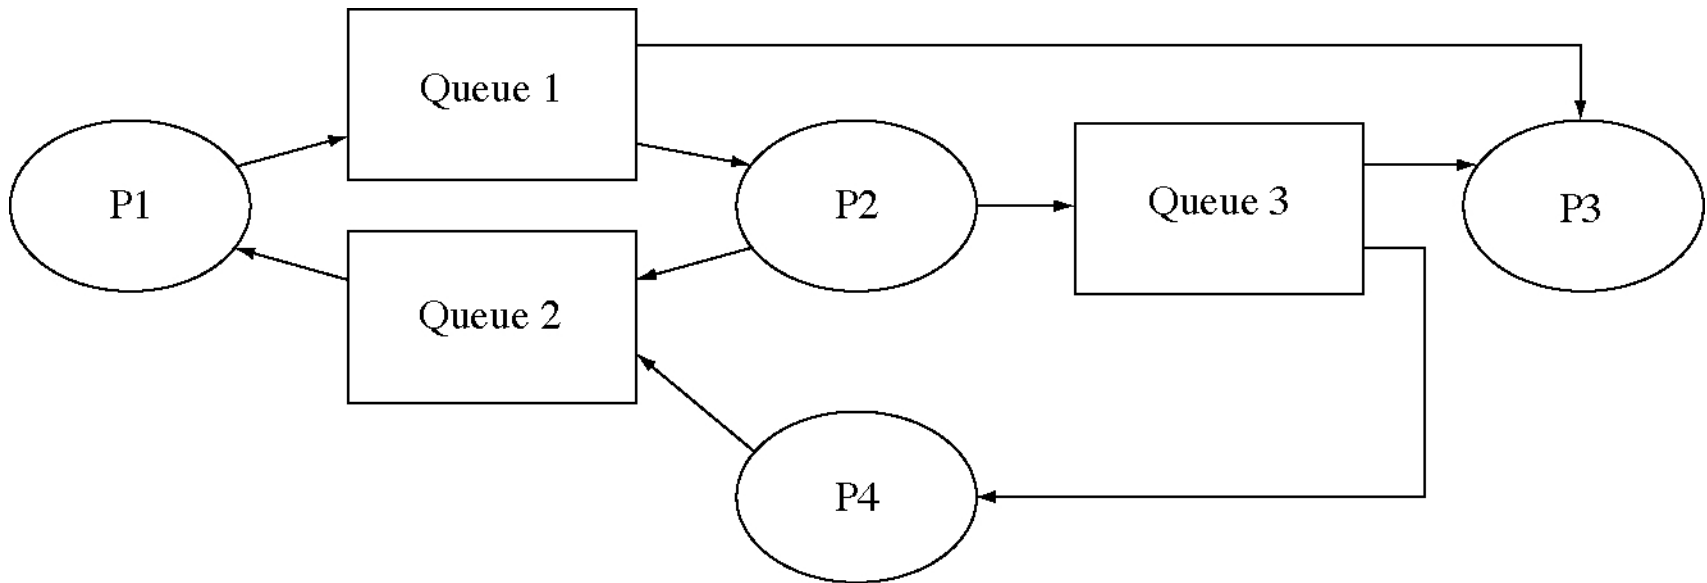
# A process hierarchy



# Interprocess communication (IPC)

- Many methods have been used: messages, pipes, sockets, remote procedure call, etc.
- Messages and message queues:
  - The most common method
  - Send messages to message queues
  - Receive messages from message queues

# Example of message passing paths



# Message passing system calls

- `int CreateMessageQueue( void )`
  - returns a message queue identifier (mqid)
- `int SendMessage(int mqid, int *msg)`
  - send to a message queue (no waiting)
- `int ReceiveMessage(int mqid, int *msg)`
  - receive from a message queue
  - wait for a message if the queue is empty
- `int DestroyMessageQueue(int mqid)`

# Message: file sender (1 of 2)

- ```
void SendMsgTo( int msg_q_id, int msg0=0, int
msg1=0, int msg2=0 ) {
    int msg[8];
    msg[0] = msg0; msg[1] = msg1; msg[2] = msg2;
    (void)SendMessage( msg_q_id, msg );
}
enum { Reading=0, Writing=1, ReadAndWrite=2 };
enum{ FileToOpen=1, SendQueue=2, ReceiveQueue=3 };
void main( int argc, char * argv[ ] ) {
    int fromFD = open( argv[FileToOpen], Reading );
    if( fromFD < 0 ) {
        cerr << "Could not open file "
              << argv[FileToOpen] << endl;
        exit( 1 );
    }
    int to_q = atoi(argv[SendQueue]);
```



# Message: file sender (2 of 2)

- ```
while( 1 ) {
    char ch;
    int n = read( fromFD, &ch, 1 );
    if( n <= 0 )
        break;
    SendMsgTo( to_q, ch );
}
close( fromFD );
SendMsgTo( to_q, 0 );
int msg[8];
int from_q = atoi(argv[ReceiveQueue]);
ReceiveMessage( from_q, msg );
cout << msg[0] << " characters\n";
exit( 0 );
}
```

# Message: file receiver

- enum{ SendQueue=1, ReceiveQueue=2 };  
void main( int argc, char \* argv[ ] ) {  
 // start the count at zero.  
 int count = 0;  
 int msg[8];  
 int from\_q = atoi(argv[SendQueue]);  
 while( 1 ) {  
 ReceiveMessage( from\_q, msg );  
 if( msg[0] == 0 )  
 break;  
 // Any message with nonzero content  
 // is a character to count.  
 ++count;  
 }  
 // Send the count back to the sender.  
 int to\_q = atoi(argv[ReceiveQueue]);  
 (void) SendMsgTo( to\_q, count );  
 exit( 0 );  
}

# Message: start processes (1 of 2)

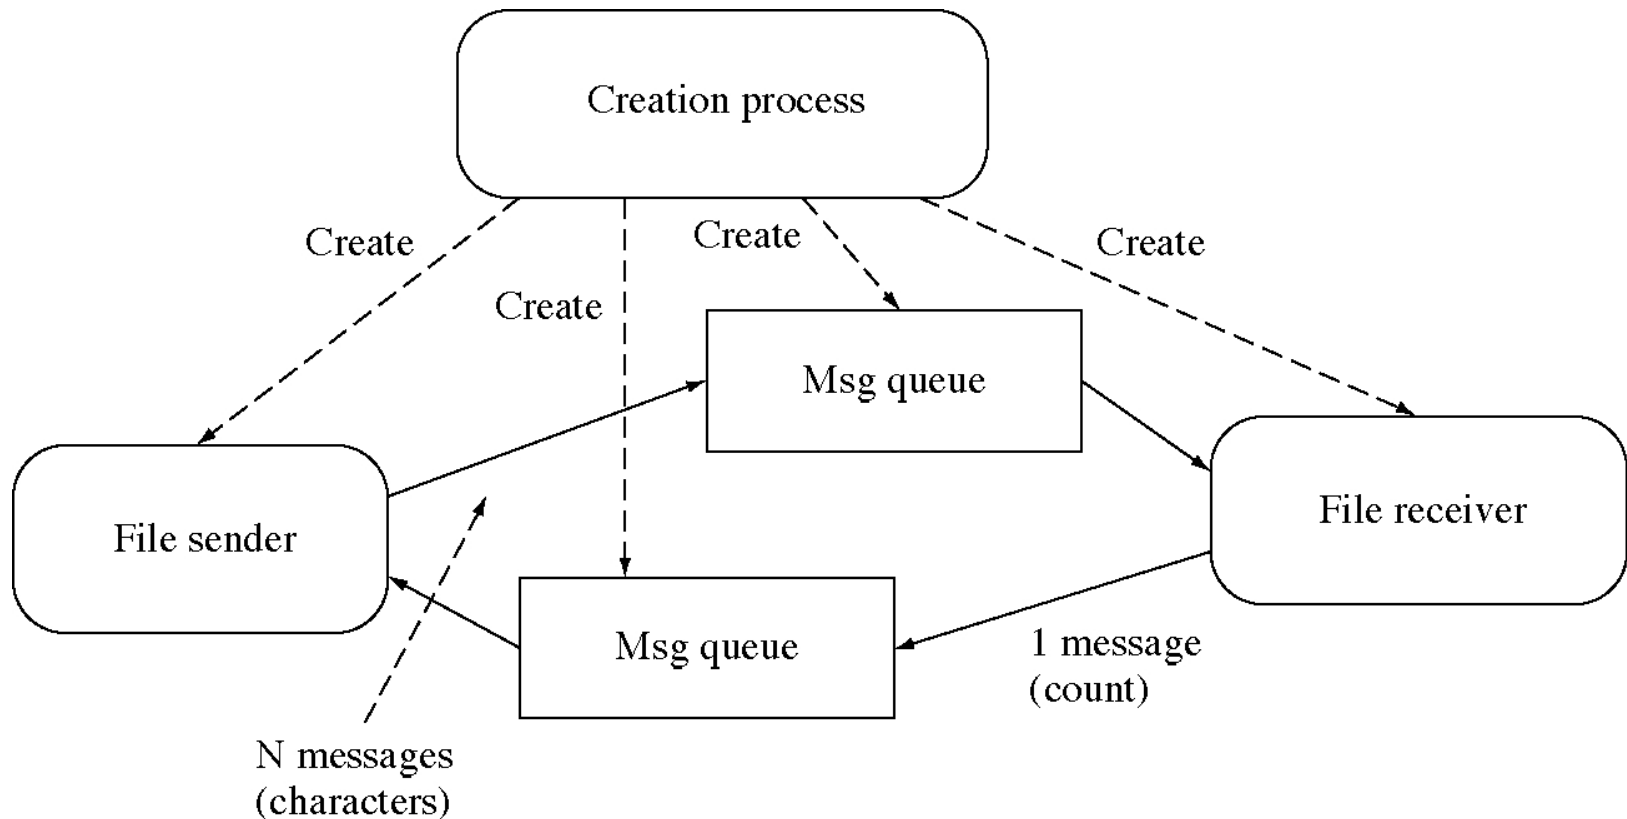
- ```
int CreateProcessWithArgs(char * prog_name,
    char * arg1=0, char * arg2=0, char * arg3=0) {
    char *args[5];
    args[0] = prog_name;
    args[1] = arg1;
    args[2] = arg2;
    args[3] = arg3;
    args[4] = 0;
    int argc = 4;
    if( arg3 == 0) --argc;
    if( arg2 == 0) --argc;
    if( arg1 == 0) --argc;
    return CreateProcess( prog_name, argc, args );
}

char * itoa( int n ) {
    char * result = new char[8];
    sprintf( result, "%d", n );
    return result;
}
```

# Message: start processes (2 of 2)

- ```
void main( int argc, char * argv[ ] ) {  
    //Create the message queues the processes will  
    use.  
    int q1 = CreateMessageQueue();  
    int q2 = CreateMessageQueue();  
    // Create the two processes, sending each the  
    // identifier for the message queues it will use.  
    int pid1 = CreateProcessWithArgs( "FileSend",  
        "FileToSend", itoa(q1), itoa(q2) );  
    int pid2 = CreateProcessWithArgs( "FileReceive",  
        itoa(q1), itoa(q2) );  
    // Wait for the two processes to complete.  
    int ret1 = wait( pid1 );  
    int ret2 = wait( pid2 );  
    // We do not use the return code ret1 and ret2  
    //   in this example.  
    // Destroy the message queues.  
    DestroyMessageQueue( q1 );  
    DestroyMessageQueue( q2 );  
    Exit( 0 );  
}
```

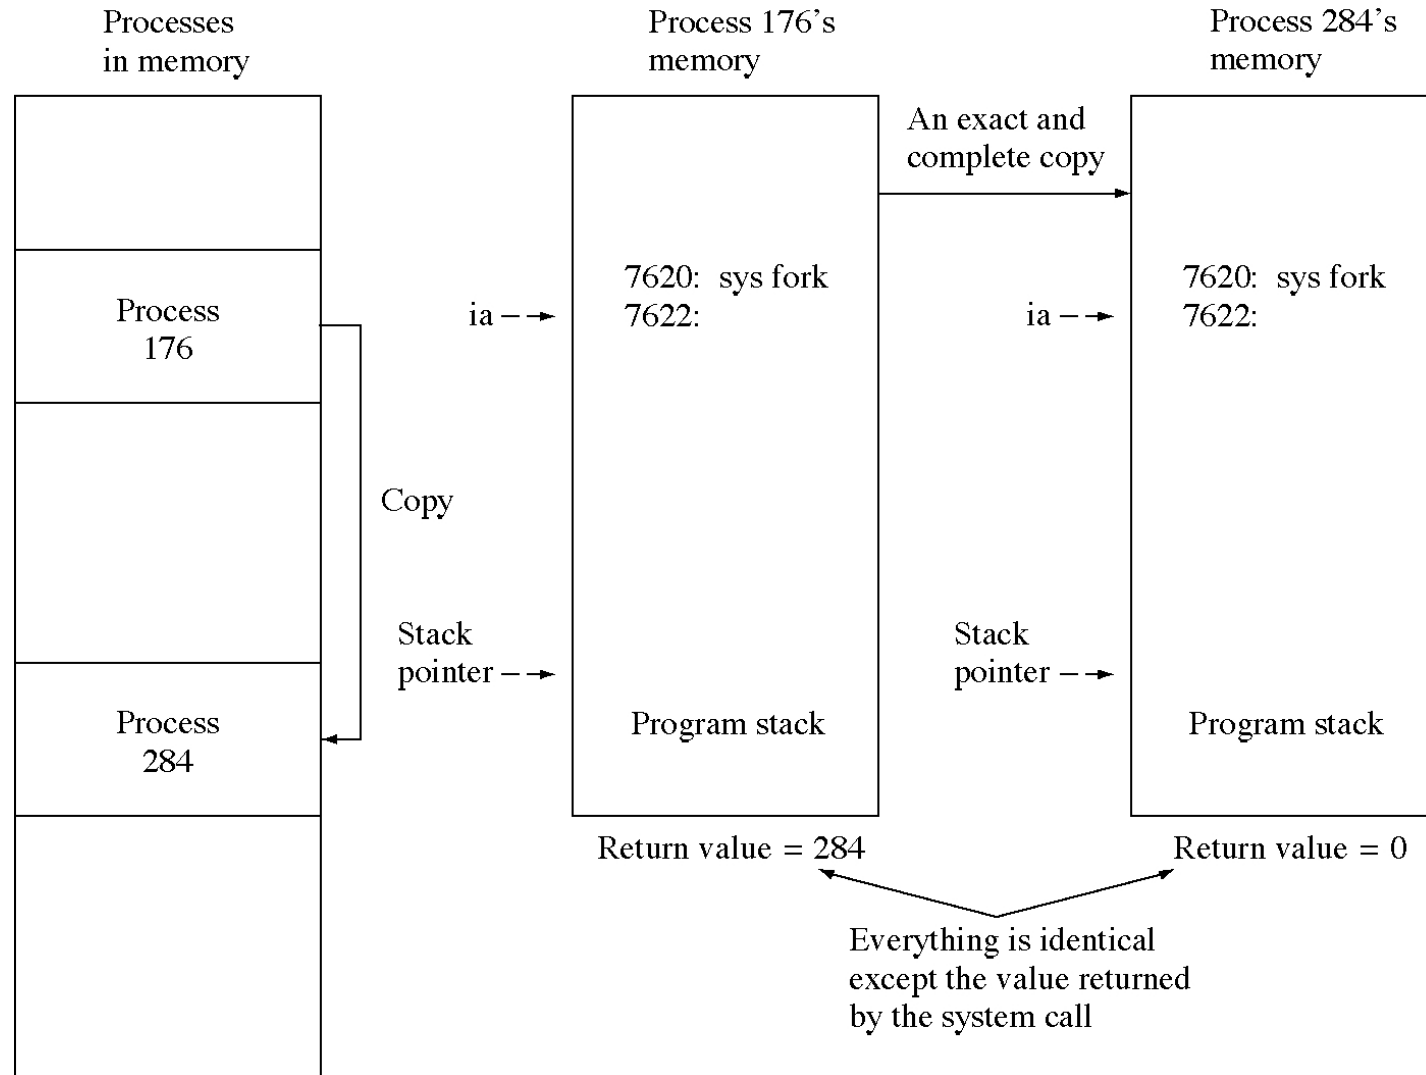
# Objects for sending a file with messages



# UNIX-style process creation

- `int fork()`
  - creates an exact copy of the calling process
- `int execv(char *progName, char *argv[ ])`
  - runs a new program in the calling process
  - destroying the old program
- `int exit(int retCode)`
  - exits the calling process
- `int wait(int *retCode)`
  - waits for any exited child, returns its pid

# UNIX fork



# Create process (UNIX-style)


- ```
void CreateProcess3( void ) {
    int pid1, pid2;
    char *argv[3] = {"compiler", "fileToCompile", 0};
    pid1 = fork();
    if( pid1 == 0 ) { // Child process code begins here
        execv( "compiler", argv ); // execute compiler
        // Child process code ends here.
        // execv does not return
    }
    // Parent executes here because pid1 != 0
    argv[0] = "editor";
    argv[1] = "fileToEdit";
    argv[2] = 0;
    if( (pid2 = fork()) == 0 )
        execv( "editor", argv );
    int reta, retb;
    int pida = wait( &reta );
    int pidb = wait( &retb );
}
```



# Standard input and output

- Most programs are filters:
  - one input stream (standard input)
  - some processing
  - one output stream (standard output)
- So the OS starts a program out with two open files, standard input and standard output
- `grep helvetica <fontList >helvList`

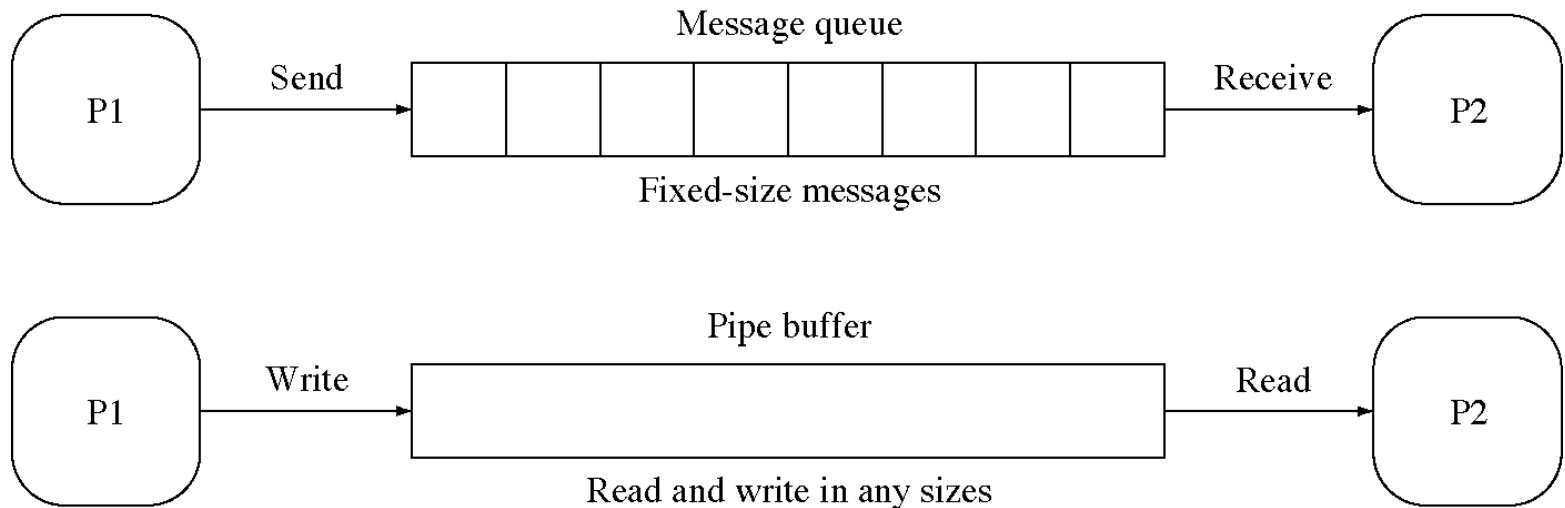
# Redirection of standard input and output



# Pipes

- Pipe: another IPC mechanism
  - uses the familiar file interface
  - not a special interface (like messages)
- Connects an open file of one process to an open file of another process
  - Often used to connect the standard output of one process to the standard input of another process

# Messages and pipes compared



# Pipe: file sender

- ```
enum { Reading=0, Writing=1, ReadAndWrite=2 };
void main( int argc, char * argv[ ] ) {
    int fromFD = open( argv[1], Reading );
    int to_pipe = open( argv[2], Writing );
    while( 1 ) {
        char ch;
        int n = read( fromFD, &ch, 1 );
        if( n == 0 ) break;
        write( to_pipe, &ch, 1 );
    }
    close( fromFD );
    close( to_pipe );
    int n, from_pipe = open( argv[3], Reading );
    // int n is four bytes long, so we read four bytes

    read( from_pipe, &n, 4 );
    close( from_pipe );
    cout << n << " characters\n";
    exit( 0 );
}
```

# Pipe: file receiver

- enum { Reading=0, Writing=1, ReadAndWrite=2 };  
void main( int argc, char \* argv[ ] ) {  
 int count = 0;  
 // The first argument is the pipe to read from.  
 int from\_pipe = open( argv[1], Reading );  
 while( 1 ) {  
 char ch;  
 int n = read( from\_pipe, &ch, 1 );  
 if( n == 0 )  
 break;  
 ++count;  
 }  
 close( from\_pipe );  
 // send the count back to the sender.  
 int to\_pipe = open( argv[2], Writing );  
 write( to\_pipe, &count, 4 );  
 close( to\_pipe );  
 exit( 0 );  
}

# Pipe: create processes

- ```
void main( int argc, char * argv[ ] ) {  
    int pid1 = CreateProcessWithArgs("FileSend",  
        "FileToSend", "PipeToReceiver", "PipeToSender");  
    int pid2 = CreateProcessWithArgs( "FileReceive",  
        "PipeToReceiver", "PipeToSender" );  
    int ret1 = wait( pid1 );  
    int ret2 = wait( pid2 );  
    exit( 0 );  
}
```

# More on naming

- We have seen three naming systems
  - *Global character names* in the file system:  
named pipes
  - *Process-local names* (file identifiers):  
anonymous pipes
  - *Global integer names* picked by the OS:  
message queues
- We can use any of the systems to name objects



# Design technique:

## Connection in protocols

- File interface is a *connection protocol*
  - open (setup), use, close
  - Best for tightly-coupled, predictable connections
- WWW interface is a *connection-less protocol*
  - Each interaction is independent
  - For loose, unpredictable connections

# OS examples

- UNIX (ATT, Bell Labs)
  - Basis of most modern OSes
- Mach (CMU)
  - Microkernel
  - Research system, now widely used
- MS/DOS (Microsoft)
  - Not a full OS

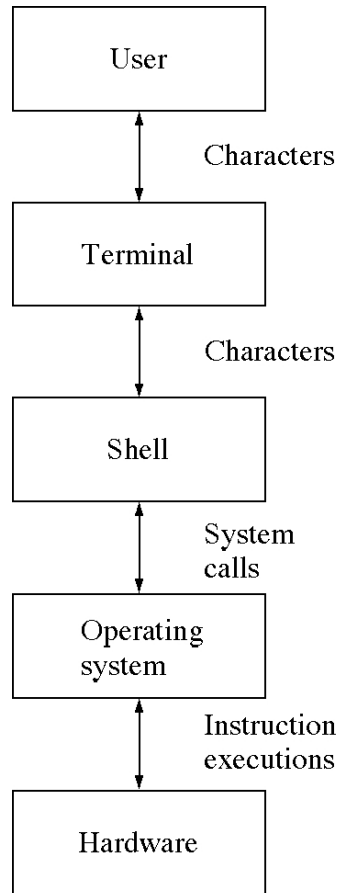
# More OS examples

- Windows NT (Microsoft)
  - Successor to MS/DOS
- OS/2 (IBM)
- Macintosh OS (Apple)
  - Innovations in the GUI
  - To be replaced by Rhapsody (Mach)


# Shell: an OS interface

- Interactive access to the OS system calls
  - copy fromFile toFile
- Contains a simple programming language
- Popularized by UNIX
  - Before UNIX: JCL, OS CLs (command languages)
  - Bourne shell, C shell (csh), Korn shell (ksh), Bourne-again shell (bash), etc.

# Two views of a shell



(a) Shell as a level



(b) Shell as a covering

# Shell: globals

- ```
#include <iostream.h>
// some constants
// maximum size of any one argument
const int ARGSIZE 50
// maximum number of arguments
const int NARGS    20

// token types returned by getWord()
const int STRING    1
const int INREDIR   2
const int OUTREDIR  3
const int NEWLINE   4

// define the argv structure
char *argv[NARGS]; // space for argv vector
char args[NARGS][ARGSIZE]; // space for arguments
```

# Shell (1 of 3)

- ```
void main( int argcount, char *arguments[ ] ) {
    int wasInRedir, wasOutRedir;
    char inRedir[ARGSIZE], outRedir[ARGSIZE];
    // each iteration will parse one command
    while( 1 ) {
        // display the prompt
        cout << "@ ";
        // So far we have not seen any redirections
        wasInRedir = 0;
        wasOutRedir = 0;
        // Set up some other variables.
        int argc = 0;
        int done = 0;
        char word[ARGSIZE];
        // Read one line from the user.
        while( !done ) {
            // getWord gets one word from the line.
            int argType = getWord(word);
            // getWord returns the type of the word it read
```

# Shell (2 of 3)

- ```
    switch( argType ) {
case INREDIR:
    wasInRedir = 1;
    (void)getWord(inRedir);
    break;
case OUTREDIR:
    wasOutRedir = 1;
    (void)getWord(outRedir);
    break;
case STRING:
    strcpy(args[argc], word);
    argv[argc] = &args[argc][0];
    ++argc;
    break;
case NEWLINE:
    done = 1;
    break;
    }
}
argv[argc] = NULL;
if( strcmp(args[0], "logout") == 0 )
    break;
```



# Shell (3 of 3)

- ```
if( fork() == 0 ) {
    if( wasInRedir ) {
        close(0); // close standard input
        open(inRedir, 0); //reopen as redirect file }
    if( wasOutRedir ) {
        close(1); // close standard output
        enum { UserWrite=0755 };
        creat(outRedir, UserWrite); }
    char cmd[60];
    strcpy(cmd, "./"); strcat(cmd, args[0]);
    execv(cmd, &argv[0]);
    strcpy(cmd, "/bin/"); strcat(cmd, args[0]);
    execv(cmd, &argv[0]);
    cout << "Child: could not exec \"" << args[0]
        << "\"\n";
    exit(1);
}
int status; (void) wait(&status);
}
cout << "Shell exiting.\n";
}
```

# Design technique: Interactive and programming interfaces

- Interactive interfaces have advantages:
  - for exploration
  - for interactive use
- Programming interfaces have advantages :
  - for detailed interactions
  - Inter-application programming
  - Scripting, COM, CORBA
- It is useful for a program to have both