

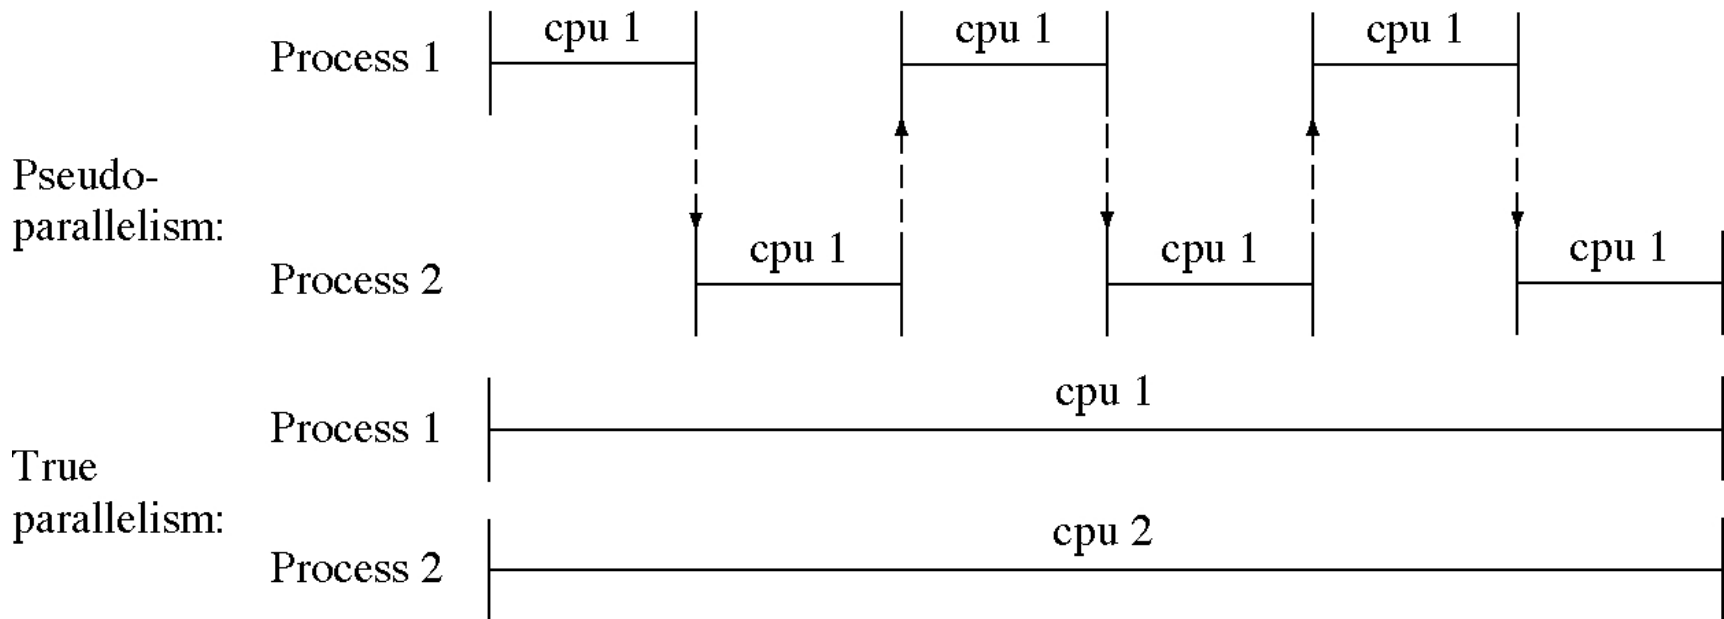
# Parallel Systems

## Chapter 6

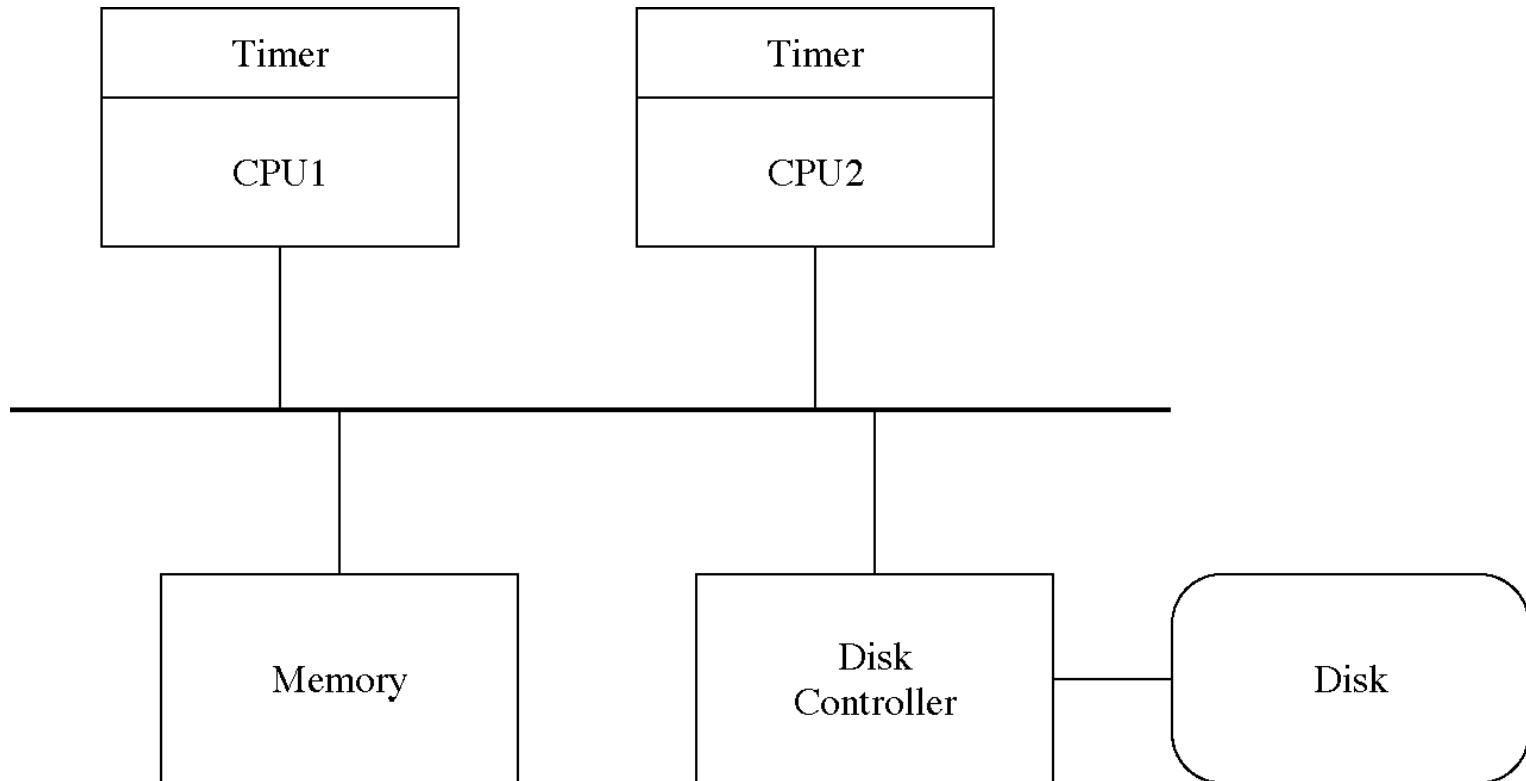
# Key concepts in chapter 6

- Parallelism
  - physical (a.k.a. true) parallelism
  - logical (a.k.a. pseudo-) parallelism
- Multiprocessor OS (MPOS)
  - race conditions
  - atomic actions
  - spin locks
- Threads
- Kernel-mode processes
- Implementation of mutual exclusion


# Parallel actions in a program



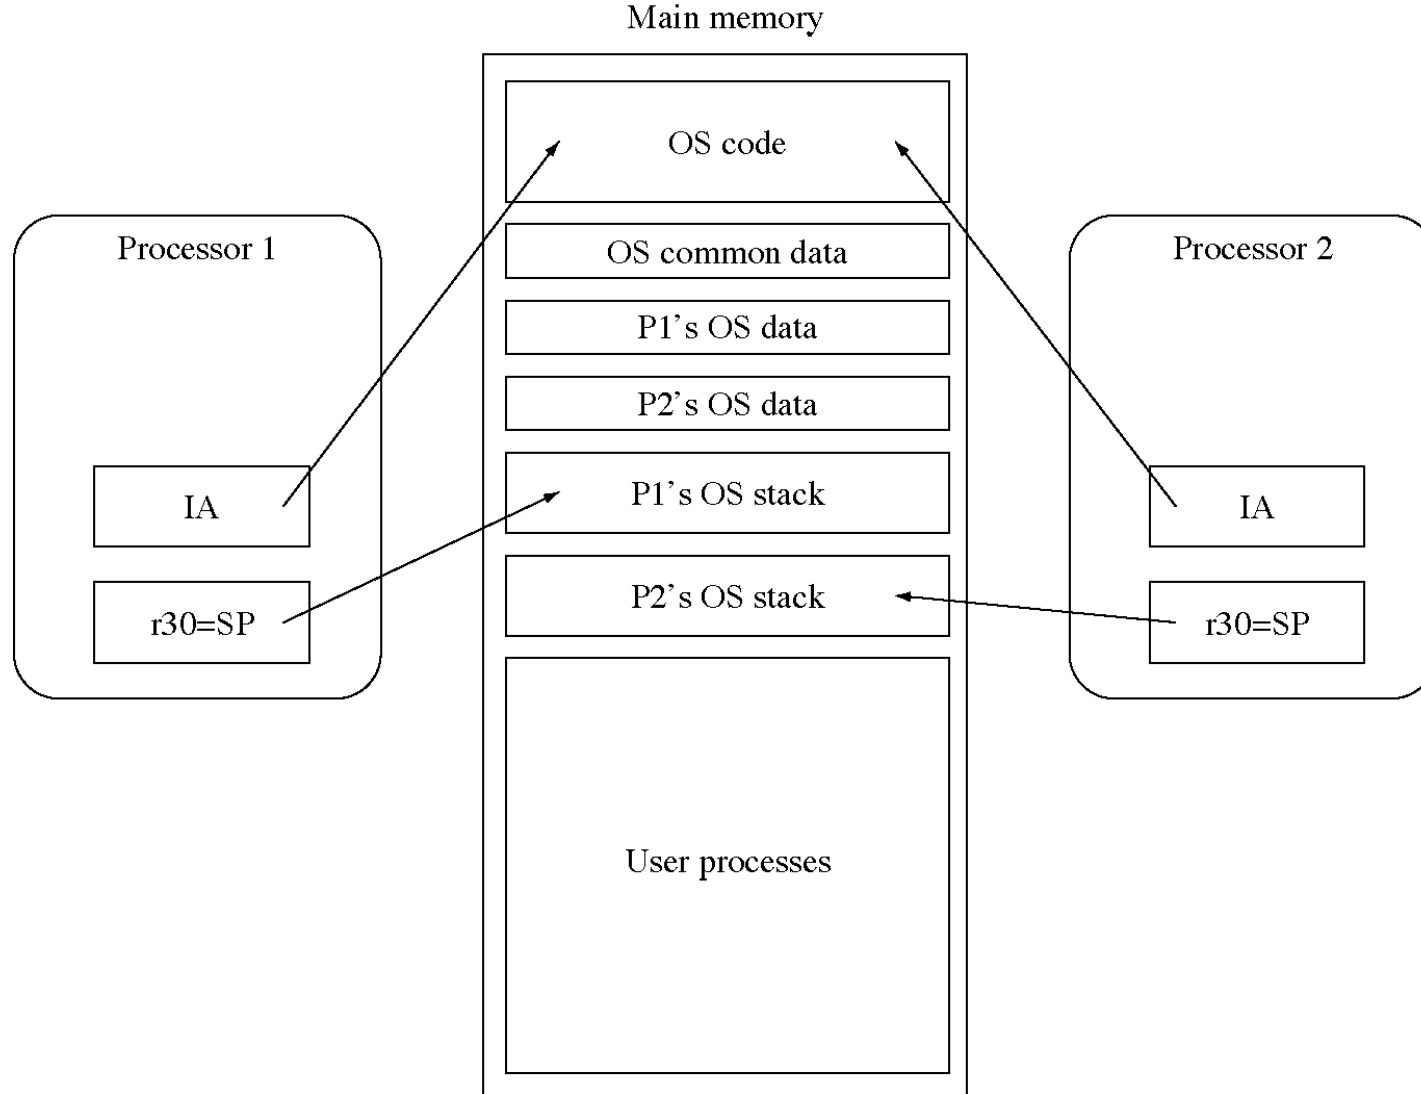
# Parallel hardware



# Two operating systems



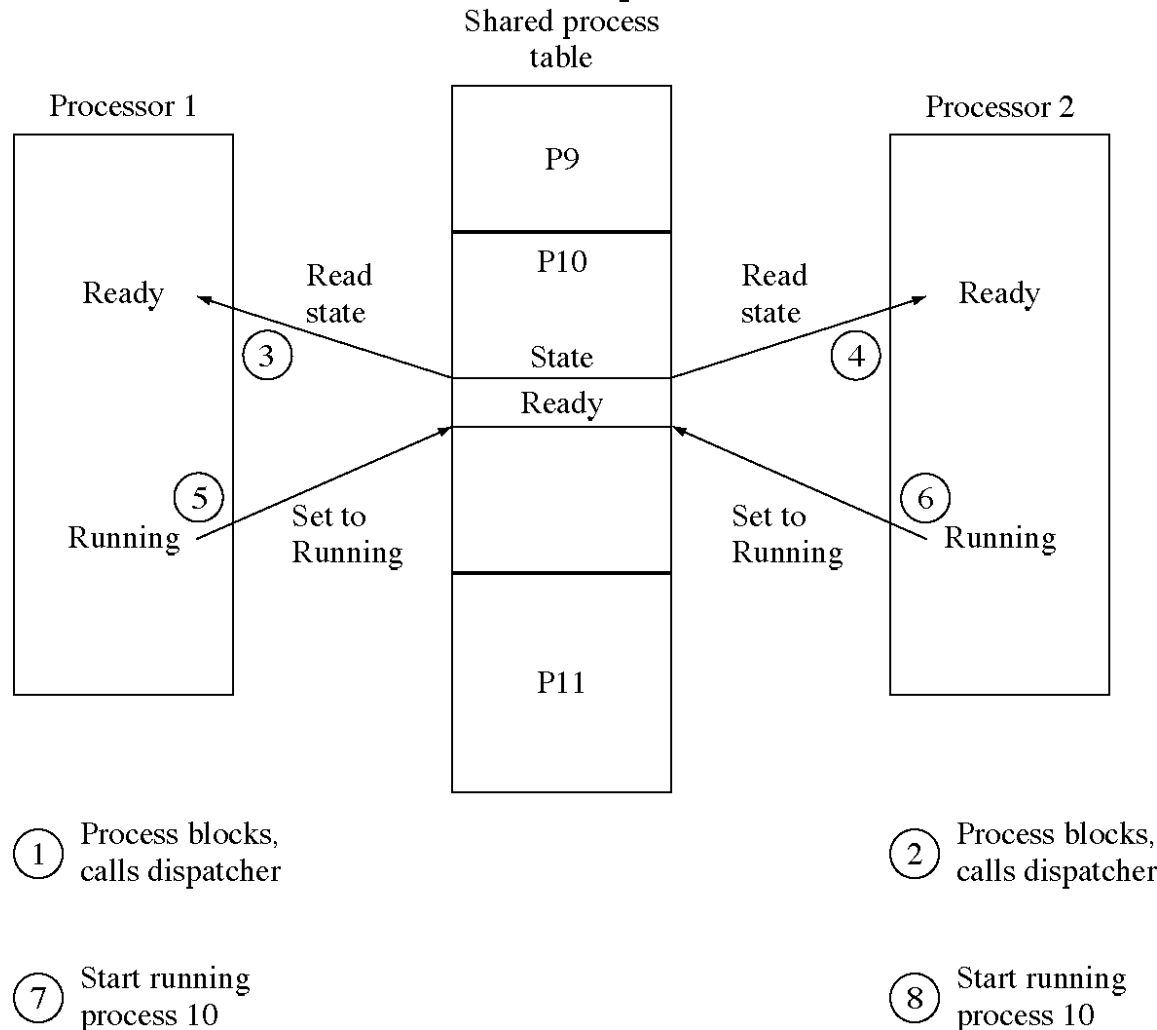
# Sharing the OS between processors



# Global data

- Shared (between the processors)
  - process\_using\_disk
  - disk queue
  - process table (pd)
  - message queues
- Duplicated (one for each processor)
  - stack
  - current\_process

# Race condition with a shared process table





# Atomic actions with the ExchangeWord instruction

Allowed	Allowed	Not allowed	Not allowed	Not allowed	Not allowed
<pre> sequenceDiagram     participant P1     participant P2     Note over P1: P1 read     Note over P1: P1 write     Note over P2: P2 read     Note over P2: P2 write             </pre>	<pre> sequenceDiagram     participant P1     participant P2     Note over P2: P2 read     Note over P2: P2 write     Note over P1: P1 read     Note over P1: P1 write             </pre>	<pre> sequenceDiagram     participant P1     participant P2     Note over P1: P1 read     Note over P2: P2 read     Note over P2: P2 write     Note over P1: P1 write             </pre>	<pre> sequenceDiagram     participant P1     participant P2     Note over P1: P1 read     Note over P2: P2 read     Note over P1: P1 write     Note over P2: P2 write             </pre>	<pre> sequenceDiagram     participant P1     participant P2     Note over P1: P1 read     Note over P2: P2 read     Note over P2: P2 write     Note over P1: P1 write             </pre>	<pre> sequenceDiagram     participant P1     participant P2     Note over P2: P2 read     Note over P1: P1 read     Note over P1: P1 write     Note over P2: P2 write             </pre>

Time  
↓

# Hardware implementation of atomic actions

- Only the “bus master” can access memory
- Solution
  - allow one processor to be the bus master for two bus cycles
  - one to read the word
  - one to write the word
  - Presto! An atomic read/write action

# MPOS: global data

- `// Each processor has a stack and a current process`  
`int p1_current_process;`  
`int p2_current_process;`  
`int p1_SystemStack[SystemStackSize];`  
`int p2_SystemStack[SystemStackSize];`  
`// Shared process table`  
`ProcessDescriptor pd[NumberOfProcesses];`
- `// Shared message queues`  
`MessageBuffer`  
`message_buffer[NumberOfMessageBuffers];`  
`int free_message_buffer;`  
`int message_queue_allocated[NumberOfMessageQueues];`  
`Queue<MessageBuffer *>`  
`* message_queue[NumberOfMessageQueues];`  
`Queue<WaitQueueItem *>`  
`* wait_queue[NumberOfMessageQueues];`  
  
`// Shared disk data`  
`int process_using_disk;`  
`Queue<DiskRequest *> disk_queue;`

# MPOS: system initialization

```
• int main( void ) { // Proc. 1 system initialization

    // Set up the system's stack
    asm{ load  p1_SystemStack+SystemStackSize,r30 }
    // set up the interrupt vectors ... as before
    // set up the pds ... as before
    // process 1 is the initialization process
    // ... as before
    // set up the pool of free message buffers
    // ... as before
    process_using_disk = 0;
    p1_current_process = 0;
    // start processor 2
    Dispatcher();
}

int main( void ) { // Proc. 2 system initialization
    // Set up the system's stack
    asm{ load  p2_SystemStack+SystemStackSize,r30 }
    p2_current_process = 0;
    Dispatcher();
```

# Protecting the process table

- ```
// A global variable to restricts access to the
// process table
int ProcessTableFreeFlag = 1;
void StartUsingProcesstTable( void ) {
    int flag;
    while( 1 ) {
        asm{
            move    r0,r1
            exchangeword r1,ProcessTableFreeFlag
            store   r1,flag
        }
        // Busy wait until we get a 1 from the flag.
        if( flag == 1 )
            break;
    }
}
void FinishUsingProcesstTable( void ) {
    ProcessTableFreeFlag = 1;
}
```

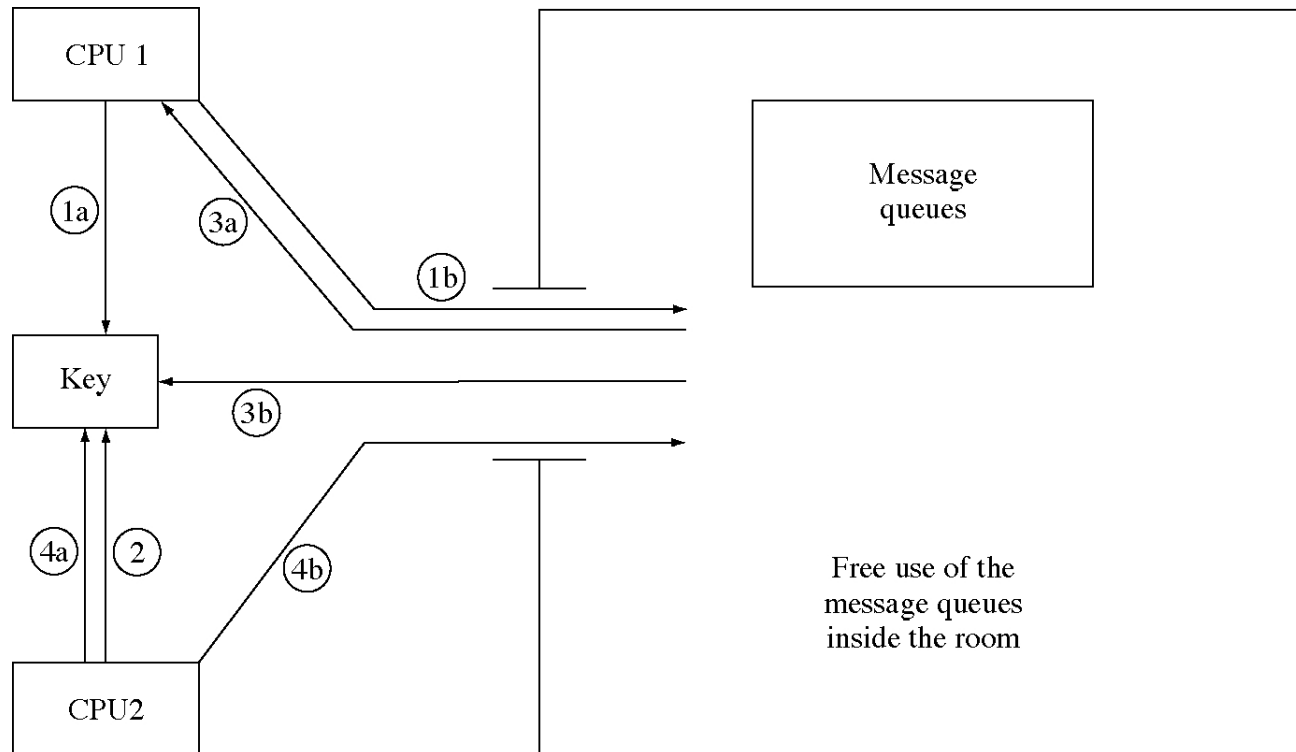
# MPOS: dispatcher

```
• int SelectProcessToRun( void ) {
    static int next_proc = NumberOfProcesses;
    int i, return_value = -1; int current_process;
    StartUsingProcessTable(); // <----- NEW CODE
    if( ProcessorExecutingThisCode() == 1 )
        current_process = p1_current_process;
    else current_process = p2_current_process;
    if( current_process > 0 && pd[current_process].state == Ready
        && pd[current_process].timeLeft > 0 ) {
        pd[current_proc].state = Running;
        return_value = current_process;
    } else {
        for( i = 1; i < NumberOfProcesses; ++i ) {
            if( ++next_proc >= NumberOfProcesses ) next_proc = 1;
            if( pd[next_proc].slotAllocated
                && pd[next_proc].state == Ready ) {
                pd[next_proc].state = Running;
                pd[next_proc].timeLeft = TimeQuantum;
                return_value = next_proc;
                break;
            }
        }
    }
    FinishUsingProcessTable(); // <----- NEW CODE
    return return_value;
}
```

# Busy waiting

- *Processes* wait by being suspended and resumed when the desired event occurs
- *Processors* wait by continually checking the lock word
  - this is called busy waiting
  - the lock is called a spin lock


# Protecting all the message queues with one lock



1. CPU1 (a) gets key and (b) enters the room.
2. CPU2 fails to get the key but keeps on trying.
3. CPU1 (a) leaves the room and (b) returns the key.
4. CPU2 (a) succeeds in getting the key and (b) enters the room.



# Protecting each message queue with a separate lock



# Protecting access to shared data

- ```
void StartUsingSharedMemory(
    int * shared_memory_flag ) {
    int flag;
    while( 1 ) {
        // check first to see if we have a chance
        while( * shared_memory_flag == 0 )
            ;
        asm{
            move    r0,r1
            exchangelword r1,* shared_memory_flag
            store   r1,flag
        }
        // Busy wait until we get a 1 from the flag.
        if( flag == 1 ) break;
    }
}

void FinishUsingSharedMemory(
    int * shared_memory_flag ) {
    *shared_memory_flag = 1;
}
```

# Protecting the message queues

- ```
int message_queue_flag[NumberOfMessageQueues]
    = {1,1,1,...,1};
case SendMessageSystemCall:
    int * user_msg; asm { store r9,user_msg }
    int to_q; asm { store r10,to_q }
    if( !message_queue_allocated[to_q] ) {
        pd[current_process].sa.reg[1] = -1; break;
    }
    int msg_no = GetMessageBuffer();
    if( msg_no == EndOfFreeList ) {
        pd[current_process].sa.reg[1] = -2; break;
    }
    CopyToSystemSpace( current_process, user_msg,
        message_buffer[msg_no], MessageSize )
    StartUsingSharedMemory(&message_queue_flag[to_q]);
    if( !wait_queue[to_q].Empty() ) {
        WaitQueueItem item = wait_queue.Remove();
        TransferMessage( msg_no, item.buffer );
        pd[item.pid].state = Ready;
    } else { message_queue[to_q].Insert( msg_no ); }
    FinishUsingSharedMemory(&message_queue_flag[to_q]);
    pd[current_process].sa.reg[1] = 0;
    break;
```

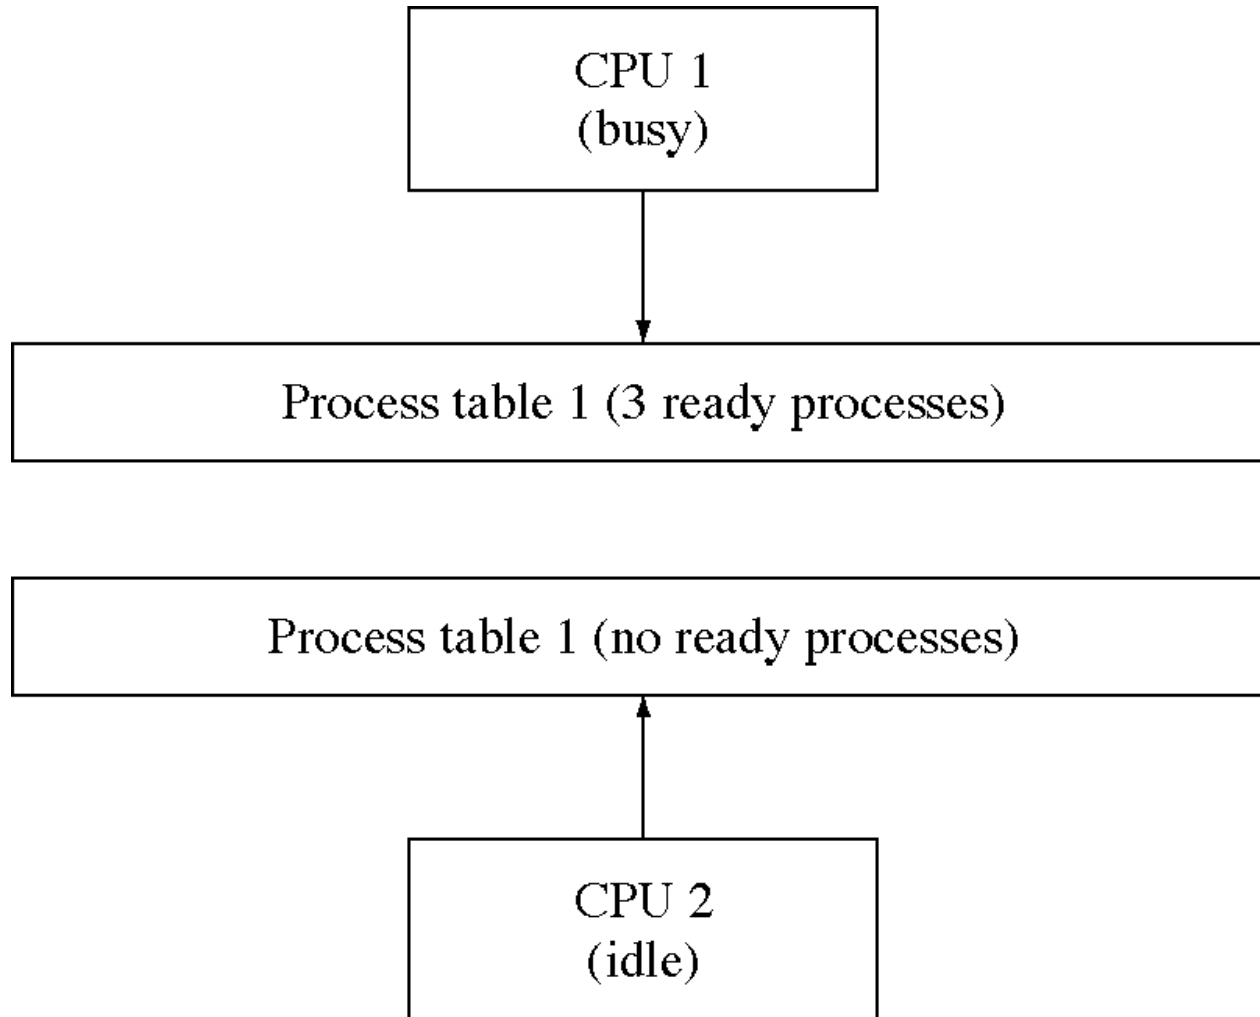
# Protecting the message buffers

- `int message_buffer_flag = 1; //`

```
int GetMessageBuffer( void ) {  
    // get the head of the free list  
    StartUsingSharedMemory( &message_buffer_flag );  
    int msg_no = free_msg_buffer;  
    if( msg_no != EndOfFreeList ) {  
        // follow the link to the next buffer  
        free_msg_buffer = message_buffer[msg_no][0];  
    }  
    FinishUsingSharedMemory( &message_buffer_flag );  
    return msg_no;  
}
```

```
void FreeMessageBuffer( int msg_no ) {  
    StartUsingSharedMemory( &message_buffer_flag );  
    message_buffer[msg_no][0] = free_msg_buffer;  
    free_msg_buffer = msg_no;  
    FinishUsingSharedMemory( &message_buffer_flag );  
}
```

# Using two process tables




# Threads


- A **process** has two characteristics
  - *resources*: a process can hold resources such as an address space and open files.
  - *dispatchability*: a process can be scheduled by the dispatcher and execute instructions.
- We can split these two characteristics into:
  - a (new) process which holds resources
  - a **thread** which can be dispatched and execute instructions.
  - several threads can exist in the same process, that is, in the same address space.

# Threads executing in two processes

Process 1's code space



Process 2's code space



# Process and thread analogies

- A process is the software version of a standalone computer
  - One processor (for each computer or process), no shared memory
  - Communicate (between computers or processes) with messages
- A thread is the software version of a processor in a multiple-processor, shared-memory computer
  - Multiple processors (threads), shared memory
  - Communicate (between processors or threads) through shared memory



# Thread-related system calls

- `int CreateThread(  
    char *startAddress, char *stackAddress)`
- `int ExitThread(int returnCode)`
- `WaitThread(int tid)`
- Basically the same as for processes

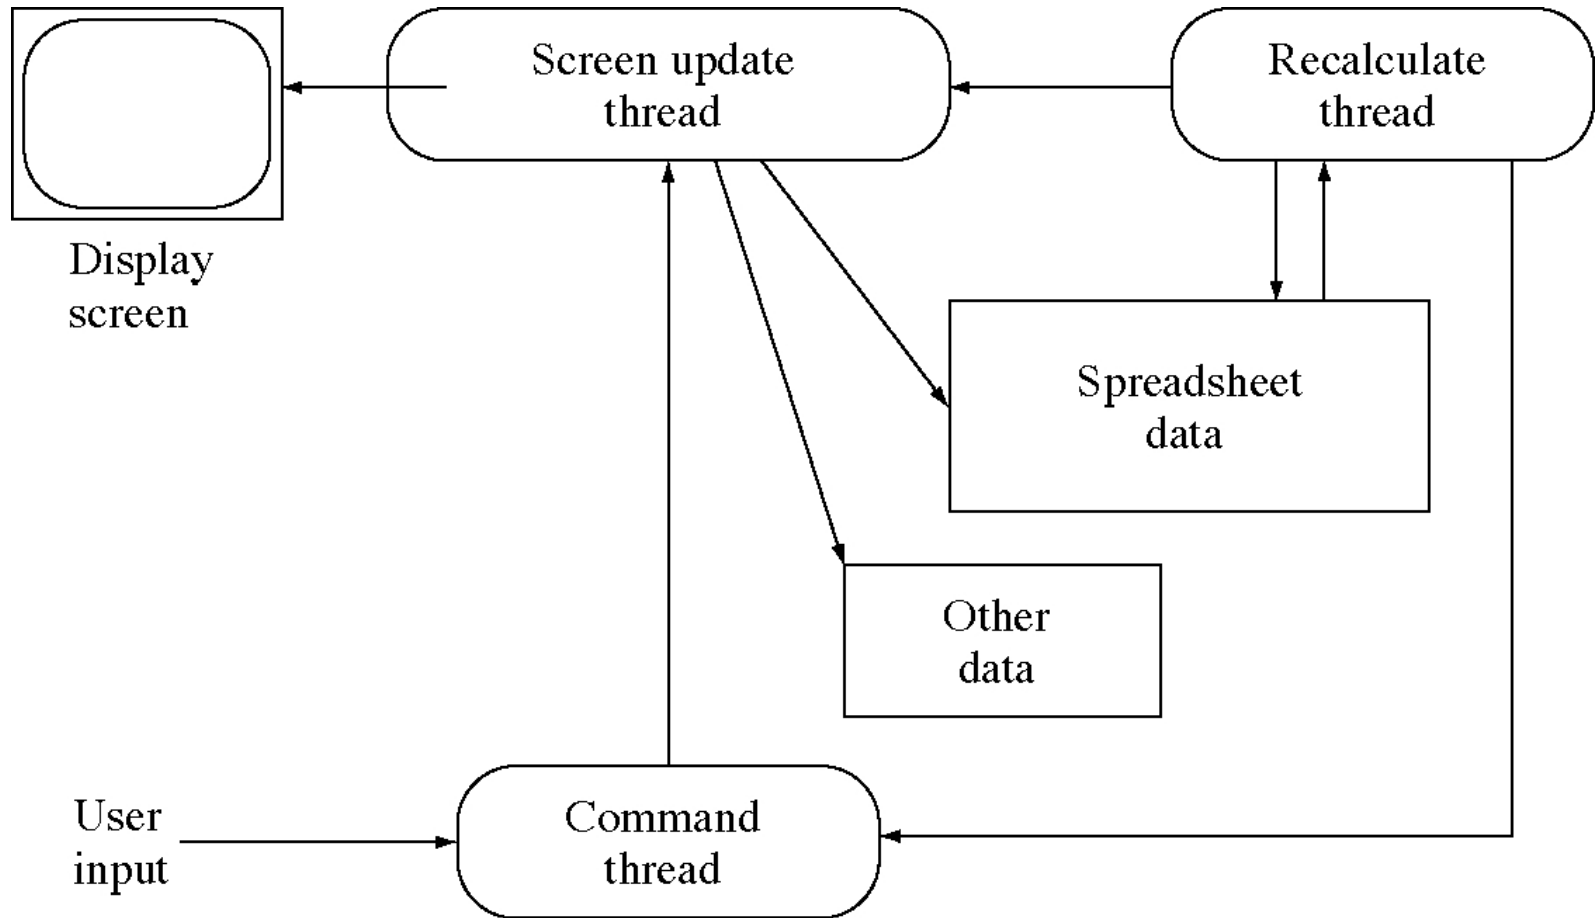
# Advantages of threads

- Threads allows parallel activity inside a single address space
  - Inter-thread communication and synchronization is very fast
- Threads are cheap to create
  - mainly because we do not need another address space
  - They are also called *lightweight processes* (LWPs)

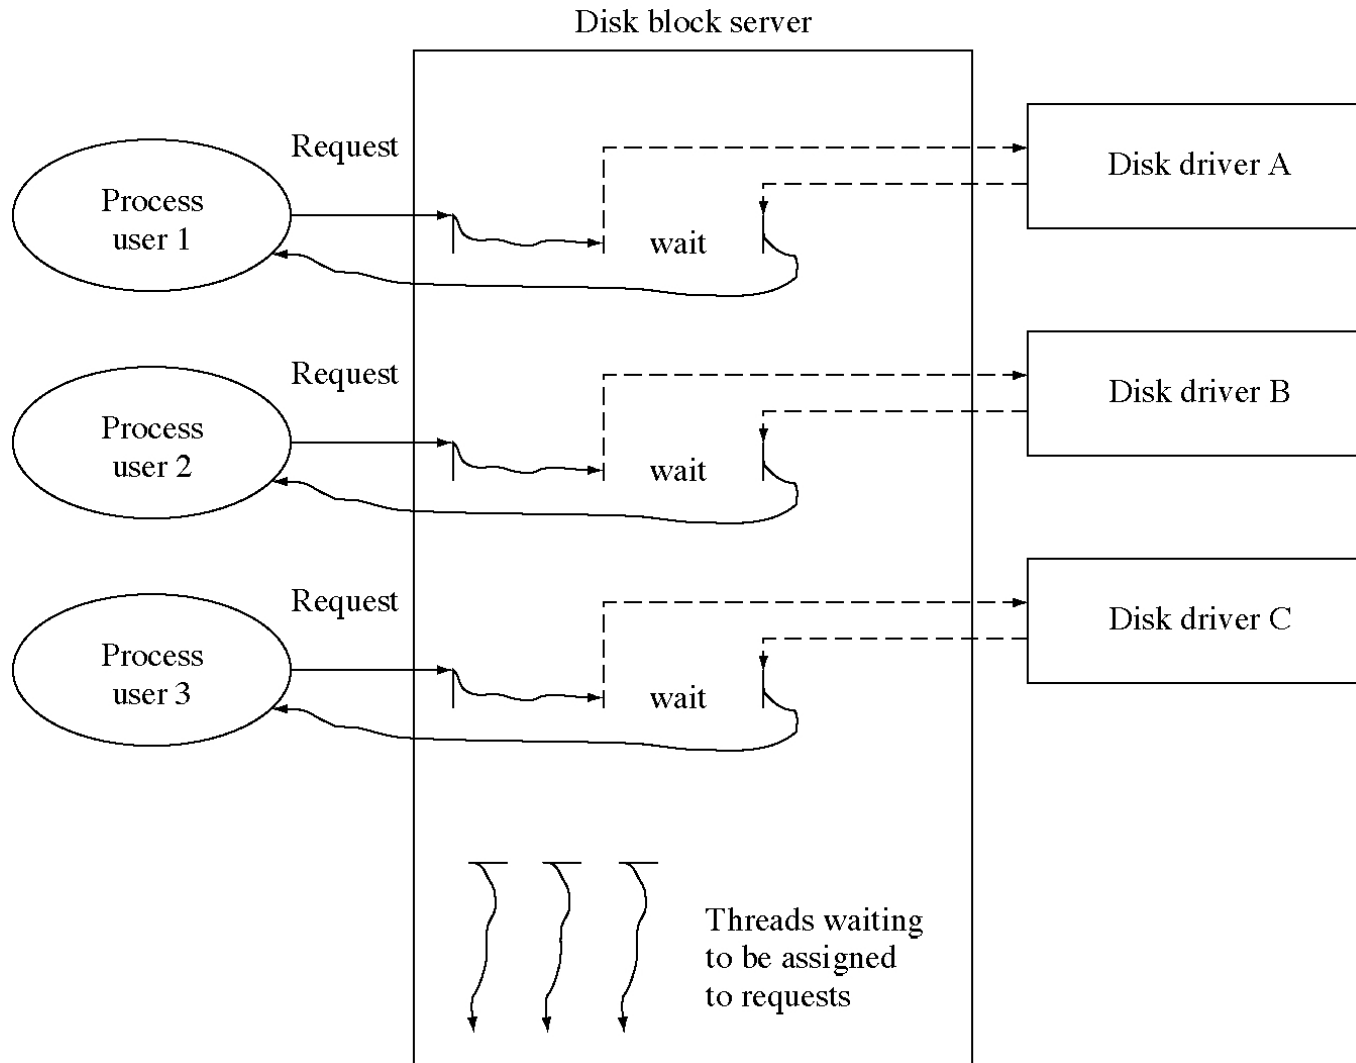
# Uses of threads

- Partition process activities
  - each part of the process has its own thread
- Waiting
  - threads can wait for events without holding up other threads
- Replicate activities
  - servers can allocate a thread for each request

# Threads in a spreadsheet program



# Disk server using threads



# Threads: global data

- ```
enum ThreadState { Ready, Running, Blocked };
struct ThreadDescriptor {
    int slotAllocated;
    int timeLeft;
    int nextThread, prevThread;
    int pid;
    ThreadState state;
    SaveArea sa;
};
struct ProcessDescriptor {
    int slotAllocated;
    int threads; void * base, bound;
};
// This is the thread currently running.
int current_thread;
// We need a process table and a thread table.
ProcessDescriptor pd[NumberOfProcesses];
// pd[0] is the systemThreadDescriptor
td[NumberOfThreads];
int thread_using_disk;
```

# Threads: system initialization

- ```
int main( void ) {  
    // Other initialization is the same ...  
    // The thread slots start out free.  
    for( i = 1; i < NumberOfThreads; ++i )  
        td[i].slotAllocated = False;  
    // Other initialization is the same ...  
}
```

# Threads: create process

```
• int CreateProcess( int first_block, int n_blocks ) {
    int pid;
    for( pid = 1; pid < NumberOfProcesses; ++pid )
        if( !(pd[pid].slotAllocated) ) break;
    if( pid >= NumberOfProcesses ) return -1;
    pd[pid].slotAllocated = True;
    pd[pid].base = pid * ProcessSize;
    pd[pid].bound = ProcessSize;
    char * addr = (char *)(pd[pid].sa.base);
    for( i = 0; i < n_blocks; ++i ) {
        while( DiskBusy() ) ;    // Busy wait for I/O.
        IssueDiskRead( first_block + i, addr, 0 );
        addr += DiskBlockSize;
    }
    int tid = CreateThread( pid, 0, ProcessSize );
    pd[pid].threads = tid;
    if( tid == -1 ) return -1;
    td[tid].nextThread = td[tid].prevThread = tid;
    return pid;
}
```



# Threads: create thread

- ```
int CreateThread(  
    int pid, char * startAddress, char *stackBegin ) {  
    // startAddress = address in the process  
    //    address space to begin execution.  
    // stackBegin = initial value of r30  
    int tid;  
    for( tid = 1; tid < NumberOfThreads; ++tid )  
        if( !(td[tid].slotAllocated) ) break;  
    if( tid >= NumberOfThreads ) {  
        return -1; }  
    td[tid].slotAllocated = True;  
    td[tid].pid = pid;  
    td[tid].state = Ready;  
    td[tid].sa.base = pd[pid].base;  
    td[tid].sa.bound = pd[pid].bound;  
    td[tid].sa.psw = 3;  
    td[tid].sa.ia = startAddress;  
    td[tid].sa.r30 = stackBegin;  
    return tid;  
}
```

# Threads: dispatcher

- ```
void Dispatcher( void ) {
    current_thread = SelectThreadToRun();
    RunThread( current_thread ); }
int SelectThreadToRun( void ) {
    static int next_thread = NumberOfThreads;
    if( current_thread > 0 && td[current_thread].slotAllocated
        && td[current_thread].state == Ready
        && td[current_thread].timeLeft > 0 ) {
        td[current_thread].state = Running;
        return current_thread; }
    for( int i = 0; i < NumberOfThreads; ++i ) {
        if( ++next_thread >= NumberOfThreads ) next_thread = 0;
        if( td[next_thread].slotAllocated
            && td[next_thread].state == Ready ) {
            td[next_thread].state = Running;
            return next_thread;
        }
    }
    return -1;    // No thread is ready to run
}
void RunThread( int tid ) {
    /* ... same except it uses tid instead of pid */}
```

# Threads: system calls (1 of 2)

- ```
void SystemCallInterruptHandler( void ) {  
    // The Exit system call is eliminated.  You exit a  
    // process by exiting from all its threads.  
    case CreateProcessSystemCall:  
        int block_number; asm { store r9,block_number }  
        int number_of_blocks;  
        asm { store r10, number_of_blocks }  
        td[current_thread].sa.reg[1]  
        = CreateProcess( block_number, number_of_blocks);  
        break;  
    case CreateThreadSystemCall:  
        int start_addr; asm { store r9,start_addr }  
        int stack_begin; asm { store r10,stack_begin }  
        int pid = td[current_thread].pid;  
        int new_thread  
        = CreateThread( pid, start_addr, stack_begin );  
        td[current_thread].sa.reg[1] = new_thread;  
        int next_thread = td[current_thread].nextThread;  
        td[new_thread].nextThread = next_thread;  
        td[new_thread].prevThread = current_thread;  
        td[current_thread].nextThread = new_thread;  
        td[next_thread].prevThread = new_thread;  
        break;
```

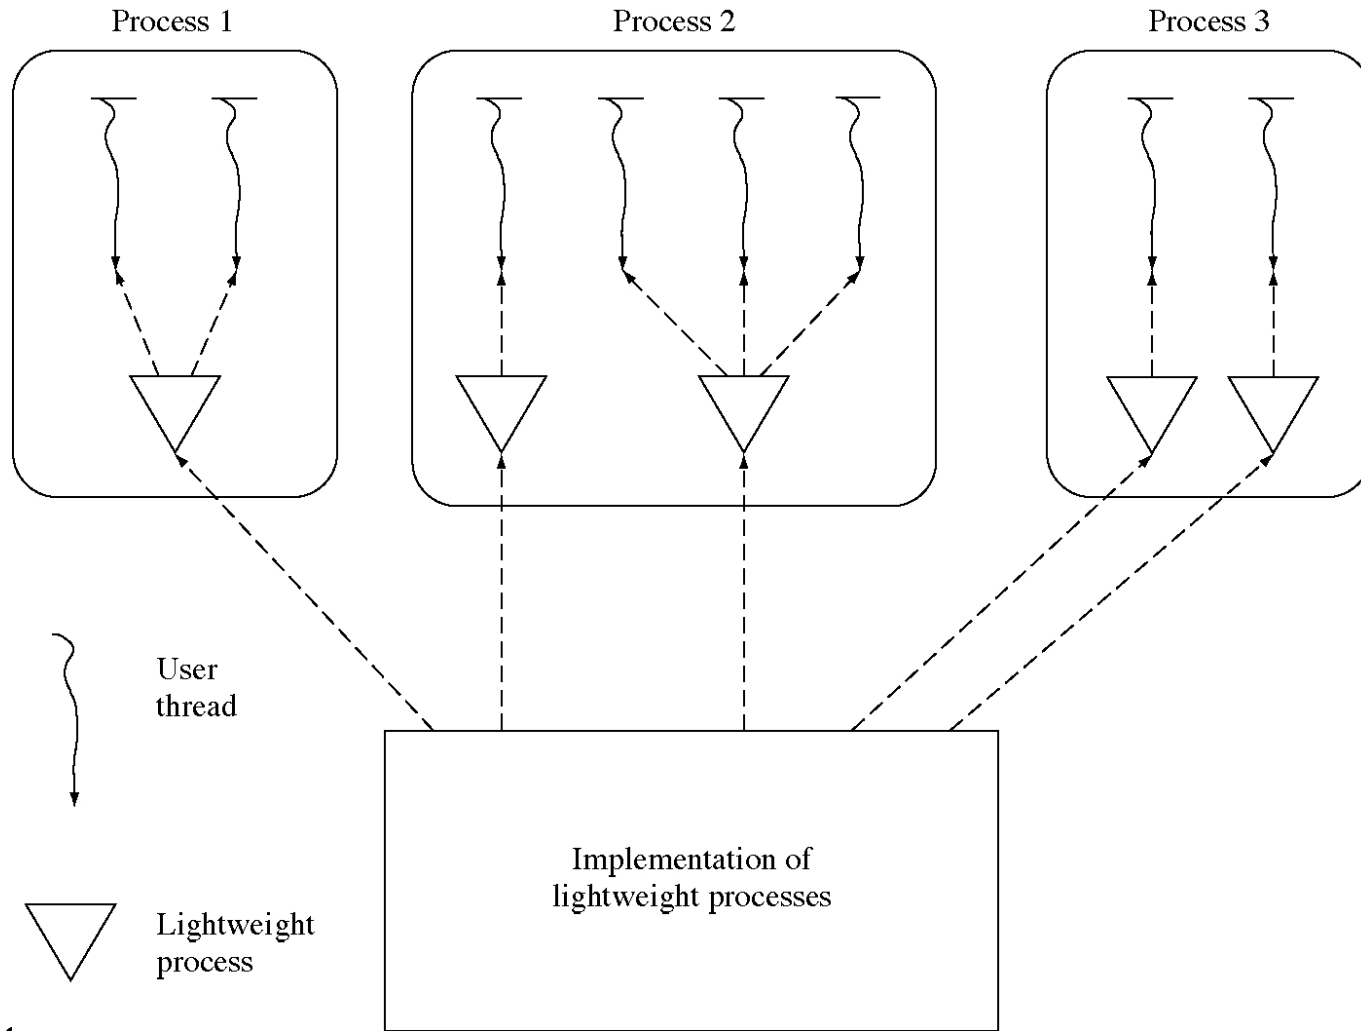
# Threads: system calls (2 of 2)

- ```
case ExitThreadSystemCall:
    td[current_thread].slotAllocated = False;
    int next_thread = td[current_thread].nextThread;
    int pid = td[current_thread].pid;
    if( next_thread == current_thread ) {
        pd[pid].slotAllocated = False;
        // No other exit process processing. In a real OS
        // we might free memory, close unclosed files, etc.
    } else {
        // Unlink it from the list and make sure pd[pid]
        // is not pointing to the deleted thread.
        int prev_thread = td[current_thread].prevThread;
        td[next_thread].prevThread = prev_thread;
        td[prev_thread].nextThread = next_thread;
        pd[pid].threads = next_thread;
    }
    break;
}
Dispatcher();
}
```

# Kinds of threads

- We have described *lightweight processes*
  - threads implemented by the OS
- A process can implement *user threads*
  - threads implemented (solely) by the user
  - these are more efficient
  - but if the OS blocks one of them, they are all blocked
- *Kernel threads* are threads that run inside the OS (a.k.a. the kernel)

# Combining user threads and lightweight processes



# Threads in real OSs

- Almost all modern OSs implement threads
- Plan 9 has a very flexible form of process creation instead of threads
- Several user thread packages are available
- Some programming languages implement threads in the language
  - Ada, Modula 3, Java

# Design technique: Separation of concepts

- We separated resource holding and dispatchability into separate concepts: process and thread
  - This allowed us to have several threads inside a single process
- If two concepts can be used separately it is often useful to separate them in the software



# Design technique: Reentrant programs

- Threads allow two threads of control to be executing in the same code at the same time
- This leads to sharing problems
  - such code must be reentrant
  - this is the same problem we had with two processor sharing OS global data

# Kernel-mode processes

- Many UNIX implementations have long allowed processes to execute inside the kernel (the OS) during system calls
  - this is actually an example of kernel threads
  - it neatly solves the problem of suspending a system call

# Kernel-mode: globals

- ```
struct ProcessDescriptor {
    int slotAllocated;
    int timeLeft; // time left from the last time slice
    ProcessState state;
    // SaveArea sa; <---- ELIMINATED
    int inSystem; // set to 0 in CreateProcess <--- NEW
    int * lastsa; // most recent save area on the stack
                // <-- NEW
    char sstack[SystemStackSize]; // system mode stack
                                //<---- NEW};
    // We don't need a common system stack any more
    // int SystemStack[SystemStackSize]; <----
    ELIMINATED
    // CHANGED: this is now a queue of ints.
    Queue<int> * wait_queue[NumberOfMessageQueues];
```

# Kernel-mode: create process

- ```
int CreateProcessSysProc(
    int first_block, int n_blocks ) {
    // ... the beginning part is the same as it was
    // before except this is added:
    pd[pid].inSystem = 0;
    // Read in the image of the process.
    char * addr = (char *)(pd[pid].sa.base);
    for( i = 0; i < n_blocks; ++i ) {
        DiskIO(DiskReadSystemCall, first_block+i, addr);
        addr += DiskBlockSize;
    }
    return pid;
}
```

# System calls (1 of 5)

- ```
void SystemCallInterruptHandler( void ) {  
    // *** BEGIN NEW CODE ***  
    //      All this initial interrupt handling is new.  
    int * savearea;  
    int saveTimer;  
    if( pd[current_process].inSystem == 0 ) {  
        // This is the first entry into system mode  
        savearea = &(pd[current_process].sstack  
                    +SystemStackSize-sizeof(SaveArea)-4);  
    } else {  
        // we were already in system mode so the system  
        // stack already has some things on it.  
        asm { store r30,savearea }  
        // allocate space on the stack for the save area  
        savearea -= sizeof(SaveArea)+4;  
    }  
}
```

# System calls (2 of 5)

- asm {  
    store    timer, saveTimer  
    load     #0, timer  
    store    iia, savearea+4  
    store    ipsw, savearea+8  
    store    base, savearea+12  
    store    bound, savearea+16  
    storeall  savearea+20  
    load     savearea, r30  
}  
pd[current\_process].timeLeft = saveTimer;  
pd[current\_process].state = Ready;  
\*savearea = pd[current\_process].lastsa;  
// link to previous save area  
pd[current\_process].lastsa = savearea;  
++(pd[current\_process].inSystem);  
// state saved, so enable interrupts  
asm { load #2, psw }  
// \*\*\* END NEW CODE \*\*\* for interrupt handling  
code

# System calls (3 of 5)

- ```
// fetch the system call number and switch on it
int system_call_number;
asm { store r8,system_call_number }
switch( system_call_number ) {
case CreateProcessSystemCall:
    // ... the same
case ExitProcessSystemCall:
    // ... the same
case CreateMessageQueueSystemCall:
    // ... the same
case SendMessageSystemCall:
    // get the arguments
    int * user_msg; asm { store r9,user_msg }
    int to_q; asm { store r10,to_q }
    // check for an invalid queue identifier
    if( !message_queue_allocated[to_q] ) {
        pd[current_process].sa.reg[1] = -1;
        break;
    }
```

# System calls (4 of 5)

- ```
int msg_no = GetMessageBuffer();
// make sure we are not out of message buffers
if( msg_no == EndOfFreeList ) {
    pd[current_process].sa.reg[1] = -2;
    break;
}
// copy message vector from the system caller's
// memory into the system's message buffer
CopyToSystemSpace( current_process, user_msg,
    message_buffer[msg_no], MessageSize );
if( !wait_queue[to_q].Empty() ) {
    // process is waiting for a message, unblock
it
    int pid = wait_queue.Remove();
    pd[pid].state = Ready;
}
// put it on the queue
message_queue[to_q].Insert( msg_no );
pd[current_process].sa.reg[1] = 0;
break;
```



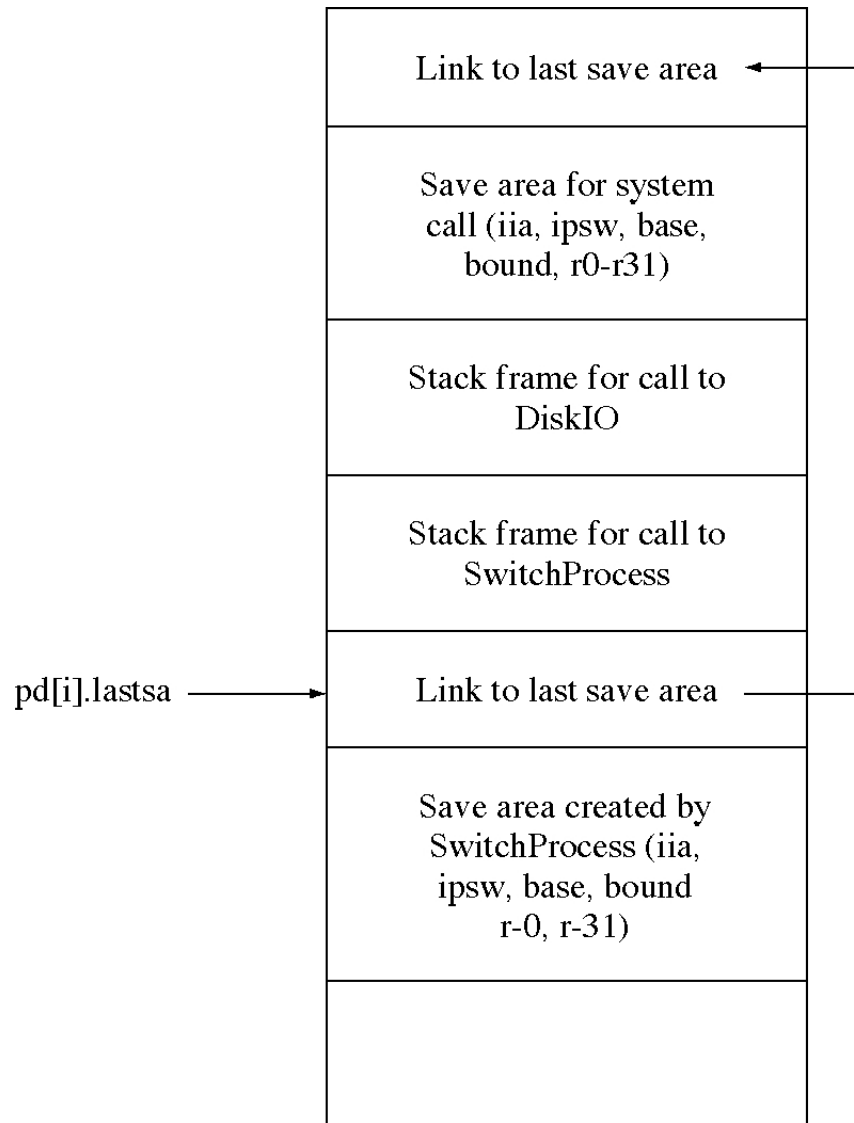
# System calls (5 of 5)

- case ReceiveMessageSystemCall:  
    int \* user\_msg; asm { store r9,user\_msg }  
    int from\_q; asm { store r10,from\_q }  
    if( !message\_queue\_allocated[from\_q] ) {  
        pd[current\_process].sa.reg[1] = -1;  
        break; }  
    if( message\_queue[from\_q].Empty() ) {  
        pd[current\_process].state = Blocked;  
        wait\_queue[from\_q].Insert( current\_process );  
        SwitchProcess(); }  
    int msg\_no = message\_queue[from\_q].Remove();  
    TransferMessage( msg\_no, user\_msg );  
    pd[current\_process].sa.reg[1] = 0;  
    break;  
case DiskReadSystemCall:  
case DiskWriteSystemCall:  
    // ... the same  
}  
Dispatcher();  
}}

# Kernel-mode: SwitchProcess

- ```
void SwitchProcess() {  
    // Called when a system mode process wants  
    // to wait for something  
    int * savearea; asm { store r30,savearea }  
    savearea -= sizeof(SaveArea)+4;  
    asm {  
        // save the registers.  
        // arrange to return from the procedure call  
        // when we return.  
        store    r31,savearea+4  
        store    psw,savearea+8  
        store    base,savearea+12  
        store    bound,savearea+16  
        storeall savearea+20  
    }  
    pd[current_process].state = Blocked;  
    Dispatcher();  
}
```

# Kernel-mode process system stack



# Kernel-mode: DiskIO

- ```
void DiskIO(int command, int disk_block,
    char * buffer ) {
    // Create a new disk request
    //   and fill in the fields.
    DiskRequest * req = new DiskRequest;
    req->command = command;
    req->disk_block = disk_block;
    req->buffer = buffer;
    req->pid = current_process;

    // Then insert it on the queue.
    disk_queue.Insert( (void *)req );
    pd[current_process].state = Blocked;

    // Wake up the disk scheduler if it is idle.
    ScheduleDisk();
    SwitchProcess(); // NEW CODE
}
```

# Kernel-mode: dispatcher

- ```
void RunProcess( int pid ) {  
    if( pid >= 0 ) {  
        asm { move    #0,psw } // Disable interrupts  
        int * savearea = pd[pid].lastsa;  
        pd[pid].lastsa = *savearea;  
        --(pd[pid].inSystem);  
        int quantum = pd[pid].timeLeft  
        asm {  
            load      savearea+4,iia  
            load      savearea+8,ipsw  
            load      savearea+12,base  
            load      savearea+16,bound  
            loadall   savearea+20  
            load      quantum,timer  
            rti  
        }  
    } else {  
        waitLoop: goto waitLoop;  
    }  
}
```

# Kernel-mode: schedule disk

- ```
void ScheduleDisk( void ) {  
    StartUsingProcessTable();  
    if( pd[scheduleDiskPid].state == Blocked )  
        pd[scheduleDiskPid].state = Ready;  
    FinishUsingProcessTable();  
}
```

# Kernel-mode: real schedule disk

```
• void RealScheduleDisk( void ) {  
    while( 1 ) { // NEW CODE  
        // If the disk is already busy, wait for it.  
        if( DiskBusy() )  
            SwitchProcess(); // NEW CODE  
        // Get the first disk request  
        // from the disk request queue.  
        DiskRequest * req = disk_queue.Remove();  
        // Wait, if there is no disk request to service.  
        if( req == 0 )  
            SwitchProcess(); // NEW CODE  
        // record which process is waiting for the disk  
        process_using_disk = req->pid;  
        // issue read or write, with interrupt enabled  
        if( req->command == DiskReadSystemCall )  
            IssueDiskRead(req->disk_block, req->buffer, 1);  
        else  
            IssueDiskWrite(req->disk_block, req->buffer, 1);  
    }  
}
```

# Kernel-mode process tradeoffs

- Waiting inside the kernel is easy
- Interrupts inside the kernel are easy
- But every process needs a kernel-mode stack
  - this adds up to a lot of memory allocated to stacks
- UNIX implementations have used kernel-mode processes from the beginning
  - but newer versions are getting away from it to save memory



# Implementation of mutual exclusion in the OS

- Three low-level solutions
  - disable interrupts
    - easy but only possible for single processor systems
  - special hardware instruction (exchangeword)
    - the preferred solution
  - software mutual exclusion
    - no special hardware required

# Using mutual exclusion

- ```
void Process1( void ) {  
    while( 1 ) {  
        DoSomeStuff();  
        EnterCriticalSection( 0 );  
        DoCriticalSectionStuff();  
        LeaveCriticalSection( 0 );  
    }  
}  
  
void Process2( void ) {  
    while( 1 ) {  
        DoSomeStuff();  
        EnterCriticalSection( 1 );  
        DoCriticalSectionStuff();  
        LeaveCriticalSection( 1 );  
    }  
}
```

# Implementing mutual exclusion

- `enum{ False = 0, True = 1 };`

```
// This is global data available to both processes
int interested[2] = {False, False};
int turn = 0;
```


```
void EnterCriticalSection( int this_process ) {
    int other_process = 1 - this_process;
    interested[this_process] = True;
    turn = this_process;
    while( turn == this_process
           && interested[other_process] ) {
        // do nothing, just wait for this loop to exit
    }
}
```

```
void LeaveCriticalSection( int this_process ) {
    interested[this_process] = False;
}
```


# Comparing the three solutions

- Disabling interrupts
  - is fast and does not involve busy waiting
  - is the best solution for a single processor
- Using ExchangeWord
  - requires hardware assistance and busy waiting
  - is the best solution for multiprocessors which share memory
- Peterson's solution
  - requires busy waiting but no special hardware
  - is the best solution for distributed systems with no central control


# Varieties of multiple processors



**(a) Multiprogramming — one CPU rapidly switched between processes**



**(b) Multiprocessing (shared memory) — multiple CPUs but one memory**



**(c) Multicomputers (no shared memory) — multiple CPU+memory units**

# Mutual exclusion (Dekker's)

- ```
// This is global data available to both processes
int interested[2] = {False, False};
int turn = 0;
void EnterCriticalSection( int this_process ) {
    int other_process = 1 - this_process;
    interested[this_process] = True;
    while( interested[other_process] ) {
        if( turn == other_process ) {
            interested[this_process] = False;
            while( turn == other_process )
                /* do nothing*/ ;
            interested[this_process] = True;
        }
    }
}
void LeaveCriticalSection( int this_process ) {
    int other_process = 1 - this_process;
    turn = other_process;
    interested[this_process] = False;
}
```