


# Processes

## Chapter 8

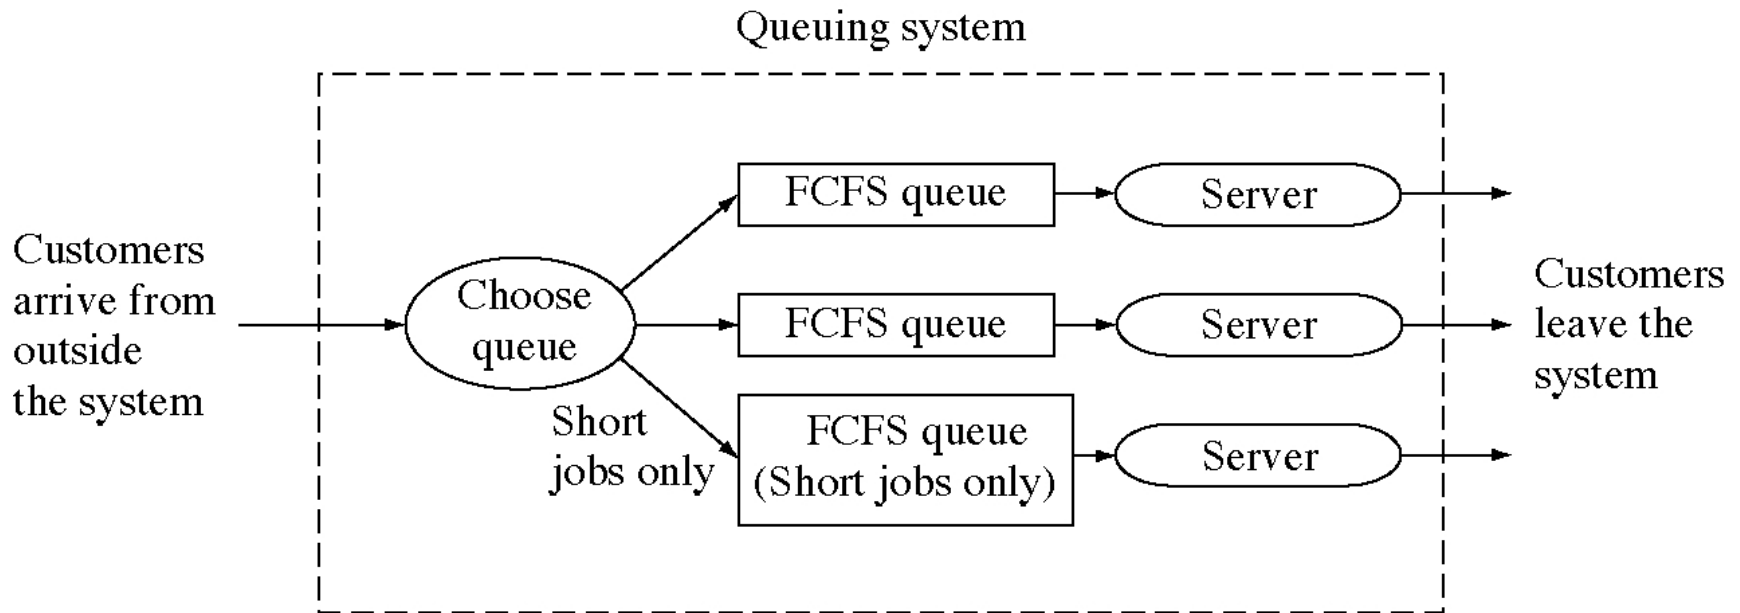
# Key concepts in chapter 8

- Non-preemptive scheduling: first-come-first-served (or first-in-first-out), shortest-job-first, priority
- Preemptive scheduling: round-robin, multiple round-robin queues
- Policy versus mechanism
- Deadlock and starvation
- Remote procedure calls
- Synchronization
- Semaphores
- Monitors

# FCFS at the supermarket




# SJF at the supermarket



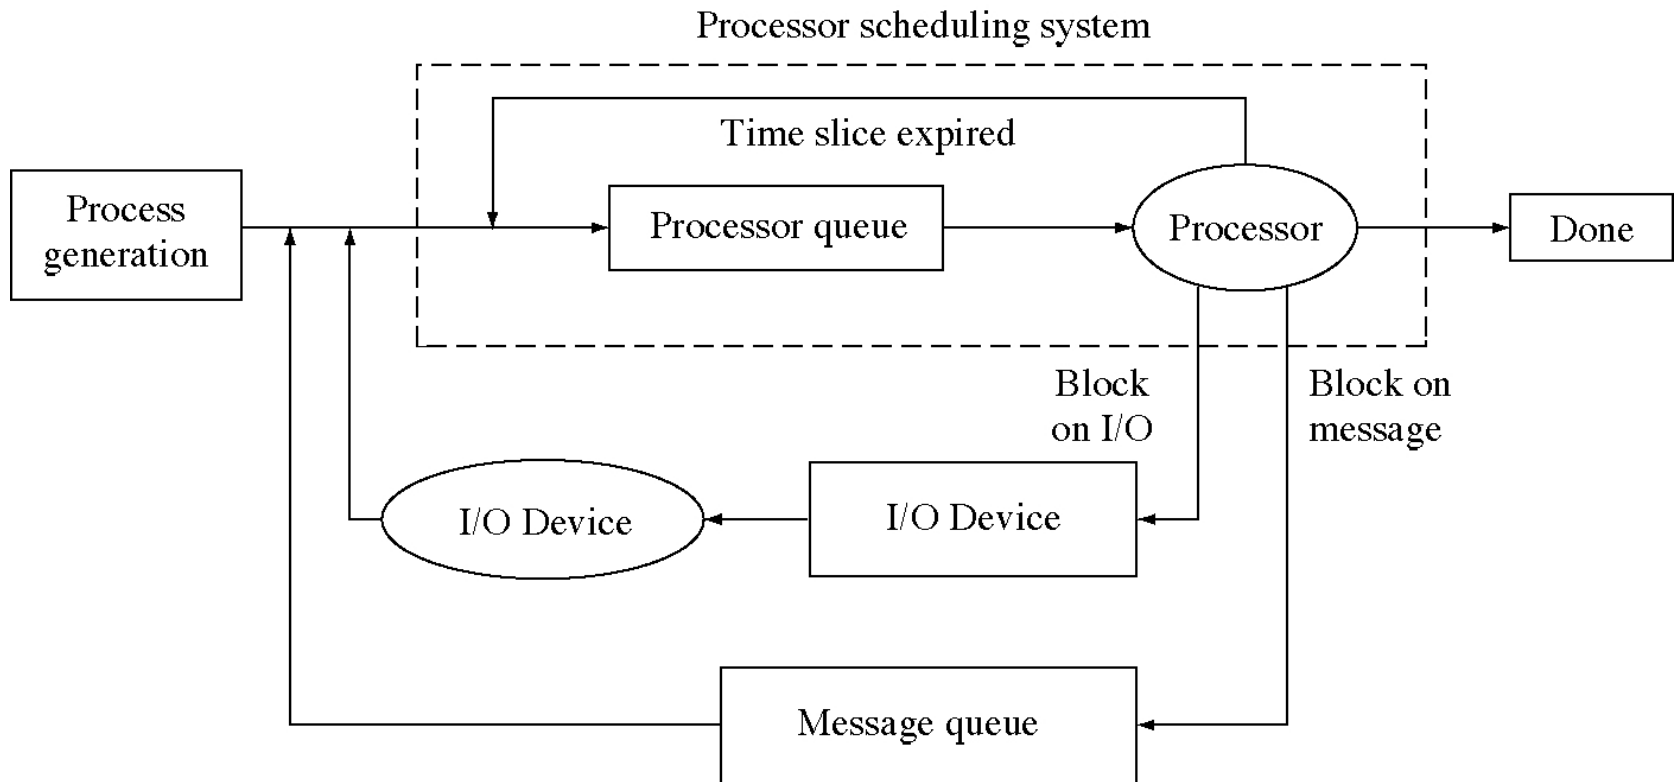
# Everyday scheduling methods

- First-come, first served
- Shorter jobs first
- Higher priority jobs first
- Job with the closest deadline first
- Round-robin

# Round-robin scheduling



# Process flow in an OS




# Round-robin scheduling

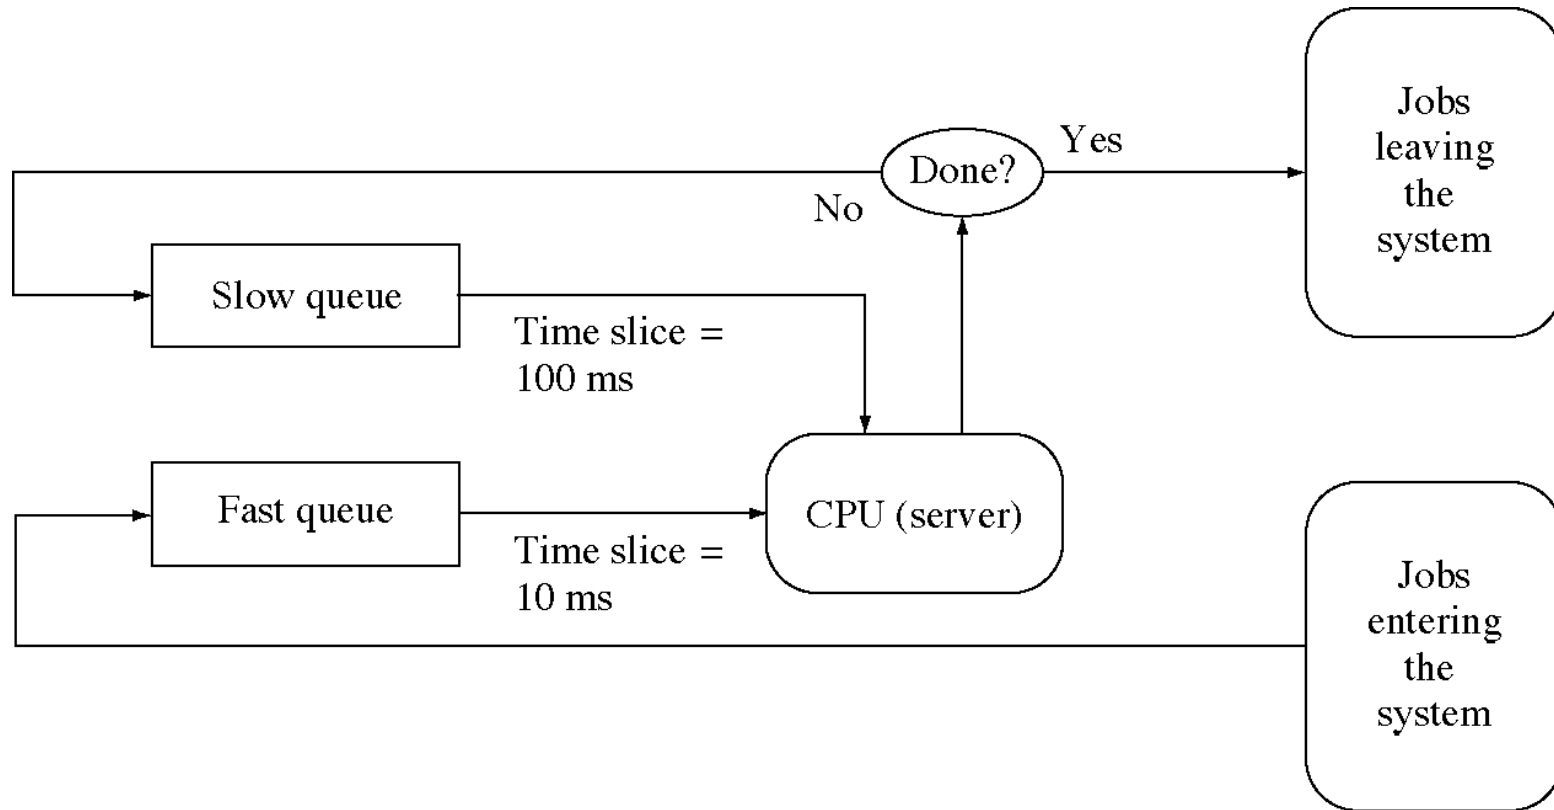
- Preemptive: processes lose the processor before they are finished using it
- Time-slice or quantum: time of each turn
  - 5 to 100 milliseconds
  - could be adaptive and change with the load
- Performs poorly if the load is heavy



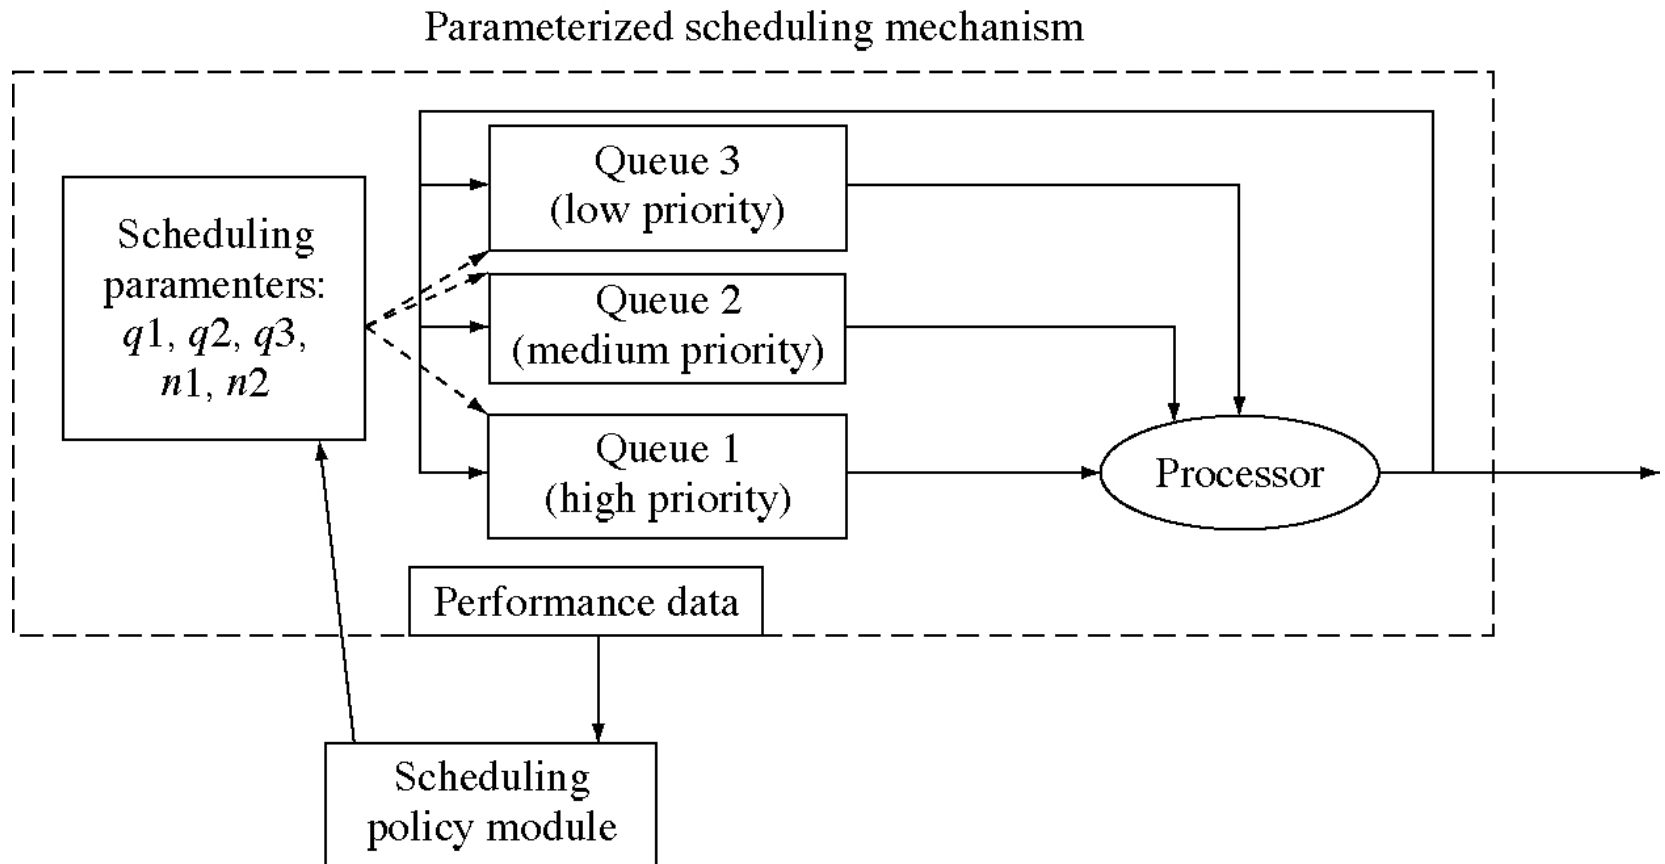
# Round-robin time slice



# Two-queue scheduling



# Three-queue scheduling




# Policy and mechanism in scheduling

- We can parameterize the three-queue scheduler
  - and make it into a range of scheduling algorithms
- Scheduling policy shifts through the day
  - so schedulers should have settable parameters
  - the scheduler is the mechanism


# A scheduling example

- Job 1, 10 seconds, priority 3
- Job 2, 2 seconds, priority 2
- Job 3, 5 seconds, priority 1
- Job 4, 3 seconds, priority 4

# A scheduling example




First-come, first-served:  $10 + 12 + 17 + 20 = 59/4 = 14.75$




Shortest-job-first:  $2 + 5 + 10 + 20 = 37/4 = 9.25$



Priority:  $5 + 7 + 17 + 20 = 49/4 = 12.25$



Round robin (time slice = 2):  $2 + 9 + 14 + 20 = 45/4 = 11.25$

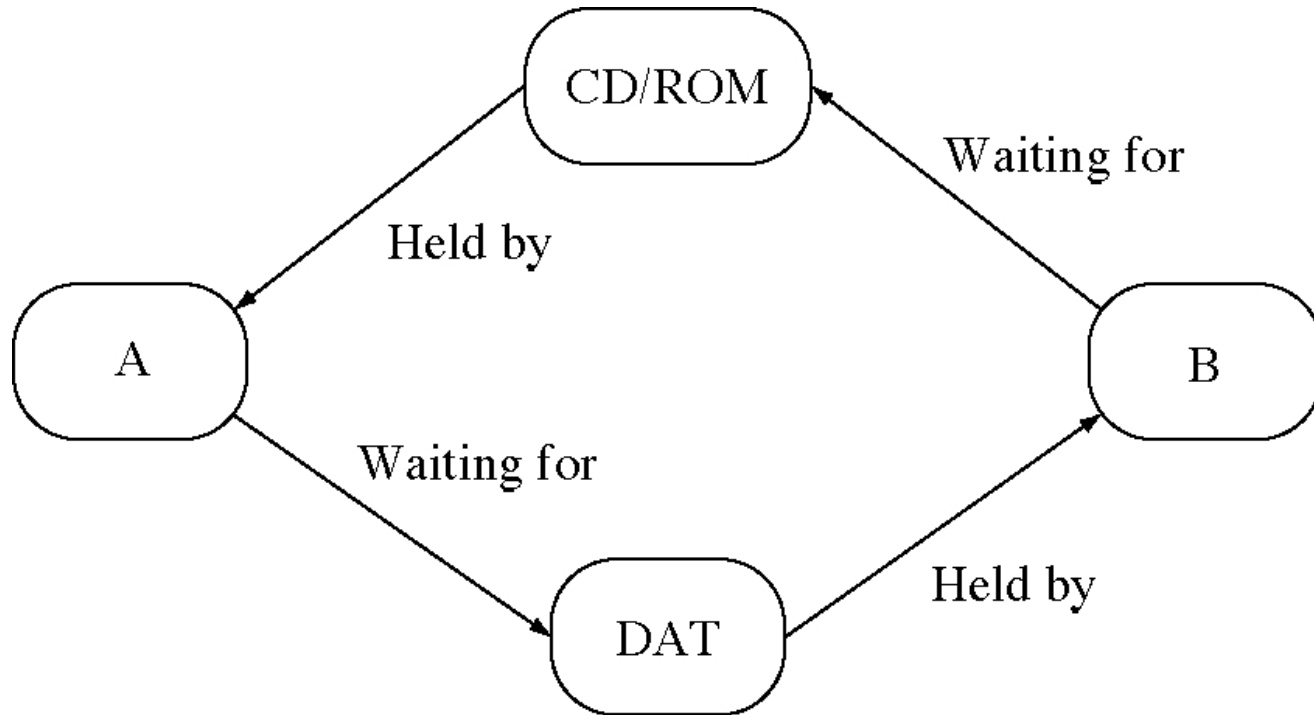


Round robin (continuous):  $8 + 11 + 15 + 20 = 54/4 = 13.50$

# Scheduling in real OSs

- All use a multiple-queue system
- UNIX SVR4: 160 levels, three classes
  - time-sharing, system, real-time
- Solaris: 170 levels, four classes
  - time-sharing, system, real-time, interrupt
- OS/2 2.0: 128 level, four classes
  - background, normal, server, real-time
- Windows NT 3.51: 32 levels, two classes
  - regular and real-time
- Mach uses “hand-off scheduling”

# Two-process deadlock





# Deadlock

- When two or more processes are waiting for each other to finish using resources
- This causes the system to come to a halt
- Deadlock is a problem if you don't take steps to prevent or detect it.
- Message deadlock:
  - ProcessA { receive(B, msg); send(B, msg): }
  - ProcessB { receive(A, msg); send(A, msg): }
  - Could result from a message being lost

# Conditions for deadlock

- Resources cannot be preempted.
- Resources cannot be shared.
- A process is holding one resource and requesting another.
- Circular waits are possible.

# Dealing with deadlock

- Three solutions
  - *Prevention*: place restrictions on resource request to make deadlock impossible
  - *Avoidance*: plan ahead to avoid deadlock.
  - *Recovery*: detect when deadlock occurs and recover from it

# Deadlock prevention

- Allow preemption
  - this is not easy
- Allow sharing
  - sometimes you can create multiple, virtual copies of the resource
- Avoid hold-and-wait
  - require all resources to be acquired in one request
- Avoid circular wait
  - require resource to be acquired in a fixed order

# Deadlock avoidance

- There are algorithms for this
  - e.g. the Banker's algorithm
- But they have problems
  - they are slow
  - they do not account for resource failure or loss

# Deadlock recovery

- The usual solution
- There are algorithms to detect deadlock
  - but they are slow
  - timeouts are another solution

# Two-phase locking

- Each process has two stages
  - Stage 1
    - acquire all resources you will need
    - if a resource is locked, then release all resources you are holding and start over
  - Stage 2
    - Use the resources
- Deadlock is prevented (no hold-and-wait)

# Starvation


- When a process is unlucky and cannot get a resource it needs
  - for example, in a random queue
  - calling in to a number that is usually busy
- Most common when you are requesting two resources together



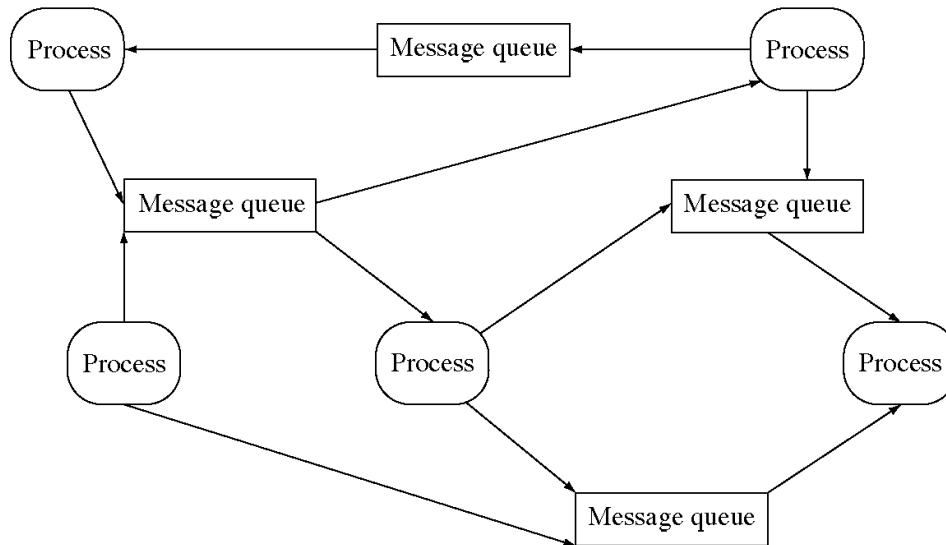
# Variations in message-passing

- Using PIDs as message addresses
  - less flexible than indirection through queues
- Non-blocking receives
  - if there is no message, receive returns an error
  - more flexible but require busy waiting for a message
- Blocking sends
  - send, receive, reply model
  - less flexible but avoids OS message buffering

# Addressing messages to processes



(a) Messages directly to processes




(b) Message-to-message queues


# Design technique: Indirection

- Indirection: going through an intermediate object to get to an object
- Examples
  - sending a message to a queue then another processes receives it from the queue
  - access a private variable in an object using set and get functions
- Advantage: you get control just before the object is accessed


# Monitoring two message queues simultaneously



(a) Use extra processes to create a common queue



(b) Poll both queues

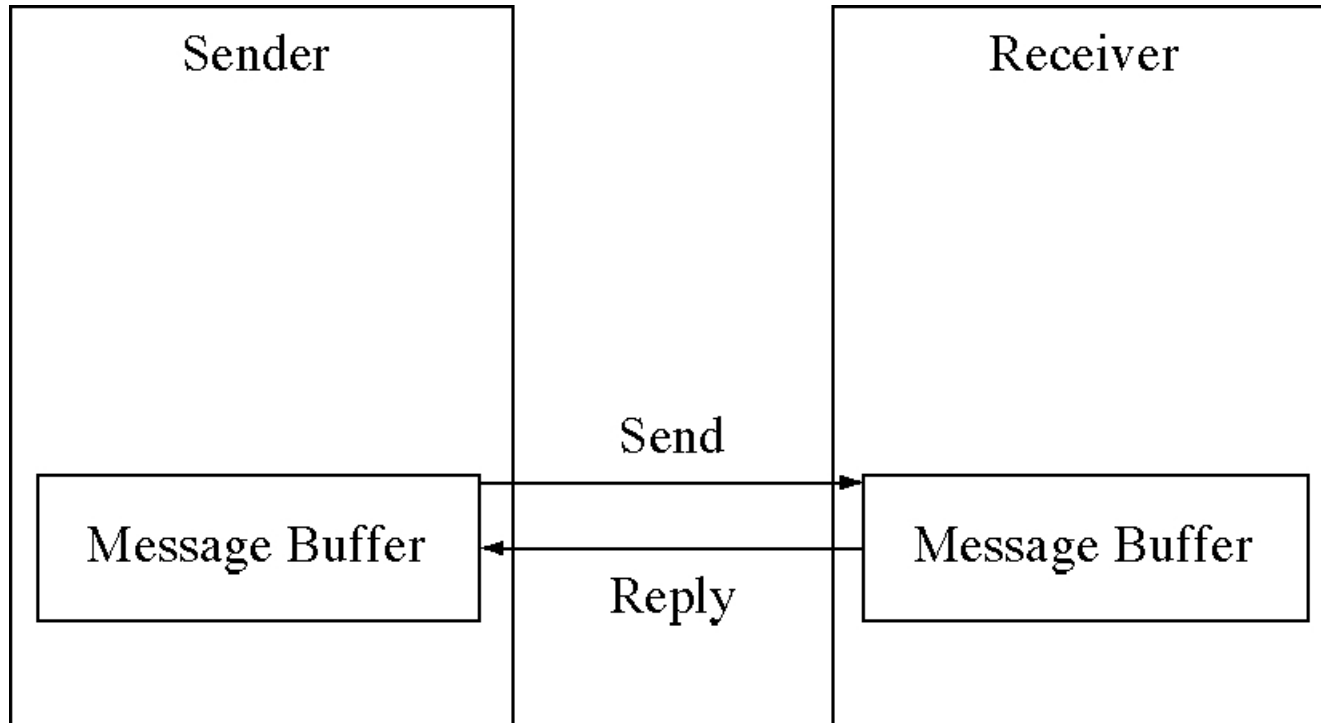


(c) Read from two queues simultaneously.

# Design technique: Adding a new facility

- Three ways to do it
  - Example: receiving from two messages queues
  - 1. use an existing facility
    - create threads to listen on each queue
  - 2. create a new low-level primitive and use it
    - add a non-blocking receive and poll with it
  - 3. create a new high-level primitive to do the job
    - add a system call to receive from two queues

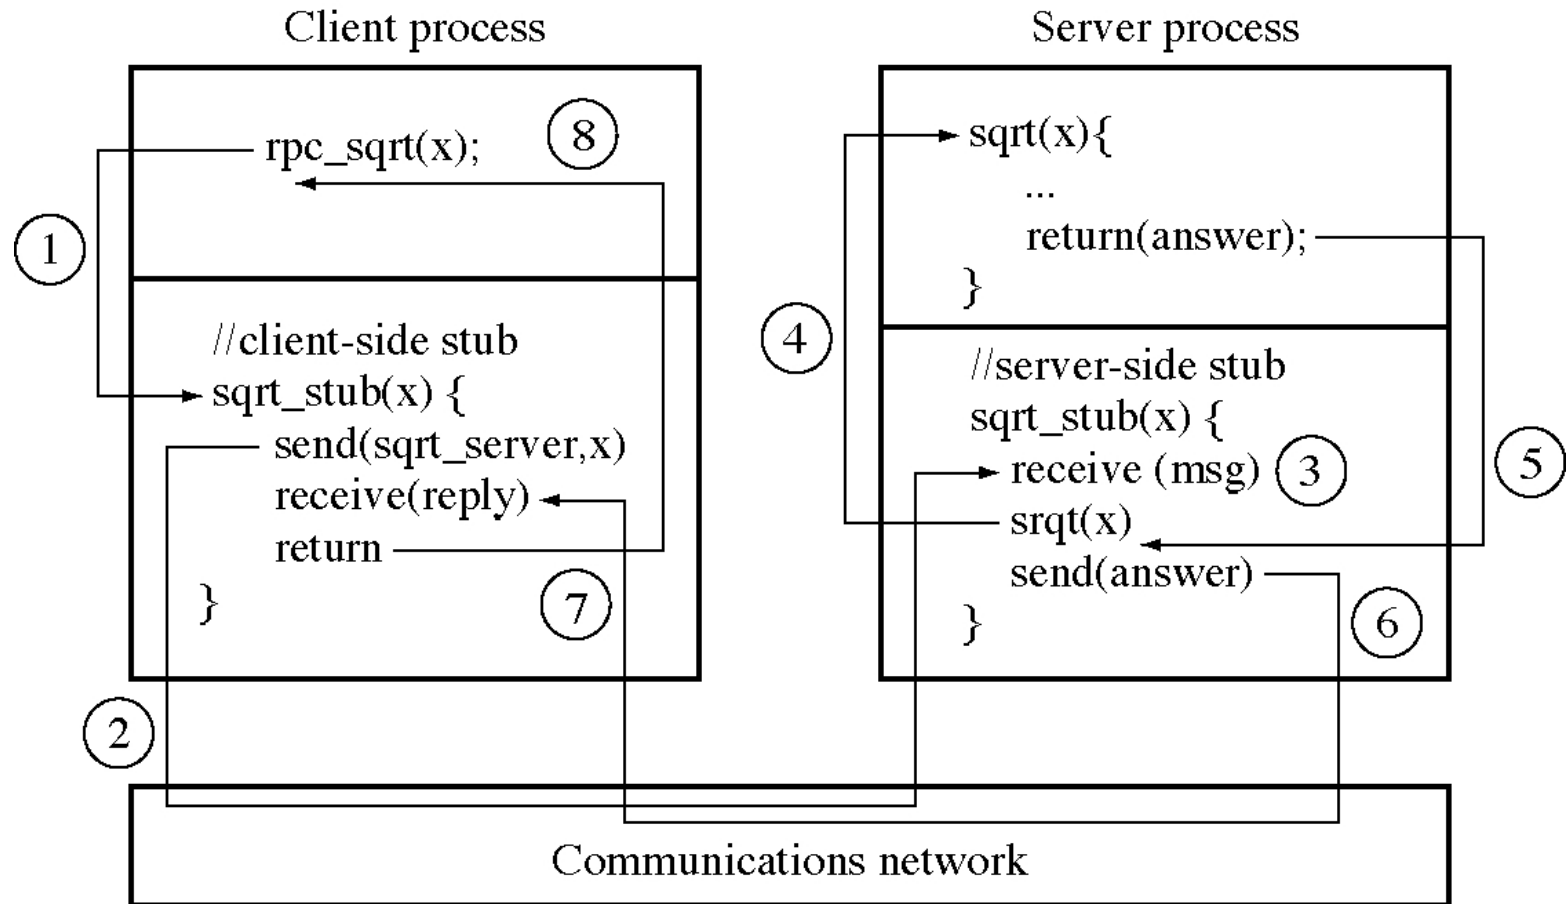
# Send-receive-reply model



# Remote procedure calls (RPCs)

- Remote procedure call (or RPC)
  - another IPC technique
  - RPC uses the familiar procedure call model
- Calling a procedure in another address space looks the same as calling a local procedure
- Implemented with stubs

# Remote procedure call flow of control





# Implementing RPCs:

## Client-side stub

- ```
retValue = RemoteService( arg1, arg2, arg3 );  
// client-side "stub" procedure  
int RemoteService( int arg1, int arg2, int arg3 ) {  
    MessageBuffer msg;  
    msg[0] = RemoteServiceCode;  
    msg[1] = arg1;  
    msg[2] = arg2;  
    msg[3] = arg3;  
    SendMessage( RemoteServer, msg );  
    ReceiveMessage( RemoteServer, msg );  
    return msg[1];  
}
```

# RPC implementation issues

- An RPC is not really procedure call
  - arguments must be by value (no pointers)
  - server must be found
  - server can fail
  - messages can be lost
- RPC libraries are provided to allow reuse of most of the RPC mechanism

# Synchronization

- *Synchronization* is processes (or threads) waiting for each other, basically signaling
- In the OS we used various methods: disabling interrupts, spin locks, kernel-mode processes
- In user processes we have seen only messages, which combine data transfer with synchronization, but they can be separated

# Semaphores

- Basic synchronization problems:
  - mutual exclusion
  - signaling
- Semaphores are a shared-memory synchronization primitive that solves these two problems

# Semaphore operations

- `int AttachSemaphore( int global_id)`
  - a global, numerical ID (`global_id`)
  - it returns a local identifier for the semaphore
- `int Wait(int local_id)`
  - `local_id` from the `AttachSemaphore`
- `int Signal(int local_id)`
- `int DetachSemaphore(int local_id)`
  - `local_id` from the `AttachSemaphore`

# Binary semaphore implementation

```
• void Wait( int sema_id ) {
    Semaphore * sema = GetSemaphoreFromId( sema_id );
    Lock( sema->lock );
    if( !sema->busy ) {
        sema->busy = True;
    } else {
        sema->queue->Insert( current_process );
        Block( current_process );
    }
    Unlock( sema->lock );
}

void Signal( int sema_id ) {
    Semaphore * sema = GetSemaphoreFromId( sema_id );
    Lock( sema->lock );
    if( sema->queue->Empty() ) {
        sema->busy = False;
    } else {
        process = sema->queue->Remove();
        Unblock( process );
    }
    Unlock( sema->lock );
}
```

# Two-process mutual exclusion

- ```
// This is the code for process A or B
void main( int argc, char * argv[ ] ) {
    int mutex_sema = AttachSemaphore( FMutexID );
    if( ThisIsProcessA() )
        Signal( mutex_sema );
    // Initialize the semaphore to not busy
    while( 1 ) {
        DoOtherThings();
        // The critical section
        Wait( mutex_sema );
        UseFileF();
        Signal( mutex_sema );
    }
    DetachSemaphore( mutex_sema );
}
```

# Two-process rendezvous

- ```
void main(int argc, char *argv[]){ // Game player A
    int a_sema = AttachSemaphore( GPAID );
    int b_sema = AttachSemaphore( GPBID );
    // Tell B that A is ready to go.
    Signal( b_sema );
    // Wait until B is ready to go.
    Wait( a_sema );
    DetachSemaphore( a_sema );
    DetachSemaphore( b_sema );
}

void main(int argc, char *argv[]){ // Game player B
    int b_sema = AttachSemaphore( GPBID );
    int a_sema = AttachSemaphore( GPAID );
    // Tell A that B is ready to go.
    Signal( a_sema );
    // Wait until A is ready to go.
    Wait( b_sema );
    DetachSemaphore( a_sema );
    DetachSemaphore( b_sema );
}
```



# Producer-consumer

- Buffer buffer;  
void main() { // The Producer  
    int empty\_buffer = AttachSemaphore(EmptyBufferID);  
    int full\_buffer = AttachSemaphore(FullBufferID);  
    while( 1 ) {  
        Wait( empty\_buffer );  
        FillBuffer( buffer );  
        Signal( full\_buffer );  
    }  
    DetachSemaphore(empty\_buffer);  
    DetachSemaphore(full\_buffer);  
}  
void main() { // The Consumer  
    int empty\_buffer = AttachSemaphore(EmptyBufferID);  
    int full\_buffer = AttachSemaphore(FullBufferID);  
    Signal( empty\_buffer );  
    while( 1 ) {  
        Wait( full\_buffer );  
        ConsumeBuffer( buffer );  
        Signal( empty\_buffer );  
    }  
    DetachSemaphore( empty\_buffer );  
    DetachSemaphore( full\_buffer );  
}

# Counting semaphore implementation

```
• void Wait( int sema_id ) {  
    Semaphore * sema = GetSemaphoreFromId( sema_id );  
    Lock( sema->lock );  
    if( sema->count > 0 ) --sema->count;  
    else {  
        sema->queue->Insert( current_process );  
        Block( current_process );  
    }  
    Unlock( sema->lock );  
}  
void Signal( int sema_id ) {  
    Semaphore * sema = GetSemaphoreFromId( sema_id );  
    Lock( sema->lock );  
    if( sema->queue->Empty() ) ++sema->busy;  
    else {  
        process = sema->queue->Remove();  
        Unblock( process );  
    }  
    Unlock( sema->lock );  
}
```

# N-buffer producer

- ```
Const int MaxBuffers = 20;
Queue<Buffer> buffer_queue;
void main() { // The Producer
    int use_buffer_queue
        = AttachSemaphore(UseBufferQueueID);
    int empty_buffer = AttachSemaphore(EmptyBufferID);
    int full_buffer = AttachSemaphore(FullBufferID);
    Buffer buffer;
    while( 1 ) {
        FillBuffer( buffer );
        Wait( empty_buffer );
        Wait( use_buffer_queue );
        buffer_queue.Insert( buffer );
        Signal( use_buffer_queue );
        Signal( full_buffer );
    }
    DetachSemaphore( use_buffer_queue );
    DetachSemaphore( empty_buffer );
    DetachSemaphore( full_buffer );
}
```

# N-buffer consumer

- ```
void main() { // The Consumer
    int use_buffer_queue
        = AttachSemaphore( UseBufferQueueID );
    int empty_buffer = AttachSemaphore(EmptyBufferID);
    int full_buffer = AttachSemaphore(FullBufferID);
    Buffer buffer;
    Signal( use_buffer_queue ); // Buffer queue free
    for( int i = 0; i < MaxBuffers; ++i )
        Signal( empty_buffer );
    while( 1 ) {
        Wait( full_buffer );
        Wait( use_buffer_queue );
        buffer = buffer_queue.Remove();
        Signal( use_buffer_queue );
        Signal( empty_buffer );
        ConsumeBuffer( buffer );
    }
    DetachSemaphore( use_buffer_queue );
    DetachSemaphore( empty_buffer );
    DetachSemaphore( full_buffer );
}
```

# Semaphores and messages

- Semaphores
  - are basically messages with no content
  - they handle synchronization only
  - data are transferred in shared memory
- Messages
  - are more appropriate when there is no shared memory

# System constants and globals

- ```
// system limits (we can change these)
const int NumberOfSemaphores = 50;

// system call numbers
const int AttachSemaphoreSystemCall = 8;
const int DetachSemaphoreSystemCall = 9;
const int SignalSemaphoreSystemCall = 10;
const int WaitSemaphoreSystemCall = 11;

// semaphore data structures
struct Semaphore {
    int allocated;
    int count;
    int use_count;
    int id;
    Queue<Pid> queue;
};

Semaphore sema[NumberOfSemaphores];
```

# System initialization

- ```
int main( void ) {  
    // ... all what we had before, plus  
    // initialize all semaphores to "not allocated"  
    for( i = 0; i < NumberOfSemaphores; ++i ) {  
        sema.allocated[i] = False;  
    }  
}
```

# System call handler (1 of 3)

- case AttachSemaphoreSystemCall:  
    int sema\_id; asm { store r9,sema\_id }  
    pd[current\_process].sa.reg[1]  
        = AttachSemaphore( sema\_id );  
    break;  
case DetachSemaphoreSystemCall:  
    int sid; asm { store r9,sid }  
    pd[current\_process].sa.reg[1]  
        = DetachSemaphore( sid );  
    break;  
case SignalSemaphoreSystemCall:  
    int sid; asm { store r9,sid }  
    pd[current\_process].sa.reg[1]  
        = SignalSemaphore( sid );  
    break;  
case WaitSemaphoreSystemCall:  
    int sid; asm { store r9,sid }  
    pd[current\_process].sa.reg[1]  
        = WaitSemaphore( sid );  
    break;



# System call handler (2 of 3)

- case SendMessageSystemCall:  
    // get the arguments  
    int \* user\_msg; asm { store r9,user\_msg }  
    int to\_q; asm { store r10,to\_q }  
    // check for an invalid queue identifier  
    if( !message\_queue\_allocated[to\_q] ) {  
        pd[current\_process].sa.reg[1] = -1;  
        break;  
    }  
    int msg\_no = GetMessageBuffer();  
    // make sure we have not run out of message  
    buffers  
    if( msg\_no == EndOfFreeList ) {  
        pd[current\_process].sa.reg[1] = -2;  
        break;  
    }  
    // copy the message vector from the system  
    caller's  
    // memory into the system's message buffer  
    CopyToSystemSpace( current\_process, user\_msg,  
    message\_buffer[msg\_no], MessageSize )

# System call handler (3 of 3)

- ```
// put it on the queue
message_queue[to_q].Insert( msg_no );
// notify any waiters that it is there
SignalSemaphore( message_semaphore[to_q] ); // NEW
pd[current_process].sa.reg[1] = 0;
break;
case ReceiveMessageSystemCall:
    int * user_msg; asm { store r9,user_msg }
    int from_q; asm { store r10,from_q }
    // check for an invalid queue identifier
    if( !message_queue_allocated[from_q] ) {
        pd[current_process].sa.reg[1] = -1;
        break;
    }
    WaitSemaphore( message_semaphore[from_q] ); // NEW
    int msg_no = message_queue[from_q].Remove();
    TransferMessage( msg_no, user_msg );
    pd[current_process].sa.reg[1] = 0;
    break;
```

# Attach semaphore

```
• int AttachSemaphore( int sema_id ) {  
    int i, free_slot = -1;  
    for( i = 0; i < NumberOfSemaphores; ++i )  
        if( sema[i].allocated ) {  
            if( sema[i].id == sema_id ) break;  
            else free_slot = i  
        }  
    if( i >= NumberOfSemaphores ) { // found sema_id?  
        if( free_slot == -1 ) // found free slot?  
            return -1; // No, so return an error code.  
        i = free_slot;  
        sema[i].allocated = True;  
        sema[i].count = 0;  
        sema[i].use_count = 0;  
        sema[i].id = sema_id;  
        sema[i].queue = new Queue<Pid>;  
    }  
    ++sema[i].use_count;  
    return i;  
}
```

# Detach semaphore

- ```
int DetachSemaphore( int sid ) {  
    if( !sema[sid].allocated ) {  
        pd[current_process].sa.reg[1] = -1;  
        break;  
    }  
    if( --sema[sid].use_count == 0 ) {  
        sema[sid].allocated = False;  
        delete sema[sid].queue;  
    }  
    return 0;  
}
```

# Signal semaphore

- ```
int SignalSemaphore( int sid ) {  
    if( !sema[sid].allocated ) {  
        pd[current_process].sa.reg[1] = -1;  
        break;  
    }  
    if( sema[sid].queue->Empty() )  
        ++sema[sid].count;  
    else {  
        int pid = sema[sid].queue->Remove();  
        pd[pid].state = Ready;  
    }  
    return 0;  
}
```

# Wait semaphore

- ```
int WaitSemaphore( int sid ) {  
    if( !sema[sid].allocated ) {  
        pd[current_process].sa.reg[1] = -1;  
        break;  
    }  
    if( sema[sid].count > 0 )  
        --sema[sid].count;  
    else {  
        sema[sid].queue->Insert(current_process);  
        pd[current_process].state = Blocked;  
    }  
    return 0;  
}
```

# System initialization

- processTableSemaphore  
    = AttachSemaphore( ProcessTableID );  
disk\_queue = AttachSemaphore( DiskQueueID);  
disk\_free = AttachSemaphore( DiskFreeID );

# Dispatcher

- ```
int SelectProcessToRun( void ) {  
    static int next_proc = NumberOfProcesses;  
    int i, return_value = -1;  
    // Get exclusive access to the process table.  
    WaitSemaphore(processTableSemaphore);    // NEW  
    // ... use process table as before  
    SignalSemaphore(processTableSemaphore); // NEW  
    return return_value;  
}
```



# Disk I/O

- ```
void DiskIO( int command, int disk_block,
             char * buffer ) {
    // Create a new disk request
    // and fill in the fields.
    DiskRequest * req = new DiskRequest;
    req->command = command;
    req->disk_block = disk_block;
    req->buffer = buffer;
    req->semaphore = pd[current_process].semaphore;
    // Then insert it on the queue.
    disk_queue.Insert( req );
    // Wake up the disk scheduler if it is idle.
    SignalSemaphore( disk_queue );
    WaitSemaphore( pd[current_process].semaphore );
}
```

# Disk scheduling

- ```
void RealScheduleDisk( void ) {  
    while( 1 ) { // NEW CODE  
        WaitSemaphore( disk_free ); // NEW CODE  
        WaitSemaphore( disk_queue ); // NEW CODE  
        // Get the first disk request  
        // from the disk request queue.  
        DiskRequest * req = disk_queue.RemoveFirst();  
        // remember which process is waiting  
        // for the disk operation  
        disk_completion_semaphore = req->semaphore;  
        // issue the read or write,  
        // with disk interrupt enabled  
        if( req->command == DiskReadSystemCall )  
            IssueDiskRead(req->disk_block, req->buffer, 1);  
        else  
            IssueDiskWrite(req->disk_block, req->buffer, 1);  
    }  
}
```

# Disk interrupt handler

- ```
void DiskInterruptHandler( void ) {  
    if( current_process > 0 ) {  
        // was there a running process?  
        // Save the processor state of the system  
        caller.  
        // ... as before  
    }  
    // Notify the waiting process that  
    // the disk transfer is complete  
    SignalSemaphore( disk_completion_semaphore );  
    // Notify any waiters that the disk is free  
    SignalSemaphore( disk_free );  
    // now run a process  
    Dispatcher();  
}
```

# Monitors

- Another synchronization primitive
  - more structured, monitors are modules
- Components of a monitor module
  - variables: of any kind
  - condition variables: for monitor signaling
  - procedures: can be called from outside the monitor, they comprise the monitor's interface

# Counter monitor

- ```
monitor Counter {  
    private:  
        int count = 0;  
    public:  
        void Increment( void ) { count = count + 1; }  
        int GetCount( void ) { return count; }  
}  
int main() { // one process  
    while( 1 ) {  
        // ... Do other things  
        Counter.Increment();  
        // ... continue on  
        int n = Counter.GetCount();  
    }  
}  
int main() { // another process  
    while( 1 ) {  
        // ... Do other things  
        Counter.Increment();  
    }  
}
```

# Signal monitor

- ```
monitor Signal {  
    private:  
        int IsSignaled = 0;  
        condition SendSignal;  
    public:  
        void SendSignal( void ) {  
            IsSignaled = 1;  
            signal( SendSignal );  
        }  
        void WaitForSignal( void ) {  
            if( !IsSignaled )  
                wait( SendSignal );  
        }  
    }  
}
```

# Using the signal monitor

- ```
int main() { // the Signal Sender
    // ... Do things up to the signal point
    Signal.SendSignal();
    // ... continue on
}
int main() { // the Signal Receiver
    // ... Do things up to the signal point
    Signal.WaitForSignal();
    // ... continue on when the signal is receive
}
```

# Bounded buffer monitor (1 of 2)

- ```
monitor BoundedBufferType {  
    private:  
        BufferItem * buffer;  
        int NumberOfBuffers;  
        int next_in, nextout;  
        int current_size;  
        condition NotEmpty, NotFull;  
    public:  
        BoundedBufferType( int size ) {  
            buffers = new BufferItem[size];  
            NumberOfBuffers = size;  
            next_in = 0; next_out = 0; current_size = 0;  
        }  
    }
```



# Bounded buffer monitor (2 of 2)

- ```
void Put( BufferItem item ) {
    if( current_size == NumberOfBuffers )
        wait( NotFull );
    buffer[next_in] = item;
    next_in = (next_in+1) % NumberOfBuffers;
    if( ++current_size == 1 )
        signal( NotEmpty );
}
BufferItem Get( void ) {
    if( current_size == 0 )
        wait( NotEmpty );
    BufferItem item = buffer[next_out];
    next_out = (next_out+1) % NumberOfBuffers;
    if( --current_size == NumberOfBuffers-1 )
        signal( NotFull );
    return item;
}
}
```

# Using a bounded buffer monitor

- `BoundedBufferType BoundedBuffer;`

```
int main() { // the Producer
    while( 1 ) {
        BufferItem item = ProduceItem();
        BoundedBuffer.Put( item );
    }
}
```

```
int main() { // the Consumer
    while( 1 ) {
        BufferItem item = BoundedBuffer.Get();
        ConsumeItem( item );
    }
}
```

# Counting semaphore monitor

- monitor Semaphore {  
  private:  
    int count = 0;  
    condition NotBusy;  
  public:  
    void Signal( void ) {  
      if( ++count > 0 )  
        signal( NotBusy );  
    }  
  
    void Wait( void ) {  
      while( count <= 0 )  
        wait( NotBusy );  
      --count;  
    }  
  }  
}

# Using a semaphore monitor

- ```
int main() {           // one process
    while( 1 ) {
        // do other stuff
        // enter critical section
        Semaphore.Wait();
        // do critical section
        Semaphore.Signal();
    }
}
```

```
int main() {           // another process
    while( 1 ) {
        // do other stuff
        // enter critical section
        Semaphore.Wait();
        // do critical section
        Semaphore.Signal();
    }
}
```

# Protected counter in Ada95

- ```
-- declare the interface to the task
task Counter is
    entry GetCount( count : out integer );
    entry Increment;
private
    count : integer;
end Counter;
-- the implementation of the protected variable
task body Counter is
    loop
        select
            accept GetCount( count_out : out integer ) do
                count_out := count;
            end;
        or
            accept Increment do
                count := count + 1;
            end;
        end select;
    end loop;
end Counter;
```

# Using a protected counter

- task body OneProcess is begin  
    loop  
        -- do other things than incrementing the counter  
        Counter.Increment;  
        -- do other things  
        Counter.GetCount( n );  
    end loop;  
end OneProcess;  
  
task body AnotherProcess is begin  
    loop  
        -- do other things than incrementing the counter  
        Counter.Increment;  
        -- do other things  
    end loop;  
end AnotherProcess;

# Signal in Ada95

- ```
task Signal is
    entry SendSignal;
    entry WaitForSignal;
private
    IsSignaled : boolean := False;
end Counter;

task body Signal is
    accept SendSignal do
        IsSignaled := True;
    end;
    accept WaitForSignal do
        null;
    end;
end Signal;
```

# Using signal in Ada95

- ```
task body SignalSender is begin
    -- get to point where event occurs
    Signal.SendSignal;
    -- go on to other things
end SignalSender;

task body SignalReceiver is begin
    -- get to the point where you need
    -- to wait for the event
    Signal.WaitForSignal;
    -- respond to event
end SignalReceiver;
```



# Bounded buffer in Ada95 (1 of 2)

- ```
task BoundedBuffer is
    entry Put( x : in Item );
    entry Get( x : out Item );
end BoundedBuffer;
task body BoundedBuffer is
    NumberOfBuffers : constant integer := 20;
    buffers : array(1 .. NumberOfBuffers) of Item;
    current_size :
        integer range 0 .. NumberOfBuffers := 0;
    next_in, next_out
        : integer range 1 .. NumberOfBuffers := 1;
```

# Bounded buffer in Ada95 (2 of 2)

```
• begin
    loop
        select
            when current_size < NumberOfBuffers =>
                accept Put( x : in item ) do
                    buffers( next_in ) := x;
                end;
                next_in := (next_in mod NumberOfBuffers) + 1;
                current_size := current_size + 1;
            or when current_size > 0 =>
                accept Get( x : out Item do
                    x := buffers(next_out);
                end;
                next_out := (next_out mod NumberOfBuffers)+1;
                current_size := current_size - 1;
            or
                terminate;
        end select;
    end loop;
end BoundedBuffer;
```

# Producer-consumer in Ada95

- ```
task body Producer is begin
    loop
        item := ProduceItem;
        BoundedBuffer.Put( item );
    end loop;
end Producer;

task body Consumer is begin
    loop
        BoundedBuffer.Get( item );
        ConsumeItem( item );
    end loop;
end Consumer;
```

# Semaphore in Ada95

- CountingSemaphore( StartCount : Integer := 1 ) is  
    entry Wait;  
    entry Signal;  
    entry Count( count\_out : out integer );  
private  
    CurrentCount : Integer := StartCount;  
end CountingSemaphore;  
task body CountingSemaphore is begin  
    loop  
        select  
            when CurrentCount > 0 =>  
                accept Wait do  
                    CurrentCount := CurrentCount - 1;  
                end;  
        or  
            accept Signal do  
                CurrentCount := CurrentCount + 1;  
            end;  
        or  
            accept Count( count\_out : out integer ) do  
                count\_out := CurrentCount;  
            end;  
        end select;  
    end loop;  
end CountingSemaphore;

# Protected counter in Ada95 using protecting variables

- -- the interface to the protected variable  
protected Counter is  
    function GetCount return integer;  
    procedure Increment;  
private  
    count : integer;  
end Counter;  
-- the implementation of the protected variable  
protected body Counter is  
    function GetCount return integer is begin  
        return count;  
    end GetCount;  
    procedure Increment is begin  
        count := count + 1;  
    end Increment;  
end Counter;

# Using a protected counter

- task body OneProcess is begin  
    loop  
        -- do other things than incrementing the counter  
        Counter.Increment;  
        -- do other things  
        n := Counter.GetCount;  
    end loop;  
end OneProcess;

```
task body AnotherProcess is begin
    loop
        -- do other things than incrementing the counter
        Counter.Increment;
        -- do other things
    end loop;
end AnotherProcess;
```

# Signaling in Ada95

## using protected variables

- ```
protected Signal is
  procedure SendSignal;
  entry WaitForSignal;
private
  IsSignaled : boolean := False;
end Signal;

protected body Signal is
  procedure SendSignal is begin
    IsSignaled := True;
  end SendSignal;
  entry WaitForSignal when IsSignaled is begin
    null;
  end WaitForSignal;
end Signal;
```

# Using signals in Ada95

- ```
task body SignalSender is begin
    -- get to point where event occurs
    Signal.SendSignal;
    -- go on to other things
end SignalSender;

task body SignalReceiver is begin
    -- get to the point where you need to wait
    -- for the event
    Signal.WaitForSignal;
    -- respond to event
end SignalReceiver;
```



# Bounded buffer in Ada95

- protected type BoundedBuffer is  
    entry Put( x : in Item );  
    entry Get( x : out Item );  
private  
    buffers : ItemArray(1..NumberOfBuffers);  
    next\_in, next\_out  
        : integer range 1..NumberOfBuffers := 1;  
    current\_size  
        : integer range 0..NumberOfBuffers := 0;  
end BoundedBuffer;  
protected body BoundedBuffer is  
    entry Put( x : in Item )  
        when current\_size < NumberOfBuffers is begin  
        buffers(next\_in) := x;  
        next\_in := (next\_in mod NumberOfBuffers) + 1;  
        current\_size := current\_size + 1;  
    end Put;  
    entry Get( x : out Item )  
        when current\_size > 0 is begin  
        x := buffers(next\_out);  
        next\_out := (next\_out mod NumberOfBuffers) + 1;  
        current\_size := current\_size - 1;  
    end Get;  
end BoundedBuffer;

# Producer-consumer in Ada95

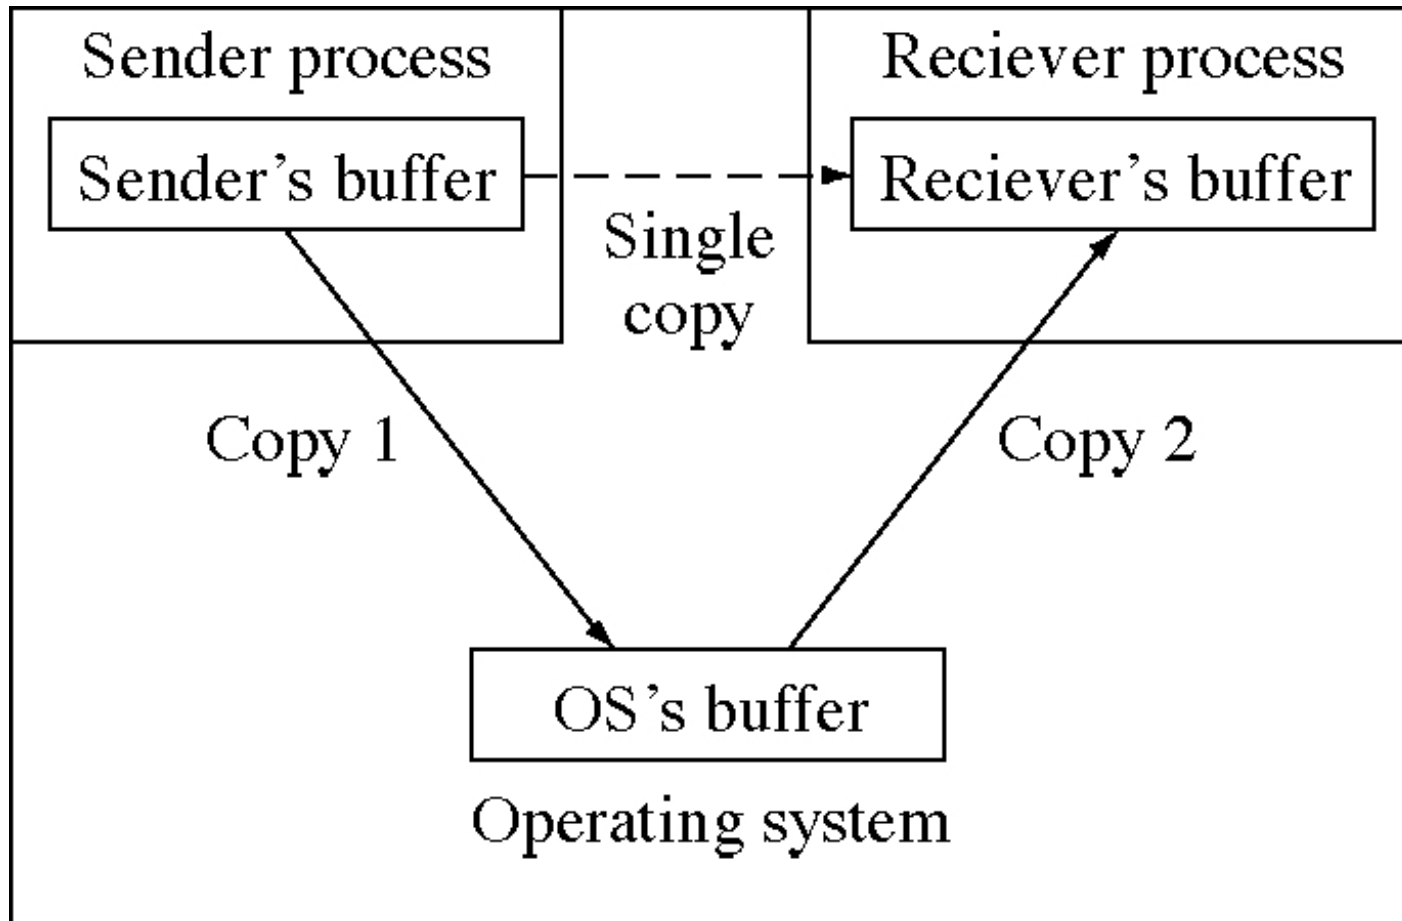
- ```
task body Producer is begin
    loop
        item := ProduceItem;
        BoundedBuffer.Put( item );
    end loop;
end Producer;

task body Consumer is begin
    loop
        BoundedBuffer.Get( item );
        ConsumeItem( item );
    end loop;
end Consumer;
```

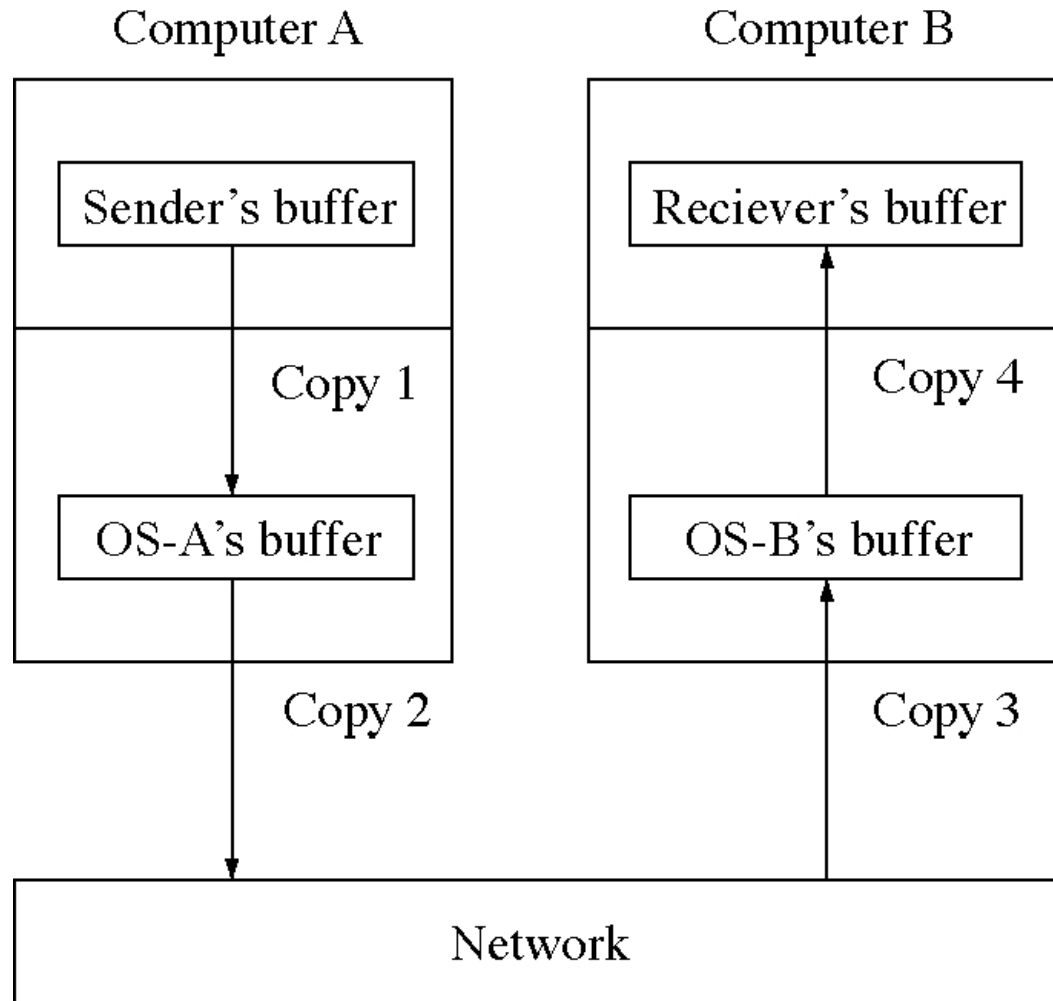
# Counting semaphore in Ada95 using protected variables

- protected type  
    CountingSemaphore( StartCount : Integer := 1 ) is  
        entry Wait;  
        procedure Signal;  
        function Count return Integer;  
private  
    CurrentCount : Integer := StartCount;  
end CountingSemaphore;  
protected body CountingSemaphore is  
    entry Wait when CurrentCount > 0 is begin  
        CurrentCount := CurrentCount - 1;  
    end Wait;  
    procedure Signal is begin  
        CurrentCount := CurrentCount + 1;  
    end Signal;  
    function Count return Integer is begin  
        return CurrentCount;  
    end Count;  
end CountingSemaphore;

# Copying messages in an OS



# Copying messages in a network



# Longer messages

- Eight words is too small
- Paging (chapter 11) will allow us to pass long messages with no copying

# IPC in Mach

- Task: Mach terminology for process
- Thread: execute in a task
- Port: a message queue
  - only one receiver
  - protected with access rights and capabilities
- Message: contains three types of information
  - data: any number of bytes, copied
  - out-of-line-data: part of the sender's address space, not copied
  - ports

# Signals

- Used in UNIX system for IPC
- An event notification
  - 30 or so fixed signal (event) types
- Basically a software interrupt
  - process defines the interrupt handler