


Virtual Memory

Chapter 11


Key concepts in chapter 11

- Fragmentation
- Virtual memory
- Paging
- File mapping

Compacting memory



Before compaction




After compaction

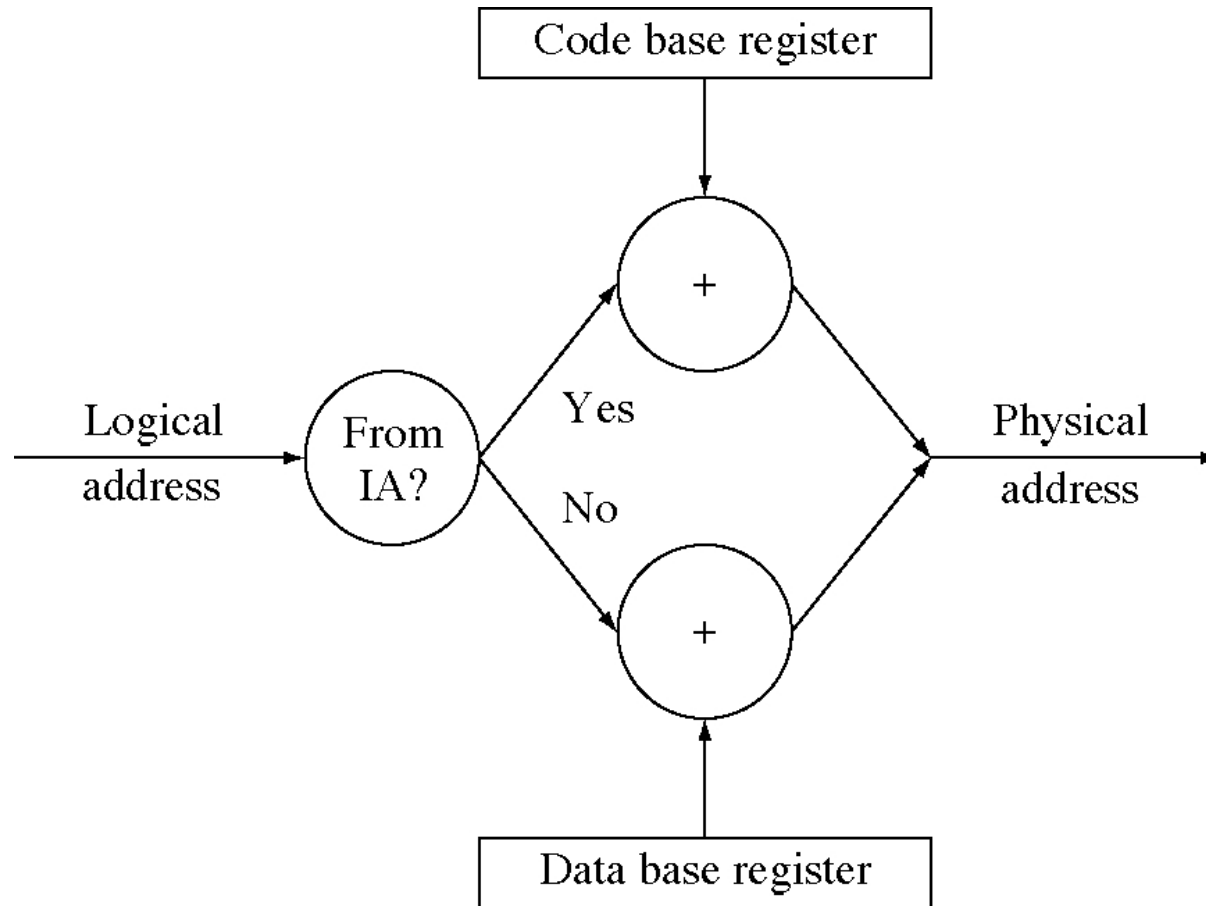
Fragmentation

- Without memory mapping, programs require physically continuous memory
- Large blocks mean large fragments
 - and wasted memory
- We need hardware memory mapping to address this problem
 - segments
 - pages
- We will look at a series of potential solutions

Separate code and data spaces




Code/data memory relocation



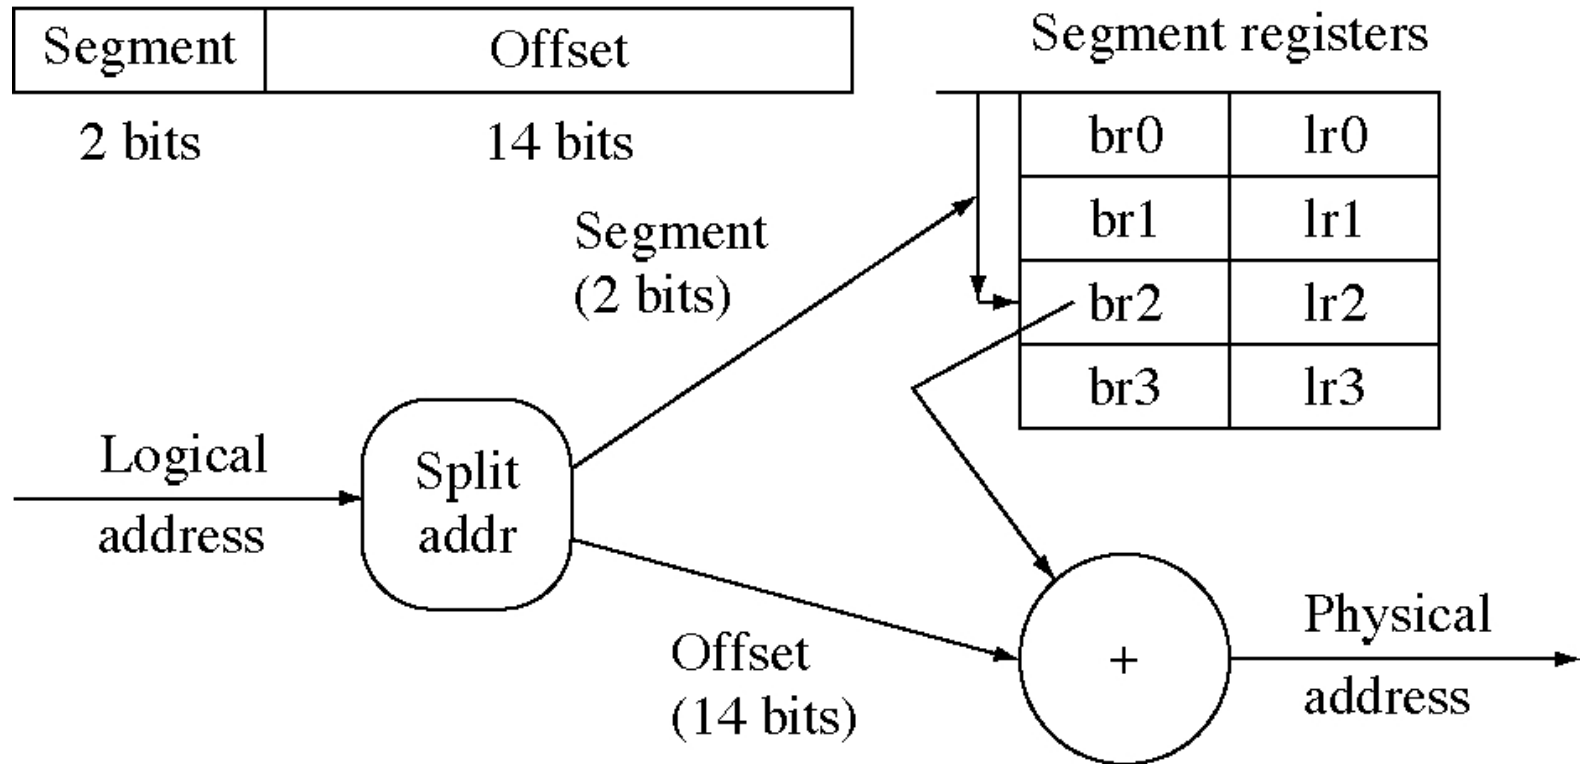
Segmentation

- Divide the logical address space into segments (variable-sized chunks of memory)
- Each segment has a base and bound register
 - and so segments do not need to be contiguous in the physical address space
 - but the logical address space is still contiguous
- DEC PDP11
 - eight segments
 - up to 8K bytes per segment


Two segmented address spaces



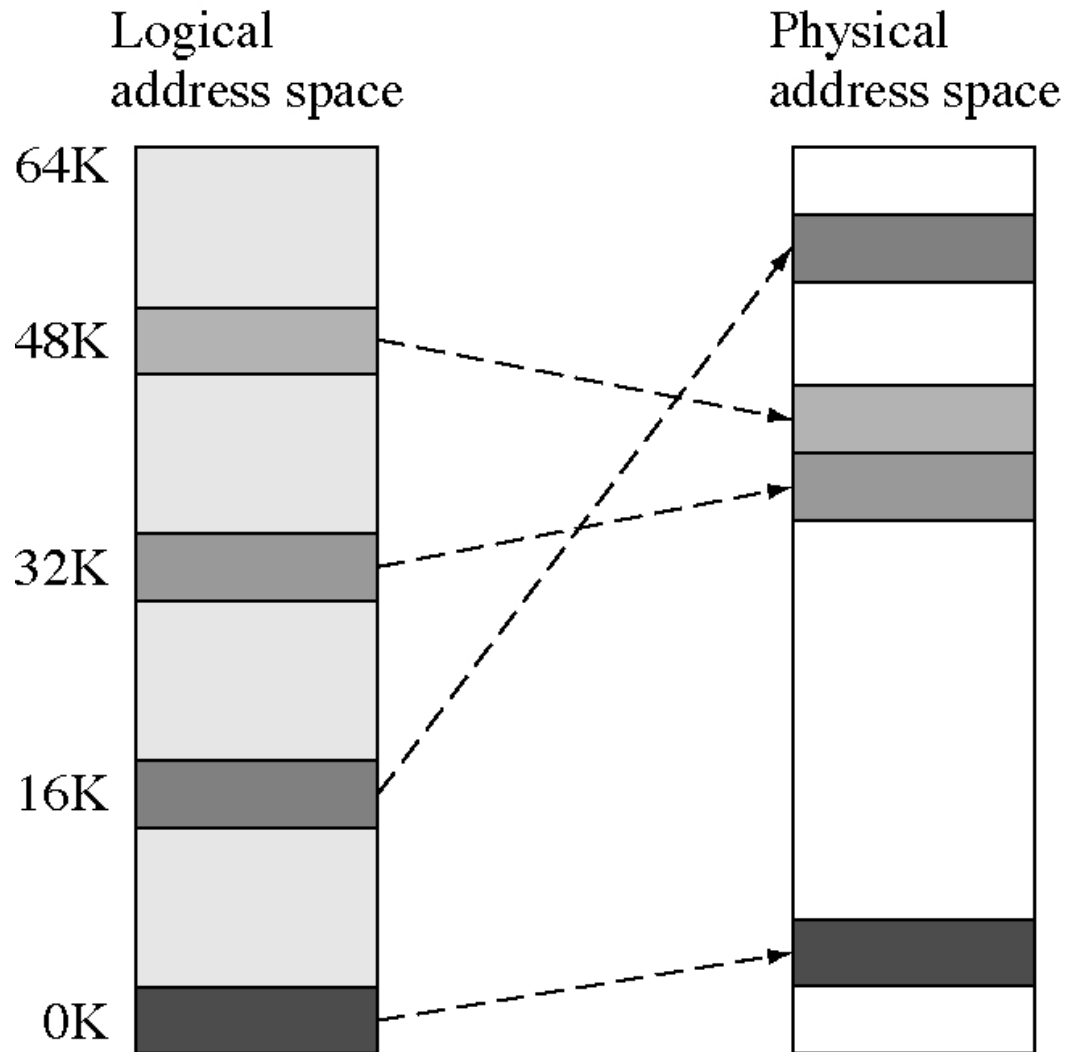
Segmentation memory mapping



Contiguous 24K logical address space



Noncontiguous 24K logical address space



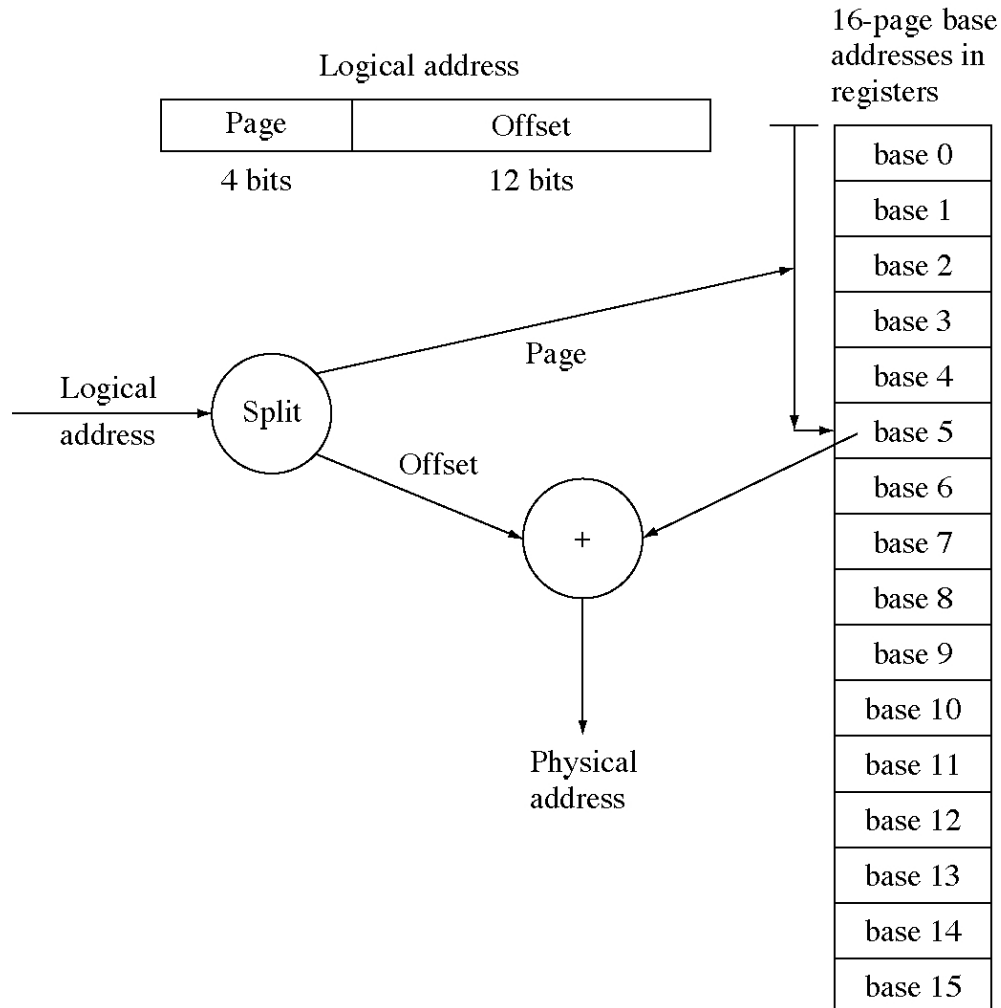
Noncontiguous logical address spaces

- This is possible with segmentation hardware
 - but is not usually a good idea
- Example segmentation map
- Segment-Base - Limit - Logical address
 - 0 100K 6K 0K-6K
 - 1 194K 6K 16K-22K
 - 2 132K 6K 32K-38K
 - 3 240K 6K 48K-54K

Segments to pages

- Large segments do not help the fragmentation problem
 - so we need small segments
- Small segments are usually full
 - so we don't need a length register
 - just make them all the same length
- Identical length segments are called *pages*
- We use page tables instead of segment tables
 - base register but no limit register

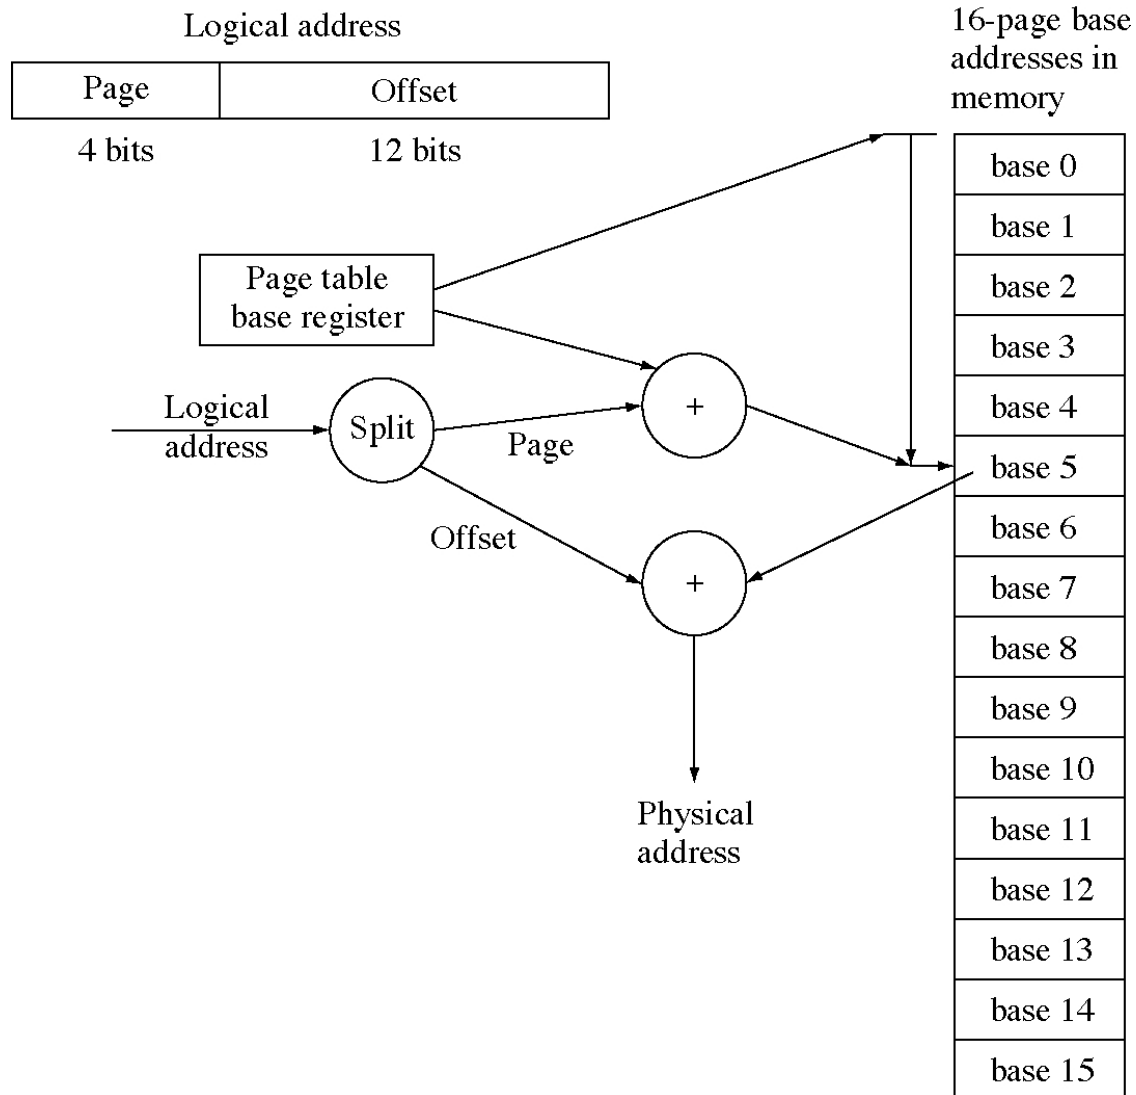
Hardware register page table




Problems with page tables in registers

- Practical limit on the number of pages
- Time to save and load page registers on context switches
- Cost of hardware registers
- *Solution*: put the page table in memory and have a single register that points to it

Page tables in memory



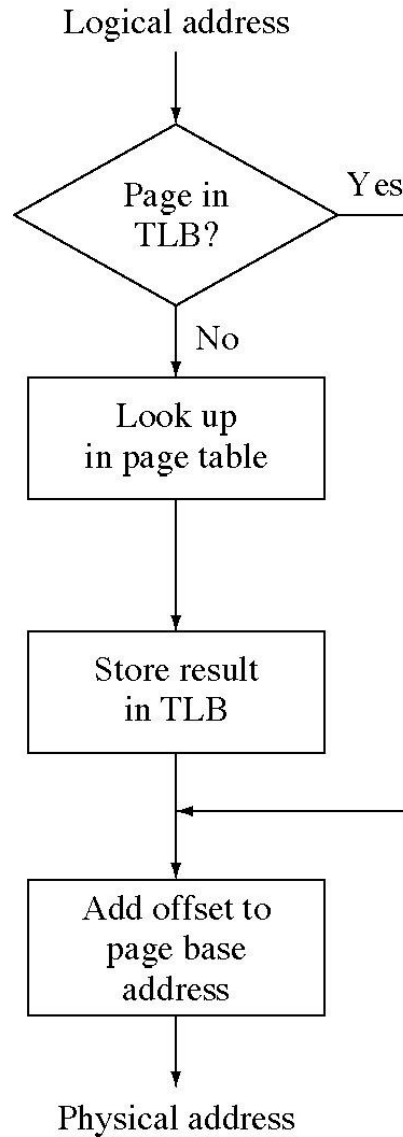
Page table mapping



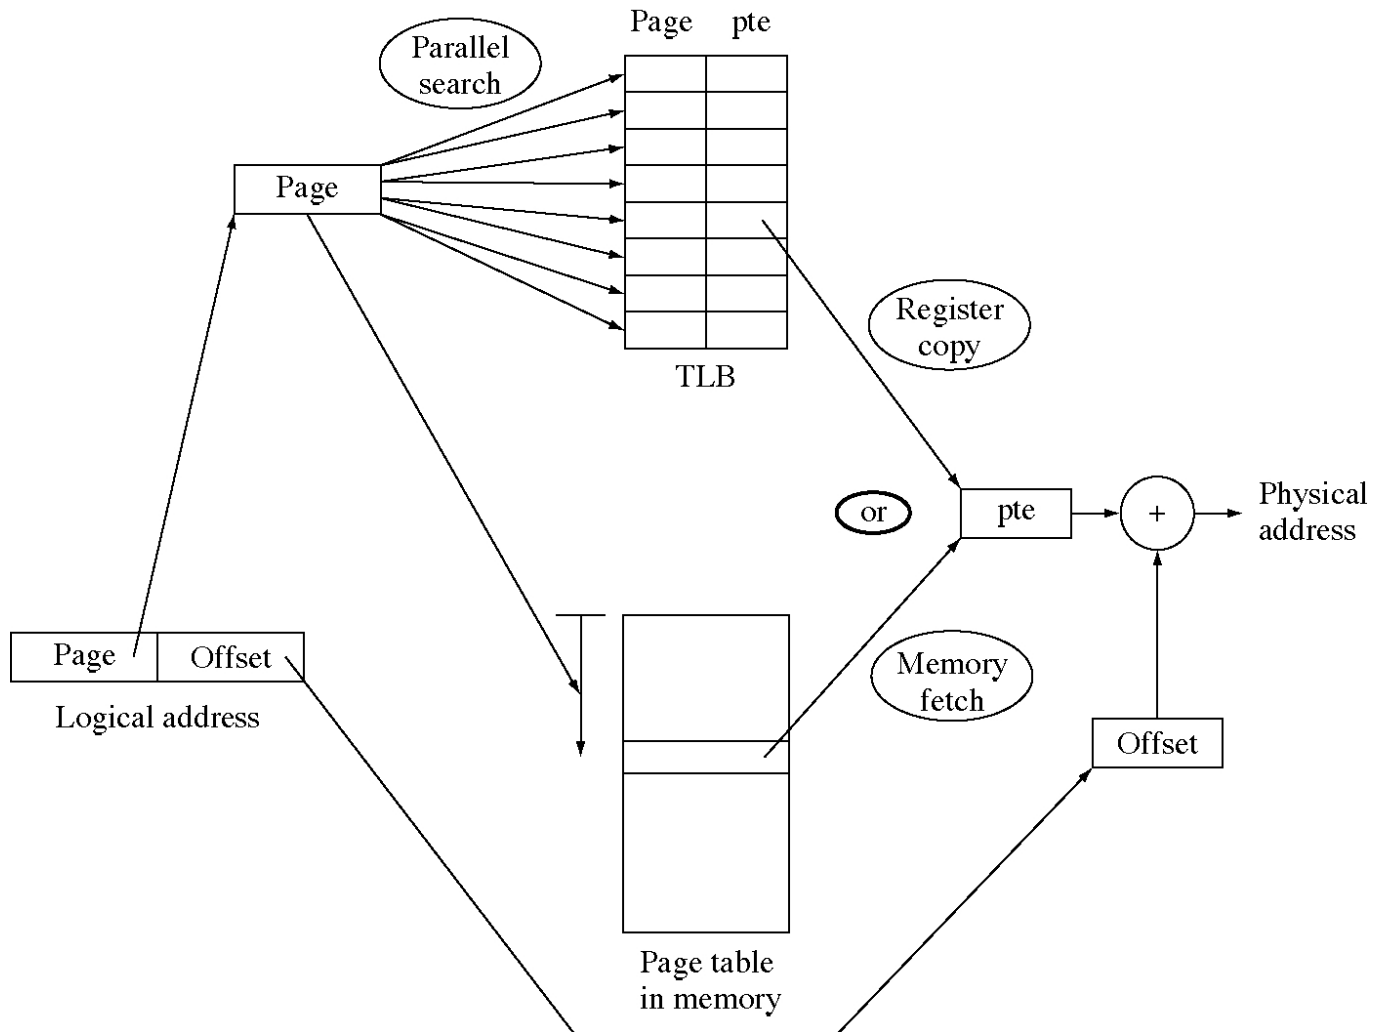
Problems with page tables in memory

- Every data memory access requires a corresponding page table memory access
 - the memory usage has doubled
 - and program speed is cut in half
- *Solution*: caching page table entries
 - called a *translation lookaside buffer*
 - or TLB

TLB flow chart



TLB lookup



Caching the page table (1 of 2)

- ```
const int PageTableCacheSize = 8;
const int pageSizeShift = 12;
const int pageSizeMask = 0xFFF;
struct CacheEntry {
 int logicalPageAddress;
 int pageBaseAddress;
} PageTableCache[PageTableCacheSize];
extern int pageTableBaseRegister;
int LeastRecentlyUsedCacheSlot(void);
```


# Caching the page table (2 of 2)

- ```
int LogicalToPhysical( int logicalAddress ) {
    int logicalPageAddress = logicalAddress & ~pageSizeMask;
    for( int i = 0; i < PageTableCacheSize; ++i ) {
        // the hardware lookup is done in parallel
        if( PageTableCache[i].logicalPageAddress
            == logicalPageAddress )
            return PageTableCache[i].pageBaseAddress;
    }
    int pteAddress = pageTableBaseRegister
                    + (logicalAddress >> pageSizeShift);
    int pageBaseAddress = MemoryFetch( pteAddress );
    // now update the cache by replacing the entry that has
    // not been used in the longest time (the least recently
    // used one) with this new entry
    i = LeastRecentlyUsedCacheSlot();
    PageTableCache[i].logicalPageAddress=logicalPageAddress;
    PageTableCache[i].pageBaseAddress = pageBaseAddress;
    return pageBaseAddress;
}
```

Why TLBs work

- Memory access is not random, that is, not all locations in the address space are equally likely to be referenced
- References are localized because
 - sequential code execution
 - loops in code
 - groups of data accessed together
 - data is accessed many times
- This property is called *locality*
- TLB hit rates are 90+% .

Good and bad cases for paging




```
for(i=0;i<8;++i)  
  sum+=a[i][0];
```

(worst case)

○ = Memory reference


□ = Page



```
for(j=0;j<8;++j)  
  sum+=a[0][j];
```

(best case)

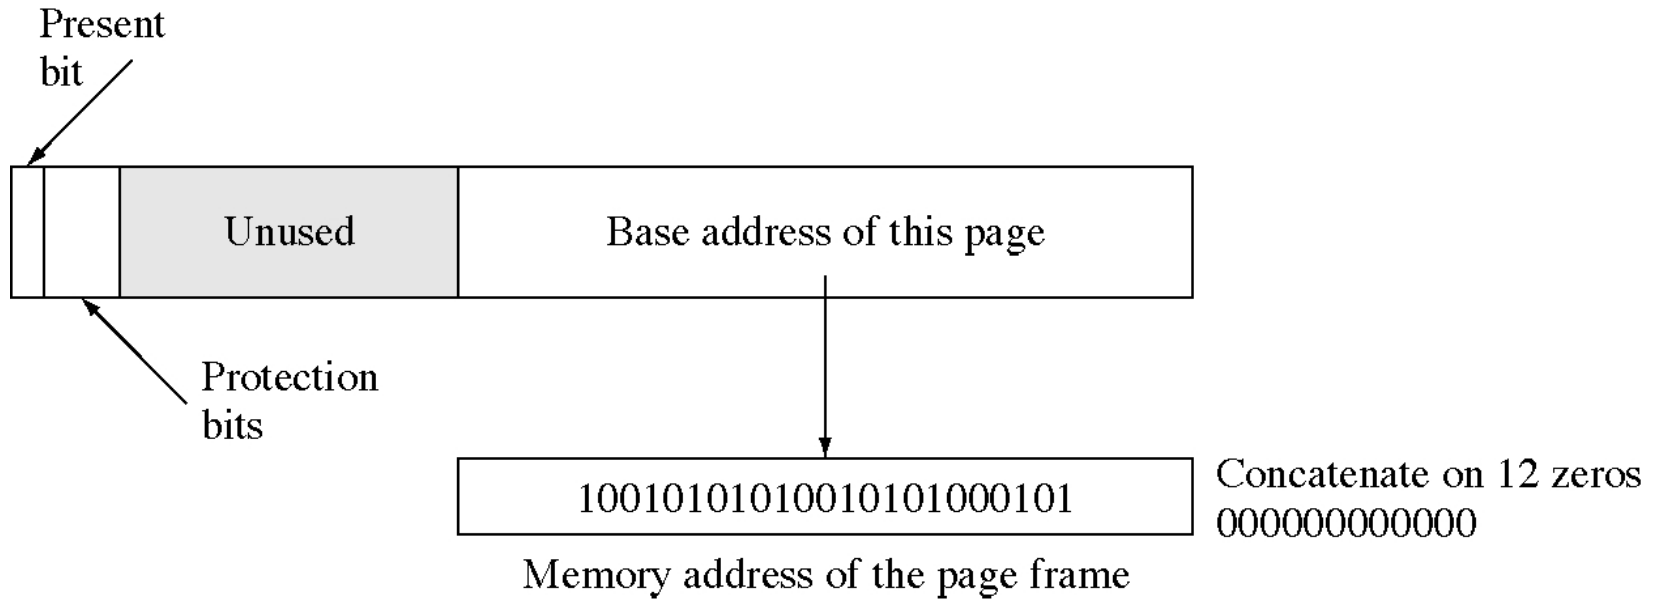
Internal and external fragmentation



Page and page frame

- Page
 - the information in the page frame
 - can be stored in memory (in a page frame)
 - can be stored on disk
 - multiple copies are possible
- Page frame
 - the physical memory that holds a page
 - a resource to be allocated

Page table entry



Page table protection

- Three bits control: read, write, execute
- Possible protection modes:
 - 000: page cannot be accessed at all
 - 001: page is read only
 - 010: page is write only
 - 100: page is execute only
 - 011: page can be read or written
 - 101: page can be read as data or executed
 - 110: write or execute, unlikely to be used
 - 111: any access is allowed

Paged memory allocator (1 of 4)

- ```
const int PageSize = 4096;
struct MemoryRequest {
 int npages;
 // size of the request in pages
 Semaphore satisfied; // signal when memory allocated
 int * pageTableArray; // store page numbers here
 MemoryRequest *next, *prev; // doubly linked list
};
// The memory request list
// keep a front and back pointer for queue discipline
MemoryRequest * RequestListFront, *RequestListBack;
// The structure for the free page list
struct FreePage {
 int pageNumber;
 FreePage *next;
};
// The free page list
FreePage * FreePageList;
int NumberOfFreePages;
```

# Paged memory allocator (2 of 4)

```
• void Initialize(int npages) {
 RequestListFront = 0; FreePageList = 0;
 NumberOfFreePages = npages;
 for(int i = 0; i < NumberOfFreePages; ++i) {
 FreePageList = new FreePage(i, FreePageList);
 }
}
// request procedure: request a piece to be allocated
void RequestBlock(int npages, Semaphore * satisfied,
 int * pageTableArray) {
 MemoryRequest * n = new MemoryRequest(npages,
 satisfied, pageTableArray, 0 , 0);
 if(RequestListFront == 0) { // list was empty
 RequestListFront = RequestListBack = n;
 } else {
 RequestListBack->next = n;
 RequestListBack = n;
 }
 TryAllocating();
}
```

# Paged memory allocator (3 of 4)

- ```
void TryAllocating( void ) {
    MemoryRequest * request = RequestListFront;
    while( request != 0 ) {
        if( CanAllocate( request ) ) {
            if( RequestListFront == RequestListBack ) {
                RequestListFront = 0;
                request = 0; // drop out of loop
            } else {
                request->prev->next = request->next;
                request->next->prev = request->prev;
                MemoryRequest * oldreq = request;
                // save the address
                request = request->next;
                delete oldreq;
            }
        } else
            request = request->next;
    }
}
```

Paged memory allocator (4 of 4)

- ```
void FreePages(int npages, int pageTable[]) {
 for(int i = 0; i < npages; ++i)
 FreePageList
 = new FreePage(pageTable[i], FreePageList);
}

int CanAllocate(MemoryRequest * request) {
 if(request->npages >= NumberOfFreePages) {
 NumberOfFreePages -= request->npages;
 int * p = request->pageTableArray;
 for(int i = 0; i < request->npages; ++i) {
 *p++ = FreePageList->pageNumber;
 FreePage * fpl = FreePageList;
 FreePageList = FreePageList->next;
 delete fpl;
 }
 return True;
 }
 return False;
}
```



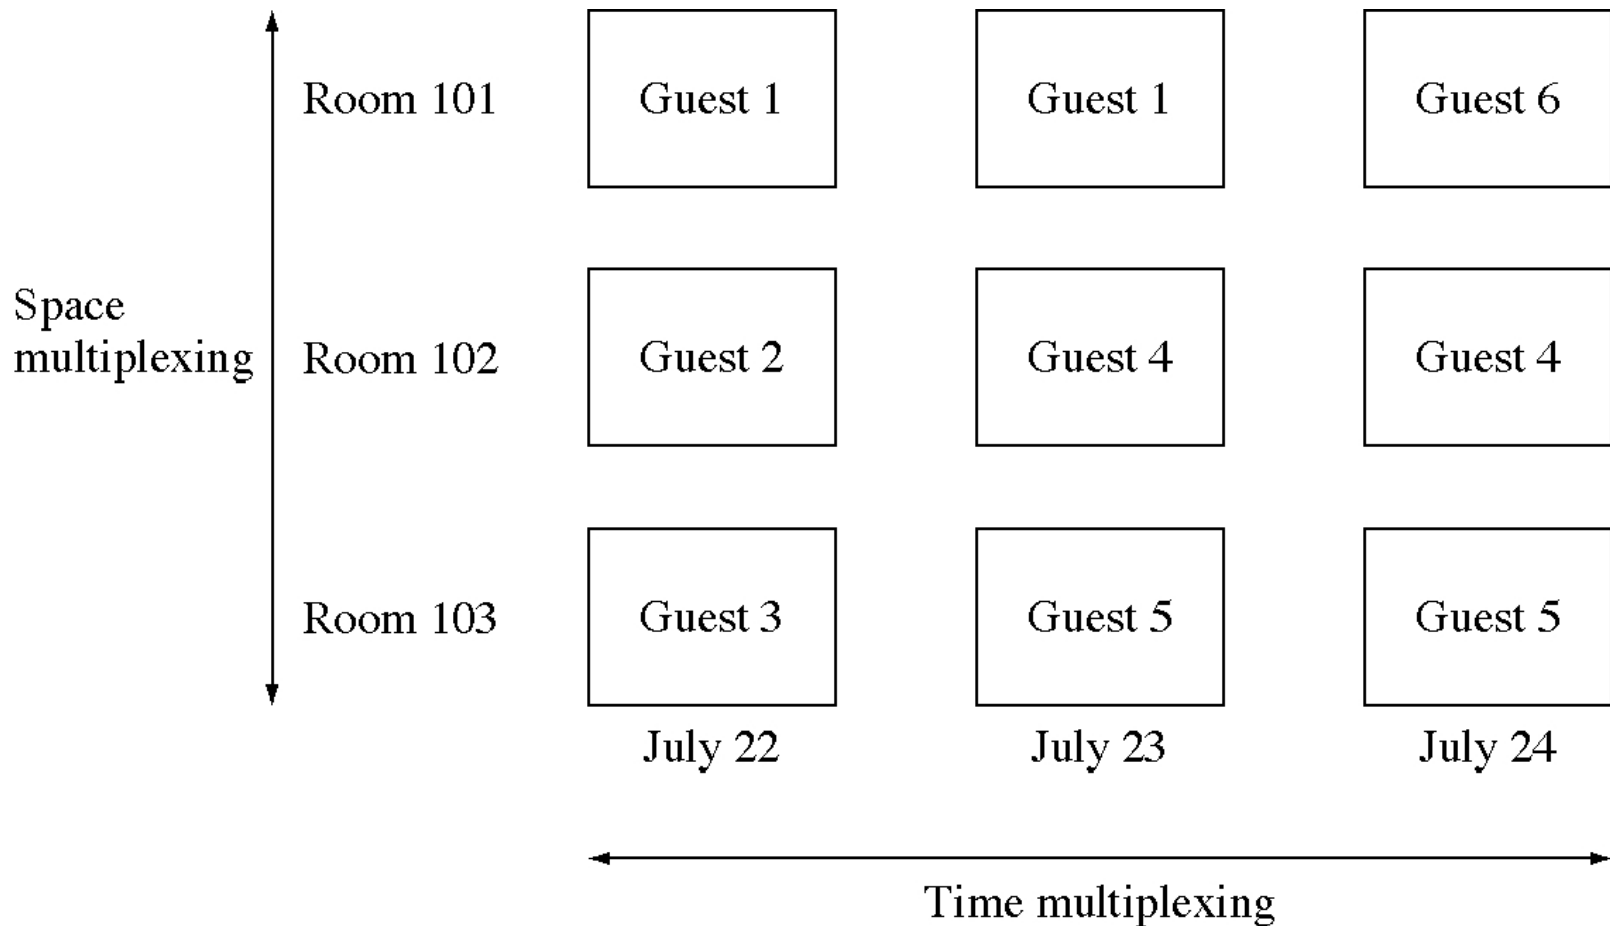
# The design process

- Evolution of solutions to the memory problem is a good example of the design process in action
  - at each stage the current solution had a problem
  - we modified the design to fix the problem
  - this created a new problem
  - we continued until the solution was good enough
- Sometimes we reused previously discarded ideas

# Time and space multiplexing

- The processor is a time resource
  - It can only be time multiplexed
- Memory is a space resource
  - We have looked at space multiplexing of memory
  - Now we will look at time multiplexing of memory

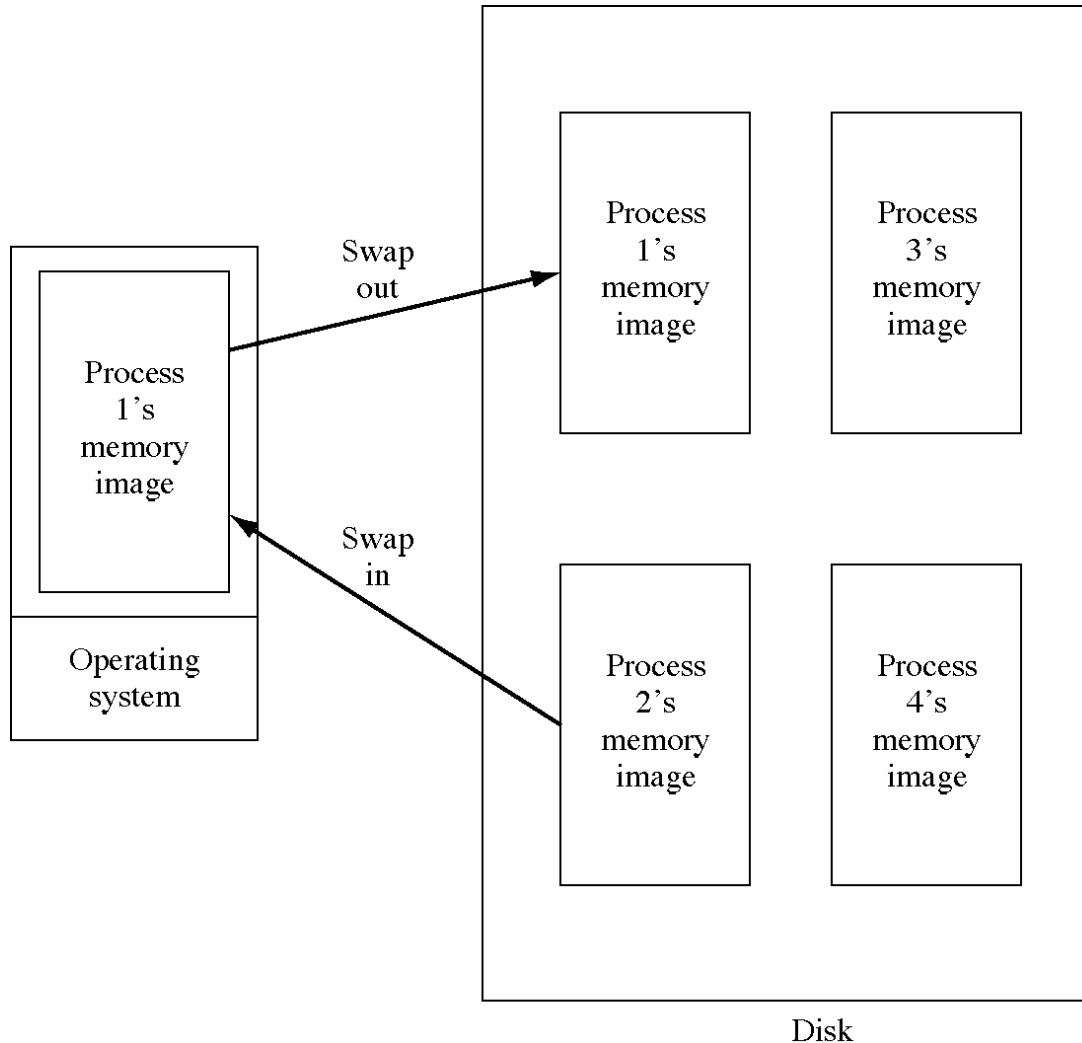
# Time and space multiplexing of hotel rooms



# Time multiplexing memory

- Swapping: move whole programs in and out of memory (to disk or tape)
  - allowed time-sharing in early OSs
- Overlays: move parts of program in and out of memory (to disk or tape)
  - allowed the running of programs that were larger than the physical memory available
  - widely used in early PC systems

# Swapping



# Design technique:

## Persistent objects

- A process was a dynamic entity in the system
  - we wanted to write it out to disk
  - and read it back in again later
  - kind of freeze and unfreeze it
- The ability to write an object to disk is called *persistence*, and it very useful
- It allows objects to live beyond the execution of the program that creates them

# How to create persistence

- Basically simple
  - write out a representation of the objects
- Problems
  - pointers and references: we must encode these in some way in order to write them out
  - following references: we also have to write out everything the object refers to.

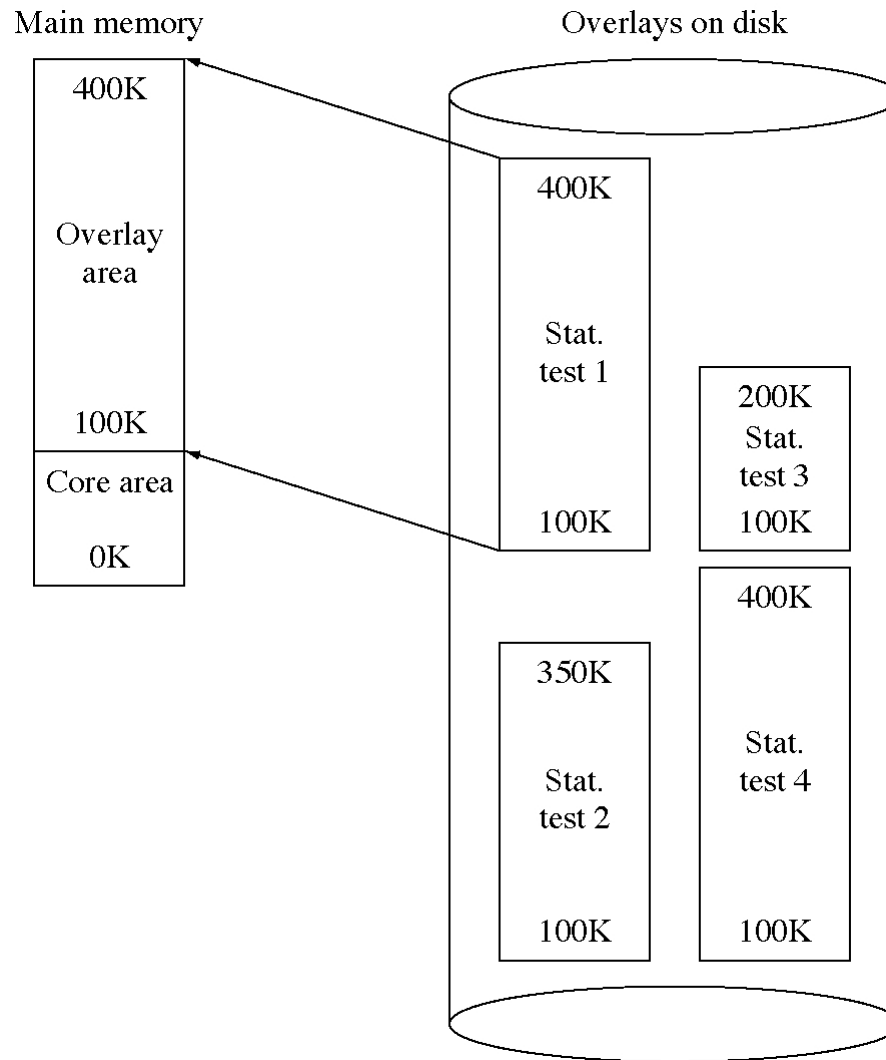
# Design stages:

## Design and implementation

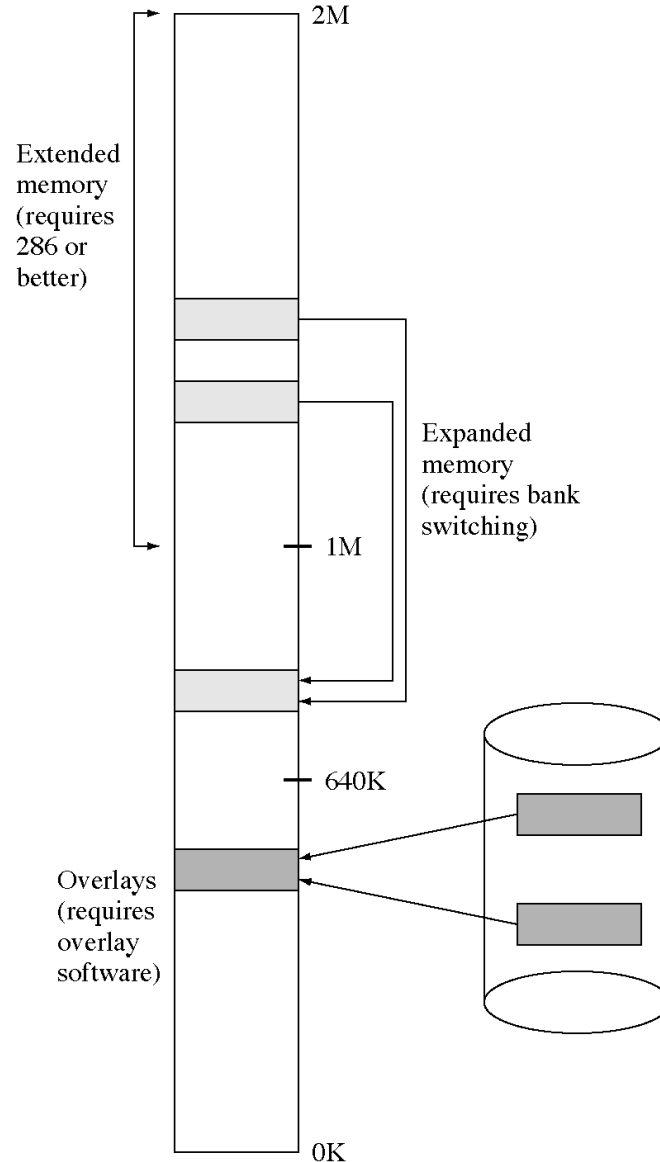
- First we decide what we want
- Then we decide how to implement it
- Each stage requires design
- Feedback is often required in order to avoid inefficient designs
- In operating systems
  - we often use useful services to the processes
  - and try to implement them efficiently



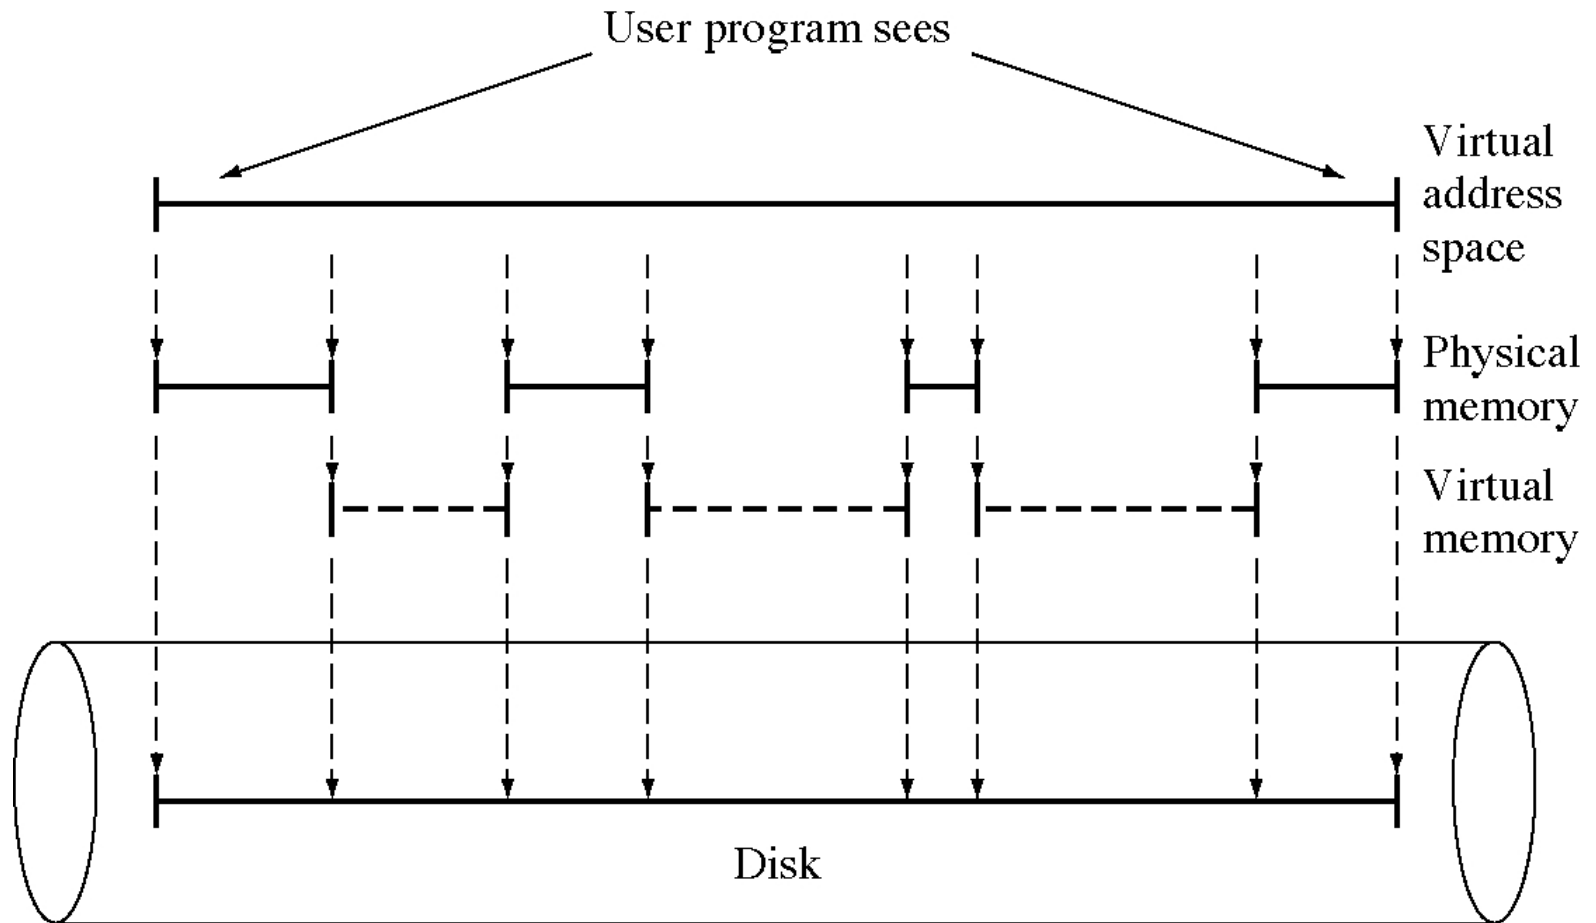
# Overlays



# Overlays in PCs



# Virtual memory



# Implementation of virtual memory

- Virtual memory allows
  - time multiplexing of memory
  - users to see a larger (virtual) address space than the physical address space
  - the operating system to split up a process in physical memory
- Implementation requires extensive hardware assistance and a lot of OS code and time
  - but it is worth it

# Virtual memory algorithm (1 of 2)

- ```
const int LogicalPages    = 1024;
const int BytesPerPage    = 4096;
const int OffsetShift     = 12;
const int OffsetMask      = 0xFFF;
const int PhysicalPages   = 512;
enum AccessType { invalid = 0, read = 1, write = 2,
execute = 3 };
struct PageTableEntry {
    int pageBase    : 9;
    int present     : 1;
    AccessType protection : 2;
    int fill        : 4; // fill to 16 bits
};
PageTableEntry UserPageTable[LogicalPages];
```


Virtual memory algorithm (2 of 2)

```
• int MemoryAccess( int logicalAddress,
    AccessType how, int dataToWrite = 0 ) {
    int page = logicalAddress >> OffsetShift;
    int offset = logicalAddress & OffsetMask;
    PageTableEntry pte = UserPageTable[page];
    if( how != pte.protection )
        if( !(how = read && pte.protection = write) ) {
            CauseInterrupt( ProtectionViolation );
            return 0;}
    if( pte.present == 0 ) {
        GenerateInterrupt( PageFault, page );
        return 0; }
    int physicalAddress
        = (pte.pageBase << OffsetShift) + offset;
    switch( how ) {
        case read: case execute:
            return PhysicalMemoryFetch(physicalAddress);
        case write:
            PhysicalMemoryStore(
                physicalAddress, dataToWrite );
            return 0;
    }
}
```

Virtual memory software

- The virtual memory (a.k.a. paging) system in the OS must respond to four events
 - process creation
 - process exit
 - process dispatch
 - page fault

Virtual memory events



Process creation actions

- 1. Compute program size (say N pages)
- 2. Allocate N page frames of swap space
- 3. Allocate a page table (in the OSs memory) for N page table entries.
- 4. Initialize the swap area
- 5. Initialize the page table: all pages are marked as not present.
- 6. Record the location in the swap area and of the page table in the process descriptor

Process exit actions

- 1. Free the memory used by the page table
- 2. Free the disk space in the swap area
- 3. Free the page frames in process was using

Process dispatch actions

- 1. Invalidate the TLB (since we are changing address spaces)
- 2. Load the hardware page table base register with the address of the page table for this process

Page fault actions

- 1. Find the faulting page (say page K)
- 2. Find an empty page frame. This will involve replacing a page.
- 3. Read in page K to this page frame
- 4. Fix up the page table entry for page K. Mark it present and set the base address.
- 5. Restart the process with the instruction that caused the page fault

Locality

- Programs do not access their address space uniformly
 - they access the same location over and over
- *Spatial locality*: processes tend to access location near to location they just accessed
 - because of sequential program execution
 - because data for a function is grouped together
- *Temporal locality*: processes tend to access data over and over again
 - because of program loops
 - because data is processed over and over again


Design technique: Locality

- Locality is almost always present
 - and often we can optimize a design by taking advantage of it.
- Caching is a common name for systems that take advantage of locality to optimize operations

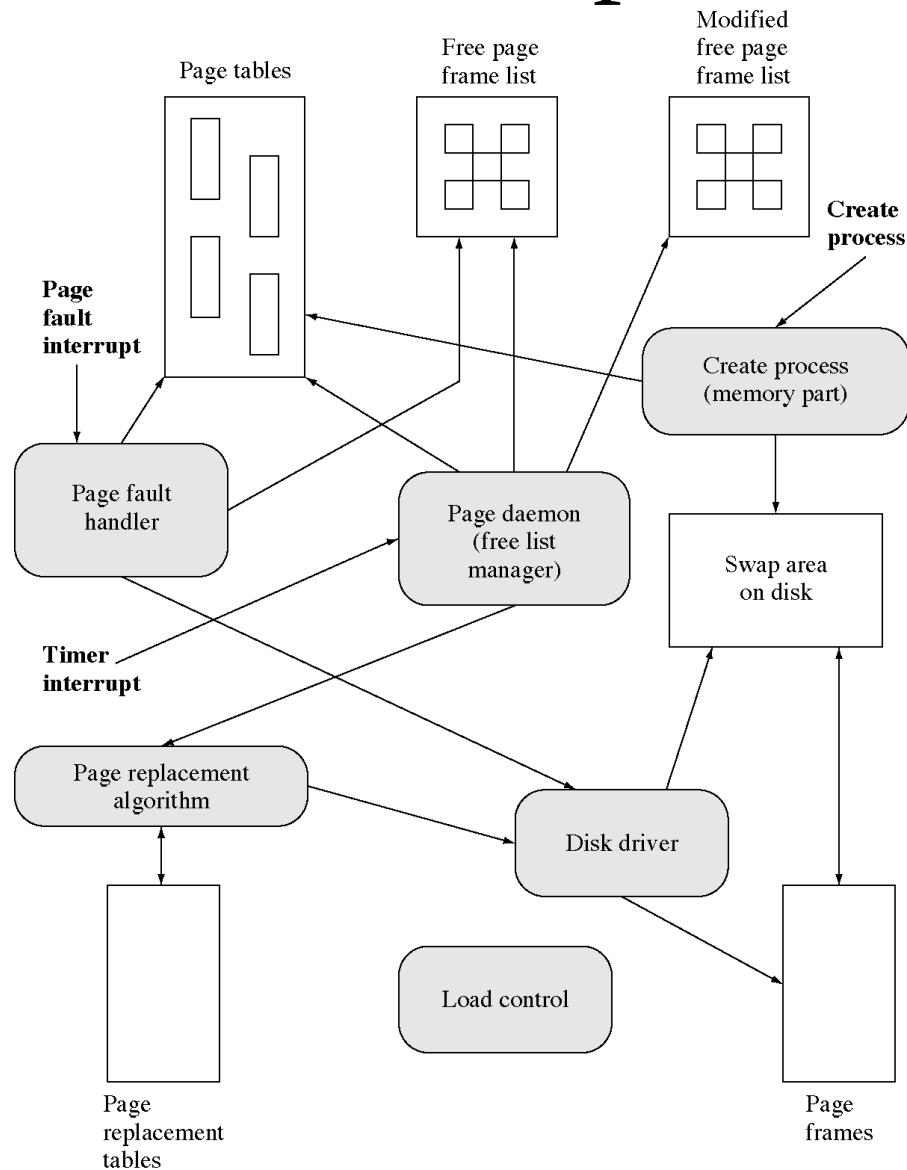
Practicality of paging

- Paging only works because of locality
 - at any one point in time programs don't need most of their pages
- Page fault rates must be very, very low for paging to be practical
 - like one page fault per 100,000 or more memory references

VM data structures



VM events and procedures



Daemons and events

- OS usually respond to events
 - but sometime they need to be proactive
- An OS daemon is a process that wakes up every so often and looks to see if it has any work to do
- It is useful to keep a pool of free pages
 - so page faults can be handled immediately
- A paging daemon wakes up every so often and keeps the free page pool large enough

Design technique:

System models and daemons

- Operating systems are essentially reactive
 - they react to interrupts
- But if we include timer interrupts
 - then operating systems are sort of doing things on their own, that is, being proactive
- This is what we call a daemon
 - a daemon wakes up and checks to see if something needs doing

Reactive and proactive user interfaces

- Graphical user interfaces are generally reactive, they wait for user actions
 - this is a good model and puts the user in control which is good psychologically
- Agents are a new user interface concept
 - agents are proactive
 - they go out and look for useful things to do

Design technique: Polling, software interrupts and hooks

- How can a process know when an event occurs? There are two approaches
- *polling*: it can check periodically
- *interrupts*: it can ask another process to interrupt it when the event occurs
 - the other process is usually the one that causes or handles the event
 - so it is not much trouble for it to inform the waiting process that the event has occurred.

Polling versus interrupts

- Daemons use polling to discover events
 - and then react to the event
- Polling is easier to set up but is less efficient than being interrupted
 - but interrupts require a process to do the interrupting
 - there might be no such process
 - or it might not be set up to provide interrupts


Hooks

- A hook is the ability to register a procedure to be called when an event occurs
- Systems that provide hooks are easy to modify
 - emacs provides hooks for many editing events
 - widget callback functions are hooks
 - a hook is basically a software interrupt

File mapping

- *File mapping* is the mapping of a file on disk into the virtual address space of a process.
- File I/O then consists of reading and writing words in the virtual address space
 - no system calls are required for read and write
- This is also called a *memory-mapped file*.
- The I/O system and the paging system both move data between disk and memory
 - so it makes sense to combine them

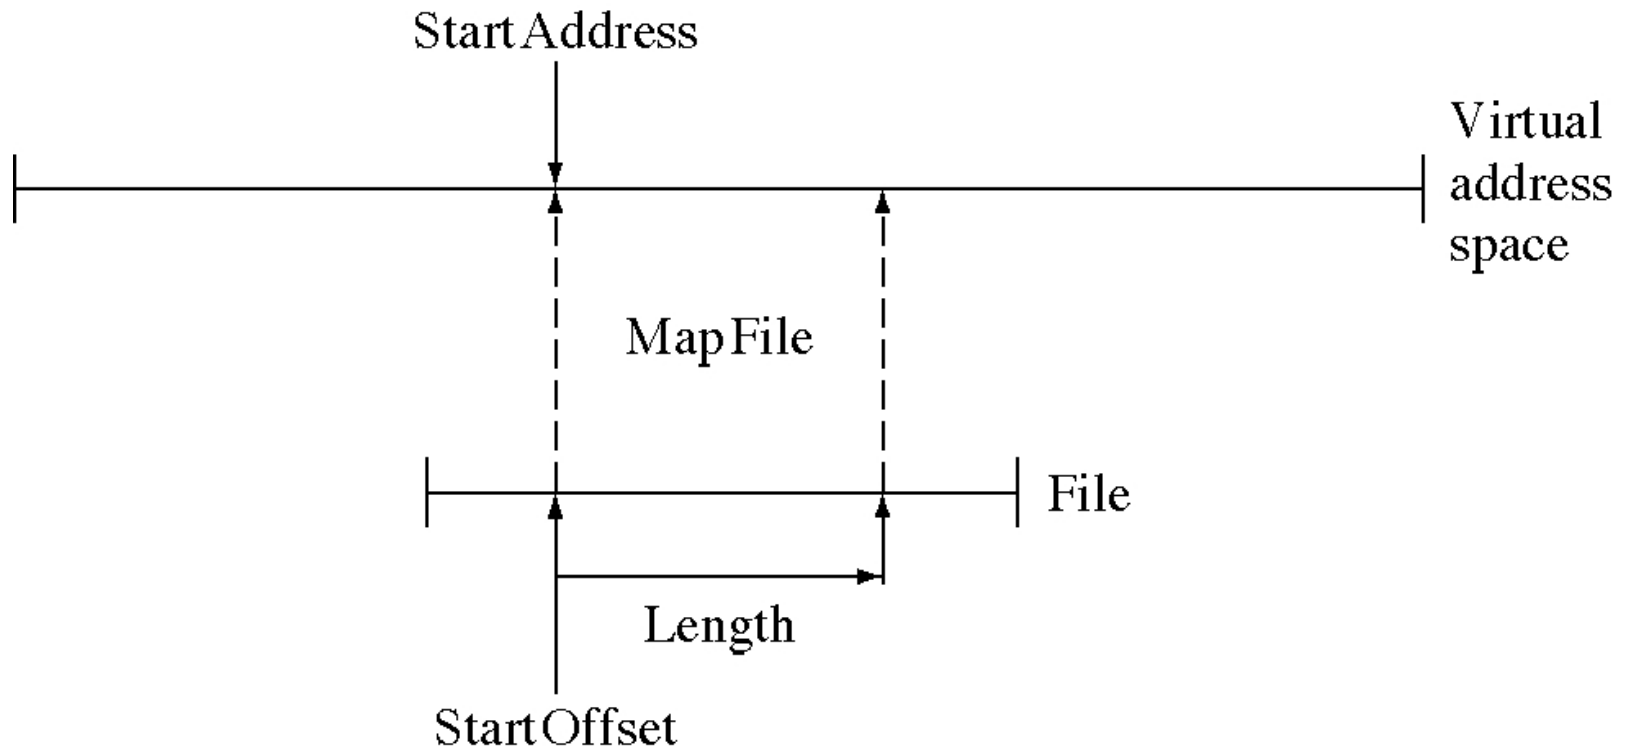
File mapping



File mapping system calls

- A system call is required to map the file into the address space (and one to remove it).
- `char * MapFile(
 int openFileId, // file already open
 char * startAddress = 0, // 0: OS does it
 int startOffset = 0, // into the file
 int length = 0); // 0: whole file`
- This call allows you to map a file in pieces
 - to save virtual address space

The MapFile system call




File mapping example code

- ```
int CountLetter(
 char *fileName, char letter) {
 int fid = open(fileName, Reading);
 char * fileArea = MapFile(fid);
 int fileLength = GetFileLength(fid);
 int letterCount = 0;
 for(int i=0; i < fileLength; ++i) {
 if(fileArea[i] == letter)
 ++letterCount;
 }
 UnMapFile(fid);
 return letterCount;
}
```


# Advantages of file mapping

- Simpler OS interface: no explicit I/O
- More efficient: system calls are not required for reading and writing files
- Reduces the number of copies of file data in memory
- Almost all modern operating system use it

# Multiple memory copies of file data



(a) Each process has a physical copy of the data



(b) The processes share a physical copy of the file data

# Virtual memory in the IBM 801

