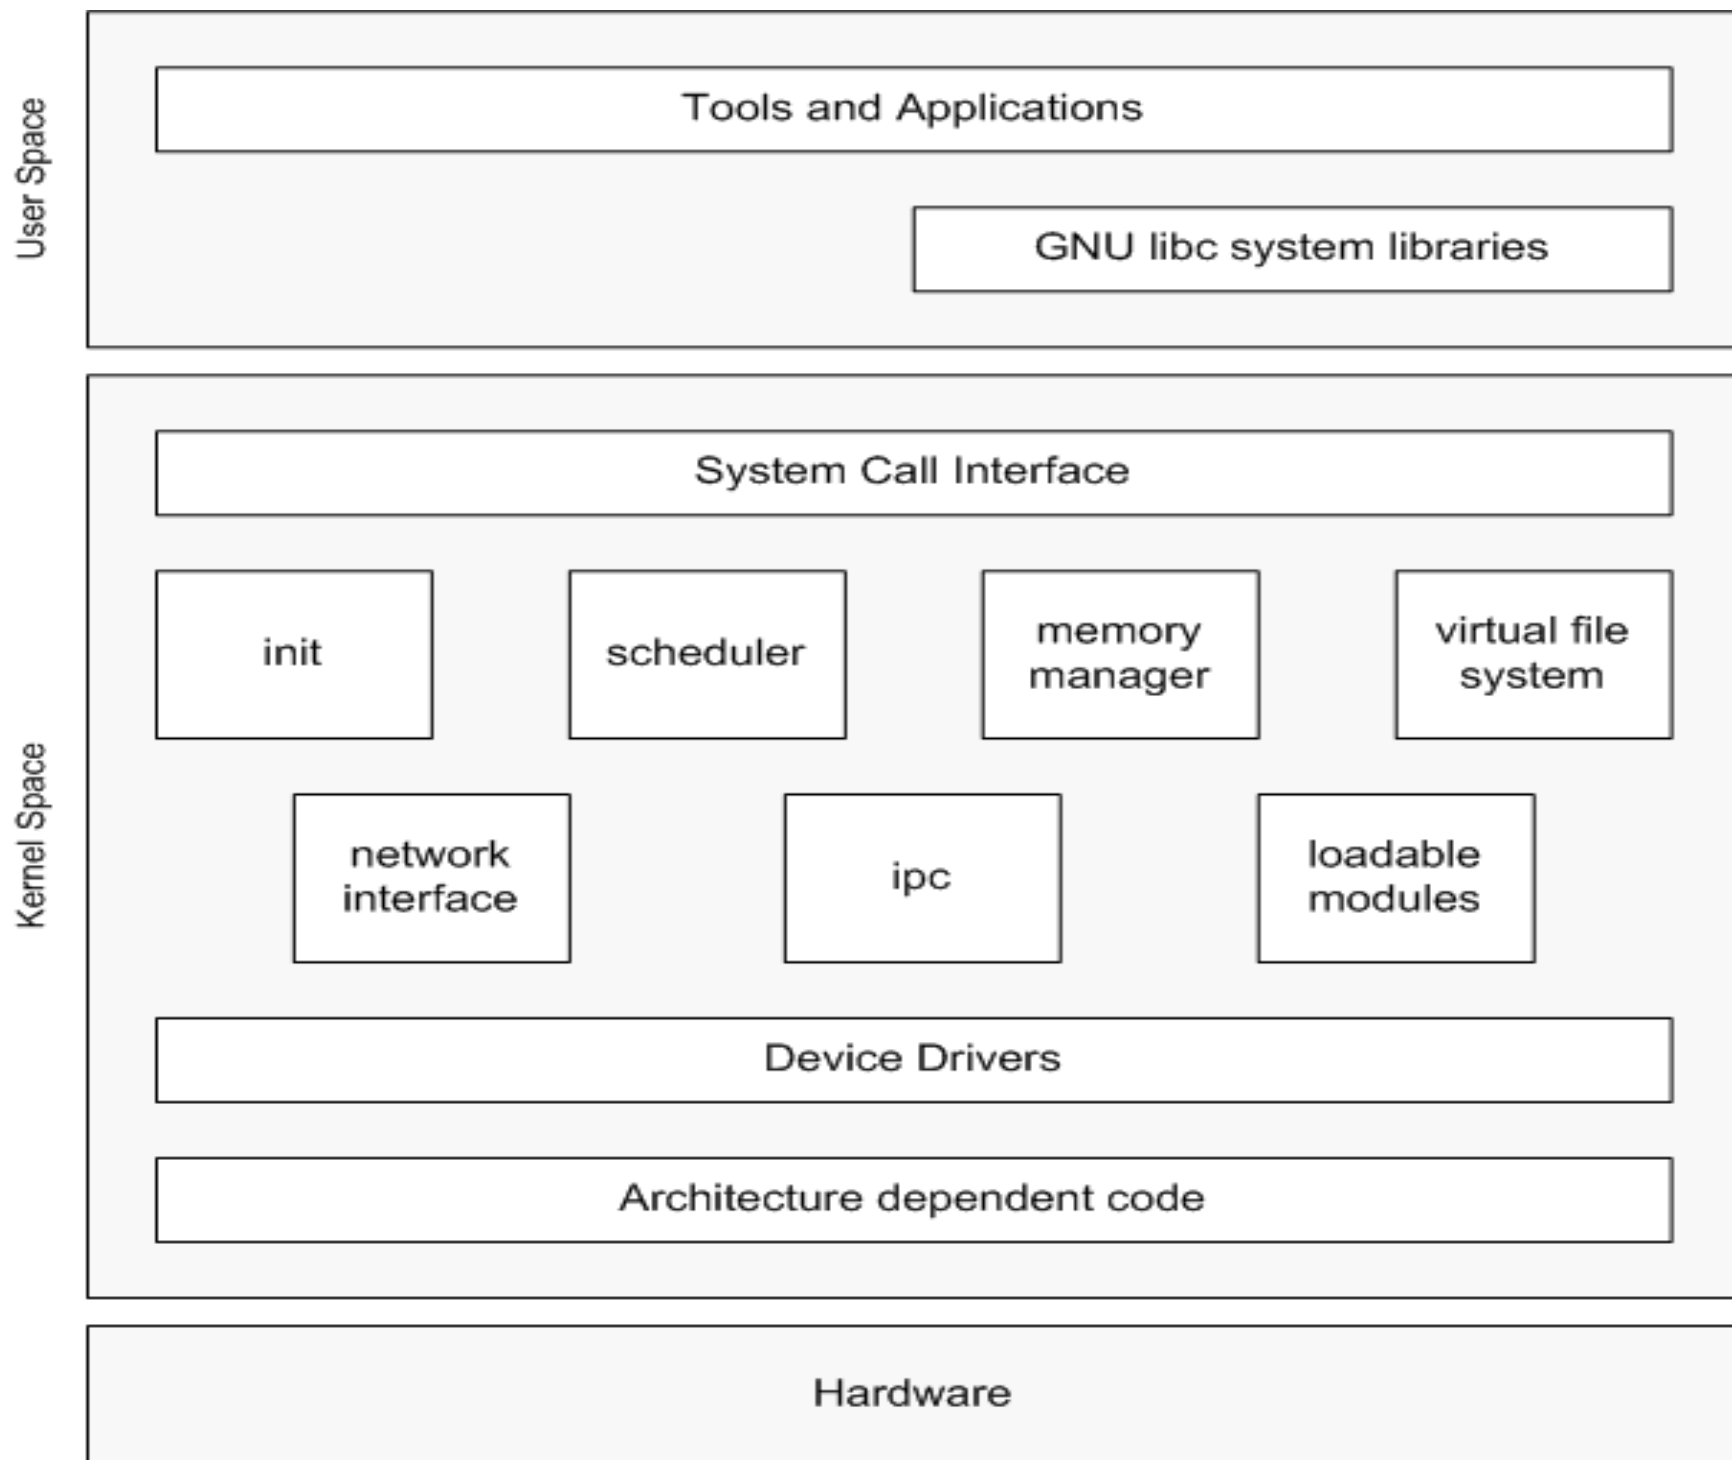


# Unix/Linux Primer

- OS knowledge assumed
- experience on GPOS/RTOS assumed
- Introduction to Unix/Linux Architecture  
will be covered in this session
- More details will be covered when we discuss  
specific interfaces
- Main focus will be on Linux

# Monolithic kernel Architecture

- It is a large kernel composed of several logically different components put together to form a large program image for the kernel
- Supports kernel module loading support for run-time extension of the kernel
- Most Unix variants are of this type
- Mac OS X is one exception



# System call wrappers

- Machine-dependent wrappers part of glibc
- Used implement services offered by the kernel to user-space applications
- Software generated interrupt
- Mode-switch to kernel-space and back to user-space
- Examples : `fork()`, `read()`, `write()`, `pipe()`, ....

# Library function calls

- Not system call wrappers
- May or may not use system call wrappers
- Examples : `printf()`, `fopen()`, `fread()`, `strncpy()`,...
- More convenient to use but less powerful
- Both system calls wrappers and Library function calls are taken by the Standards discussed in the previous sections

# System call wrappers/Interrupts/Exceptions

- All of them use similar mechanisms  
(subtle differences exist)
- All them end up calling the kernel
- Kernel executes in the context of the process that is currently executing when the event occurs
- Exceptions take care of FATAL program errors  
(invalid memory access, illegal instruction,.....)

# Process model

- Process descriptors(PD) – also known as task structures in Linux represent a process in the system
- PD contains only the most essential information
- Most resources allocated to the process are represented through other add-on data-structures
- Multi-tasking
- Independent and isolated (virtual)address spaces

# Process model

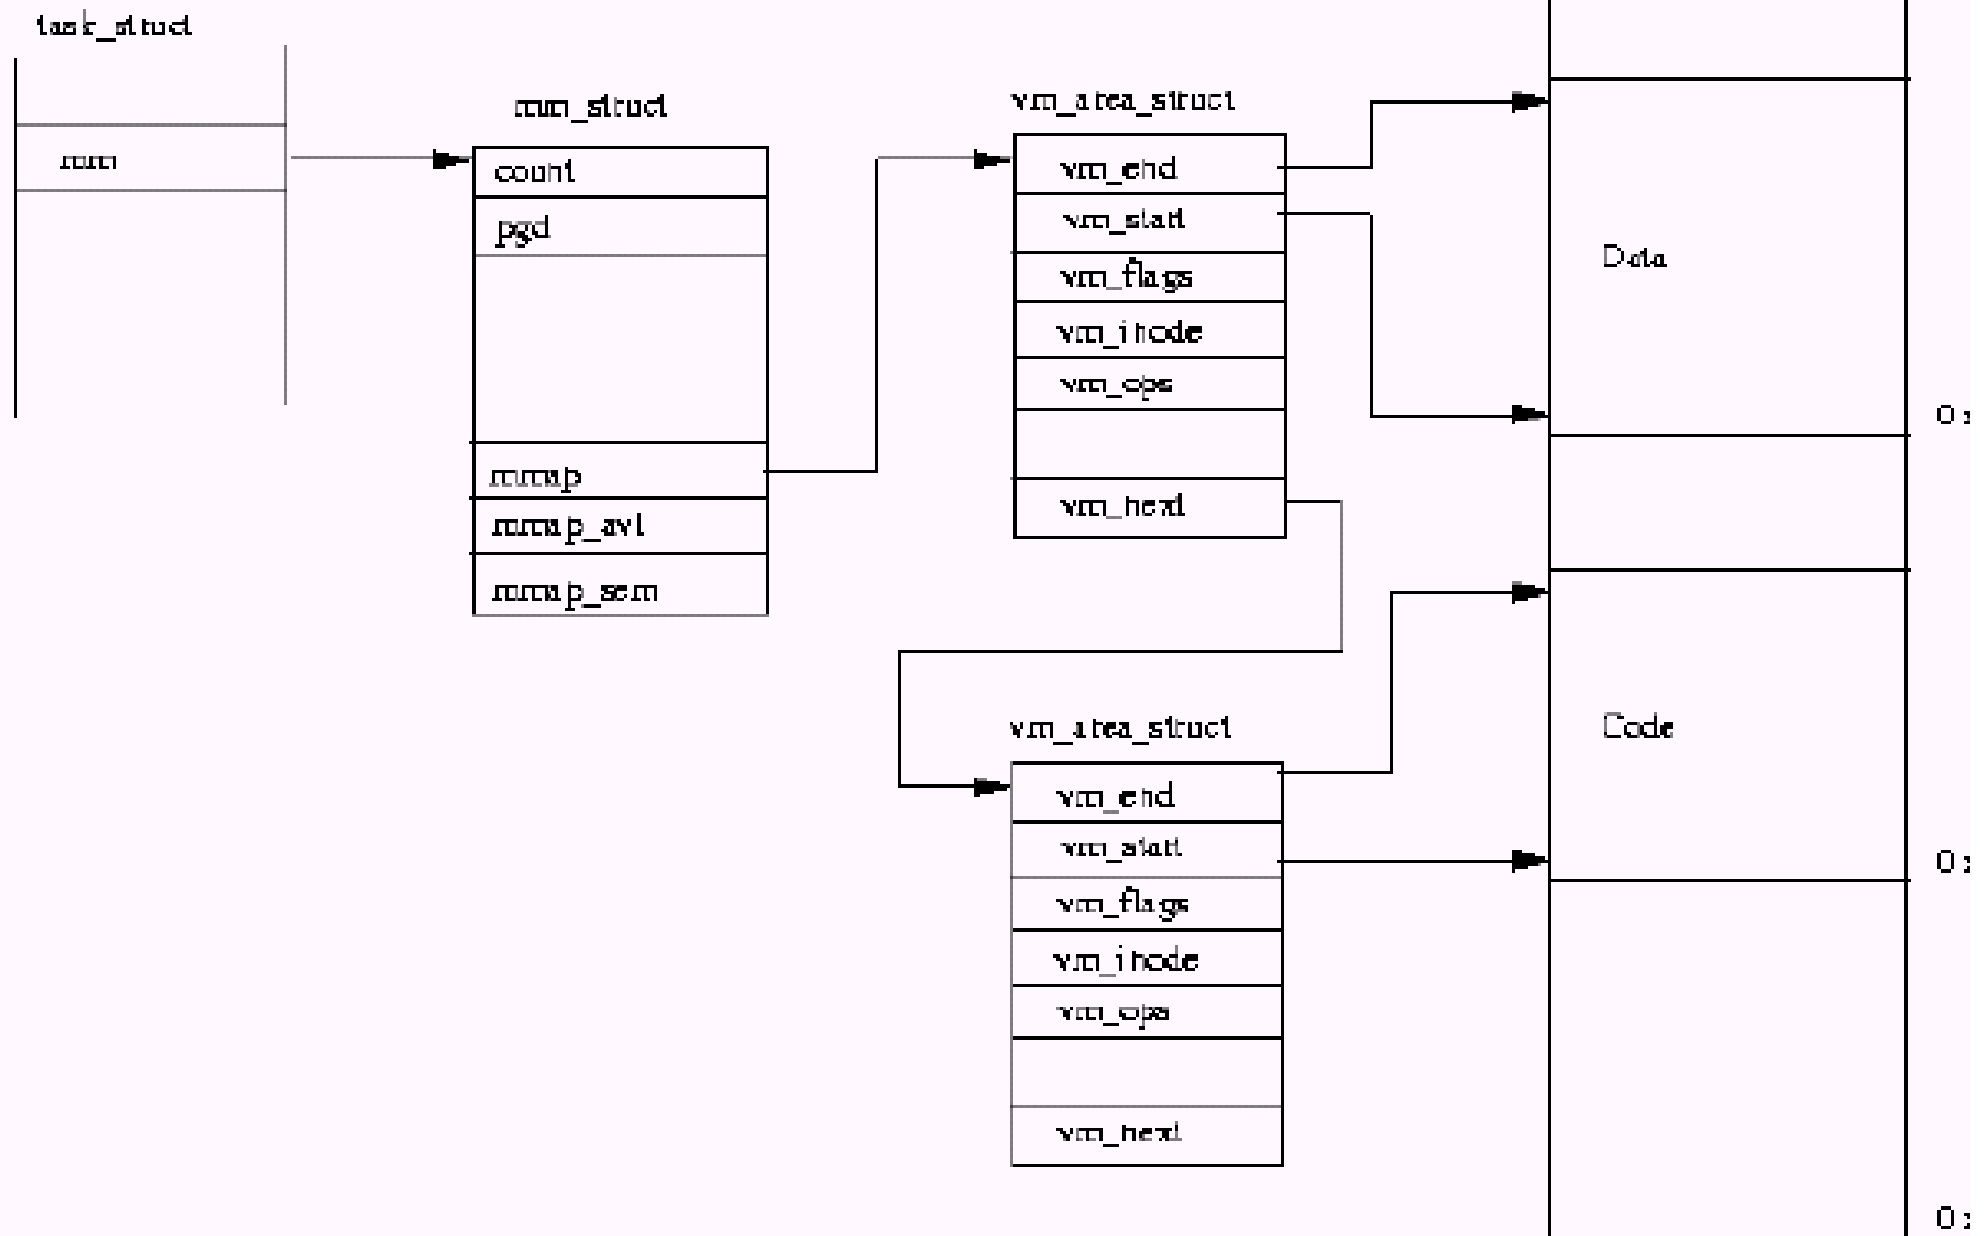
- Each process is uniquely identified by PID
- Parent – child relationship is maintained
- Parent – child relationship and its importance
- Process states
- Init process with  $PID = 1$
- Zombie processes
- Orphaned processes and Init process

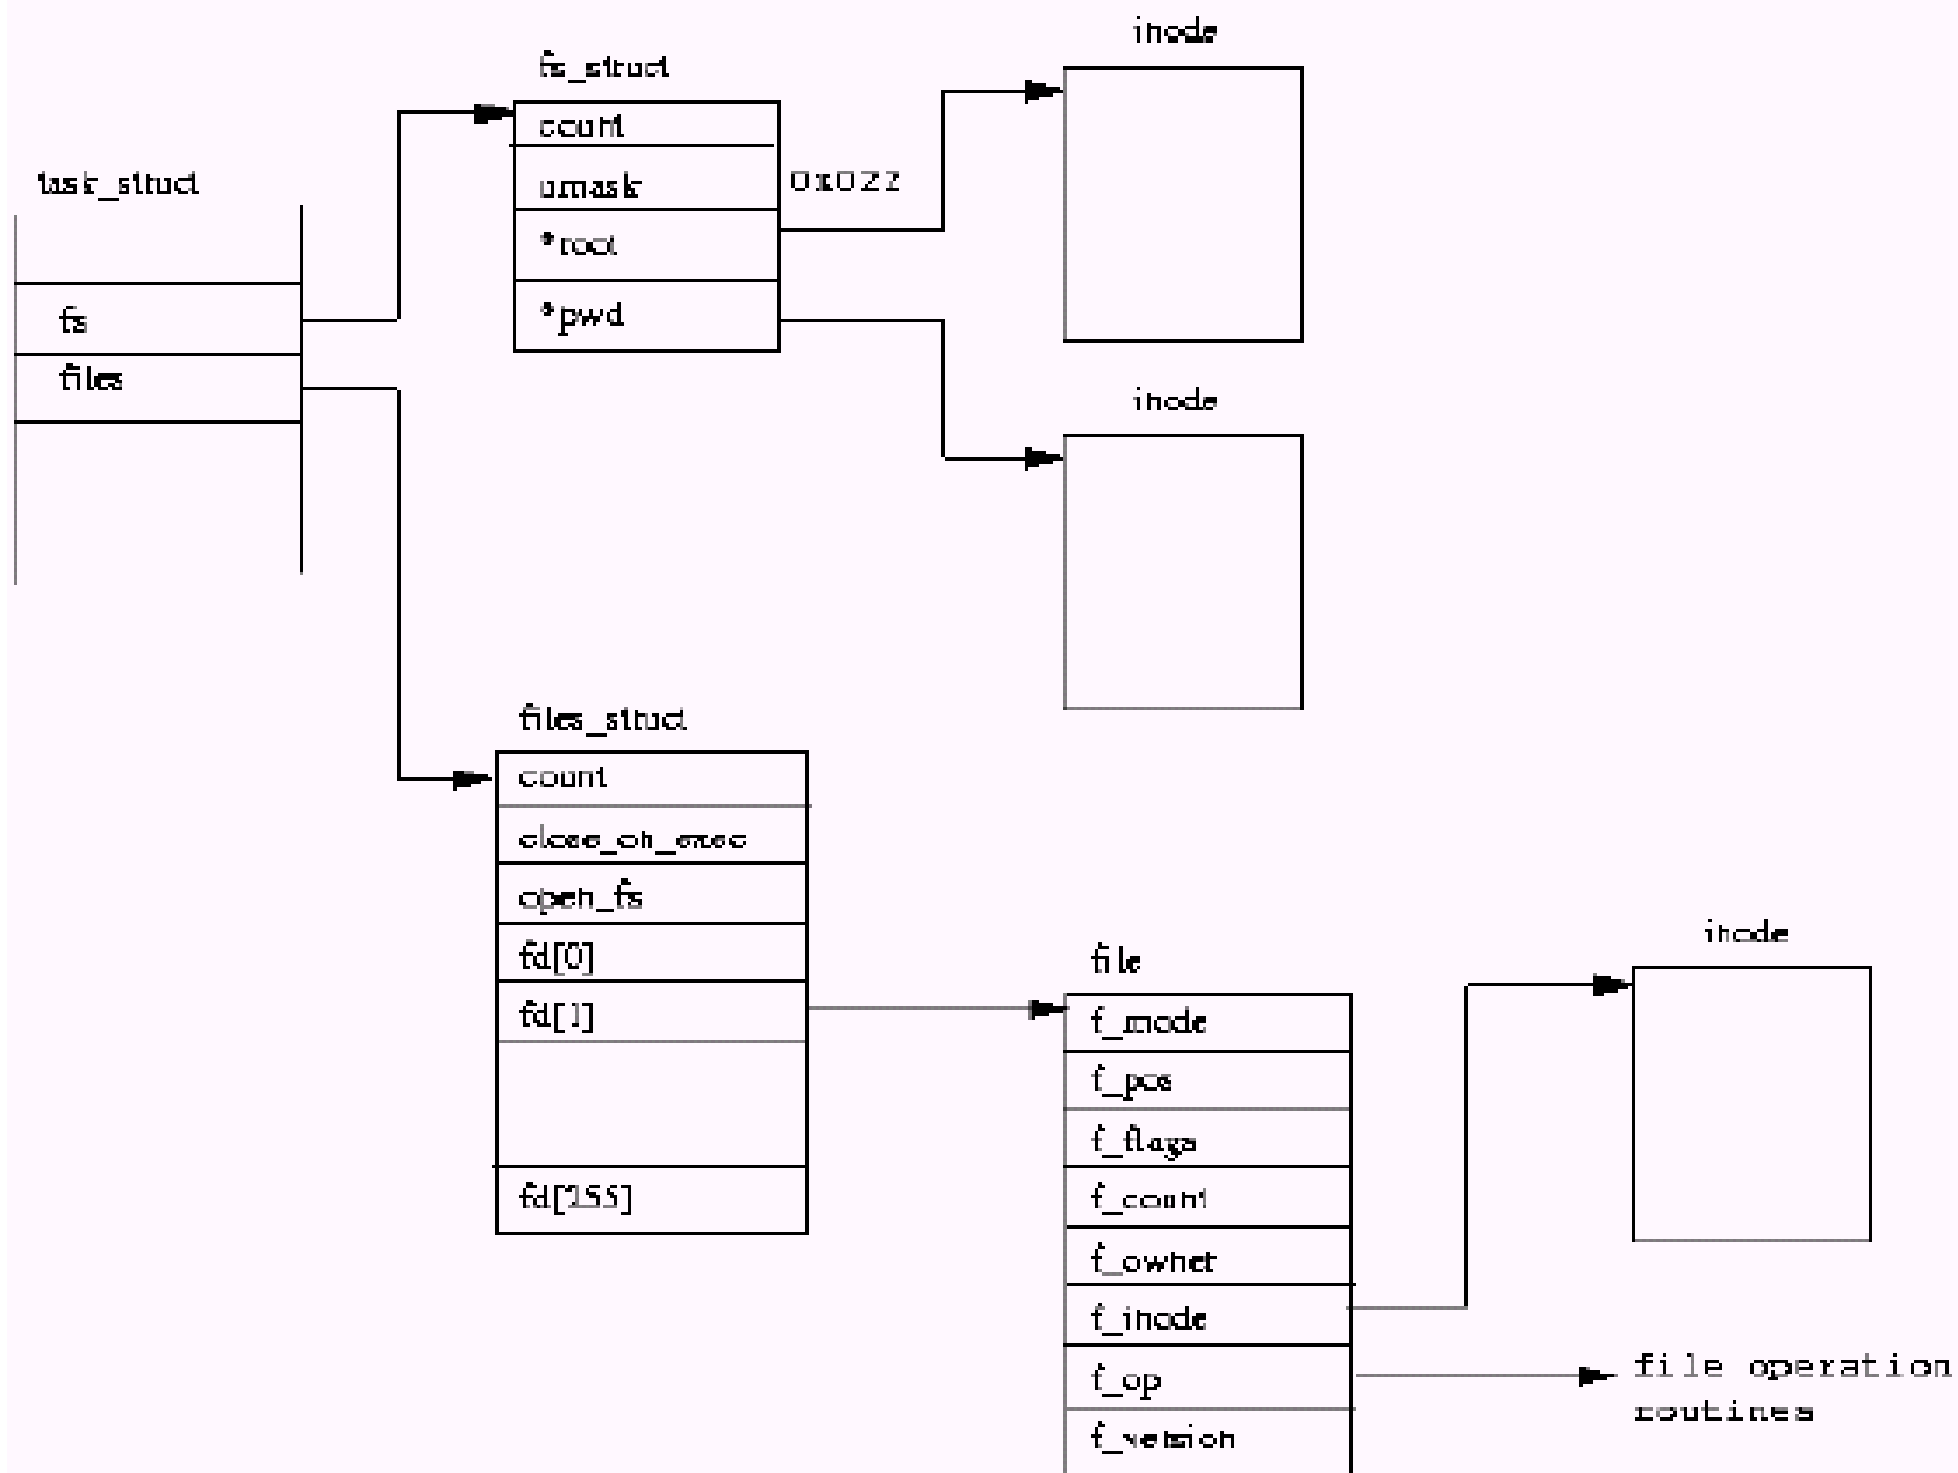


```

struct task_struct{
    volatile long
    long
    struct task_struct
    struct task_struct
    int
    int
    struct task_struct
    struct wait_queue
    struct task_struct
    unsigned long
    struct tms
    unsigned long
    unsigned short
    struct thread_struct
    struct files_struct
    struct mm_struct
    struct signal_struct
    sigset_t
    state;
    counter, priority;
    *next_task, *prev_task;
    *next_run, *prev_run;
    exit_code, exit_signal;
    pid;
    *p_opptr, *p_cpctr;
    *wait_chldexit;
    *pidhash_next;
    policy;
    times;
    start_time;
    uid, gid;
    tss;
    *files;
    *mm;
    *sig;
    signal, blocked;
};

```





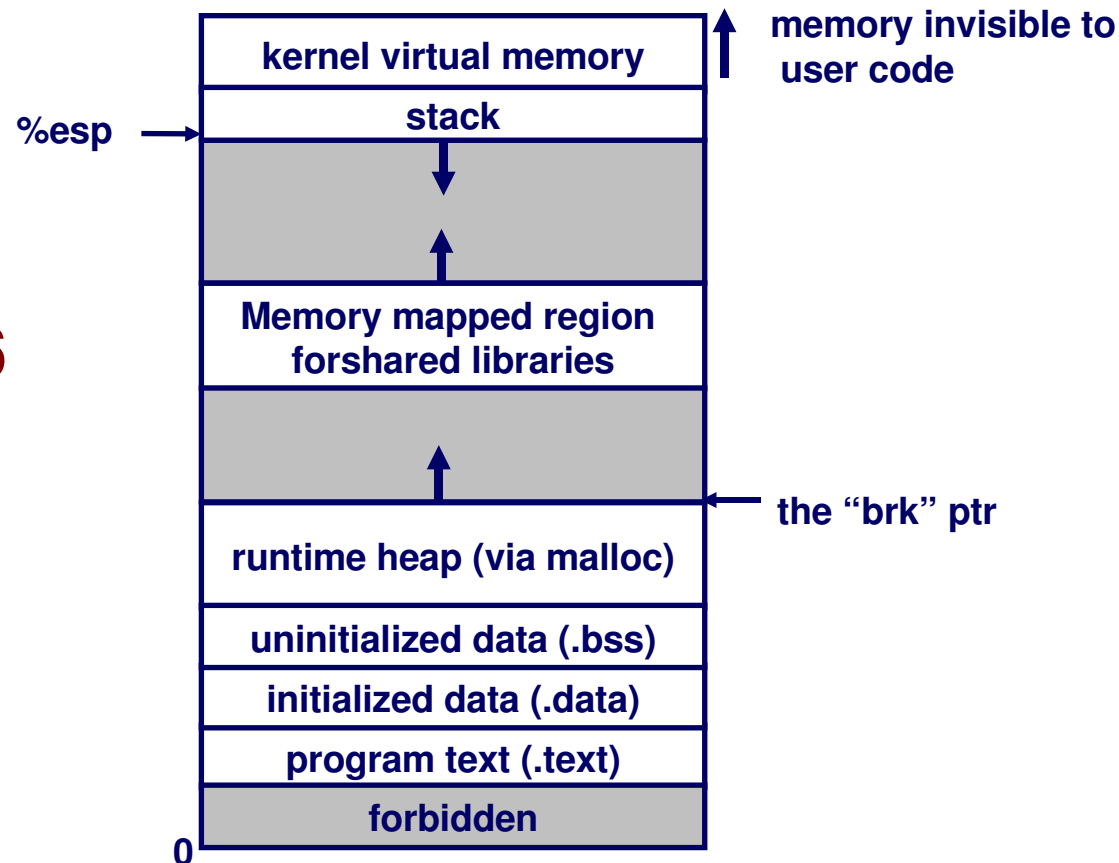
# Process Address Space

Multiple processes can reside in physical memory.

How do we resolve address conflicts?

- what if two processes access something at the same address?

**Linux/x86  
process  
memory  
image**



UNINTERRUPTIBLE  
INTERRUPTIBLE

ZOMBIE

terminate

wait

signal

preempt

event

schedule






create


signal

STOPPED

RUNNING

CPU





# CPU scheduling policies

- FAIR-SHARE / POSIX SCHED\_OTHER
- FIXED-PRIORITY REAL-TIME SCHEDULER/  
POSIX SCHED\_FIFO
- FIXED-PRIORITY RR REAL-TIME  
SCHEDULER/ POSIX SCHED\_RR
- The schedulers are preemptive
- The kernel is preemptive in both user-space/  
kernel-space

# CPU scheduling policies

- Real-time priorities – 1-99
- Non real-time priorities – 100-139
- System calls can be used to adjust the policies/priorities
- A good candidate for soft real-time OS
- Timer handler runs every 1ms (used to be 10ms)

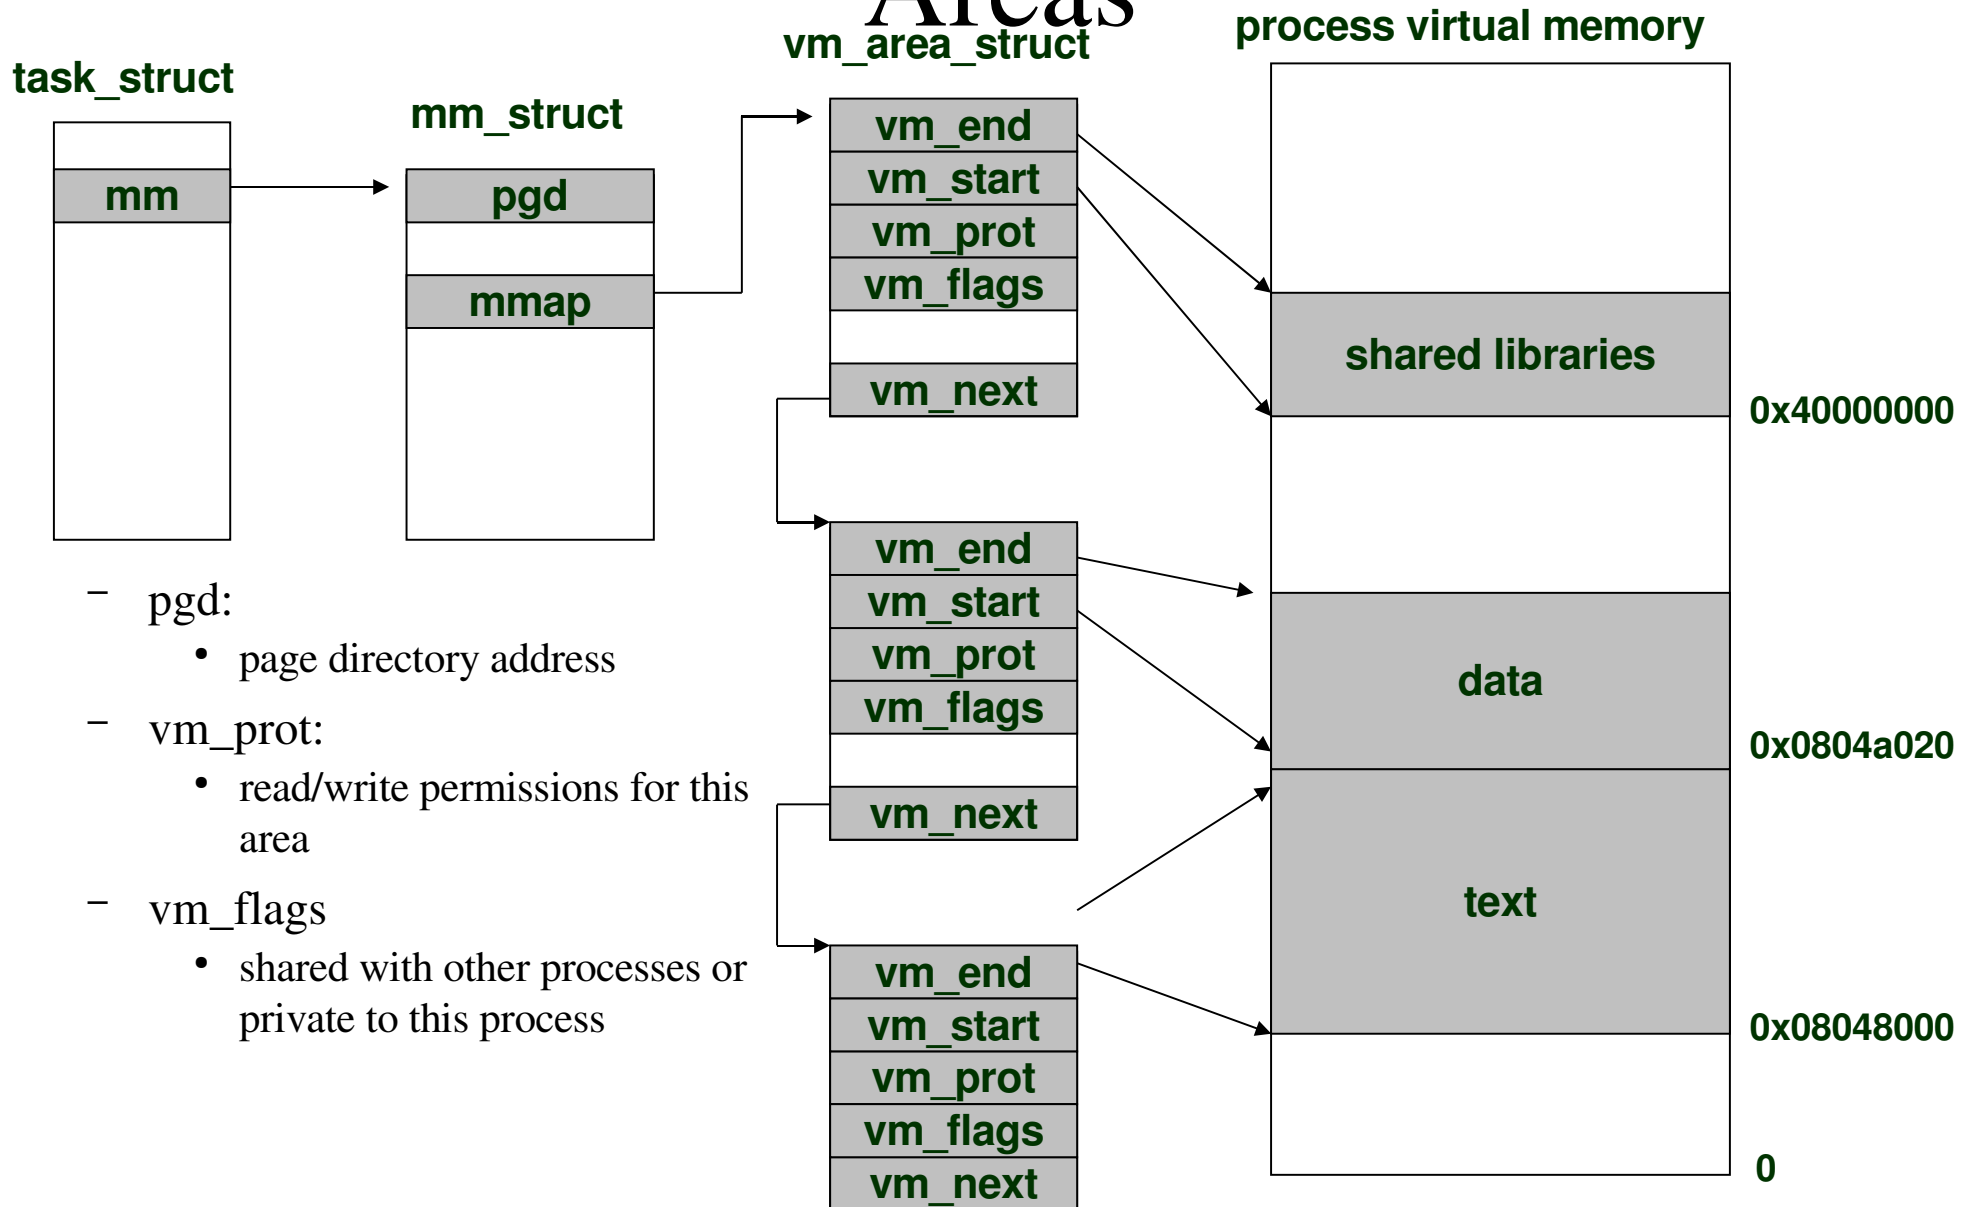


# Virtual memory model

- Private address spaces
- May have shared memory areas (shared memory) with other processes
- Demand-paged virtual memory scheme
- Copy-on-write is used during process creation
- Memory area structures and page-tables describe VM for a process
- Swap partitions exist

# Linux Organizes VM as Collection of

“Areas”



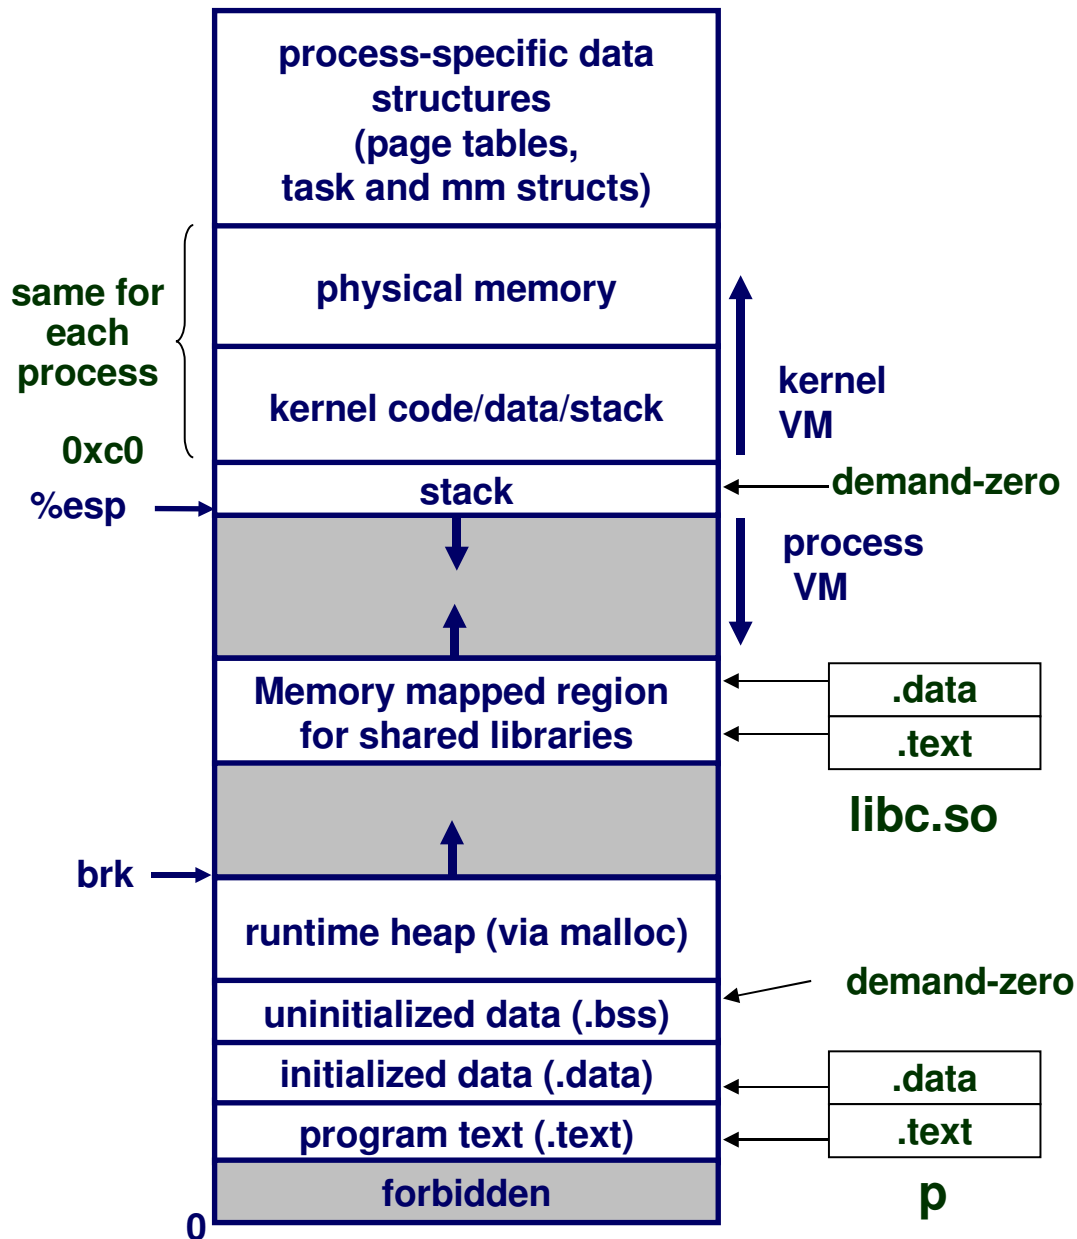
# Process model revisited

- Process duplication using `fork()`
- Process image overlay using `exec()`
- Linux uses Unix mechanism for creating processes – using `fork()`
- Changing the currently associated program with the process – using `exec()`
- We will see the advantages of this during the process system calls

# To create a new process using `fork ( )`:

- make copies of the old process's `mm_struct`, `vm_area_struct`'s, and page tables.
  - at this point the two processes are sharing all of their pages.
  - How to get separate spaces without copying all the virtual pages from one space to another?
    - “copy on write” technique.
- copy-on-write
  - make pages of writeable areas read-only
  - flag `vm_area_struct`'s for these areas as private “copy-on-write”.
  - writes by either process to these pages will cause page faults.
    - fault handler recognizes copy-on-write, makes a copy of the page, and restores write permissions.
- Net result:
  - copies are deferred until absolutely necessary (i.e., when one


# Exec() Revisited




To run a new program `p` in the current process using `exec()`:

- free `vm_area_struct`'s and page tables for old areas.
- create new `vm_area_struct`'s and page tables for new areas.
  - stack, bss, data, text, shared libs.
  - text and data backed by ELF executable object file.
  - bss and stack initialized to zero.
- set PC to entry point in `.text`
  - Linux will swap in code and data pages as needed.

# Linux File System Hierarchy



# Regular Files and Directories



# Linux File system

- Each file is uniquely represented by an i-node
- Regular files vs directory files
- Directory entry in the file system
- Hard-links
- Soft-links
- Hard-links vs Soft-links
- Open files and associated data-structures
- VFS – support for multiple file-systems



# Linux i-node

## inode

### file information:

size (in bytes)  
owner UID and GID  
relevant times (3)  
link and block counts  
permissions


direct pointers  
to beginning file blocks

single indirect pointer

double indirect pointer

triple indirect pointer


pointers to  
next file  
blocks



# Links

- A *link* is an association between a filename and an inode.
- UNIX has two types of links: hard and symbolic (also called soft).
- Directory entries are called hard links because they directly link filenames to inodes.
- Each inode contains a count of the number of hard links to the inode.
- When a file is created, a new directory entry is created and a new inode is assigned.
- Additional hard links can be created with  
    ln newname oldname or with the link system call.
- A new hard link to an existing file creates a new directory entry but assigns no other additional disk space.
- A new hard link increments the link count in the inode.
- A hard link can be removed with the rm command or the unlink system call.
- These decrement the link count.
- The inode and associated disk space are freed when the count is decremented to 0.

directory entry in /dirA



directory entry in /dirA


inode	name
-------	------

12345	name1
-------	-------

directory entry in /dirB

inode	name
-------	------

12345	name2
-------	-------



# Symbolic links or soft links

- A symbolic link is a special type of file that contains the name of another file.
- A reference to the name of a symbolic link causes the operating system to use the name stored in the file, rather than the name itself.
- Symbolic links are created with the command:  
`ln -s newname oldname` or `symlink()` system call
- Symbolic links do not affect the link count in the inode.
- Unlike hard links, symbolic links can span filesystems.

# Symbolic links or soft links..

directory entry in /dirA

inode	name
12345	name1

directory entry in /dirB

inode	name
13579	name2

inode 12345

⋮
1
⋮
23567
⋮

block 23567

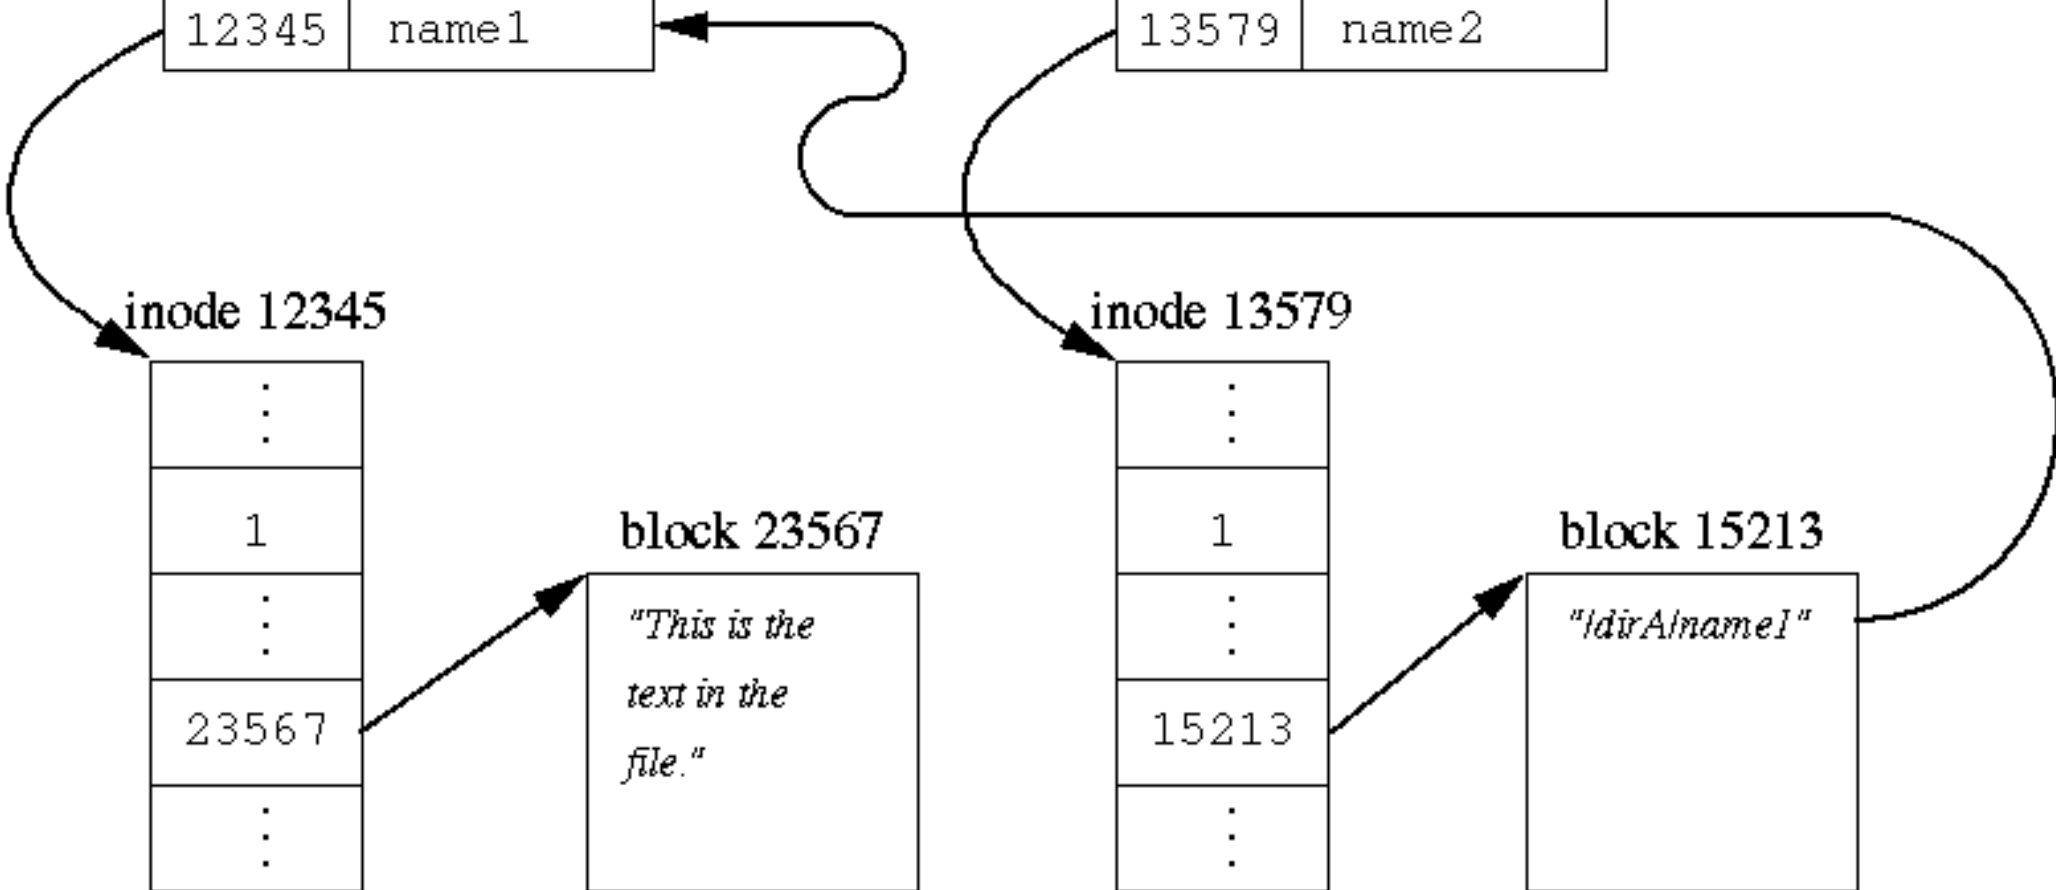
*"This is the  
text in the  
file."*

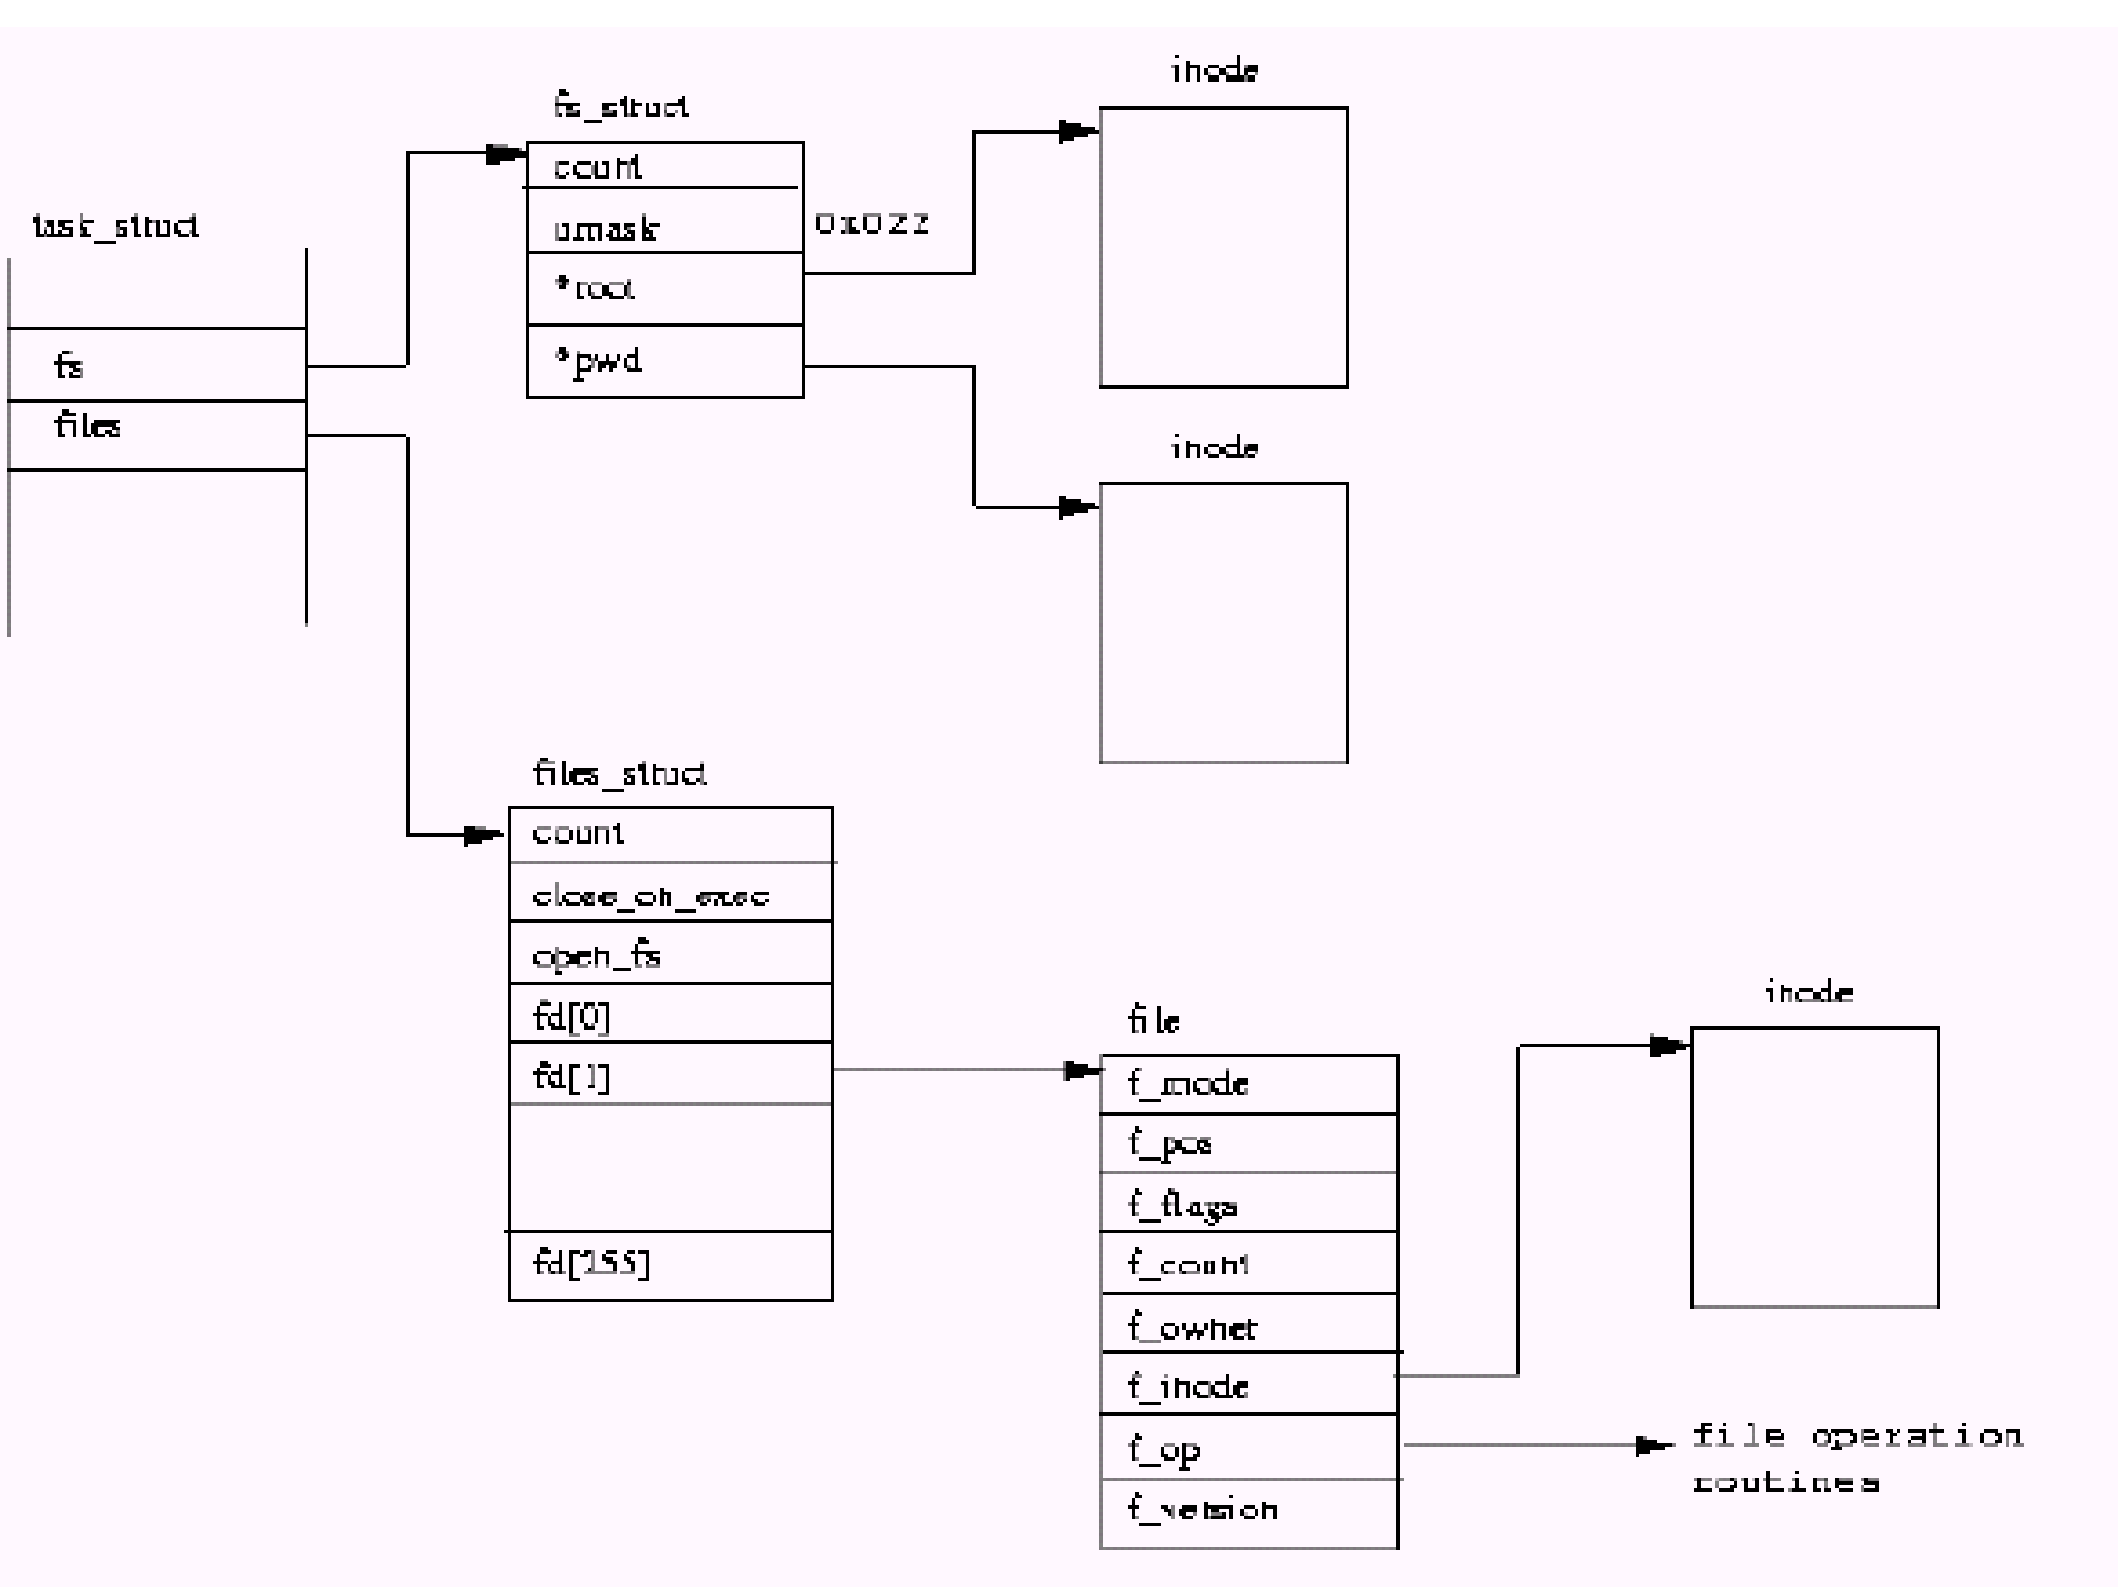
inode 13579

⋮
1
⋮
15213
⋮

block 15213

*"/dirA/name1"*





# Virtual File System (VFS)

Filesystems

ext2

smbfs


ramfs

...

reiserfs

proc

Device Drivers





# IPCS

- Signals
- Pipes
- FIFOs
- Message queues
- Semaphores
- Shared memory

## Signals :

A signal is *generated* when the event that causes the signal occurs.

- A signal is *delivered* when the process takes action based on the signal.
- The *lifetime* of a signal is the interval between its generation and delivery.
- A signal that has been generated but not yet delivered is *pending*.
- A process *catches* a signal if it executes a *signal handler* when the signal is delivered.
- Alternatively, a process can *ignore* a signal when it is delivered, that is to take no action.
- The function `sigaction` is used to specify what is to happen to a signal when it is delivered.
- The *signal mask* determines the action to be taken when the signal is generated. It contains a list of signals to be *blocked*.
- A *blocked* signal is not delivered to a process until it is unblocked.
- The function `sigprocmask` is used to modify the signal mask.
- Each signal has a *default action* which is usually to terminate.

```
#define SIGINT    2  /* interrupt, generated from terminal */
#define SIGILL    4  /* illegal instruction */
#define SIGABRT   6  /* abort process */
#define SIGFPE    8  /* floating point exception */
#define SIGKILL   9  /* kill a process */
#define SIGUSR1  10  /* user defined signal 1 */
#define SIGSEGV  11  /* segmentation violation */
#define SIGUSR2  12  /* user defined signal 2 */
#define SIGALRM  14  /* alarm clock timeout */
#define SIGCHLD  17  /* sent to parent on child exit */
#define SIGXCPU  24  /* cpu time limit exceeded */
```

# Sending signals from command line

- You can send a signal to a process from the command line using kill
- kill -l will list the signals the system understands
- kill [-signal] pid will send a signal to a process.
- The optional argument may be a name or a number. The default is SIGTERM.
- To unconditionally kill a process, use:  
kill -9 pid which is  
kill -KILL pid.

# Pipes

- Pipes are kernel buffers used to communicate between related processes
- They have a pipe data-structure and a kernel buffer of fixed-size
- Unstructured / Stream-oriented
- Typically, uni-directional (exceptions do occur)
- Uses the open file model described earlier
- System call `pipe()` creates a pipe and provides 2 file descriptors for read and write

# FIFOs

- FIFOs are kernel buffers used to communicate between related and unrelated processes
- They have a pipe data-structure and a kernel buffer of fixed-size
- Unstructured / Stream-oriented
- Typically, uni-directional
- Uses the open file model described earlier
- System call `mkfifo()` creates a FIFO file on-disk and must be explicitly opened for read and write

# SEMAPHORE


beginning section

WAIT (Guard)


critical section

SIGNAL (Guard)


remainder section




(a)



(b)



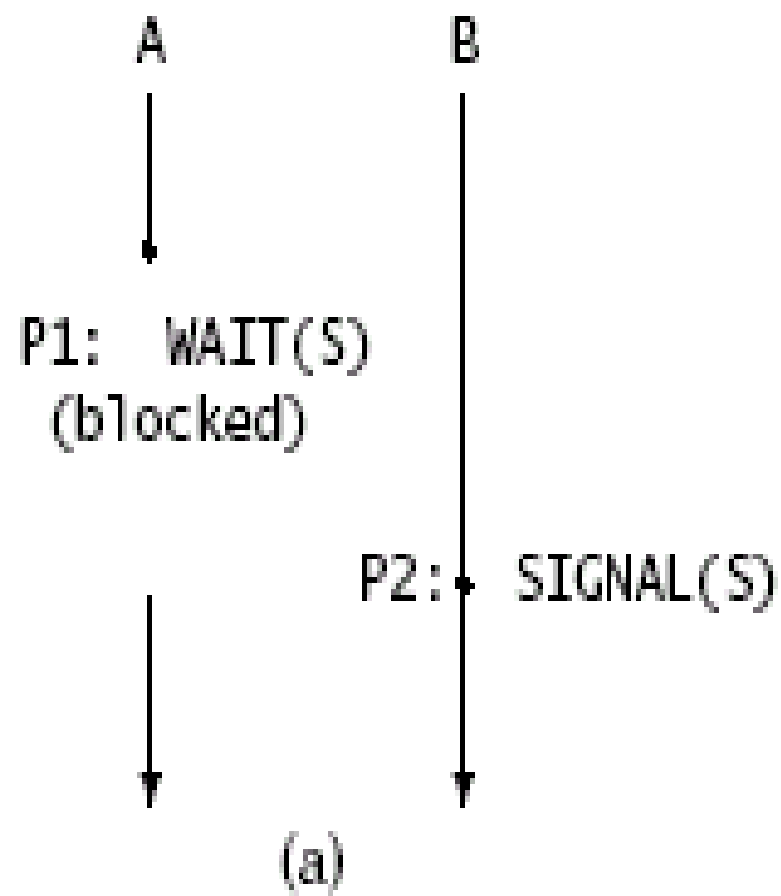



(c)



(d)







(a)

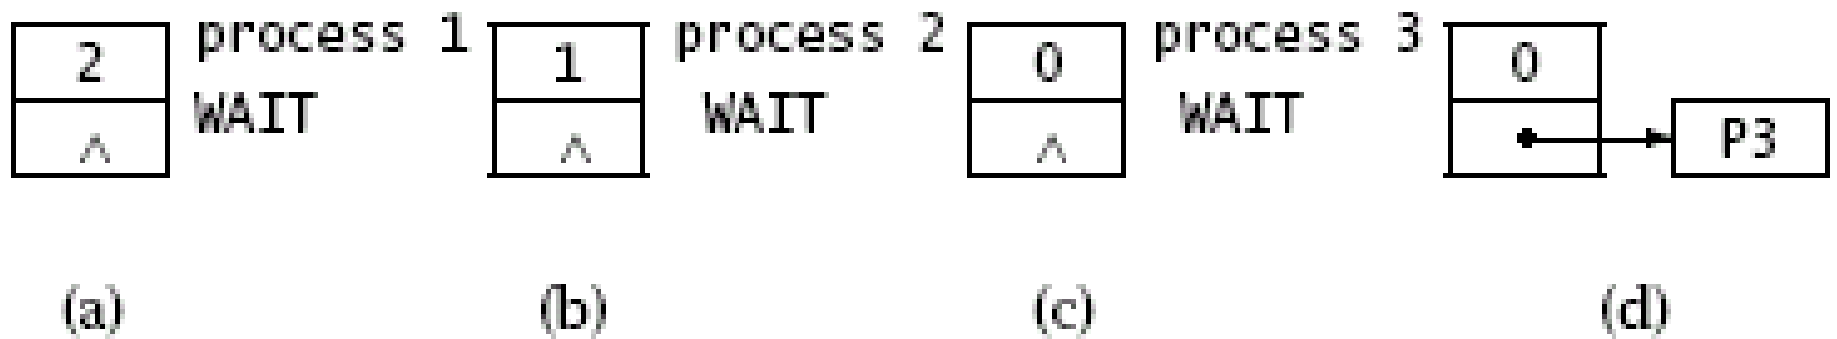
0
^

(b)

1
^

(c)

# COUNTING SEMAPHORE



# System V semaphore

```
struct semid_ds{  
    struct ipc_perm  sem_perm;  
    struct sem       *sem_base;  
    struct sem_queue *sem_pending;  
};
```

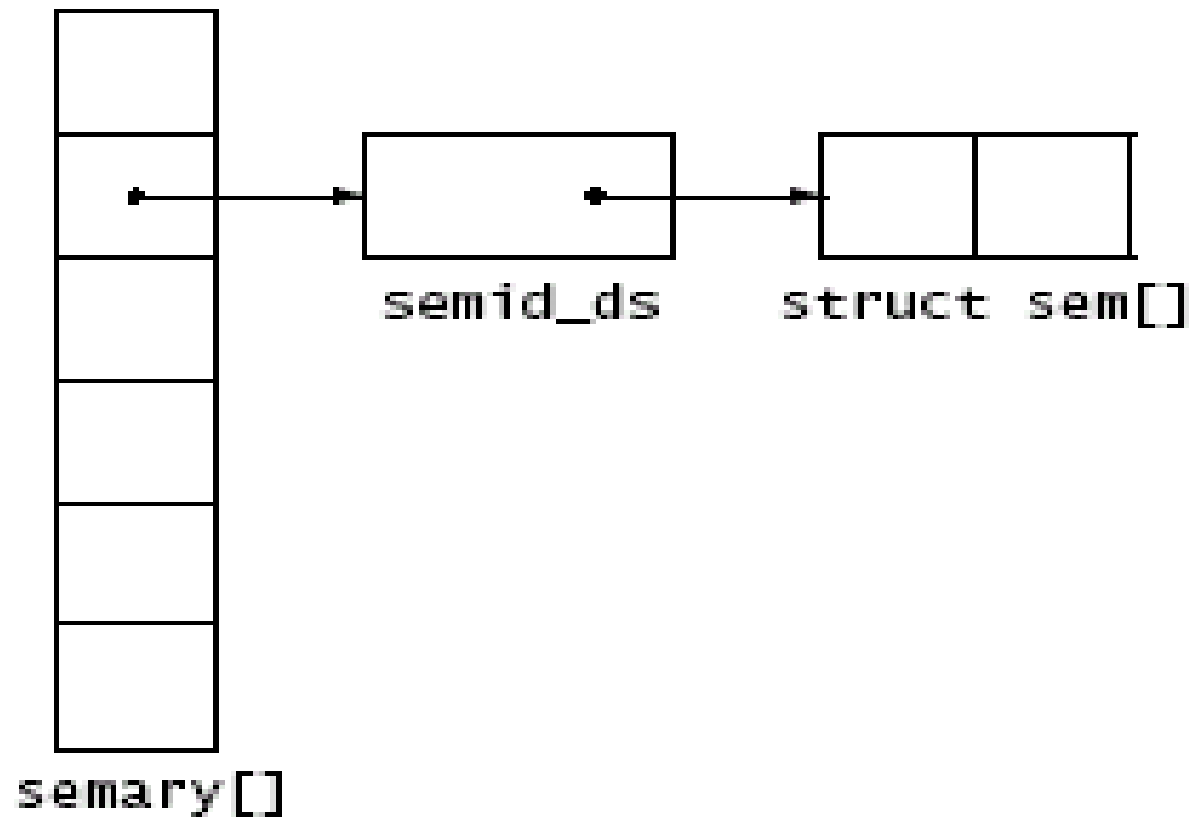
# System V semaphores ..

```
struct ipc_perm{  
    kernel_key_t  key;  /* user supplied key */  
    kernel_uid_t  uid;  /* owner's user id   */  
    kernel_gid_t  gid;  /* owner's group id  */  
    kernel_mode_t mode; /* access modes      */  
};
```

## System V semaphore ...

```
struct sem{  
    int semval; /* current value */  
    int sempid; /* process which last operated on sem */  
};
```

# System V semaphore ....

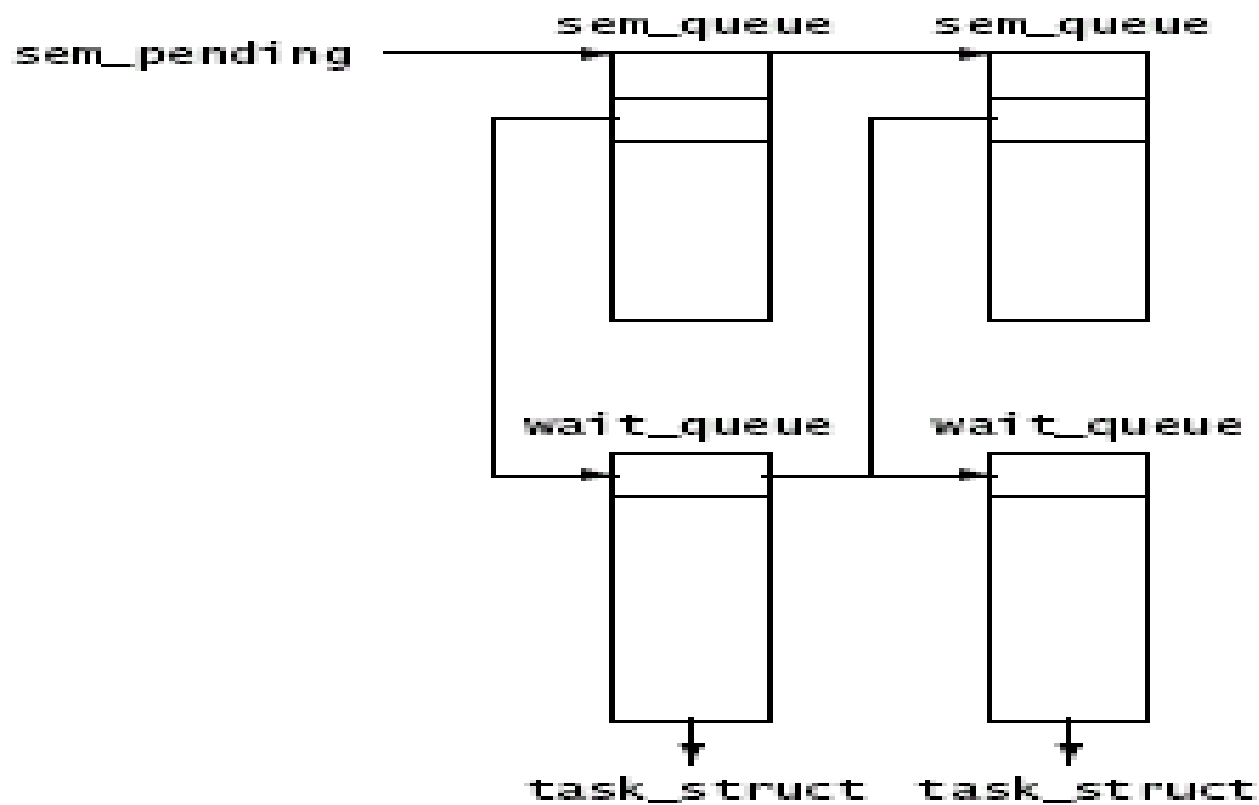


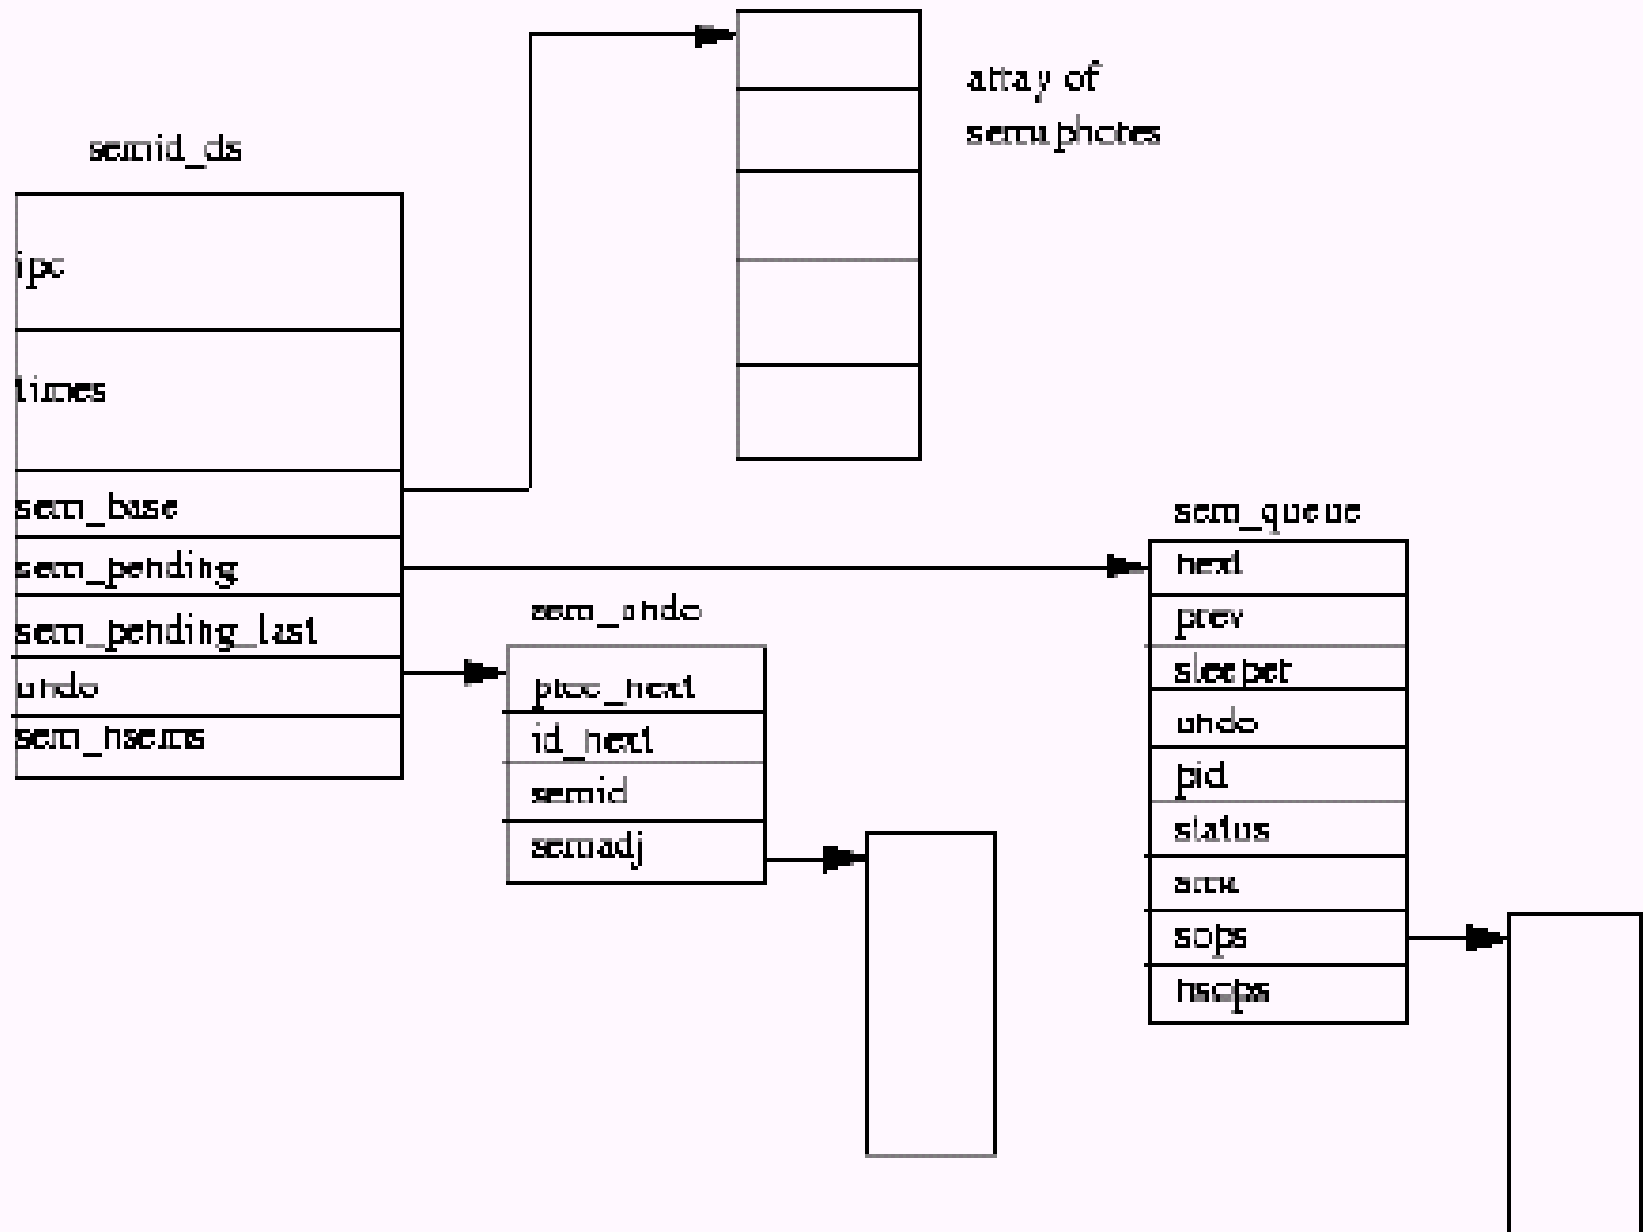
# System V semaphore .....

```
struct sem_queue{  
    struct sem_queue  *next;  
    struct wait_queue *sleepers;  
    struct semid_ds   *sma;  
    struct sembuf      *sops;  
    int                nsops;  
};
```




# System V semaphore .....





# System V message queues



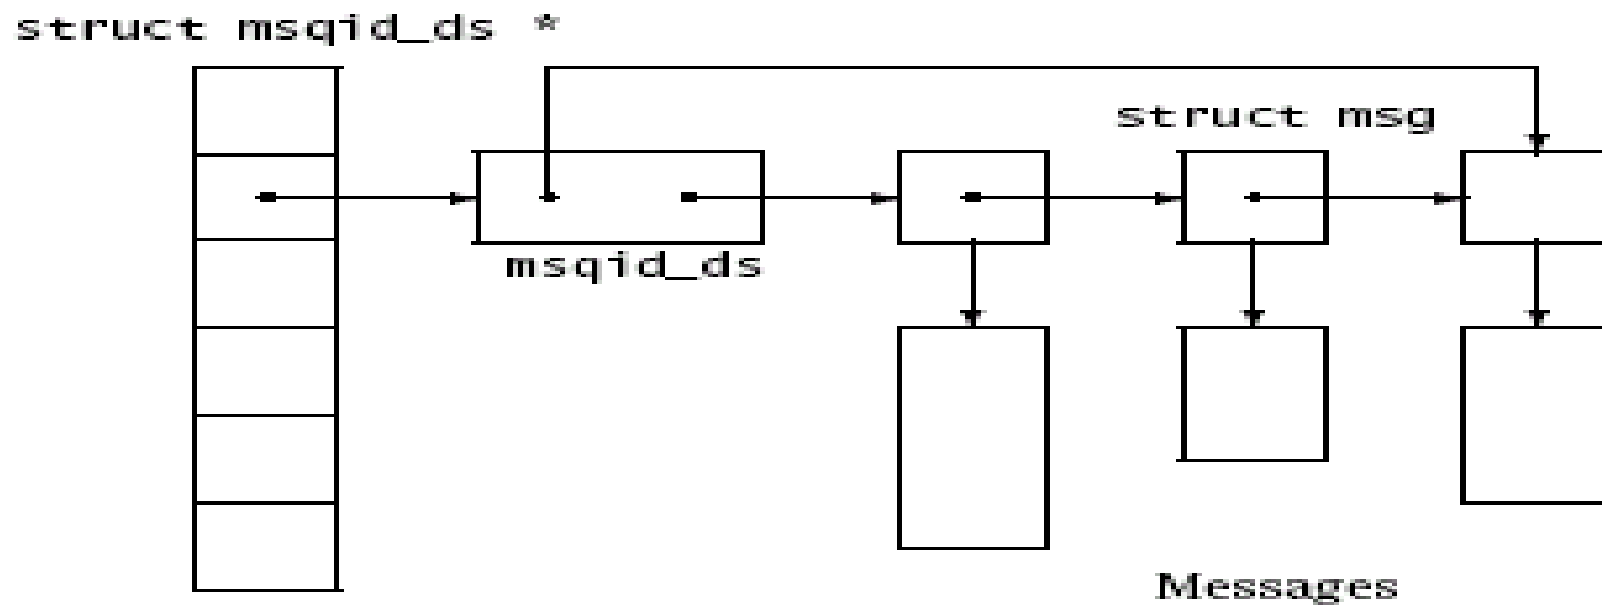
# System V message queues..

```
msqid_ds{  
    struct ipc_perm  msg_perm;  /* access permissions      */  
    struct msg        *msg_first; /* first message on queue */  
    struct msg        *msg_last; /* last message on queue  */  
    struct wait_queue *wwait;    /* blocked writing threads */  
    struct wait_queue *rwait;    /* blocked reading threads */  
    unsigned short    msg_qnum;  /* number of messages on queue */  
};
```

# System v message queues ...

```
struct msg{  
    struct msg *msg_next; /* next message on queue */  
    long      msg_type; /* as specified by sender */  
    char      *msg_spot; /* message text address */  
    time_t    msg_stime; /* msgsnd time */  
    short     msg_ts;    /* message text size */  
};
```

# System V message queues ....



# System V message queues .....

```
struct msgbuf{  
    long mtype; /* message type */  
    char mtext[]; /* message contents */  
};
```

msgid\_ds

ipc
*msg_last
*msg_first
times
*wwait
*rwait
msg_qhwm

msg

*msg_head
msg_type
*msg_appl
msg_stime
msg_ls
message

msg

*msg_head
msg_type
*msg_appl
msg_stime
msg_ls
message

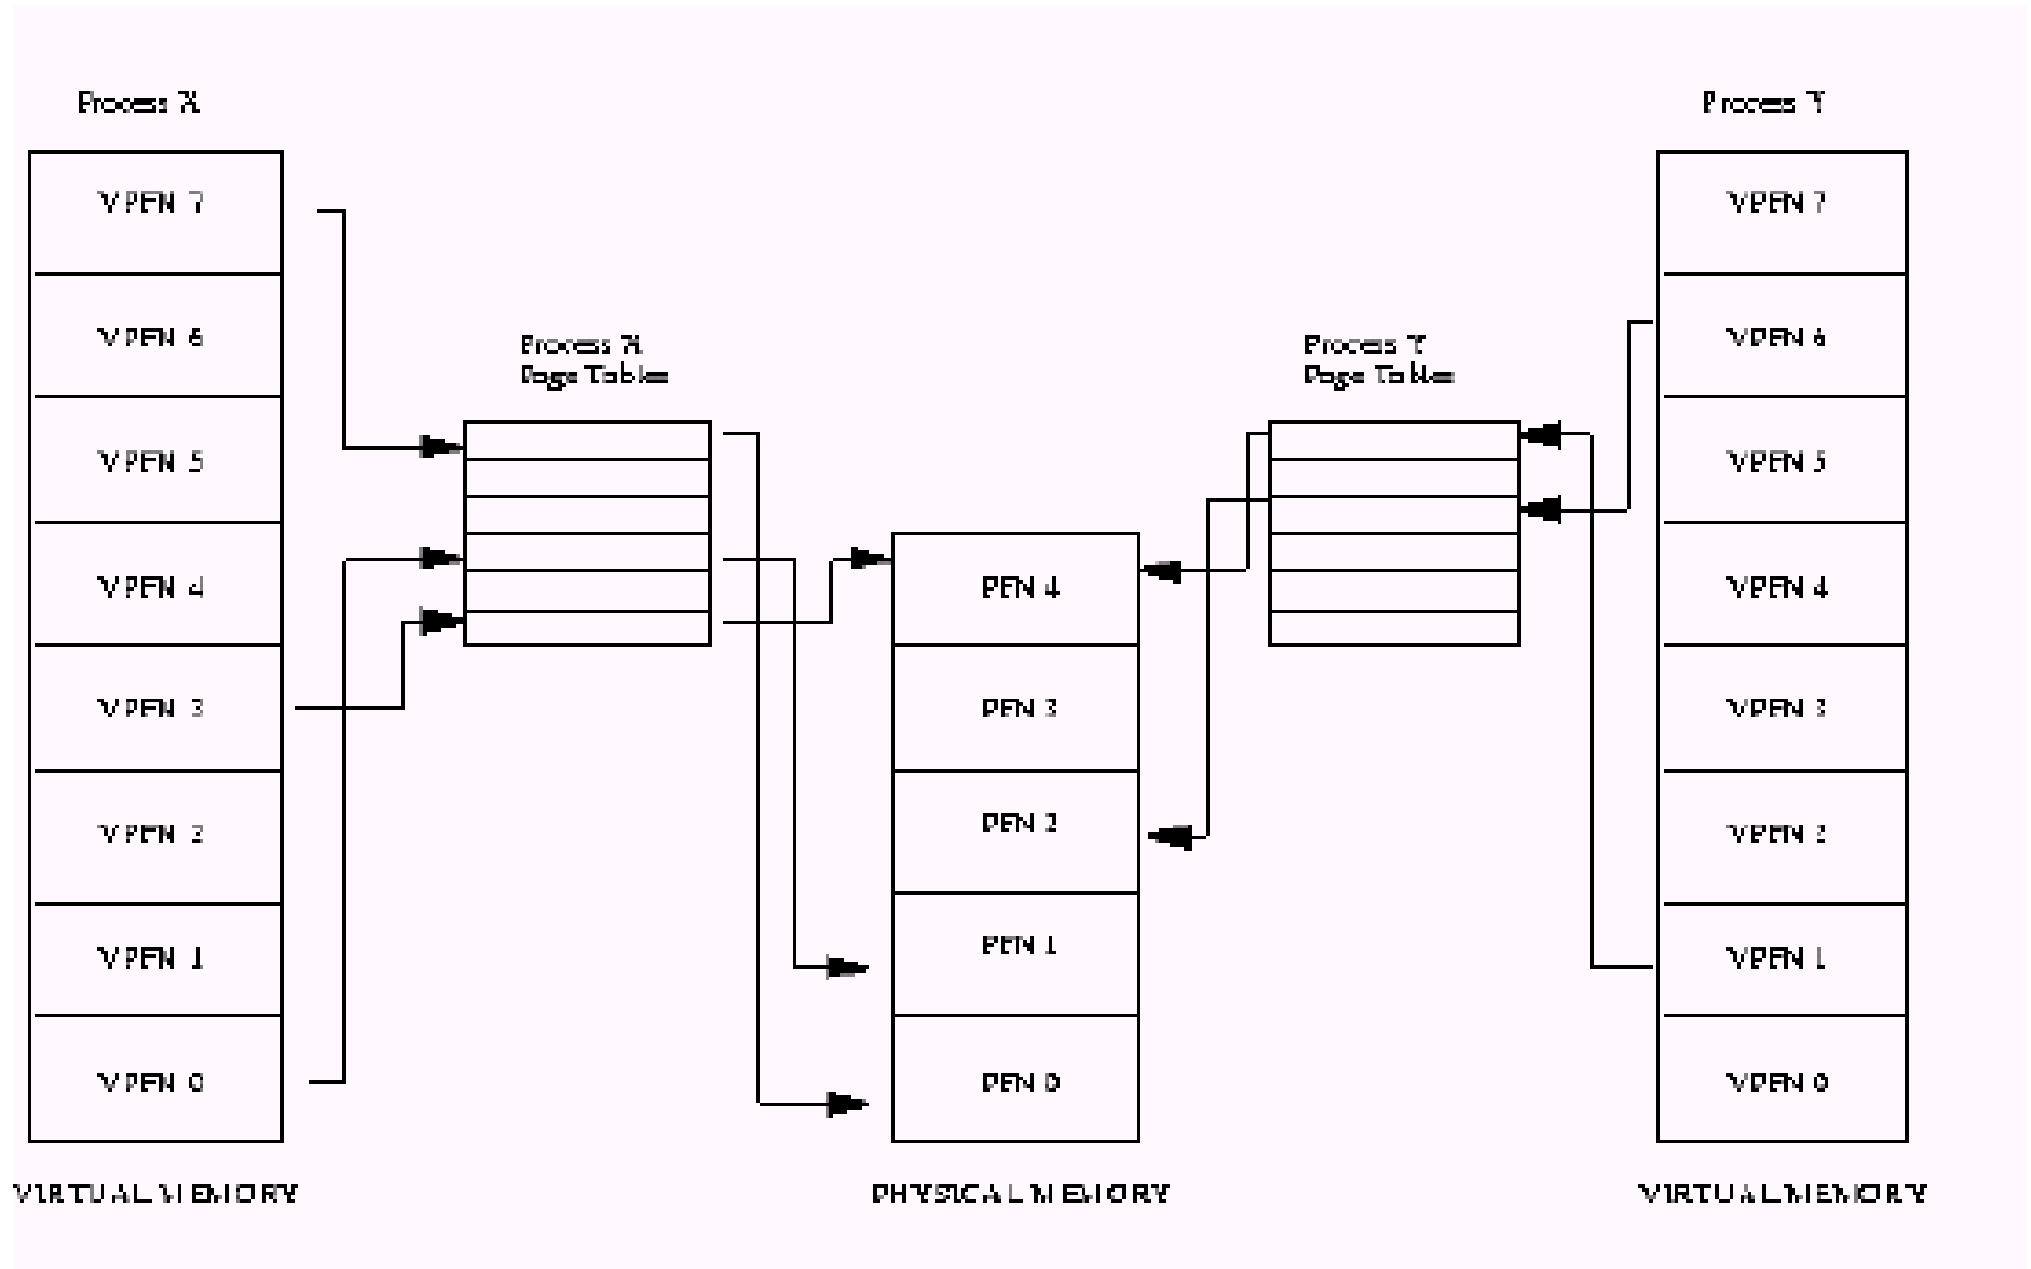
↑  
msg\_ls  
↓

↑  
msg\_ls  
↓

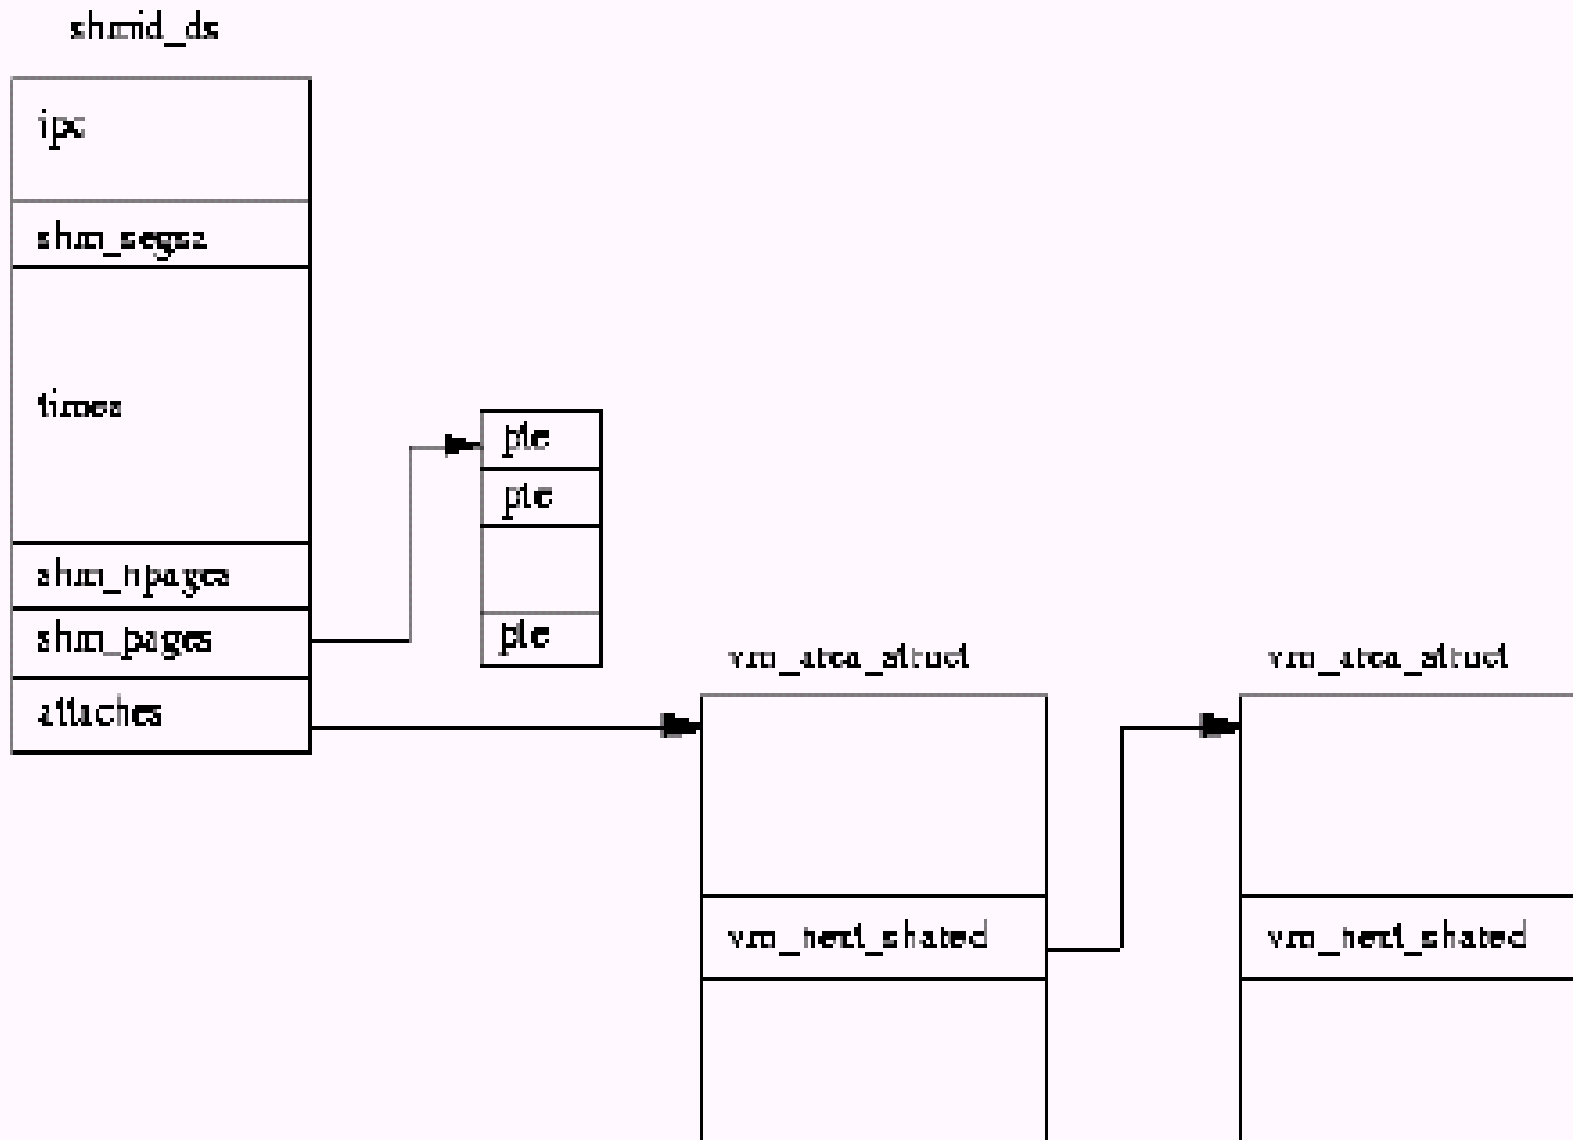
← msg\_qhwm →



# An Abstract Model of Virtual Memory



# System V



# Threading Model

<b>Per process items</b>	<b>Per thread items</b>
Address space Global variables Open files Child processes Pending alarms Signals and signal handlers Accounting information	Program counter Registers Stack State

# Threading Model

- Kernel-level threading with 1:1 model
- PD also acts as TD/LWP
- NPTL library and kernel support for  
POSIX Threading Standard
- CPU scheduling discussed earlier applies to  
multi-threading model